

BSD Unix 2.11 man entries
6/9/2019 - The ShadowTron Blog

Generated using the simh PDP-11/70 emulator with the PiDP11 Front Panel

PiDP11 - <https://obsolescence.wixsite.com/obsolescence/pidp-11>
SimH - <http://simh.trailing-edge.com/>
ShadowTronBlog - https://www.youtube.com/channel/UCtUiwjYcRS_u6Egc8iTkHNg
<http://shadowtron.com>
shadowtronblog@gmail.com

Manual Area covered

=====

--> 1 Commands and Application Programs
2 System Calls
3 C Library Subroutines
3F Fortran Library
4 Special Files
5 File Formats
6 Games
7 Miscellaneous
8 System Maintenance

**** Manual 1 - Commands and Application Programs ****

intro	introduction to commands
adb	debugger
addbib	create or extend bibliographic database
apply	apply a command to a set of arguments
apropos	locate commands by keyword lookup
ar	archive and library maintainer
as	VAX-11 assembler
at	execute commands at a later time
atq	print the queue of jobs waiting to be run
atrm	remove jobs spooled by at
awk	pattern scanning and processing language
basename	strip filename affixes
bc	arbitrary-precision arithmetic language
bib	bibliographic formatter; list bibliographic reference items
biff	be notified if mail arrives and who it is from
binmail	send or receive mail among users
cal	print calendar
calendar	reminder service
cat	catenate and print
cb	C program beautifier
cc	C compiler
cd	change working directory
checknr	check nroff/troff files
chfn	change password file information
chgrp	change group
chmod	change mode
chsh	change password file information
clear	clear terminal screen
cmp	compare two files
col	filter reverse line feeds
colcrt	filter nroff output for CRT previewing
colrm	remove columns from a file
comm	select or reject lines common to two sorted files
compress	compress and expand data
cp	copy
crypt	encode/decode
csh	a shell (command interpreter) with C-like syntax
ctags	create a tags file
date	print and set the date
dbx	debugger
dc	desk calculator
dd	convert and copy a file
deroff	remove nroff, troff, tbl and eqn constructs
df	disk free
diction	print wordy sentences; thesaurus for diction
diff	differential file and directory comparator
diff3	3-way differential file comparison
du	summarize disk usage
echo	echo arguments
ed	text editor
efl	Extended Fortran Language
eqn	typeset mathematics
error	analyze and disperse compiler error messages
ex	text editor
expand	expand tabs to spaces, and vice versa
expr	evaluate arguments as an expression

f77	Fortran 77 compiler
false	provide truth values
file	determine file type
find	find files
finger	ser information lookup program
fmt	simple text formatter
fold	fold long lines for finite width output device
fp	Functional Programming language compiler/interpreter
fpr	print Fortran file
from	who is my mail from?
fsplit	split a multi-routine Fortran file into individual files
ftp	ARPANET file transfer program
gcore	get core images of running processes
gprof	display call graph profile data
graph	draw a graph
grep	search a file for a pattern
groups	show group memberships
head	give first few lines
hostid	set or print identifier of current host system
hostname	set or print name of current host system
indent	indent and format C program source
install	install binaries
iostat	report I/O statistics
join	relational database operator
jove	an interactive display-oriented text editor
kill	terminate a process with extreme prejudice
last	indicate last logins of users and teletypes
lastcomm	show last commands executed in reverse order
ld	link editor
learn	computer aided instruction about UNIX
leave	remind you when you have to leave
lex	generator of lexical analysis programs
lint	a C program verifier
lisp	lisp interpreter
liszt	compile a Franz Lisp program
ln	make links
lock	reserve a terminal
logger	make entries in the system log
login	sign on
look	find lines in a sorted list
lookbib	build inverted index for a bibliography, find references in a bibliography
lorder	find ordering relation for an object library
lpq	spool queue examination program
lpr	off line print
lprm	remove jobs from the line printer spooling queue
lptest	generate lineprinter ripple pattern
ls	list contents of directory
lxref	lisp cross reference program
m4	macro processor
mail	send and receive mail
make	maintain program groups
man	find manual information by keywords; print out the manual
mesg	permit or deny messages
mh	Message Handler
mkdir	make a directory
mkstr	create an error message file by massaging C source
more	file perusal filter for crt viewing
mset	retrieve ASCII to IBM 3270 keyboard map
msgsg	system messages and junk mail program

mt	magnetic tape manipulating program
mv	move or rename files
netstat	show network status
newaliases	rebuild the data base for the mail aliases file
nice	run a command at low priority (\fIsh\fR only)
nm	print name list
notes	a news system
nroff	text formatting
od	octal, decimal, hex, ascii dump
pagesize	print system page size
passwd	change password file information
patch	a program for applying a diff file to an original
pc	Pascal compiler
pdx	pascal debugger
pi	Pascal interpreter code translator
pix	Pascal interpreter and executor
plot	graphics filters
pmerge	pascal file merger
pr	print file
printenv	print out the environment
prof	display profile data
ps	process status
ptx	permuted index
pwd	working directory name
px	Pascal interpreter
pxp	Pascal execution profiler
pxref	Pascal cross-reference program
quota	display disc usage and limits
ranlib	convert archives to random libraries
ratfor	rational Fortran dialect
rcp	remote file copy
rcsintro	introduction to RCS commands
rdist	remote file distribution program
readnews	read news articles
refer	find and insert literature references in documents
rev	reverse lines of a file
rlogin	remote login
rm	remove (unlink) files or directories
rmail	handle remote mail received via uucp
rmdir	remove (unlink) directories or files
roffbib	run off bibliographic database
rsh	remote shell
ruptime	show host status of local machines
rwho	who's logged in on local machines
sccs	front end for the SCCS subsystem
script	make typescript of terminal session
sed	stream editor
sendbug	mail a system bug report to 4bsd-bugs
sh	command language
size	size of an object file
sleep	suspend execution for an interval
soelim	eliminate \&.so's from nroff input
sort	sort or merge files
sortbib	sort bibliographic database
spell	find spelling errors
spline	interpolate smooth curve
split	split a file into pieces
strings	find the printable strings in a object, or other binary, file
strip	remove symbols and relocation bits

struct	structure Fortran programs
stty	set terminal options
style	analyze surface characteristics of a document
su	substitute user id temporarily
sum	sum and count blocks in a file
symorder	rearrange name list
sysline	display system status on status line of a terminal
systat	display system statistics on a crt
tabs	set terminal tabs
tail	deliver the last part of a file
talk	talk to another user
tar	tape archiver
tbl	format tables for nroff or troff
tc	phototypesetter simulator
tcopy	copy a mag tape
tee	pipe fitting
telnet	user interface to the TELNET protocol
test	condition command
tftp	trivial file transfer program
time	time a command
tip	connect to a remote system
tk	paginator for the Tektronix 4014
tn3270	full-screen remote login to IBM VM/CMS
touch	update date last modified of a file
tp	manipulate tape archive
tr	translate characters
troff	text formatting and typesetting
true	provide truth values
tset	terminal dependent initialization
tsort	topological sort
tty	get terminal name
ul	do underlining
unifdef	remove ifdef'ed lines
uniq	report repeated lines in a file
units	conversion program
uptime	show how long system has been up
users	compact list of users who are on the system
uucp	unix to unix copy
uuencode	encode/decode a binary file for transmission via mail
uulog	display UUCP log files
uuname	list names of UUCP hosts
uuq	examine or manipulate the uucp queue
uusend	send a file to a remote host
uux	unix to unix command execution
vacation	return ``I am on vacation'' indication
vgrind	grind nice listings of programs
vi	screen oriented (visual) display editor based on ex
vlp	Format Lisp programs to be printed with nroff, vtroff, or troff
vmstat	report virtual memory statistics
vnews	read news articles
vwidth	make troff width table for a font
w	who is on and what they are doing
wait	await completion of process
wall	write to all users
wc	word count
what	show what versions of object modules were used to construct a file
whatis	describe what a command is
whereis	locate source, binary, and or manual for program

which	locate a program file including aliases and paths (\fIcsh\fR only)
who	who is on the system
whoami	print effective current user id
whois	DARPA Internet user name directory service
window	window environment
write	write to another user
xsend	secret mail
xstr	extract strings from C programs to implement shared strings
yacc	yet another compiler-compiler
yes	be repetitively affirmative

NAME

intro - introduction to commands

This section describes publicly accessible commands in alphabetic order. Certain distinctions of purpose are made in the headings:

- (1) Commands of general utility.
- (1C) Commands for communication with other systems.
- (1G) Commands used primarily for graphics and computer-aided design.

N.B.: Commands related to system maintenance used to appear in section 1 manual pages and were distinguished by (1M) at the top of the page. These manual pages now appear in section 8.

SEE ALSO

Section (6) for computer games.

How to get started, in the Introduction.

DIAGNOSTICS

Upon termination each command returns two bytes of status, one supplied by the system giving the cause for termination, and (in the case of 'normal' termination) one supplied by the program, see wait and exit(2). The former byte is 0 for normal termination, the latter is customarily 0 for successful execution, nonzero to indicate troubles such as erroneous parameters, bad or inaccessible data, or other inability to cope with the task at hand. It is called variously 'exit code', 'exit status' or 'return code', and is described only where special conventions are involved.

NAME

adb - debugger (2BSD)

SYNOPSIS

adb [-w] [-k] [-Idir] [objfil [corfil]]

DESCRIPTION

Adb is a general purpose debugging program. It may be used to examine files and to provide a controlled environment for the execution of UNIX programs.

Objfil is normally an executable program file, preferably containing a symbol table; if not then the symbolic features of adb cannot be used although the file can still be examined. The default for objfil is a.out. Corfil is assumed to be a core image file produced after executing objfil; the default for corfil is core.

Requests to adb are read from the standard input and responses are to the standard output. If the -w flag is present then both objfil and corfil are created if necessary and opened for reading and writing so that files can be modified using adb.

The -k option makes adb do UNIX kernel memory mapping; it should be used when core is a UNIX crash dump or /dev/mem.

The -I option specifies a directory where files to be read with \$< or \$<< (see below) will be sought; the default is /usr/share/adb.

Adb ignores QUIT; INTERRUPT causes return to the next adb command.

In general requests to adb are of the form

```
[address] [, count] [command] [;]
```

If address is present then dot is set to address. Initially dot is set to 0. For most commands count specifies how many times the command will be executed. The default count is 1. Address and count are expressions.

The interpretation of an address depends on the context it is used in. If a subprocess is being debugged then addresses are interpreted in the usual way in the address space of the subprocess. For further details of address mapping see ADDRESSES.

EXPRESSIONS

. The value of dot.

`+` The value of dot incremented by the current increment.

`^` The value of dot decremented by the current increment.

`"` The last address typed.

`integer`

An octal number if integer begins with a 0; a hexadecimal number if preceded by #; otherwise a decimal number. This default interpretation of integers may be changed via the \$o and \$d commands.

`integer.fraction`

A 32 bit floating point number.

`'cccc'` The ASCII value of up to 4 characters. \ may be used to escape a '.

`< name` The value of name, which is either a variable name or a register name. Adb maintains a number of variables (see VARIABLES) named by single letters or digits. If name is a register name then the value of the register is obtained from the system header in corfil. The register names are those printed by the \$r command.

`symbol` A symbol is a sequence of upper or lower case letters, underscores or digits, not starting with a digit. The backslash character \ may be used to escape other characters. The value of the symbol is taken from the symbol table in objfil. An initial _ or ~ will be prepended to symbol if needed. If the symbol is a text symbol and objfil is an overlay type, the default is ~symbol, which is the subroutine itself, not the entry interface in the base segment.

`_ symbol`

In C, the 'true name' of an external symbol begins with _. It may be necessary to utter this name to distinguish it from internal or hidden variables of a program. For a subroutine in an overlay, ~symbol is the actual subroutine, and _symbol is the entry point in the base segment (the "thunk"). Note that to explicitly specify the local name, the ~ must be preceded by a backslash, since ~ is the bitwise complement operator: \~symbol must be typed.

`routine.name`

The address of the variable name in the specified C routine. Both routine and name are symbols. If name

is omitted the value is the address of the most recently activated C stack frame corresponding to routine.

(exp) The value of the expression exp.

Monadic operators

*exp The contents of the location addressed by exp in corfil.

@exp The contents of the location addressed by exp in objfil.

-exp Integer negation.

~exp Bitwise complement.

Dyadic operators are left associative and are less binding than monadic operators.

e1+e2 Integer addition.

e1-e2 Integer subtraction.

e1*e2 Integer multiplication.

e1%e2 Integer division.

e1&e2 Bitwise conjunction.

e1|e2 Bitwise disjunction.

e1#e2 E1 rounded up to the next multiple of e2.

COMMANDS

Most commands consist of a verb followed by a modifier or list of modifiers. The following verbs are available. (The commands '?' and '/' may be followed by '*'; see ADDRESSES for further details.)

?f Locations starting at address in objfil are printed according to the format f. dot is incremented by the sum of the increments for each format letter (q.v.).

/f Locations starting at address in corfil are printed according to the format f and dot is incremented as for '?'.

=f The value of address itself is printed in the styles indicated by the format f. (For i format '?' is printed for the parts of the instruction that reference

subsequent words.)

A format consists of one or more characters that specify a style of printing. Each format character may be preceded by a decimal integer that is a repeat count for the format character. While stepping through a format dot is incremented temporarily by the amount given for each format letter. If no format is given then the last format is used. The format letters available are as follows.

- o 2 Print 2 bytes in octal. All octal numbers output by adb are preceded by 0.
- O 4 Print 4 bytes in octal.
- q 2 Print in signed octal.
- Q 4 Print long signed octal.
- d 2 Print in decimal.
- D 4 Print long decimal.
- x 2 Print 2 bytes in hexadecimal.
- X 4 Print 4 bytes in hexadecimal.
- u 2 Print as an unsigned decimal number.
- U 4 Print long unsigned decimal.
- f 4 Print the 32 bit value as a floating point number.
- F 8 Print double floating point.
- b 1 Print the addressed byte in octal.
- c 1 Print the addressed character.
- C 1 Print the addressed character using the following escape convention. the standard escape convention where control characters are printed as ^X and the delete character is printed as ^?.
- s n Print the addressed characters until a zero character is reached.
- S n Print a string using the ^X escape convention (see C above). n is the length of the string including its zero terminator.
- Y 4 Print 4 bytes in date format (see ctime(3)).
- i n Print as machine instructions. n is the number of bytes occupied by the instruction. This style of printing causes variables 1 and 2 to be set to the offset parts of the source and destination respectively.
- a 0 Print the value of dot in symbolic form. Symbols are checked to ensure that they have an appropriate type as indicated below.
 - / local or global data symbol
 - ? local or global text symbol
 - = local or global absolute symbol
- p 2 Print the addressed value in symbolic form using the same rules for symbol lookup as a.
- t 0 When preceded by an integer tabs to the next appropriate tab stop. For example, 8t moves to

```

    the next 8-space tab stop.
r 0  Print a space.
n 0  Print a newline.
"... " 0
    Print the enclosed string.
^    Dot is decremented by the current increment.
    Nothing is printed.
+    Dot is incremented by 1.  Nothing is printed.
-    Dot is decremented by 1.  Nothing is printed.

```

newline

If the previous command temporarily incremented dot, make the increment permanent. Repeat the previous command with a count of 1.

[?/]l value mask

Words starting at dot are masked with mask and compared with value until a match is found. If L is used then the match is for 4 bytes at a time instead of 2. If no match is found then dot is unchanged; otherwise dot is set to the matched location. If mask is omitted then -1 is used.

[?/]w value ...

Write the 2-byte value into the addressed location. If the command is W, write 4 bytes. Odd addresses are not allowed when writing to the subprocess address space.

[?/]m b1 e1 f1[?/]

New values for (b1, e1, f1) are recorded. If less than three expressions are given then the remaining map parameters are left unchanged. If the '?' or '/' is followed by '*' then the last segment (b2,e2,f2) of the mapping is changed. If the list is terminated by '?' or '/' then the file (objfil or corfil respectively) is used for subsequent requests. (So that, for example, '/m?' will cause '/' to refer to objfil.)

>name

Dot is assigned to the variable or register named.

! A shell (/bin/sh) is called to read the rest of the line following '!'.

\$modifier

Miscellaneous commands. The available modifiers are:

```

<f  Read commands from the file f.  If this command is
    executed in a file, further commands in the file
    are not seen.  If f is omitted, the current input
    stream is terminated.  If a count is given, and is
    zero, the command will be ignored.  The value of

```

the count will be placed in variable 9 before the first command in f is executed.

<<f Similar to < except it can be used in a file of commands without causing the file to be closed. Variable 9 is saved during the execution of this command, and restored when it completes. There is a (small) finite limit to the number of << files that can be open at once.

>f Append output to the file f, which is created if it does not exist. If f is omitted, output is returned to the terminal.

? Print process id, the signal which caused stoppage or termination, as well as the registers as \$r. This is the default if modifier is omitted.

r Print the general registers and the instruction addressed by pc. Dot is set to pc.

f Print the floating registers in single or double length. If the floating point status of ps is set to double (0200 bit) then double length is used anyway.

b Print all breakpoints and their associated counts and commands.

a ALGOL 68 stack backtrace. If address is given then it is taken to be the address of the current frame (instead of r4). If count is given then only the first count frames are printed.

c C stack backtrace. If address is given then it is taken as the address of the current frame instead of the contents of the frame-pointer register. If C is used then the names and (16 bit) values of all automatic and static variables are printed for each active function. If count is given then only the first count frames are printed.

e The names and values of external variables are printed.

w Set the page width for output to address (default 80).

s Set the limit for symbol matches to address (default 255).

o All integers input are regarded as octal.

d Reset integer input as described in EXPRESSIONS.

q Exit from adb.

v Print all non zero variables in octal.

m Print the address map.

:modifier

Manage a subprocess. Available modifiers are:

bc Set breakpoint at address. If objfil is overlaid and address is in the overlay region, the breakpoint is set in the overlay of the symbol mentioned if address was given symbolically,

otherwise it is set in the current overlay (that in variable c). The breakpoint is executed count-1 times before causing a stop. Each time the breakpoint is encountered the command c is executed. If this command is omitted or sets dot to zero then the breakpoint causes a stop.

- d Delete breakpoint at address.
- r Run objfil as a subprocess. If address is given explicitly then the program is entered at this point; otherwise the program is entered at its standard entry point. count specifies how many breakpoints are to be ignored before stopping. Arguments to the subprocess may be supplied on the same line as the command. An argument starting with < or > causes the standard input or output to be established for the command.
- cs The subprocess is continued with signal s, see sigvec(2). If address is given then the subprocess is continued at this address. If no signal is specified then the signal that caused the subprocess to stop is sent. Breakpoint skipping is the same as for r.
- ss As for c except that the subprocess is single stepped count times. If there is no current subprocess then objfil is run as a subprocess as for r. In this case no signal can be sent; the remainder of the line is treated as arguments to the subprocess.
- k The current subprocess, if any, is terminated.

VARIABLES

Adb provides a number of variables. Named variables are set initially by adb but are not used subsequently (except for c). Numbered variables are reserved for communication as follows.

- 0 The last value printed.
- 1 The last offset part of an instruction source.
- 2 The previous value of variable 1.
- 9 The count on the last \$< or \$<< command.

On entry the following are set from the system header in the corfil. If corfil does not appear to be a core file then these values are set from objfil.

- b The base address of the data segment.
- c The current overlay. This is set from corfil, and in

turn sets the overlay map. This variable and the map may be changed by referring to other overlays symbolically (e.g. by examining text in subroutines in other overlays), using the \$r command (which resets c from the core file), or explicitly (with the command >c).

d The data segment size.
 e The entry point.
 m The 'magic' number (0405, 0407, 0410, 0411, 0430 or 0431). The object is overlaid if m is 0430 or 0431.
 o The sum of the overlay segment sizes (if overlaid).
 s The stack segment size.
 t The text segment size.

ADDRESSES

The address in a file associated with a written address is determined by a mapping associated with that file. Each mapping is represented by two or three triples (b1, e1, f1), (bo, eo, fo) for overlaid object files, and (b2, e2, f2) and the file address corresponding to a written address is calculated as follows.

b1<address<e1 => file address=address+f1-b1, otherwise,
 if overlaid,
 bo<address<eo => file address=address+fo-bo, otherwise,
 b2<address<e2 => file address=address+f2-b2,

otherwise, the requested address is not legal. In some cases (e.g. for programs with separated I and D space) the two segments for a file may overlap. If a ? or / is followed by an * then only the last triple is used.

The initial setting of both mappings is suitable for normal a.out and core files. If the object file is not of the kind expected then, for both files, b1 is set to 0, e1 is set to the maximum file size and f1 is set to 0; in this way the whole file can be examined with no address translation.

So that adb may be used on large files all appropriate values are kept as signed 32 bit integers.

FILES

a.out
 core

SEE ALSO

cc(1), ptrace(2), a.out(5), core(5)
 ADB - A Tutorial, S. R. Bourne

DIAGNOSTICS

`Adb' when there is no current command or format. Comments

about inaccessible files, syntax errors, abnormal termination of commands, etc. Exit status is 0, unless last command failed or returned nonzero status.

BUGS

When single stepping, system calls do not count as an executed instruction.

Local variables whose names are the same as an external variable may foul up the accessing of the external.

Local variables cannot be accessed in routines that are in overlays.

Since no shell is invoked to interpret the arguments of the `:r` command, the customary wild-card and variable expansions cannot occur.

NAME

addbib - create or extend bibliographic database

SYNOPSIS

addbib [-p promptfile] [-a] database

DESCRIPTION

When this program starts up, answering ``y'' to the initial ``Instructions?'' prompt yields directions; typing ``n'' or RETURN skips them. Addbib then prompts for various bibliographic fields, reads responses from the terminal, and sends output records to a database. A null response (just RETURN) means to leave out that field. A minus sign (-) means to go back to the previous field. A trailing backslash allows a field to be continued on the next line. The repeating ``Continue?'' prompt allows the user either to resume by typing ``y'' or RETURN, to quit the current session by typing ``n'' or ``q'', or to edit the database with any system editor (vi, ex, edit, ed).

The -a option suppresses prompting for an abstract; asking for an abstract is the default. Abstracts are ended with a CTRL-d. The -p option causes addbib to use a new prompting skeleton, defined in promptfile. This file should contain prompt strings, a tab, and the key-letters to be written to the database.

The most common key-letters and their meanings are given below. Addbib insulates you from these key-letters, since it gives you prompts in English, but if you edit the bibliography file later on, you will need to know this information.

%A	Author's name
%B	Book containing article referenced
%C	City (place of publication)
%D	Date of publication
%E	Editor of book containing article referenced
%F	Footnote number or label (supplied by refer)
%G	Government order number
%H	Header commentary, printed before reference
%I	Issuer (publisher)
%J	Journal containing article
%K	Keywords to use in locating reference
%L	Label field used by -k option of refer
%M	Bell Labs Memorandum (undefined)
%N	Number within volume
%O	Other commentary, printed at end of reference
%P	Page number(s)
%Q	Corporate or Foreign Author (unreversed)
%R	Report, paper, or thesis (unpublished)
%S	Series title

%T Title of article or book
%V Volume number
%X Abstract - used by roffbib, not by refer
%Y,Z ignored by refer

Except for `A', each field should be given just once. Only relevant fields should be supplied. An example is:

%A Bill Tuthill
%T Refer - A Bibliography System
%I Computing Services
%C Berkeley
%D 1982
%O UNX 4.3.5.

FILES

promptfile optional file to define prompting

SEE ALSO

refer(1), sortbib(1), roffbib(1), indxbib(1), lookbib(1)

AUTHORS

Al Stangenberger, Bill Tuthill

NAME

apply - apply a command to a set of arguments

SYNOPSIS

apply [-ac] [-n] command args ...

DESCRIPTION

Apply runs the named command on each argument arg in turn. Normally arguments are chosen singly; the optional number n specifies the number of arguments to be passed to command. If n is zero, command is run without arguments once for each arg. Character sequences of the form %d in command, where d is a digit from 1 to 9, are replaced by the d'th following unused arg. If any such sequences occur, n is ignored, and the number of arguments passed to command is the maximum value of d in command. The character '%' may be changed by the -a option.

Examples:

 apply echo *
is similar to ls(1);
 apply -2 cmp a1 b1 a2 b2 ...
compares the 'a' files to the 'b' files;
 apply -0 who 1 2 3 4 5
runs who(1) 5 times; and
 apply 'ln %1 /usr/joe' *
links all files in the current directory to the directory /usr/joe.

SEE ALSO

sh(1)

AUTHOR

Rob Pike

BUGS

Shell metacharacters in command may have bizarre effects; it is best to enclose complicated commands in single quotes ' '.

There is no way to pass a literal '%2' if '%' is the argument expansion character.

NAME

apropos - locate commands by keyword lookup

SYNOPSIS

apropos keyword ...

DESCRIPTION

Apropos shows which manual sections contain instances of any of the given keywords in their title. Each word is considered separately and case of letters is ignored. Words which are part of other words are considered; thus, when looking for compile, apropos will find all instances of 'compiler' also. Try

apropos password

and

apropos editor

If the line starts 'name(section) ...' you can do 'man section name' to get the documentation for it. Try 'apropos format' and then 'man 3s printf' to get the manual on the subroutine printf.

Apropos is actually just the -k option to the man(1) command.

FILES

/usr/man/whatis data base

SEE ALSO

man(1), whatis(1), catman(8)

AUTHOR

William Joy

NAME

ar - create and maintain library archives

SYNOPSIS

```
ar -d [-Tv] archive file ...
ar -m [-Tv] archive file ...
ar -m [-abiTv] position archive file ...
ar -p [-Tv] archive [file ...]
ar -q [-cTv] archive file ...
ar -r [-cuTv] archive file ...
ar -r [-abciuTv] position archive file ...
ar -t [-Tv] archive [file ...]
ar -x [-ouTv] archive [file ...]
```

DESCRIPTION

The ar utility creates and maintains groups of files combined into an archive. Once an archive has been created, new files can be added and existing files can be extracted, deleted, or replaced.

Files are named in the archive by a single component, i.e., if a file referenced by a path containing a slash ('/') is archived it will be named by the last component of that path. When matching paths listed on the command line against file names stored in the archive, only the last component of the path will be compared.

All informational and error messages use the path listed on the command line, if any was specified, otherwise the name in the archive is used. If multiple files in the archive have the same name, and paths are listed on the command line to ``select'' archive files for an operation, only the first file with a matching name will be selected.

The normal use of ar is for the creation and maintenance of libraries suitable for use with the loader (see ld(1)) although it is not restricted to this purpose. The options are as follows:

- a A positioning modifier used with the options -r and -m. The files are entered or moved after the archive member position, which must be specified.
- b A positioning modifier used with the options -r and -m. The files are entered or moved before the archive member position, which must be specified.
- c Whenever an archive is created, an informational message to that effect is written to standard error. If the -c option is specified, ar creates the archive silently.

- d Delete the specified archive files.
- i Identical to the -b option.
- m Move the specified archive files within the archive. If one of the options -a, -b or -i are specified, the files are moved before or after the position file in the archive. If none of those options are specified, the files are moved to the end of the archive.
- o Set the access and modification times of extracted files to the modification time of the file when it was entered into the archive. This will fail if the user is not the owner of the extracted file or the super-user.
- p Write the contents of the specified archive files to the standard output. If no files are specified, the contents of all the files in the archive are written in the order they appear in the archive.
- q (Quickly) append the specified files to the archive. If the archive does not exist a new archive file is created. Much faster than the -r option, when creating a large archive piece-by-piece, as no checking is done to see if the files already exist in the archive.
- r Replace or add the specified files to the archive. If the archive does not exist a new archive file is created. Files that replace existing files do not change the order of the files within the archive. New files are appended to the archive unless one of the options -a, -b or -i is specified.
- T Select and/or name archive members using only the first fifteen characters of the archive member or command line file name. The historic archive format had sixteen bytes for the name, but some historic archiver and loader implementations were unable to handle names that used the entire space. This means that file names that are not unique in their first fifteen characters can subsequently be confused. A warning message is printed to the standard error output if any file names are truncated. (See ar(5) for more information.)
- t List the specified files in the order in which they appear in the archive, each on a separate line. If no files are specified, all files in the archive are listed.
- u Update files. When used with the -r option, files in the archive will be replaced only if the disk file has

a newer modification time than the file in the archive. When used with the `-x` option, files in the archive will be extracted only if the archive file has a newer modification time than the file on disk.

- `-v` Provide verbose output. When used with the `-d`, `-m`, `-q` or `-x` options, `ar` gives a file-by-file description of the archive modification. This description consists of three, white-space separated fields: the option letter, a dash (``-'`) and the file name. When used with the `-r` option, `ar` displays the description as above, but the initial letter is an ``a'` if the file is added to the archive and an ``r'` if the file replaces a file already in the archive.

When used with the `-p` option, the name of each printed file is written to the standard output before the contents of the file, preceded by a single newline character, and followed by two newline characters, enclosed in less-than (``<'`) and greater-than (``>'`) characters.

When used with the `-t` option, `ar` displays an ``ls -l'` style listing of information about the members of the archive. This listing consists of eight, white-space separated fields: the file permissions (see `strmode(3)`), the decimal user and group ID's, separated by a single slash (``/'`), the file size (in bytes), the file modification time (in the `date(1)` format ``%b %e %H:%M %Y'`), and the name of the file.

- `-x` Extract the specified archive members into the files named by the command line arguments. If no members are specified, all the members of the archive are extracted into the current directory.

If the file does not exist, it is created; if it does exist, the owner and group will be unchanged. The file access and modification times are the time of the extraction (but see the `-o` option). The file permissions will be set to those of the file when it was entered into the archive; this will fail if the user is not the owner of the extracted file or the super-user.

The `ar` utility exits 0 on success, and >0 if an error occurs.

ENVIRONMENT

TMPDIR

The pathname of the directory to use when creating temporary files.

FILES

/tmp	default temporary file directory
ar.XXXXXX	temporary file names

COMPATIBILITY

By default, ar writes archives that may be incompatible with historic archives, as the format used for storing archive members with names longer than fifteen characters has changed. This implementation of ar is backward compatible with previous versions of ar in that it can read and write (using the -T option) historic archives. The -T option is provided for compatibility only, and will be deleted in a future release. See ar(5) for more information.

STANDARDS

The ar utility is expected to offer a superset of the POSIX 1003.2 functionality.

SEE ALSO

ld(1), ranlib(1), strmode(3), ar(5)

NAME

as - assembler

SYNOPSIS

as [-u] [-V] [-o objfile] file ...

DESCRIPTION

As assembles the concatenation of the named files. The options are:

- u Treat all undefined symbols in the assembly as external globals.
- V Produce an object suitable for loading into an automatic text overlaid program.
- o Use objfil for the name of the resultant object. If this is omitted, a.out is used. If no errors occurred during the assembly and if there were no unresolved external references, it is made executable.

The special file name -- serves two purposes. It signals the end of all options and causes stdin to be read for input. Thus it is now possible to pipe data to the assembler:

```
/lib/cpp -E foo.s | sed -e '^#;/;' | as -o foo.o --
```

The file name -- may be placed between normal files, when EOF is detected on stdin the next file in the argument list is opened and read.

If no input files are specified then stdin is read.

FILES

/tmp/atml temporary
a.out object

SEE ALSO

adb(1), ld(1), nm(1), a.out(5)
UNIX Assembler Manual by D. M. Ritchie

DIAGNOSTICS

When an input file cannot be read, its name followed by a question mark is typed and assembly ceases. When syntactic or semantic errors occur, a single-character diagnostic is typed out together with the line number and the file name in which it occurred. Errors in pass 1 cause cancellation of pass 2. The possible errors are:

) Parentheses error
] Parentheses error

<	String not terminated properly
*	Indirection used illegally
.	Illegal assignment to `.'
a	Error in address
b	Branch instruction is odd or too remote
e	Error in expression
f	Error in local (`f' or `b') type symbol
g	Garbage (unknown) character
i	End of file inside an if
m	Multiply defined symbol as label
o	Word quantity assembled at odd address
p	`.' different in pass 1 and 2
r	Relocation error
u	Undefined symbol
x	Syntax error

BUGS

Syntax errors can cause incorrect line numbers in following diagnostics.

NAME

at - execute commands at a later time

SYNOPSIS

at [-c] [-s] [-m] time [day] [file]

DESCRIPTION

At spools away a copy of the named file to be used as input to sh(1) or csh(1). If the -c flag (for (csh(1))) or the -s flag (for (sh(1))) is specified, then that shell will be used to execute the job; if no shell is specified, the current environment shell is used. If no file name is specified, at prompts for commands from standard input until a ^D is typed.

If the -m flag is specified, mail will be sent to the user after the job has been run. If errors occur during execution of the job, then a copy of the error diagnostics will be sent to the user. If no errors occur, then a short message is sent informing the user that no errors occurred.

The format of the spool file is as follows: A four line header that includes the owner of the job, the name of the job, the shell used to run the job, and whether mail will be set after the job is executed. The header is followed by a cd command to the current directory and a umask command to set the modes on any files created by the job. Then at copies all relevant environment variables to the spool file. When the script is run, it uses the user and group ID of the creator of the spool file.

The time is 1 to 4 digits, with an optional following 'A', 'P', 'N' or 'M' for AM, PM, noon or midnight. One and two digit numbers are taken to be hours, three and four digits to be hours and minutes. If no letters follow the digits, a 24 hour clock time is understood.

The optional day is either (1) a month name followed by a day number, or (2) a day of the week; if the word 'week' follows, invocation is moved seven days further off. Names of months and days may be recognizably truncated. Examples of legitimate commands are

```
at 8am jan 24
at -c -m 1530 fr week
at -s -m 1200n week
```

At programs are executed by periodic execution of the command /usr/libexec/atrun from cron(8). The granularity of at depends upon the how often atrun is executed.

Error output is lost unless redirected or the `-m` flag is requested, in which case a copy of the errors is sent to the user via `mail(1)`.

FILES

<code>/usr/spool/at</code>	spooling area
<code>/usr/spool/at/yy.ddd.hhhh.*</code>	job file
<code>/usr/spool/at/past</code>	directory where jobs are executed from
<code>/usr/spool/at/lasttimedone</code>	last time atrun was run
<code>/usr/libexec/atrun</code>	executor (run by <code>cron(8)</code>)

SEE ALSO

`atq(1)`, `atrm(1)`, `calendar(1)`, `sleep(1)`, `cron(8)`

DIAGNOSTICS

Complains about various syntax errors and times out of range.

BUGS

Due to the granularity of the execution of `/usr/libexec/atrun`, there may be bugs in scheduling things almost exactly 24 hours into the future.

If the system crashes, mail is not sent to the user informing them that the job was not completed.

Sometimes old spool files are not removed from the directory `/usr/spool/at/past`. This is usually due to a system crash, and requires that they be removed by hand.

NAME

atq - print the queue of jobs waiting to be run

SYNOPSIS

atq [-c] [-n] [name ...]

DESCRIPTION

Atq prints the queue of jobs that are waiting to be run at a later date. These jobs were created with the at(1) command. With no flags, the queue is sorted in the order that the jobs will be executed.

If the -c flag is used, the queue is sorted by the time that the at command was given.

The -n flag prints only the total number of files that are currently in the queue.

If a name(s) is provided, only those files belonging to that user(s) are displayed.

FILES

/usr/spool/at spool area

SEE ALSO

at(1), atrm(1), cron(8)

NAME

atrm - remove jobs spooled by at

SYNOPSIS

atrm [-f] [-i] [-] [[job #] [name]...]

DESCRIPTION

Atrm removes jobs that were created with the at(1) command. With the - flag, all jobs belonging to the person invoking atrm are removed. If a job number(s) is specified, atrm attempts to remove only that job number(s).

If the -f flag is used, all information regarding the removal of the specified jobs is suppressed. If the -i flag is used, atrm asks if a job should be removed; a response of 'y' causes the job to be removed.

If a user(s) name is specified, all jobs belonging to that user(s) are removed. This form of invoking atrm is useful only to the super-user.

FILES

/usr/spool/at spool area

SEE ALSO

at(1), atq(1), cron(8)

NAME

awk - pattern scanning and processing language

SYNOPSIS

awk [-Fc] [prog] [file] ...

DESCRIPTION

Awk scans each input file for lines that match any of a set of patterns specified in prog. With each pattern in prog there can be an associated action that will be performed when a line of a file matches the pattern. The set of patterns may appear literally as prog, or in a file specified as -f file.

Files are read in order; if there are no files, the standard input is read. The file name '-' means the standard input. Each line is matched against the pattern portion of every pattern-action statement; the associated action is performed for each matched pattern.

An input line is made up of fields separated by white space. (This default can be changed by using FS, vide infra.) The fields are denoted \$1, \$2, ... ; \$0 refers to the entire line.

A pattern-action statement has the form

```
pattern { action }
```

A missing { action } means print the line; a missing pattern always matches.

An action is a sequence of statements. A statement can be one of the following:

```
if ( conditional ) statement [ else statement ]
while ( conditional ) statement
for ( expression ; conditional ; expression ) statement
break
continue
{ [ statement ] ... }
variable = expression
print [ expression-list ] [ >expression ]
printf format [ , expression-list ] [ >expression ]
next # skip remaining patterns on this input line
exit # skip the rest of the input
```

Statements are terminated by semicolons, newlines or right braces. An empty expression-list stands for the whole line. Expressions take on string or numeric values as appropriate, and are built using the operators +, -, *, /, %, and concatenation (indicated by a blank). The C operators ++, --,

`+=`, `-=`, `*=`, `/=`, and `%=` are also available in expressions. Variables may be scalars, array elements (denoted `x[i]`) or fields. Variables are initialized to the null string. Array subscripts may be any string, not necessarily numeric; this allows for a form of associative memory. String constants are quoted "...".

The `print` statement prints its arguments on the standard output (or on a file if `>file` is present), separated by the current output field separator, and terminated by the output record separator. The `printf` statement formats its expression list according to the format (see `printf(3S)`).

The built-in function `length` returns the length of its argument taken as a string, or of the whole line if no argument. There are also built-in functions `exp`, `log`, `sqrt`, and `int`. The last truncates its argument to an integer. `substr(s, m, n)` returns the `n`-character substring of `s` that begins at position `m`. The function `sprintf(fmt, expr, expr, ...)` formats the expressions according to the `printf(3S)` format given by `fmt` and returns the resulting string.

Patterns are arbitrary Boolean combinations (`!`, `||`, `&&`, and parentheses) of regular expressions and relational expressions. Regular expressions must be surrounded by slashes and are as in `egrep`. Isolated regular expressions in a pattern apply to the entire line. Regular expressions may also occur in relational expressions.

A pattern may consist of two patterns separated by a comma; in this case, the action is performed for all lines between an occurrence of the first pattern and the next occurrence of the second.

A relational expression is one of the following:

```
expression matchop regular-expression
expression relop expression
```

where a `relop` is any of the six relational operators in C, and a `matchop` is either `~` (for contains) or `!~` (for does not contain). A conditional is an arithmetic expression, a relational expression, or a Boolean combination of these.

The special patterns `BEGIN` and `END` may be used to capture control before the first input line is read and after the last. `BEGIN` must be the first pattern, `END` the last.

A single character `c` may be used to separate the fields by starting the program with


```
BEGIN { FS = "c" }
```

or by using the -Fc option.

Other variable names with special meanings include NF, the number of fields in the current record; NR, the ordinal number of the current record; FILENAME, the name of the current input file; OFS, the output field separator (default blank); ORS, the output record separator (default newline); and OFMT, the output format for numbers (default "%.6g").

EXAMPLES

Print lines longer than 72 characters:

```
length > 72
```

Print first two fields in opposite order:

```
{ print $2, $1 }
```

Add up first column, print sum and average:

```
    { s += $1 }  
END  { print "sum is", s, " average is", s/NR }
```

Print fields in reverse order:

```
{ for (i = NF; i > 0; --i) print $i }
```

Print all lines between start/stop pairs:

```
/start/, /stop/
```

Print all lines whose first field is different from previous one:

```
$1 != prev { print; prev = $1 }
```

SEE ALSO

lex(1), sed(1)

A. V. Aho, B. W. Kernighan, P. J. Weinberger, Awk - a pattern scanning and processing language

BUGS

There are no explicit conversions between numbers and strings. To force an expression to be treated as a number add 0 to it; to force it to be treated as a string concatenate "" to it.

NAME

basename - strip filename affixes

SYNOPSIS

basename string [suffix]

DESCRIPTION

Basename deletes any prefix ending in '/' and the suffix, if present in string, from string, and prints the result on the standard output. It is normally used inside substitution marks `` in shell procedures.

This shell procedure invoked with the argument /usr/src/bin/cat.c compiles the named file and moves the output to cat in the current directory:

```
cc $1
mv a.out `basename $1 .c`
```

SEE ALSO

sh(1)

NAME

bc - arbitrary-precision arithmetic language

SYNOPSIS

bc [-c] [-l] [file ...]

DESCRIPTION

Bc is an interactive processor for a language which resembles C but provides unlimited precision arithmetic. It takes input from any files given, then reads the standard input. The -l argument stands for the name of an arbitrary precision math library. The syntax for bc programs is as follows; L means letter a-z, E means expression, S means statement.

Comments

are enclosed in /* and */.

Names

simple variables: L
array elements: L [E]
The words `ibase', `obase', and `scale'

Other operands

arbitrarily long numbers with optional sign and decimal point.
(E)
sqrt (E)
length (E) number of significant decimal digits
scale (E) number of digits right of decimal point
L (E , ... , E)

Operators

+ - * / % ^ (% is remainder; ^ is power)
++ -- (prefix and postfix; apply to names)
== <= >= != < >
= += -= *= /= %= ^=

Statements

E
{ S ; ... ; S }
if (E) S
while (E) S
for (E ; E ; E) S
null statement
break
quit

Function definitions

define L (L , ... , L) {
 auto L, ... , L
 S; ... S
}

```

    return ( E )
}

```

Functions in -l math library

```

s(x) sine
c(x) cosine
e(x) exponential
l(x) log
a(x) arctangent
j(n,x)    Bessel function

```

All function arguments are passed by value.

The value of a statement that is an expression is printed unless the main operator is an assignment. Either semicolons or newlines may separate statements. Assignment to scale influences the number of digits to be retained on arithmetic operations in the manner of dc(1). Assignments to ibase or obase set the input and output number radix respectively.

The same letter may be used as an array, a function, and a simple variable simultaneously. All variables are global to the program. 'Auto' variables are pushed down during function calls. When using arrays as function arguments or defining them as automatic variables empty square brackets must follow the array name.

For example

```

scale = 20
define e(x){
    auto a, b, c, i, s
    a = 1
    b = 1
    s = 1
    for(i=1; 1==1; i++){
        a = a*x
        b = b*i
        c = a/b
        if(c == 0) return(s)
        s = s+c
    }
}

```

defines a function to compute an approximate value of the exponential function and

```

    for(i=1; i<=10; i++) e(i)

```

prints approximate values of the exponential function of the first ten integers.

Bc is actually a preprocessor for dc(1), which it invokes automatically, unless the -c (compile only) option is present. In this case the dc input is sent to the standard output instead.

FILES

/usr/share/misc/lib.b mathematical library
dc(1) desk calculator proper

SEE ALSO

dc(1)
L. L. Cherry and R. Morris, BC - An arbitrary precision desk-calculator language

BUGS

No &&, ||, or ! operators.
For statement must have all three E's.
Quit is interpreted when read, not when executed.

NAME

biff - be notified if mail arrives and who it is from

SYNOPSIS

biff [yn]

DESCRIPTION

Biff informs the system whether you want to be notified when mail arrives during the current terminal session. The command

biff y

enables notification; the command

biff n

disables it. When mail notification is enabled, the header and first few lines of the message will be printed on your screen whenever mail arrives. A ``biff y'' command is often included in the file .login or .profile to be executed at each login.

Biff operates asynchronously. For synchronous notification use the MAIL variable of sh(1) or the mail variable of csh(1).

SEE ALSO

csh(1), sh(1), mail(1), comsat(8C)

NAME

binmail - send or receive mail among users

SYNOPSIS

```
/bin/mail [ + ] [ -i ] [ person ] ...  
/bin/mail [ + ] [ -i ] -f file
```

DESCRIPTION

Note: This is the old version 7 UNIX system mail program. The default mail command is described in Mail(1), and its binary is in the directory /usr/ucb.

mail with no argument prints a user's mail, message-by-message, in last-in, first-out order; the optional argument + displays the mail messages in first-in, first-out order. For each message, it reads a line from the standard input to direct disposition of the message.

newline

Go on to next message.

d Delete message and go on to the next.

p Print message again.

- Go back to previous message.

s [file] ...

Save the message in the named files ('mbox' default).

w [file] ...

Save the message, without a header, in the named files ('mbox' default).

m [person] ...

Mail the message to the named persons (yourself is default).

EOT (control-D)

Put unexamined mail back in the mailbox and stop.

q Same as EOT.

!command

Escape to the Shell to do command.

* Print a command summary.

An interrupt normally terminates the mail command; the mail file is unchanged. The optional argument -i tells mail to continue after interrupts.

When persons are named, mail takes the standard input up to an end-of-file (or a line with just `.`) and adds it to each person's `mail' file. The message is preceded by the sender's name and a postmark. Lines that look like postmarks are prepended with `>'. A person is usually a user name recognized by login(1). To denote a recipient on a remote system, prefix person by the system name and exclamation mark (see uucp(1C)).

The -f option causes the named file, for example, `mbox', to be printed as if it were the mail file.

When a user logs in he is informed of the presence of mail.

FILES

/etc/passwd	to identify sender and locate persons
/usr/spool/mail/*	incoming mail for user *
mbox	saved mail
/tmp/ma*	temp file
/usr/spool/mail/*.lock	lock for mail directory
dead.letter	unmailable text

SEE ALSO

Mail(1), write(1), uucp(1C), uux(1C), xsend(1), sendmail(8)

BUGS

Race conditions sometimes result in a failure to remove a lock file.

Normally anybody can read your mail, unless it is sent by xsend(1). An installation can overcome this by making mail a set-user-id command that owns the mail directory.

NAME

cal - print calendar

SYNOPSIS

cal [month] year

DESCRIPTION

Cal prints a calendar for the specified year. If a month is also specified, a calendar just for that month is printed. Year can be between 1 and 9999. The month is a number between 1 and 12. The calendar produced is that for England and her colonies.

Try September 1752.

BUGS

The year is always considered to start in January even though this is historically naive. Beware that `cal 78' refers to the early Christian era, not the 20th century.

NAME

calendar - reminder service

SYNOPSIS

calendar [-]

DESCRIPTION

Calendar consults the file ``calendar'` in the current directory and prints out lines that contain today's or tomorrow's date anywhere in the line. Most reasonable month-day dates such as ``Dec. 7,'` ``december 7,'` ``12/7,'` etc., are recognized, but not ``7 December'` or ``7/12'`. If you give the month as ```*''` with a date, i.e. ```* 1''`, that day in any month will do. On weekends ``tomorrow'` extends through Monday.

When an argument is present, calendar does its job for every user who has a file ``calendar'` in his login directory and sends him any positive results by mail(1). Normally this is done daily in the wee hours under control of cron(8).

The file ``calendar'` is first run through the ```C''` preprocessor, `/lib/cpp`, to include any other calendar files specified with the usual ```#include''` syntax. Included calendars will usually be shared by all users, maintained and documented by the local administration.

FILES

calendar
/usr/libexec/calendar to figure out today's and tomorrow's dates
/etc/passwd
/tmp/cal*
/lib/cpp, egrep, sed, mail as subprocesses

SEE ALSO

at(1), cron(8), mail(1)

BUGS

Calendar's extended idea of ``tomorrow'` doesn't account for holidays.

NAME

cat - concatenate and print

SYNOPSIS

cat [-u] [-n] [-s] [-v] file ...

DESCRIPTION

Cat reads each file in sequence and displays it on the standard output. Thus

```
cat file
```

displays the file on the standard output, and

```
cat file1 file2 >file3
```

concatenates the first two files and places the result on the third.

If no input file is given, or if the argument '-' is encountered, cat reads from the standard input file. Output is buffered in the block size recommended by stat(2) unless the standard output is a terminal, when it is line buffered. The -u option makes the output completely unbuffered.

The -n option displays the output lines preceded by lines numbers, numbered sequentially from 1. Specifying the -b option with the -n option omits the line numbers from blank lines.

The -s option crushes out multiple adjacent empty lines so that the output is displayed single spaced.

The -v option displays non-printing characters so that they are visible. Control characters print like ^X for control-x; the delete character (octal 0177) prints as ^?. Non-ascii characters (with the high bit set) are printed as M- (for meta) followed by the character of the low 7 bits. A -e option may be given with the -v option, which displays a '\$' character at the end of each line. Specifying the -t option with the -v option displays tab characters as ^I.

SEE ALSO

cp(1), ex(1), more(1), pr(1), tail(1)

BUGS

Beware of 'cat a b >a' and 'cat a b >b', which destroy the input files before reading them.

NAME

cb - C program beautifier

SYNOPSIS

cb

DESCRIPTION

Cb places a copy of the C program from the standard input on the standard output with spacing and indentation that displays the structure of the program.

NAME

cc - C compiler (2BSD)

SYNOPSIS

cc [option] ... file ...

DESCRIPTION

Cc is the UNIX C compiler. Cc accepts several types of arguments:

Arguments whose names end with ``.c'` are taken to be C source programs; they are compiled, and each object program is left on the file whose name is that of the source with ``.o'` substituted for ``.c'`. The ``.o'` file is normally deleted, however, if a single C program is compiled and loaded all at one go.

In the same way, arguments whose names end with ``.s'` are taken to be assembly source programs and are assembled, producing a ``.o'` file.

The following options are interpreted by cc. See `ld(1)` for load-time options.

- c Suppress the loading phase of the compilation, and force an object file to be produced even if only one program is compiled.
- w Suppress warning diagnostics.
- p Arrange for the compiler to produce code which counts the number of times each routine is called. If loading takes place, replace the standard startup routine by one which automatically calls `monitor(3)` at the start and arranges to write out a `mon.out` file at normal termination of execution of the object program. An execution profile can then be generated by use of `prof(1)`.
- O Invoke an object-code improver.
- S Compile the named C programs, and leave the assembler-language output on corresponding files suffixed ``.s'`.
- M Run only the macro preprocessor on the named C programs, requesting it to generate Makefile dependencies and send the result to the standard output.
- E Run only the macro preprocessor on the named C programs, and send the result to the standard output.

- C prevent the macro preprocessor from eliding comments.
- o output
 Name the final output file output. If this option is used the file 'a.out' will be left undisturbed.
- Dname=def
-Dname Define the name to the preprocessor, as if by '#define'. If no definition is given, the name is defined as "1".
- Uname Remove any initial definition of name.
- Idir '#include' files whose names do not begin with '/' are always sought first in the directory of the file argument, then in directories named in -I options, then in directories on a standard list.
- Ldir Library archives are sought first in directories named in -L options, then in directories on a standard list.
- Bstring
 Find substitute compiler passes in the files named string with the suffixes cpp, c0, c1 and c2. If string is empty, use a standard backup version.
- t[p012]
 Find only the designated compiler passes in the files whose names are constructed by a -B option. In the absence of a -B option, the string is taken to be '/usr/c/'.

Other arguments are taken to be either loader option arguments, or C-compatible object programs, typically produced by an earlier cc run, or perhaps libraries of C-compatible routines. These programs, together with the results of any compilations specified, are loaded (in the order given) to produce an executable program with name a.out.

FILES

file.c	input file
file.o	object file
a.out	loaded output
/tmp/ctm?	temporary
/lib/cpp	preprocessor
/lib/c[01]	compiler
/lib/c2	optional optimizer
/lib/crt0.o	runtime startoff
/lib/mcrt0.o	startoff for profiling
/lib/libc.a	standard library, see intro(3)

/usr/lib/libc_p.aprofiling library, see intro(3)
/usr/include standard directory for '#include' files
mon.out file produced for analysis by prof(1)

SEE ALSO

B. W. Kernighan and D. M. Ritchie, The C Programming
Language, Prentice-Hall, 1978
B. W. Kernighan, Programming in C-a tutorial
D. M. Ritchie, C Reference Manual
monitor(3), prof(1), adb(1), ld(1), as(1)

DIAGNOSTICS

The diagnostics produced by C itself are intended to be
self-explanatory. Occasional messages may be produced by
the assembler or loader.

BUGS

The compiler currently ignores advice to put char, unsigned
char, long, float, or double variables in registers.

NAME

cd - change working directory

SYNOPSIS

cd directory

DESCRIPTION

Directory becomes the new working directory. The process must have execute (search) permission in directory.

Because a new process is created to execute each command, cd would be ineffective if it were written as a normal command. It is therefore recognized and executed by the shells. In csh(1) you may specify a list of directories in which directory is to be sought as a subdirectory if it is not a subdirectory of the current directory; see the description of the cdp_{ath} variable in csh(1).

SEE ALSO

csh(1), sh(1), pwd(1), chdir(2)

NAME

checknr - check nroff/troff files

SYNOPSIS

```
checknr [ -s ] [ -f ] [ -a.x1.y1.x2.y2. ... .xn.yn ] [
-c.x1.x2.x3 ... .xn ] [ file ... ]
```

DESCRIPTION

Checknr checks a list of nroff(1) or troff(1) input files for certain kinds of errors involving mismatched opening and closing delimiters and unknown commands. If no files are specified, checknr checks the standard input. Delimiters checked are:

- (1) Font changes using \fx ... \fP.
- (2) Size changes using \sx ... \s0.
- (3) Macros that come in open ... close forms, for example, the .TS and .TE macros which must always come in pairs.

Checknr knows about the ms(7) and me(7) macro packages.

Additional pairs of macros can be added to the list using the -a option. This must be followed by groups of six characters, each group defining a pair of macros. The six characters are a period, the first macro name, another period, and the second macro name. For example, to define a pair .BS and .ES, use -a.BS.ES

The -c option defines commands which would otherwise be complained about as undefined.

The -f option requests checknr to ignore \f font changes.

The -s option requests checknr to ignore \s size changes.

Checknr is intended to be used on documents that are prepared with checknr in mind, much the same as lint. It expects a certain document writing style for \f and \s commands, in that each \fx must be terminated with \fP and each \sx must be terminated with \s0. While it will work to directly go into the next font or explicitly specify the original font or point size, and many existing documents actually do this, such a practice will produce complaints from checknr. Since it is probably better to use the \fP and \s0 forms anyway, you should think of this as a contribution to your document preparation style.

SEE ALSO

nroff(1), troff(1), checkeq(1), ms(7), me(7)

DIAGNOSTICS

Complaints about unmatched delimiters.

Complaints about unrecognized commands.

Various complaints about the syntax of commands.

BUGS

There is no way to define a 1 character macro name using -a.

Does not correctly recognize certain reasonable constructs,
such as conditionals.

NAME

chpass - add or change user database information

SYNOPSIS

chpass [-a list] [user]

DESCRIPTION

Chpass allows editing of the user database information associated with user or, by default, the current user. The information is formatted and supplied to an editor for changes. The vi editor will be used unless the environmental variable EDITOR selects an alternate editor. When the editor terminates, the information is re-read and used to update the user database itself. Only the user, or the super-user, may edit the information associated with the user.

Only the information that the user is allowed to change is displayed.

Possible display items are as follows:

Login:	user's login name
Password:	user's encrypted password
Uid:	user's id
Gid:	user's login group id
Change:	password change time
Expire:	account expiration time
Class:	user's general classification
Home Directory:	user's home directory
Shell:	user's login shell
Full Name:	user's real name
Location:	user's normal location
Home Phone:	user's home phone
Office Phone:	user's office phone

The login field is the user name used to access the computer account.

The password field contains the encrypted form of the user's password.

The uid field is the number associated with the login field. Both of these fields should be unique across the system (and often across a group of systems) as they control file access.

While it is possible to have multiple entries with identical login names and/or identical user id's, it is usually a mistake to do so. Routines that manipulate these files will often return only one of the multiple entries, and that one by random selection.

The group field is the group that the user will be placed in upon login. Since this system supports multiple groups (see groups(1)) this field currently has little special meaning. This field may be filled in with either a number or a group name (see group(5)).

The change field is the date by which the password must be changed.

The expire field is the date on which the account expires.

Both the change and expire fields should be entered in the form ``month day year'' where month is the month name (the first three characters are sufficient), day is the day of the month, and year is the year.

The class field is currently unused. In the near future it will be a key to a termcap(5) style database of user attributes.

The user's home directory is the full UNIX path name where the user will be placed on login.

The shell field is the command interpreter the user prefers. If the shell field is empty, the Bourne shell (/bin/sh) is assumed.

When altering a login shell, and not the super-user, the user must select an approved shell from the list in /etc/shells.

The last four fields are for storing the user's full name, office location, and home and work telephone numbers.

The super-user is also allowed to directly supply a user database entry, in the format specified by passwd(5), as an argument to the -a option. This argument must be a colon (``:``) separated list of all the user database fields, although they may be empty.

Once the information has been verified, chpasswd uses mkpasswd(8) to update the user database. This is run in the background, and, at very large sites could take several minutes. Until this update is completed, the password file is unavailable for other updates and the new information will not be available to programs.

FILES

/etc/master.passwd	The user database
/etc/shells	The list of approved shells

SEE ALSO

login(1), finger(1), getusershell(3), passwd(5),
mkpasswd(8), vipw(8)

Robert Morris and Ken Thompson, UNIX password security

BUGS

User information should (and eventually will) be stored
elsewhere.

NAME

chgrp - change group

SYNOPSIS

chgrp [-f -R] group file ...

DESCRIPTION

Chgrp changes the group-ID of the files to group. The group may be either a decimal GID or a group name found in the group-ID file.

The user invoking chgrp must belong to the specified group and be the owner of the file, or be the super-user.

No errors are reported when the -f (force) option is given.

When the -R option is given, chgrp recursively descends its directory arguments setting the specified group-ID. When symbolic links are encountered, their group is changed, but they are not traversed.

FILES

/etc/group

SEE ALSO

chown(2), passwd(5), group(5)

NAME

chmod - change mode

SYNOPSIS

chmod [-Rf] mode file ...

DESCRIPTION

The mode of each named file is changed according to mode, which may be absolute or symbolic. An absolute mode is an octal number constructed from the OR of the following modes:

4000	set user ID on execution
2000	set group ID on execution
1000	sticky bit, see chmod(2)
0400	read by owner
0200	write by owner
0100	execute (search in directory) by owner
0070	read, write, execute (search) by group
0007	read, write, execute (search) by others

A symbolic mode has the form:

[who] op permission [op permission] ...

The who part is a combination of the letters u (for user's permissions), g (group) and o (other). The letter a stands for all, or ugo. If who is omitted, the default is a but the setting of the file creation mask (see umask(2)) is taken into account.

Op can be + to add permission to the file's mode, - to take away permission and = to assign permission absolutely (all other bits will be reset).

Permission is any combination of the letters r (read), w (write), x (execute), X (set execute only if file is a directory or some other execute bit is set), s (set owner or group id) and t (save text - sticky). Letters u, g, or o indicate that permission is to be taken from the current mode. Omitting permission is only useful with = to take away all permissions.

When the -R option is given, chmod recursively descends its directory arguments setting the mode for each file as described above. When symbolic links are encountered, their mode is not changed and they are not traversed.

If the -f option is given, chmod will not complain if it fails to change the mode on a file.

EXAMPLES

The first example denies write permission to others, the second makes a file executable by all if it is executable by anyone:

```
chmod o-w file
chmod +X file
```

Multiple symbolic modes separated by commas may be given. Operations are performed in the order specified. The letter *s* is only useful with *u* or *g*.

Only the owner of a file (or the super-user) may change its mode.

SEE ALSO

`ls(1)`, `chmod(2)`, `stat(2)`, `umask(2)`, `chown(8)`

NAME

chpass - add or change user database information

SYNOPSIS

chpass [-a list] [user]

DESCRIPTION

Chpass allows editing of the user database information associated with user or, by default, the current user. The information is formatted and supplied to an editor for changes. The vi editor will be used unless the environmental variable EDITOR selects an alternate editor. When the editor terminates, the information is re-read and used to update the user database itself. Only the user, or the super-user, may edit the information associated with the user.

Only the information that the user is allowed to change is displayed.

Possible display items are as follows:

Login:	user's login name
Password:	user's encrypted password
Uid:	user's id
Gid:	user's login group id
Change:	password change time
Expire:	account expiration time
Class:	user's general classification
Home Directory:	user's home directory
Shell:	user's login shell
Full Name:	user's real name
Location:	user's normal location
Home Phone:	user's home phone
Office Phone:	user's office phone

The login field is the user name used to access the computer account.

The password field contains the encrypted form of the user's password.

The uid field is the number associated with the login field. Both of these fields should be unique across the system (and often across a group of systems) as they control file access.

While it is possible to have multiple entries with identical login names and/or identical user id's, it is usually a mistake to do so. Routines that manipulate these files will often return only one of the multiple entries, and that one by random selection.

The group field is the group that the user will be placed in upon login. Since this system supports multiple groups (see groups(1)) this field currently has little special meaning. This field may be filled in with either a number or a group name (see group(5)).

The change field is the date by which the password must be changed.

The expire field is the date on which the account expires.

Both the change and expire fields should be entered in the form ``month day year'' where month is the month name (the first three characters are sufficient), day is the day of the month, and year is the year.

The class field is currently unused. In the near future it will be a key to a termcap(5) style database of user attributes.

The user's home directory is the full UNIX path name where the user will be placed on login.

The shell field is the command interpreter the user prefers. If the shell field is empty, the Bourne shell (/bin/sh) is assumed.

When altering a login shell, and not the super-user, the user must select an approved shell from the list in /etc/shells.

The last four fields are for storing the user's full name, office location, and home and work telephone numbers.

The super-user is also allowed to directly supply a user database entry, in the format specified by passwd(5), as an argument to the -a option. This argument must be a colon (``:') separated list of all the user database fields, although they may be empty.

Once the information has been verified, chpasswd uses mkpasswd(8) to update the user database. This is run in the background, and, at very large sites could take several minutes. Until this update is completed, the password file is unavailable for other updates and the new information will not be available to programs.

FILES

/etc/master.passwd	The user database
/etc/shells	The list of approved shells

SEE ALSO

login(1), finger(1), getusershell(3), passwd(5),
mkpasswd(8), vipw(8)

Robert Morris and Ken Thompson, UNIX password security

BUGS

User information should (and eventually will) be stored
elsewhere.

NAME

clear - clear terminal screen

SYNOPSIS

clear

DESCRIPTION

Clear clears your screen if this is possible. It looks in the environment for the terminal type and then in /etc/termcap to figure out how to clear the screen.

FILES

/etc/termcap terminal capability data base

NAME

cmp - compare two files

SYNOPSIS

cmp [-l] [-s] file1 file2

DESCRIPTION

The two files are compared. (If file1 is '-', the standard input is used.) Under default options, cmp makes no comment if the files are the same; if they differ, it announces the byte and line number at which the difference occurred. If one file is an initial subsequence of the other, that fact is noted.

Options:

- l Print the byte number (decimal) and the differing bytes (octal) for each difference.
- s Print nothing for differing files; return codes only.

SEE ALSO

diff(1), comm(1)

DIAGNOSTICS

Exit code 0 is returned for identical files, 1 for different files, and 2 for an inaccessible or missing argument.

NAME

col - filter reverse line feeds

SYNOPSIS

col [-bfh]

DESCRIPTION

Col reads the standard input and writes the standard output. It performs the line overlays implied by reverse line feeds (ESC-7 in ASCII) and by forward and reverse half line feeds (ESC-9 and ESC-8). Col is particularly useful for filtering multicolumn output made with the `.rt` command of `nroff` and output resulting from use of the `tbl(1)` preprocessor.

Although `col` accepts half line motions in its input, it normally does not emit them on output. Instead, text that would appear between lines is moved to the next lower full line boundary. This treatment can be suppressed by the `-f` (fine) option; in this case the output from `col` may contain forward half line feeds (ESC-9), but will still never contain either kind of reverse line motion.

If the `-b` option is given, `col` assumes that the output device in use is not capable of backspacing. In this case, if several characters are to appear in the same place, only the last one read will be taken.

The control characters SO (ASCII code 017), and SI (016) are assumed to start and end text in an alternate character set. The character set (primary or alternate) associated with each printing character read is remembered; on output, SO and SI characters are generated where necessary to maintain the correct treatment of each character.

If the `-h` option is given, `col` converts white space to tabs to shorten printing time.

All control characters are removed from the input except space, backspace, tab, return, newline, ESC (033) followed by one of 7, 8, 9, SI, SO, and VT (013). This last character is an alternate form of full reverse line feed, for compatibility with some other hardware conventions. All other non-printing characters are ignored.

SEE ALSO

`troff(1)`, `tbl(1)`

BUGS

Can't back up more than 128 lines.
No more than 800 characters, including backspaces, on a line.

NAME

colcrt - filter nroff output for CRT previewing

SYNOPSIS

colcrt [-] [-2] [file ...]

DESCRIPTION

Colcrt provides virtual half-line and reverse line feed sequences for terminals without such capability, and on which overstriking is destructive. Half-line characters and underlining (changed to dashing '-') are placed on new lines in between the normal output lines.

The optional - suppresses all underlining. It is especially useful for previewing allboxed tables from tbl(1).

The option -2 causes all half-lines to be printed, effectively double spacing the output. Normally, a minimal space output format is used which will suppress empty lines. The program never suppresses two consecutive empty lines, however. The -2 option is useful for sending output to the line printer when the output contains superscripts and subscripts which would otherwise be invisible.

A typical use of colcrt would be

```
tbl exum2.n | nroff -ms | colcrt - | more
```

SEE ALSO

nroff/troff(1), col(1), more(1), ul(1)

BUGS

Should fold underlines onto blanks even with the '-' option so that a true underline character would show; if we did this, however, colcrt wouldn't get rid of cu'd underlining completely.

Can't back up more than 102 lines.

General overstriking is lost; as a special case '|' overstruck with '-' or underline becomes '+'.

Lines are trimmed to 132 characters.

Some provision should be made for processing superscripts and subscripts in documents which are already double-spaced.

NAME

colrm - remove columns from a file

SYNOPSIS

colrm [startcol [endcol]]

DESCRIPTION

Colrm removes selected columns from a file. Input is taken from standard input. Output is sent to standard output.

If called with one parameter the columns of each line will be removed starting with the specified column. If called with two parameters the columns from the first column to the last column will be removed.

Column numbering starts with column 1.

SEE ALSO

expand(1)

NAME

comm - select or reject lines common to two sorted files

SYNOPSIS

comm [- [123]] file1 file2

DESCRIPTION

Comm reads file1 and file2, which should be ordered in ASCII collating sequence, and produces a three column output: lines only in file1; lines only in file2; and lines in both files. The filename '-' means the standard input.

Flags 1, 2, or 3 suppress printing of the corresponding column. Thus comm -12 prints only the lines common to the two files; comm -23 prints only lines in the first file but not in the second; comm -123 is a no-op.

SEE ALSO

cmp(1), diff(1), uniq(1)

NAME

compress, uncompress, zcat - compress and expand data

SYNOPSIS

```
compress [ -f ] [ -v ] [ -c ] [ -b bits ] [ name ... ]
uncompress [ -f ] [ -v ] [ -c ] [ name ... ]
zcat [ name ... ]
```

DESCRIPTION

Compress reduces the size of the named files using adaptive Lempel-Ziv coding. Whenever possible, each file is replaced by one with the extension .Z, while keeping the same ownership modes, access and modification times. If no files are specified, the standard input is compressed to the standard output. Compressed files can be restored to their original form using uncompress or zcat.

The -f option will force compression of name, even if it does not actually shrink or the corresponding name.Z file already exists. Except when run in the background under /bin/sh, if -f is not given the user is prompted as to whether an existing name.Z file should be overwritten.

The -c ('`cat`') option makes compress/uncompress write to the standard output; no files are changed. The nondestructive behavior of zcat is identical to that of uncompress -c.

Compress uses the modified Lempel-Ziv algorithm popularized in "A Technique for High Performance Data Compression", Terry A. Welch, IEEE Computer, vol. 17, no. 6 (June 1984), pp. 8-19. Common substrings in the file are first replaced by 9-bit codes 257 and up. When code 512 is reached, the algorithm switches to 10-bit codes and continues to use more bits until the limit specified by the -b flag is reached (default 16). Bits must be between 9 and 16. The default can be changed in the source to allow compress to be run on a smaller machine.

After the bits limit is attained, compress periodically checks the compression ratio. If it is increasing, compress continues to use the existing code dictionary. However, if the compression ratio decreases, compress discards the table of substrings and rebuilds it from scratch. This allows the algorithm to adapt to the next "block" of the file.

Note that the -b flag is omitted for uncompress, since the bits parameter specified during compression is encoded within the output, along with a magic number to ensure that neither decompression of random data nor recompression of compressed data is attempted.

The amount of compression obtained depends on the size of the input, the number of bits per code, and the distribution of common substrings. Typically, text such as source code or English is reduced by 50-60%. Compression is generally much better than that achieved by Huffman coding (as used in pack), or adaptive Huffman coding (compact), and takes less time to compute.

The -v option causes the printing of the percentage reduction of each file.

If an error occurs, exit status is 1, else if the last file was not compressed because it became larger, the status is 2; else the status is 0.

DIAGNOSTICS

Usage: compress [-fvc] [-b maxbits] [file ...]
Invalid options were specified on the command line.
Missing maxbits
Maxbits must follow -b.
file: not in compressed format
The file specified to uncompress has not been compressed.
file: compressed with xx bits, can only handle yy bits
File was compressed by a program that could deal with more bits than the compress code on this machine. Recompress the file with smaller bits.
file: already has .Z suffix -- no change
The file is assumed to be already compressed. Rename the file and try again.
file: filename too long to tack on .Z
The file cannot be compressed because its name is longer than 12 characters. Rename and try again. This message does not occur on BSD systems.
file already exists; do you wish to overwrite (y or n)?
Respond "y" if you want the output file to be replaced; "n" if not.
uncompress: corrupt input
A SIGSEGV violation was detected which usually means that the input file is corrupted.
Compression: xx.xx%
Percentage of the input saved by compression. (Relevant only for -v.)
-- not a regular file: unchanged
When the input file is not a regular file, (e.g. a directory), it is left unaltered.
-- has xx other links: unchanged
The input file has links; it is left unchanged. See ln(1) for more information.
-- file unchanged
No savings is achieved by compression. The input remains virgin.

BUGS

Although compressed files are compatible between machines with large memory, `-b12` should be used for file transfer to architectures with a small process data space (64KB or less, as exhibited by the DEC PDP series, the Intel 80286, etc.)

`compress` should be more flexible about the existence of the ``.Z'` suffix.

NAME

cp - copy

SYNOPSIS

cp [-ip] file1 file2

cp [-ipr] file ... directory

DESCRIPTION

File1 is copied onto file2. By default, the mode and owner of file2 are preserved if it already existed; otherwise the mode of the source file modified by the current umask(2) is used. The -p option causes cp to attempt to preserve (duplicate) in its copies the modification times and modes of the source files, ignoring the present umask.

In the second form, one or more files are copied into the directory with their original file-names.

Cp refuses to copy a file onto itself.

If the -i option is specified, cp will prompt the user with the name of the file whenever the copy will cause an old file to be overwritten. An answer of 'y' will cause cp to continue. Any other answer will prevent it from overwriting the file.

If the -r option is specified and any of the source files are directories, cp copies each subtree rooted at that name; in this case the destination must be a directory.

SEE ALSO

cat(1), mv(1), rcp(1C)

NAME

csch - a shell (command interpreter) with C-like syntax

SYNOPSIS

csch [-cefinstvVxX] [arg ...]

DESCRIPTION

Csh is a first implementation of a command language interpreter incorporating a history mechanism (see History Substitutions), job control facilities (see Jobs), interactive file name and user name completion (see File Name Completion), and a C-like syntax. So as to be able to use its job control facilities, users of csh must (and automatically) use the new tty driver fully described in tty(4). This new tty driver allows generation of interrupt characters from the keyboard to tell jobs to stop. See stty(1) for details on setting options in the new tty driver.

An instance of csh begins by executing commands from the file ``.cshrc'` in the home directory of the invoker. If this is a login shell then it also executes commands from the file ``.login'` there. It is typical for users on crt's to put the command ```stty crt''` in their `.login` file, and to also invoke `tset(1)` there.

In the normal case, the shell will then begin reading commands from the terminal, prompting with ``% '.` Processing of arguments and the use of the shell to process files containing command scripts will be described later.

The shell then repeatedly performs the following actions: a line of command input is read and broken into words. This sequence of words is placed on the command history list and then parsed. Finally each command in the current line is executed.

When a login shell terminates it executes commands from the file ``.logout'` in the users home directory.

Lexical structure

The shell splits input lines into words at blanks and tabs with the following exceptions. The characters ``&' `|' `;'`
`<' `>' `(' `)'` form separate words. If doubled in `&&'`
`||'`, `<<'` or `>>'` these pairs form single words. These
parser metacharacters may be made part of other words, or
prevented their special meaning, by preceding them with `\`.
A newline preceded by a `\` is equivalent to a blank.`

In addition strings enclosed in matched pairs of quotations, ``''`, ``'` or ````, form parts of a word; metacharacters in
these strings, including blanks and tabs, do not form`

separate words. These quotations have semantics to be described subsequently. Within pairs of ``'`` or ``"``` characters a newline preceded by a `\\`` gives a true newline character.

When the shell's input is not a terminal, the character ``#`` introduces a comment which continues to the end of the input line. It is prevented this special meaning when preceded by `\\`` and in quotations using ``'``, ``"```, and ``"```.

Commands

A simple command is a sequence of words, the first of which specifies the command to be executed. A simple command or a sequence of simple commands separated by ``|`` characters forms a pipeline. The output of each command in a pipeline is connected to the input of the next. Sequences of pipelines may be separated by ``;``, and are then executed sequentially. A sequence of pipelines may be executed without immediately waiting for it to terminate by following it with an ``&``.

Any of the above may be placed in ``(' `)'`` to form a simple command (which may be a component of a pipeline, etc.) It is also possible to separate pipelines with ``||`` or ``&&`` indicating, as in the C language, that the second is to be executed only if the first fails or succeeds respectively. (See Expressions.)

Jobs

The shell associates a job with each pipeline. It keeps a table of current jobs, printed by the jobs command, and assigns them small integer numbers. When a job is started asynchronously with ``&``, the shell prints a line which looks like:

```
[1] 1234
```

indicating that the job which was started asynchronously was job number 1 and had one (top-level) process, whose process id was 1234.

If you are running a job and wish to do something else you may hit the key `^Z` (control-Z) which sends a STOP signal to the current job. The shell will then normally indicate that the job has been ``Stopped``, and print another prompt. You can then manipulate the state of this job, putting it in the background with the `bg` command, or run some other commands and then eventually bring the job back into the foreground with the foreground command `fg`. A `^Z` takes effect immediately and is like an interrupt in that pending output and

unread input are discarded when it is typed. There is another special key `^Y` which does not generate a STOP signal until a program attempts to read(2) it. This can usefully be typed ahead when you have prepared some commands for a job which you wish to stop after it has read them.

A job being run in the background will stop if it tries to read from the terminal. Background jobs are normally allowed to produce output, but this can be disabled by giving the command ```stty tostop''`. If you set this tty option, then background jobs will stop when they try to produce output like they do when they try to read input.

There are several ways to refer to jobs in the shell. The character ``%'` introduces a job name. If you wish to refer to job number 1, you can name it as ``%1'`. Just naming a job brings it to the foreground; thus ``%1'` is a synonym for ``fg %1'`, bringing job 1 back into the foreground. Similarly saying ``%1 &'` resumes job 1 in the background. Jobs can also be named by prefixes of the string typed in to start them, if these prefixes are unambiguous, thus ``%ex'` would normally restart a suspended `ex(1)` job, if there were only one suspended job whose name began with the string ``ex'`. It is also possible to say ``%?string'` which specifies a job whose text contains string, if there is only one such job.

The shell maintains a notion of the current and previous jobs. In output pertaining to jobs, the current job is marked with a ``+'` and the previous job with a ``-'`. The abbreviation ``%+'` refers to the current job and ``%-'` refers to the previous job. For close analogy with the syntax of the history mechanism (described below), ``%%'` is also a synonym for the current job.

Status reporting

This shell learns immediately whenever a process changes state. It normally informs you whenever a job becomes blocked so that no further progress is possible, but only just before it prints a prompt. This is done so that it does not otherwise disturb your work. If, however, you set the shell variable `notify`, the shell will notify you immediately of changes of status in background jobs. There is also a shell command `notify` which marks a single process so that its status changes will be immediately reported. By default `notify` marks the current process; simply say ``notify'` after starting a background job to mark it.

When you try to leave the shell while jobs are stopped, you will be warned that ``You have stopped jobs.'` You may use the `jobs` command to see what they are. If you do this or immediately try to exit again, the shell will not warn you a

second time, and the suspended jobs will be terminated.

File Name Completion

When the file name completion feature is enabled by setting the shell variable `filec` (see `set`), `cs`h will interactively complete file names and user names from unique prefixes, when they are input from the terminal followed by the escape character (the escape key, or `control-[]`). For example, if the current directory looks like

DSC.OLD	bin	cmd	lib	xmpl.c
DSC.NEW	chaosnet	cmtest	mail	xmpl.o
bench	class	dev	mbx	xmpl.out

and the input is

```
% vi ch<escape>
```

`cs`h will complete the prefix ```ch``` to the only matching file name ```chaosnet```, changing the input line to

```
% vi chaosnet
```

However, given

```
% vi D<escape>
```

`cs`h will only expand the input to

```
% vi DSC.
```

and will sound the terminal bell to indicate that the expansion is incomplete, since there are two file names matching the prefix ```D```.

If a partial file name is followed by the end-of-file character (usually `control-D`), then, instead of completing the name, `cs`h will list all file names matching the prefix. For example, the input

```
% vi D<control-D>
```

causes all files beginning with ```D``` to be listed:

```
DSC.NEW    DSC.OLD
```

while the input line remains unchanged.

The same system of escape and end-of-file can also be used to expand partial user names, if the word to be completed (or listed) begins with the character ```~```. For example, typing

```
cd ~ro<control-D>
```

may produce the expansion

```
cd ~root
```

The use of the terminal bell to signal errors or multiple matches can be inhibited by setting the variable `nobeep`.

Normally, all files in the particular directory are candidates for name completion. Files with certain suffixes can be excluded from consideration by setting the variable `figignore` to the list of suffixes to be ignored. Thus, if `figignore` is set by the command

```
% set figignore = (.o .out)
```

then typing
 % vi x<escape>
would result in the completion to
 % vi xmpl.c
ignoring the files "xmpl.o" and "xmpl.out". However, if the only completion possible requires not ignoring these suffixes, then they are not ignored. In addition, fignore does not affect the listing of file names by control-D. All files are listed regardless of their suffixes.

Substitutions

We now describe the various transformations the shell performs on the input in the order in which they occur.

History substitutions

History substitutions place words from previous command input as portions of new commands, making it easy to repeat commands, repeat arguments of a previous command in the current command, or fix spelling mistakes in the previous command with little typing and a high degree of confidence. History substitutions begin with the character `!' and may begin anywhere in the input stream (with the proviso that they do not nest.) This `!' may be preceded by an `\' to prevent its special meaning; for convenience, a `!' is passed unchanged when it is followed by a blank, tab, new-line, `=' or `(' . (History substitutions also occur when an input line begins with `^'. This special abbreviation will be described later.) Any input line which contains history substitution is echoed on the terminal before it is executed as it could have been typed without history substitution.

Commands input from the terminal which consist of one or more words are saved on the history list. The history substitutions reintroduce sequences of words from these saved commands into the input stream. The size of which is controlled by the history variable; the previous command is always retained, regardless of its value. Commands are numbered sequentially from 1.

For definiteness, consider the following output from the history command:

```
9  write michael
10 ex write.c
11 cat oldwrite.c
12 diff *write.c
```

The commands are shown with their event numbers. It is not usually necessary to use event numbers, but the current event number can be made part of the prompt by placing an

`!' in the prompt string.

With the current event 13 we can refer to previous events by event number `!11', relatively as in `!-2' (referring to the same event), by a prefix of a command word as in `!d' for event 12 or `!wri' for event 9, or by a string contained in a word in the command as in `!?mic?' also referring to event 9. These forms, without further modification, simply reintroduce the words of the specified events, each separated by a single blank. As a special case `!!' refers to the previous command; thus `!!' alone is essentially a redo.

To select words from an event we can follow the event specification by a `:' and a designator for the desired words. The words of an input line are numbered from 0, the first (usually command) word being 0, the second word (first argument) being 1, etc. The basic word designators are:

```
0    first (command) word
n    n'th argument
^    first argument, i.e. `1'
$    last argument
%    word matched by (immediately preceding) ?s? search
x-y  range of words
-y   abbreviates `0-y'
*    abbreviates `^-$', or nothing if only 1 word in event
x*   abbreviates `x-$'
x-   like `x*' but omitting word `$'
```

The `:' separating the event specification from the word designator can be omitted if the argument selector begins with a `^', `\$', `*' or `-' or `%'. After the optional word designator can be placed a sequence of modifiers, each preceded by a `:'. The following modifiers are defined:

```
h    Remove a trailing pathname component, leaving the head.
r    Remove a trailing `.xxx' component, leaving the root name.
e    Remove all but the extension `.xxx' part.
s/l/r/ Substitute l for r
t    Remove all leading pathname components, leaving the tail.
&    Repeat the previous substitution.
g    Apply the change globally, prefixing the above, e.g. `g&'.
p    Print the new command but do not execute it.
q    Quote the substituted words, preventing further substitutions.
x    Like q, but break into words at blanks, tabs and newlines.
```

Unless preceded by a `g' the modification is applied only to the first modifiable word. With substitutions, it is an error for no word to be applicable.

The left hand side of substitutions are not regular expressions in the sense of the editors, but rather strings. Any

character may be used as the delimiter in place of ``/'`; a ``\`` quotes the delimiter into the `l` and `r` strings. The character ``&`` in the right hand side is replaced by the text from the left. A ``\`` quotes ``&`` also. A null `l` uses the previous string either from a `l` or from a contextual scan string `s` in ``!?s``. The trailing delimiter in the substitution may be omitted if a newline follows immediately as may the trailing ``?`` in a contextual scan.

A history reference may be given without an event specification, e.g. ``!$``. In this case the reference is to the previous command unless a previous history reference occurred on the same line in which case this form repeats the previous reference. Thus ``!?foo?^ !$`` gives the first and last arguments from the command matching ``?foo``.

A special abbreviation of a history reference occurs when the first non-blank character of an input line is a ``^``. This is equivalent to ``!:s^`` providing a convenient shorthand for substitutions on the text of the previous line. Thus ``^lb^lib`` fixes the spelling of ``lib`` in the previous command. Finally, a history substitution may be surrounded with ``{`` and ``}`` if necessary to insulate it from the characters which follow. Thus, after ``ls -ld ~paul`` we might do ``!{l}a`` to do ``ls -ld ~paula``, while ``!la`` would look for a command starting ``la``.

Quotations with ``` and `"`

The quotation of strings by ``'`` and ``"``` can be used to prevent all or some of the remaining substitutions. Strings enclosed in ``'`` are prevented any further interpretation. Strings enclosed in ``"``` may be expanded as described below.

In both cases the resulting text becomes (all or part of) a single word; only in one special case (see Command Substitution below) does a ``"``` quoted string yield parts of more than one word; ``'`` quoted strings never do.

Alias substitution

The shell maintains a list of aliases which can be established, displayed and modified by the `alias` and `unalias` commands. After a command line is scanned, it is parsed into distinct commands and the first word of each command, left-to-right, is checked to see if it has an alias. If it does, then the text which is the alias for that command is reread with the history mechanism available as though that command were the previous input line. The resulting words replace the command and argument list. If no reference is made to the history list, then the argument list is left unchanged.

Thus if the alias for ``ls'` is ``ls -l'` the command ``ls /usr'` would map to ``ls -l /usr'`, the argument list here being undisturbed. Similarly if the alias for ``lookup'` was ``grep !^ /etc/passwd'` then ``lookup bill'` would map to ``grep bill /etc/passwd'`.

If an alias is found, the word transformation of the input text is performed and the aliasing process begins again on the reformed input line. Looping is prevented if the first word of the new text is the same as the old by flagging it to prevent further aliasing. Other loops are detected and cause an error.

Note that the mechanism allows aliases to introduce parser metasyntax. Thus we can ``alias print 'pr \!* | lpr''` to make a command which `pr`'s its arguments to the line printer.

Variable substitution

The shell maintains a set of variables, each of which has as value a list of zero or more words. Some of these variables are set by the shell or referred to by it. For instance, the `argv` variable is an image of the shell's argument list, and words of this variable's value are referred to in special ways.

The values of variables may be displayed and changed by using the `set` and `unset` commands. Of the variables referred to by the shell a number are toggles; the shell does not care what their value is, only whether they are set or not. For instance, the `verbose` variable is a toggle which causes command input to be echoed. The setting of this variable results from the `-v` command line option.

Other operations treat variables numerically. The ``@'` command permits numeric calculations to be performed and the result assigned to a variable. Variable values are, however, always represented as (zero or more) strings. For the purposes of numeric operations, the null string is considered to be zero, and the second and subsequent words of multiword values are ignored.

After the input line is aliased and parsed, and before each command is executed, variable substitution is performed keyed by ``$'` characters. This expansion can be prevented by preceding the ``$'` with a ``\'` except within ``"'`s where it always occurs, and within ``''`s where it never occurs. Strings quoted by ``''` are interpreted later (see Command substitution below) so ``$'` substitution does not occur there until later, if at all. A ``$'` is passed unchanged if followed by a blank, tab, or end-of-line.

Input/output redirections are recognized before variable expansion, and are variable expanded separately. Otherwise, the command name and entire argument list are expanded together. It is thus possible for the first (command) word to this point to generate more than one word, the first of which becomes the command name, and the rest of which become arguments.

Unless enclosed in `"' or given the `:q' modifier the results of variable substitution may eventually be command and filename substituted. Within `"', a variable whose value consists of multiple words expands to a (portion of) a single word, with the words of the variables value separated by blanks. When the `:q' modifier is applied to a substitution the variable will expand to multiple words with each word separated by a blank and quoted to prevent later command or filename substitution.

The following metasequences are provided for introducing variable values into the shell input. Except as noted, it is an error to reference a variable which is not set.

`$name`

`${name}`

Are replaced by the words of the value of variable name, each separated by a blank. Braces insulate name from following characters which would otherwise be part of it. Shell variables have names consisting of up to 20 letters and digits starting with a letter. The underscore character is considered a letter.

If name is not a shell variable, but is set in the environment, then that value is returned (but : modifiers and the other forms given below are not available in this case).

`$name[selector]`

`${name[selector]}`

May be used to select only some of the words from the value of name. The selector is subjected to `\$' substitution and may consist of a single number or two numbers separated by a `-' . The first word of a variables value is numbered `1' . If the first number of a range is omitted it defaults to `1' . If the last member of a range is omitted it defaults to `\$#name' . The selector `*' selects all words. It is not an error for a range to be empty if the second argument is omitted or in range.

`$#name`

`${#name}`

Gives the number of words in the variable. This is useful for later use in a `[selector]' .

\$0

Substitutes the name of the file from which command input is being read. An error occurs if the name is not known.

\$number

\${number}

Equivalent to `\${argv[number]}'.

\$*

Equivalent to `\${argv[*]}'.

The modifiers `:h', `:t', `:r', `:q' and `:x' may be applied to the substitutions above as may `:gh', `:gt' and `:gr'. If braces `{ '}' appear in the command form then the modifiers must appear within the braces. The current implementation allows only one `:' modifier on each `\$' expansion.

The following substitutions may not be modified with `:' modifiers.

\$?name

\${?name}

Substitutes the string `1' if name is set, `0' if it is not.

\$?0

Substitutes `1' if the current input filename is known, `0' if it is not.

\$\$

Substitute the (decimal) process number of the (parent) shell.

\$<

Substitutes a line from the standard input, with no further interpretation thereafter. It can be used to read from the keyboard in a shell script.

Command and filename substitution

The remaining substitutions, command and filename substitution, are applied selectively to the arguments of builtin commands. This means that portions of expressions which are not evaluated are not subjected to these expansions. For commands which are not internal to the shell, the command name is substituted separately from the argument list. This occurs very late, after input-output redirection is performed, and in a child of the main shell.

Command substitution

Command substitution is indicated by a command enclosed in ``'``. The output from such a command is normally broken into separate words at blanks, tabs and newlines, with null words being discarded, this text then replacing the original string. Within ``'``s, only newlines force new words; blanks and tabs are preserved.

In any case, the single final newline does not force a new word. Note that it is thus possible for a command substitution to yield only part of a word, even if the command outputs a complete line.

Filename substitution

If a word contains any of the characters ``*'``, ``?'``, ``['`` or ``{'`` or begins with the character ``~'``, then that word is a candidate for filename substitution, also known as ``globbing'``. This word is then regarded as a pattern, and replaced with an alphabetically sorted list of file names which match the pattern. In a list of words specifying filename substitution it is an error for no pattern to match an existing file name, but it is not required for each pattern to match. Only the metacharacters ``*'``, ``?'`` and ``['`` imply pattern matching, the characters ``~'`` and ``{'`` being more akin to abbreviations.

In matching filenames, the character ``.'`` at the beginning of a filename or immediately following a ``/'``, as well as the character ``/'`` must be matched explicitly. The character ``*'`` matches any string of characters, including the null string. The character ``?'`` matches any single character. The sequence ``[...]'`` matches any one of the characters enclosed. Within ``[...]'``, a pair of characters separated by ``-'`` matches any character lexically between the two.

The character ``~'`` at the beginning of a filename is used to refer to home directories. Standing alone, i.e. ``~'`` it expands to the invokers home directory as reflected in the value of the variable `home`. When followed by a name consisting of letters, digits and ``-'`` characters the shell searches for a user with that name and substitutes their home directory; thus ``~ken'`` might expand to ``/usr/ken'`` and ``~ken/chmach'`` to ``/usr/ken/chmach'``. If the character ``~'`` is followed by a character other than a letter or ``/'`` or appears not at the beginning of a word, it is left undisturbed.

The metanotation ``a{b,c,d}e'`` is a shorthand for ``abe ace ade'``. Left to right order is preserved, with results of matches being sorted separately at a low level to preserve

this order. This construct may be nested. Thus ``~source/s1/{oldls,ls}.c'` expands to ``/usr/source/s1/oldls.c` `/usr/source/s1/ls.c` whether or not these files exist without any chance of error if the home directory for ``source'` is ``/usr/source'`. Similarly ``../{memo,*box}'` might expand to ``../memo ../box ../mbox'`. (Note that ``memo'` was not sorted with the results of matching ``*box'`.) As a special case ``{'`, ``}'` and ``{'` are passed undisturbed.

Input/output

The standard input and standard output of a command may be redirected with the following syntax:

< name

Open file name (which is first variable, command and filename expanded) as the standard input.

<< word

Read the shell input up to a line which is identical to word. Word is not subjected to variable, filename or command substitution, and each input line is compared to word before any substitutions are done on this input line. Unless a quoting ``\'`, ``\"`, ``\'` or ``\'` appears in word variable and command substitution is performed on the intervening lines, allowing ``\'` to quote ``$'`, ``\'` and ``\'`. Commands which are substituted have all blanks, tabs, and newlines preserved, except for the final newline which is dropped. The resultant text is placed in an anonymous temporary file which is given to the command as standard input.

> name

>! name

>& name

>&! name

The file name is used as standard output. If the file does not exist then it is created; if the file exists, its is truncated, its previous contents being lost.

If the variable `noclobber` is set, then the file must not exist or be a character special file (e.g. a terminal or ``/dev/null'`) or an error results. This helps prevent accidental destruction of files. In this case the ``!'` forms can be used and suppress this check.

The forms involving ``&'` route the diagnostic output into the specified file as well as the standard output. Name is expanded in the same way as ``<'` input filenames are.

```
>> name
>>& name
>>! name
>>&! name
```

Uses file name as standard output like `>' but places output at the end of the file. If the variable noclobber is set, then it is an error for the file not to exist unless one of the `!' forms is given. Otherwise similar to `>'.

A command receives the environment in which the shell was invoked as modified by the input-output parameters and the presence of the command in a pipeline. Thus, unlike some previous shells, commands run from a file of shell commands have no access to the text of the commands by default; rather they receive the original standard input of the shell. The `<<' mechanism should be used to present inline data. This permits shell command scripts to function as components of pipelines and allows the shell to block read its input. Note that the default standard input for a command run detached is not modified to be the empty file `/dev/null'; rather the standard input remains as the original standard input of the shell. If this is a terminal and if the process attempts to read from the terminal, then the process will block and the user will be notified (see Jobs above).

Diagnostic output may be directed through a pipe with the standard output. Simply use the form `|&' rather than just `|'.

Expressions

A number of the builtin commands (to be described subsequently) take expressions, in which the operators are similar to those of C, with the same precedence. These expressions appear in the @, exit, if, and while commands. The following operators are available:

```
||  &&  |  ^  &  ==  !=  =~  !~  <=  >=  <  >  <<  >>
+   -   *   /   %   !   ~   (   )
```

Here the precedence increases to the right, `==' `!=' `=~' and `!~', `<=' `>=' `<' and `>', `<<' and `>>', `+' and `-', `*' `/' and `%` being, in groups, at the same level. The `==' `!=' `=~' and `!~' operators compare their arguments as strings; all others operate on numbers. The operators `=~' and `!~' are like `!=' and `==' except that the right hand side is a pattern (containing, e.g. `*'s, `?'s and instances of `[...]') against which the left hand operand is matched. This reduces the need for use of the switch statement in shell scripts when all that is really needed is pattern

matching.

Strings which begin with ``0'` are considered octal numbers. Null or missing arguments are considered ``0'`. The result of all expressions are strings, which represent decimal numbers. It is important to note that no two components of an expression can appear in the same word; except when adjacent to components of expressions which are syntactically significant to the parser (``&'`|'`<'`>'`('`')'`) they should be surrounded by spaces.

Also available in expressions as primitive operands are command executions enclosed in ``{'` and ``}'` and file enquiries of the form ``-l name'` where `l` is one of:

r	read access
w	write access
x	execute access
e	existence
o	ownership
z	zero size
f	plain file
d	directory

The specified name is command and filename expanded and then tested to see if it has the specified relationship to the real user. If the file does not exist or is inaccessible then all enquiries return false, i.e. ``0'`. Command executions succeed, returning true, i.e. ``1'`, if the command exits with status 0, otherwise they fail, returning false, i.e. ``0'`. If more detailed status information is required then the command should be executed outside of an expression and the variable status examined.

Control flow

The shell contains a number of commands which can be used to regulate the flow of control in command files (shell scripts) and (in limited but useful ways) from terminal input. These commands all operate by forcing the shell to reread or skip in its input and, due to the implementation, restrict the placement of some of the commands.

The `foreach`, `switch`, and `while` statements, as well as the `if-then-else` form of the `if` statement require that the major keywords appear in a single simple command on an input line as shown below.

If the shell's input is not seekable, the shell buffers up input whenever a loop is being read and performs seeks in this internal buffer to accomplish the rereading implied by the loop. (To the extent that this allows, backward goto's

will succeed on non-seekable inputs.)

Builtin commands

Builtin commands are executed within the shell. If a builtin command occurs as any component of a pipeline except the last then it is executed in a subshell.

alias

alias name

alias name wordlist

The first form prints all aliases. The second form prints the alias for name. The final form assigns the specified wordlist as the alias of name; wordlist is command and filename substituted. Name is not allowed to be alias or unalias.

alloc

Shows the amount of dynamic memory acquired, broken down into used and free memory. With an argument shows the number of free and used blocks in each size category. The categories start at size 8 and double at each step. This command's output may vary across system types, since systems other than the VAX may use a different memory allocator.

bg

bg %job ...

Puts the current or specified jobs into the background, continuing them if they were stopped.

break

Causes execution to resume after the end of the nearest enclosing foreach or while. The remaining commands on the current line are executed. Multi-level breaks are thus possible by writing them all on one line.

breaksw

Causes a break from a switch, resuming after the endsw.

case label:

A label in a switch statement as discussed below.

cd

cd name

chdir

chdir name

Change the shell's working directory to directory name. If no argument is given then change to the home directory of the user.

If name is not found as a subdirectory of the current directory (and does not begin with `/', `./' or `../'),

then each component of the variable `cdpath` is checked to see if it has a subdirectory name. Finally, if all else fails but `name` is a shell variable whose value begins with ``/'`, then this is tried to see if it is a directory.

`continue`

Continue execution of the nearest enclosing `while` or `foreach`. The rest of the commands on the current line are executed.

`default:`

Labels the default case in a switch statement. The default should come after all case labels.

`dirs`

Prints the directory stack; the top of the stack is at the left, the first directory in the stack being the current directory.

`echo wordlist`

`echo -n wordlist`

The specified words are written to the shells standard output, separated by spaces, and terminated with a new-line unless the `-n` option is specified.

`else`

`end`

`endif`

`endsw`

See the description of the `foreach`, `if`, `switch`, and `while` statements below.

`eval arg ...`

(As in `sh(1)`.) The arguments are read as input to the shell and the resulting command(s) executed in the context of the current shell. This is usually used to execute commands generated as the result of command or variable substitution, since parsing occurs before these substitutions. See `tset(1)` for an example of using `eval`.

`exec command`

The specified command is executed in place of the current shell.

`exit`

`exit(expr)`

The shell exits either with the value of the status variable (first form) or with the value of the specified `expr` (second form).

fg

fg %job ...

Brings the current or specified jobs into the foreground, continuing them if they were stopped.

foreach name (wordlist)

...

end

The variable name is successively set to each member of wordlist and the sequence of commands between this command and the matching end are executed. (Both foreach and end must appear alone on separate lines.)

The builtin command continue may be used to continue the loop prematurely and the builtin command break to terminate it prematurely. When this command is read from the terminal, the loop is read up once prompting with '?' before any statements in the loop are executed. If you make a mistake typing in a loop at the terminal you can rub it out.

glob wordlist

Like echo but no '\\' escapes are recognized and words are delimited by null characters in the output. Useful for programs which wish to use the shell to filename expand a list of words.

goto word

The specified word is filename and command expanded to yield a string of the form 'label'. The shell rewinds its input as much as possible and searches for a line of the form 'label:' possibly preceded by blanks or tabs. Execution continues after the specified line.

hashstat

Print a statistics line indicating how effective the internal hash table has been at locating commands (and avoiding exec's). An exec is attempted for each component of the path where the hash function indicates a possible hit, and in each component which does not begin with a '/'.

history

history n

history -r n

history -h n

Displays the history event list; if n is given only the n most recent events are printed. The -r option reverses the order of printout to be most recent first rather than oldest first. The -h option causes the history list to be printed without leading numbers. This is used to produce files suitable for sourcing

using the -h option to source.

if (expr) command

If the specified expression evaluates true, then the single command with arguments is executed. Variable substitution on command happens early, at the same time it does for the rest of the if command. Command must be a simple command, not a pipeline, a command list, or a parenthesized command list. Input/output redirection occurs even if expr is false, when command is not executed (this is a bug).

if (expr) then

...

else if (expr2) then

...

else

...

endif

If the specified expr is true then the commands to the first else are executed; otherwise if expr2 is true then the commands to the second else are executed, etc. Any number of else-if pairs are possible; only one endif is needed. The else part is likewise optional. (The words else and endif must appear at the beginning of input lines; the if must appear alone on its input line or after an else.)

jobs

jobs -l

Lists the active jobs; given the -l options lists process id's in addition to the normal information.

kill %job

kill -sig %job ...

kill pid

kill -sig pid ...

kill -l

Sends either the TERM (terminate) signal or the specified signal to the specified jobs or processes. Signals are either given by number or by names (as given in /usr/include/signal.h, stripped of the prefix ``SIG''). The signal names are listed by ``kill -l''. There is no default, saying just `kill' does not send a signal to the current job. If the signal being sent is TERM (terminate) or HUP (hangup), then the job or process will be sent a CONT (continue) signal as well.

limit

limit resource

limit resource maximum-use

limit -h

limit -h resource

limit -h resource maximum-use

Limits the consumption by the current process and each process it creates to not individually exceed maximum-use on the specified resource. If no maximum-use is given, then the current limit is printed; if no resource is given, then all limitations are given. If the -h flag is given, the hard limits are used instead of the current limits. The hard limits impose a ceiling on the values of the current limits. Only the super-user may raise the hard limits, but a user may lower or raise the current limits within the legal range.

Resources controllable currently include cputime (the maximum number of cpu-seconds to be used by each process), filesize (the largest single file which can be created), datasize (the maximum growth of the data+stack region via sbrk(2) beyond the end of the program text), stacksize (the maximum size of the automatically-extended stack region), and coredumpsize (the size of the largest core dump that will be created).

The maximum-use may be given as a (floating point or integer) number followed by a scale factor. For all limits other than cputime the default scale is 'k' or 'kilobytes' (1024 bytes); a scale factor of 'm' or 'megabytes' may also be used. For cputime the default scaling is 'seconds', while 'm' for minutes or 'h' for hours, or a time of the form 'mm:ss' giving minutes and seconds may be used.

For both resource names and scale factors, unambiguous prefixes of the names suffice.

login

Terminate a login shell, replacing it with an instance of /bin/login. This is one way to log off, included for compatibility with sh(1).

logout

Terminate a login shell. Especially useful if ignoreeof is set.

nice

nice +number

nice command

nice +number command

The first form sets the scheduling priority for this shell to 4. The second form sets the priority to the

given number. The final two forms run command at priority 4 and number respectively. The greater the number, the less cpu the process will get. The super-user may specify negative priority by using `nice -number ...'. Command is always executed in a sub-shell, and the restrictions placed on commands in simple if statements apply.

nohup

nohup command

The first form can be used in shell scripts to cause hangups to be ignored for the remainder of the script. The second form causes the specified command to be run with hangups ignored. All processes detached with `&' are effectively nohup'ed.

notify

notify %job ...

Causes the shell to notify the user asynchronously when the status of the current or specified jobs changes; normally notification is presented before a prompt. This is automatic if the shell variable notify is set.

onintr

onintr -

onintr label

Control the action of the shell on interrupts. The first form restores the default action of the shell on interrupts which is to terminate shell scripts or to return to the terminal command input level. The second form `onintr -' causes all interrupts to be ignored. The final form causes the shell to execute a `goto label' when an interrupt is received or a child process terminates because it was interrupted.

In any case, if the shell is running detached and interrupts are being ignored, all forms of onintr have no meaning and interrupts continue to be ignored by the shell and all invoked commands.

popd

popd +n

Pops the directory stack, returning to the new top directory. With an argument `+n' discards the nth entry in the stack. The elements of the directory stack are numbered from 0 starting at the top.

pushd

pushd name

pushd +n

With no arguments, pushd exchanges the top two elements of the directory stack. Given a name argument, pushd

changes to the new directory (ala cd) and pushes the old current working directory (as in csw) onto the directory stack. With a numeric argument, rotates the nth argument of the directory stack around to be the top element and changes to it. The members of the directory stack are numbered from the top starting at 0.

rehash

Causes the internal hash table of the contents of the directories in the path variable to be recomputed. This is needed if new commands are added to directories in the path while you are logged in. This should only be necessary if you add commands to one of your own directories, or if a systems programmer changes the contents of one of the system directories.

repeat count command

The specified command which is subject to the same restrictions as the command in the one line if statement above, is executed count times. I/O redirections occur exactly once, even if count is 0.

set

set name

set name=word

set name[index]=word

set name=(wordlist)

The first form of the command shows the value of all shell variables. Variables which have other than a single word as value print as a parenthesized word list. The second form sets name to the null string. The third form sets name to the single word. The fourth form sets the index'th component of name to word; this component must already exist. The final form sets name to the list of words in wordlist. In all cases the value is command and filename expanded.

These arguments may be repeated to set multiple values in a single set command. Note however, that variable expansion happens for all arguments before any setting occurs.

setenv

setenv name value

setenv name

The first form lists all current environment variables. The last form sets the value of environment variable name to be value, a single string. The second form sets name to an empty string. The most commonly used environment variable USER, TERM, and PATH are automatically imported to and exported from the csh variables

user, term, and path; there is no need to use setenv for these.

shift

shift variable

The members of argv are shifted to the left, discarding argv[1]. It is an error for argv not to be set or to have less than one word as value. The second form performs the same function on the specified variable.

source name

source -h name

The shell reads commands from name. Source commands may be nested; if they are nested too deeply the shell may run out of file descriptors. An error in a source at any level terminates all nested source commands. Normally input during source commands is not placed on the history list; the -h option causes the commands to be placed in the history list without being executed.

stop

stop %job ...

Stops the current or specified job which is executing in the background.

suspend

Causes the shell to stop in its tracks, much as if it had been sent a stop signal with ^Z. This is most often used to stop shells started by su(1).

switch (string)

case str1:

...

breaksw

...

default:

...

breaksw

endsw

Each case label is successively matched, against the specified string which is first command and filename expanded. The file metacharacters '*', '?' and '['...' may be used in the case labels, which are variable expanded. If none of the labels match before a 'default' label is found, then the execution begins after the default label. Each case label and the default label must appear at the beginning of a line. The command breaksw causes execution to continue after the endsw. Otherwise control may fall through case labels and default labels as in C. If no label matches and there is no default, execution continues after the endsw.

time**time command**

With no argument, a summary of time used by this shell and its children is printed. If arguments are given the specified simple command is timed and a time summary as described under the time variable is printed. If necessary, an extra shell is created to print the time statistic when the command completes.

umask**umask value**

The file creation mask is displayed (first form) or set to the specified value (second form). The mask is given in octal. Common values for the mask are 002 giving all access to the group and read and execute access to others or 022 giving all access except no write access for users in the group or others.

unalias pattern

All aliases whose names match the specified pattern are discarded. Thus all aliases are removed by ``unalias *'`. It is not an error for nothing to be unaliased.

unhash

Use of the internal hash table to speed location of executed programs is disabled.

unlimit**unlimit resource****unlimit -h****unlimit -h resource**

Removes the limitation on resource. If no resource is specified, then all resource limitations are removed. If `-h` is given, the corresponding hard limits are removed. Only the super-user may do this.

unset pattern

All variables whose names match the specified pattern are removed. Thus all variables are removed by ``unset *'`; this has noticeably distasteful side-effects. It is not an error for nothing to be unset.

unsetenv pattern

Removes all variables whose name match the specified pattern from the environment. See also the `setenv` command above and `printenv(1)`.

wait

All background jobs are waited for. If the shell is interactive, then an interrupt can disrupt the wait, at which time the shell prints names and job numbers of all jobs known to be outstanding.

```
while (expr)
```

```
...
```

```
end
```

While the specified expression evaluates non-zero, the commands between the while and the matching end are evaluated. Break and continue may be used to terminate or continue the loop prematurely. (The while and end must appear alone on their input lines.) Prompting occurs here the first time through the loop as for the foreach statement if the input is a terminal.

```
%job
```

Brings the specified job into the foreground.

```
%job &
```

Continues the specified job in the background.

```
@
```

```
@ name = expr
```

```
@ name[index] = expr
```

The first form prints the values of all the shell variables. The second form sets the specified name to the value of expr. If the expression contains '<', '>', '&' or '|' then at least this part of the expression must be placed within '(' ')'. The third form assigns the value of expr to the index'th argument of name. Both name and its index'th component must already exist.

The operators '*=', '+=', etc are available as in C. The space separating the name from the assignment operator is optional. Spaces are, however, mandatory in separating components of expr which would otherwise be single words.

Special postfix '++' and '--' operators increment and decrement name respectively, i.e. '@ i++'.

Pre-defined and environment variables

The following variables have special meaning to the shell. Of these, argv, cwd, home, path, prompt, shell and status are always set by the shell. Except for cwd and status this setting occurs only at initialization; these variables will not then be modified unless this is done explicitly by the user.

This shell copies the environment variable USER into the variable user, TERM into term, and HOME into home, and copies these back into the environment whenever the normal shell variables are reset. The environment variable PATH is likewise handled; it is not necessary to worry about its setting other than in the file .cshrc as inferior csh

processes will import the definition of path from the environment, and re-export it if you then change it.

argv	Set to the arguments to the shell, it is from this variable that positional parameters are substituted, i.e. '\$1' is replaced by '\$argv[1]', etc.
cdpath	Gives a list of alternate directories searched to find subdirectories in chdir commands.
cwd	The full pathname of the current directory.
echo	Set when the -x command line option is given. Causes each command and its arguments to be echoed just before it is executed. For non-builtin commands all expansions occur before echoing. Builtin commands are echoed before command and filename substitution, since these substitutions are then done selectively.
filec	Enable file name completion.
histchars	Can be given a string value to change the characters used in history substitution. The first character of its value is used as the history substitution character, replacing the default character !. The second character of its value replaces the character ^ in quick substitutions.
history	Can be given a numeric value to control the size of the history list. Any command which has been referenced in this many events will not be discarded. Too large values of history may run the shell out of memory. The last executed command is always saved on the history list.
home	The home directory of the invoker, initialized from the environment. The filename expansion of '~' refers to this variable.
ignoreeof	If set the shell ignores end-of-file from input devices which are terminals. This prevents shells from accidentally being killed by control-D's.
mail	The files where the shell checks for mail. This is done after each command completion

which will result in a prompt, if a specified interval has elapsed. The shell says 'You have new mail.' if the file exists with an access time not greater than its modify time.

If the first word of the value of mail is numeric it specifies a different mail checking interval, in seconds, than the default, which is 10 minutes.

If multiple mail files are specified, then the shell says 'New mail in name' when there is mail in the file name.

- noclobber** As described in the section on Input/output, restrictions are placed on output redirection to insure that files are not accidentally destroyed, and that '>>' redirections refer to existing files.
- noglob** If set, filename expansion is inhibited. This is most useful in shell scripts which are not dealing with filenames, or after a list of filenames has been obtained and further expansions are not desirable.
- nonomatch** If set, it is not an error for a filename expansion to not match any existing files; rather the primitive pattern is returned. It is still an error for the primitive pattern to be malformed, i.e. 'echo [' still gives an error.
- notify** If set, the shell notifies asynchronously of job completions. The default is to rather present job completions just before printing a prompt.
- path** Each word of the path variable specifies a directory in which commands are to be sought for execution. A null word specifies the current directory. If there is no path variable then only full path names will execute. The usual search path is '.', '/bin' and '/usr/bin', but this may vary from system to system. For the super-user the default search path is '/bin', '/sbin', '/usr/sbin', and '/usr/bin'. A shell which is given neither the -c nor the -t option will normally hash the contents of the directories in the path variable after reading .cshrc, and each time the path variable is reset. If new

commands are added to these directories while the shell is active, it may be necessary to give the rehash or the commands may not be found.

- prompt** The string which is printed before each command is read from an interactive terminal input. If a `!' appears in the string it will be replaced by the current event number unless a preceding `\' is given. Default is `% ', or `# ' for the super-user.
- savehist** is given a numeric value to control the number of entries of the history list that are saved in ~/.history when the user logs out. Any command which has been referenced in this many events will be saved. During start up the shell sources ~/.history into the history list enabling history to be saved across logins. Too large values of savehist will slow down the shell during start up.
- shell** The file in which the shell resides. This is used in forking shells to interpret files which have execute bits set, but which are not executable by the system. (See the description of Non-builtin Command Execution below.) Initialized to the (system-dependent) home of the shell.
- status** The status returned by the last command. If it terminated abnormally, then 0200 is added to the status. Builtin commands which fail return exit status `1', all other builtin commands set status `0'.
- time** Controls automatic timing of commands. If set, then any command which takes more than this many cpu seconds will cause a line giving user, system, and real times and a utilization percentage which is the ratio of user plus system times to real time to be printed when it terminates.
- verbose** Set by the -v command line option, causes the words of each command to be printed after history substitution.

Non-builtin command execution

When a command to be executed is found to not be a builtin command the shell attempts to execute the command via

execve(2). Each word in the variable path names a directory from which the shell will attempt to execute the command. If it is given neither a -c nor a -t option, the shell will hash the names in these directories into an internal table so that it will only try an exec in a directory if there is a possibility that the command resides there. This greatly speeds command location when a large number of directories are present in the search path. If this mechanism has been turned off (via unhash), or if the shell was given a -c or -t argument, and in any case for each directory component of path which does not begin with a '/', the shell concatenates with the given command name to form a path name of a file which it then attempts to execute.

Parenthesized commands are always executed in a subshell. Thus `(cd ; pwd) ; pwd` prints the home directory; leaving you where you were (printing this after the home directory), while `cd ; pwd` leaves you in the home directory. Parenthesized commands are most often used to prevent chdir from affecting the current shell.

If the file has execute permissions but is not an executable binary to the system, then it is assumed to be a file containing shell commands and a new shell is spawned to read it.

If there is an alias for shell then the words of the alias will be prepended to the argument list to form the shell command. The first word of the alias should be the full path name of the shell (e.g. `\$shell'). Note that this is a special, late occurring, case of alias substitution, and only allows words to be prepended to the argument list without modification.

Argument list processing

If argument 0 to the shell is '-' then this is a login shell. The flag arguments are interpreted as follows:

- b This flag forces a ``break'' from option processing, causing any further shell arguments to be treated as non-option arguments. The remaining arguments will not be interpreted as shell options. This may be used to pass options to a shell script without confusion or possible subterfuge. The shell will not run a set-user ID script without this option.
- c Commands are read from the (single) following argument which must be present. Any remaining arguments are placed in argv.
- e The shell exits if any invoked command terminates

abnormally or yields a non-zero exit status.

- f The shell will start faster, because it will neither search for nor execute commands from the file ``.cshrc'` in the invoker's home directory.
- i The shell is interactive and prompts for its top-level input, even if it appears to not be a terminal. Shells are interactive without this option if their inputs and outputs are terminals.
- n Commands are parsed, but not executed. This aids in syntactic checking of shell scripts.
- s Command input is taken from the standard input.
- t A single line of input is read and executed. A ``\'` may be used to escape the newline at the end of this line and continue onto another line.
- v Causes the verbose variable to be set, with the effect that command input is echoed after history substitution.
- x Causes the echo variable to be set, so that commands are echoed immediately before execution.
- V Causes the verbose variable to be set even before ``.cshrc'` is executed.
- X Is to -x as -V is to -v.

After processing of flag arguments, if arguments remain but none of the -c, -i, -s, or -t options was given, the first argument is taken as the name of a file of commands to be executed. The shell opens this file, and saves its name for possible resubstitution by ``$0'`. Since many systems use either the standard version 6 or version 7 shells whose shell scripts are not compatible with this shell, the shell will execute such a 'standard' shell if the first character of a script is not a ``#'`, i.e. if the script does not start with a comment. Remaining arguments initialize the variable `argv`.

Signal handling

The shell normally ignores quit signals. Jobs running detached (either by ``&'` or the `bg` or `%... &` commands) are immune to signals generated from the keyboard, including hangups. Other signals have the values which the shell inherited from its parent. The shells handling of interrupts and terminate signals in shell scripts can be

controlled by onintr. Login shells catch the terminate signal; otherwise this signal is passed on to children from the state in the shell's parent. In no case are interrupts allowed when a login shell is reading the file ``.logout'`.

AUTHOR

William Joy. Job control and directory stack features first implemented by J.E. Kulp of I.I.A.S.A, Laxenburg, Austria, with different syntax than that used now. File name completion code written by Ken Greer, HP Labs.

FILES

<code>~/.cshrc</code>	Read at beginning of execution by each shell.
<code>~/.login</code>	Read by login shell, after <code>`.cshrc'</code> at login.
<code>~/.logout</code>	Read by login shell, at logout.
<code>/bin/sh</code>	Standard shell, for shell scripts not starting with a <code>`#'</code> .
<code>/tmp/sh*</code>	Temporary file for <code>`<<'</code> .
<code>/etc/passwd</code>	Source of home directories for <code>~name'</code> .

LIMITATIONS

Words can be no longer than 1024 characters. The system limits argument lists to 10240 characters. The number of arguments to a command which involves filename expansion is limited to 1/6'th the number of characters allowed in an argument list. Command substitutions may substitute no more characters than are allowed in an argument list. To detect looping, the shell restricts the number of alias substitutions on a single line to 20.

SEE ALSO

`sh(1)`, `access(2)`, `execve(2)`, `fork(2)`, `killpg(2)`, `pipe(2)`, `sigvec(2)`, `umask(2)`, `setrlimit(2)`, `wait(2)`, `tty(4)`, `a.out(5)`, `environ(7)`, `'An introduction to the C shell'`

BUGS

When a command is restarted from a stop, the shell prints the directory it started in if this is different from the current directory; this can be misleading (i.e. wrong) as the job may have changed directories internally.

Shell builtin functions are not stoppable/restartable. Command sequences of the form ``a ; b ; c'` are also not handled gracefully when stopping is attempted. If you suspend ``b'`, the shell will then immediately execute ``c'`. This is especially noticeable if this expansion results from an alias. It suffices to place the sequence of commands in `()'s` to force it to a subshell, i.e. ``(a ; b ; c)'`.

Control over tty output after processes are started is primitive; perhaps this will inspire someone to work on a good virtual terminal interface. In a virtual terminal interface much more interesting things could be done with output

control.

Alias substitution is most often used to clumsily simulate shell procedures; shell procedures should be provided rather than aliases.

Commands within loops, prompted for by '?', are not placed in the history list. Control structure should be parsed rather than being recognized as built-in commands. This would allow control commands to be placed anywhere, to be combined with '|', and to be used with '&' and ';' metasyntax.

It should be possible to use the ':' modifiers on the output of command substitutions. All and more than one ':' modifier should be allowed on '\$' substitutions.

The way the filec facility is implemented is ugly and expensive.

NAME

ctags - create a tags file

SYNOPSIS

ctags [-BFatuwx] [-f tagsfile] name ...

DESCRIPTION

Ctags makes a tags file for ex(1) from the specified C, Pascal, Fortran, YACC, lex, and lisp sources. A tags file gives the locations of specified objects (in this case functions and typedefs) in a group of files. Each line of the tags file contains the object name, the file in which it is defined, and an address specification for the object definition. Functions are searched with a pattern, typedefs with a line number. Specifiers are given in separate fields on the line, separated by blanks or tabs. Using the tags file, ex can quickly find these objects definitions.

If the -x flag is given, ctags produces a list of object names, the line number and file name on which each is defined, as well as the text of that line and prints this on the standard output. This is a simple index which can be printed out as an off-line readable function index.

If the -v flag is given, an index of the form expected by vgrind(1) is produced on the standard output. This listing contains the function name, file name, and page number (assuming 64 line pages). Since the output will be sorted into lexicographic order, it may be desired to run the output through sort -f. Sample use:

```
ctags -v files | sort -f > index
vgrind -x index
```

Normally ctags places the tag descriptions in a file called tags; this may be overridden with the -f option.

Files whose names end in .c or .h are assumed to be C source files and are searched for C routine and macro definitions. Files whose names end in .y are assumed to be YACC source files. Files whose names end in .l are assumed to be either lisp files if their first non-blank character is `;', `(', or `[', or lex files otherwise. Other files are first examined to see if they contain any Pascal or Fortran routine definitions; if not, they are processed again looking for C definitions.

Other options are:

-F use forward searching patterns (/.../) (default).

-B use backward searching patterns (?...?).

- a append to tags file.
- t create tags for typedefs.
- w suppressing warning diagnostics.
- u causing the specified files to be updated in tags, that is, all references to them are deleted, and the new values are appended to the file. (Beware: this option is implemented in a way which is rather slow; it is usually faster to simply rebuild the tags file.)

The tag main is treated specially in C programs. The tag formed is created by prepending M to the name of the file, with a trailing .c removed, if any, and leading pathname components also removed. This makes use of ctags practical in directories with more than one program.

FILES

tags output tags file

SEE ALSO

ex(1), vi(1)

AUTHOR

Ken Arnold; FORTRAN added by Jim Kleckner; Bill Joy added Pascal and -x, replacing cxref; C typedefs added by Ed Pelegri-Llopart.

BUGS

Recognition of functions, subroutines and procedures for FORTRAN and Pascal is done in a very simpleminded way. No attempt is made to deal with block structure; if you have two Pascal procedures in different blocks with the same name you lose.

The method of deciding whether to look for C or Pascal and FORTRAN functions is a hack.

Does not know about #ifdefs.

Should know about Pascal types. Relies on the input being well formed to detect typedefs. Use of -tx shows only the last line of typedefs.

NAME

date - print and set the date

SYNOPSIS

date [-nu] [-d dst] [-t timezone] [yymmddhhmm [.ss]]

DESCRIPTION

If no arguments are given, the current date and time are printed. Providing an argument will set the desired date; only the superuser can set the date. The -d and -t flags set the kernel's values for daylight savings time and minutes west of GMT. If dst is non-zero, future calls to `gettimeofday(2)` will return a non-zero `tz_dsttime`. Timezone provides the number of minutes returned by future calls to `gettimeofday(2)` in `tz_minuteswest`. The -u flag is used to display or set the date in GMT (universal) time. yy represents the last two digits of the year; the first mm is the month number; dd is the day number; hh is the hour number (24 hour system); the second mm is the minute number; .ss is optional and represents the seconds. For example:

```
date 8506131627
```

sets the date to June 13 1985, 4:27 PM. The year, month and day may be omitted; the default values will be the current ones. The system operates in GMT. Date takes care of the conversion to and from local standard and daylight-saving time.

If `timed(8)` is running to synchronize the clocks of machines in a local area network, date sets the time globally on all those machines unless the -n option is given.

FILES

/usr/adm/wtmp to record time-setting. In /usr/adm/messages, date records the name of the user setting the time.

SEE ALSO

`gettimeofday(2)`, `utmp(5)`, `timed(8)`,
TSP: The Time Synchronization Protocol for UNIX 4.3BSD, R. Gusella and S. Zatti

DIAGNOSTICS

Exit status is 0 on success, 1 on complete failure to set the date, and 2 on successfully setting the local date but failing globally.

Occasionally, when `timed` synchronizes the time on many hosts, the setting of a new time value may require more than a few seconds. On these occasions, date prints: 'Network time being set'. The message 'Communication error with timed' occurs when the communication between date and timed

fails.

BUGS

The system attempts to keep the date in a format closely compatible with VMS. VMS, however, uses local time (rather than GMT) and does not understand daylight-saving time. Thus, if you use both UNIX and VMS, VMS will be running on GMT.

NAME

dc - desk calculator

SYNOPSIS

dc [file]

DESCRIPTION

Dc is an arbitrary precision arithmetic package. Ordinarily it operates on decimal integers, but one may specify an input base, output base, and a number of fractional digits to be maintained. The overall structure of dc is a stacking (reverse Polish) calculator. If an argument is given, input is taken from that file until its end, then from the standard input. The following constructions are recognized:

number

The value of the number is pushed on the stack. A number is an unbroken string of the digits 0-9. It may be preceded by an underscore `_` to input a negative number. Numbers may contain decimal points.

`+ - / * % ^`

The top two values on the stack are added (+), subtracted (-), multiplied (*), divided (/), remaindered (%), or exponentiated (^). The two entries are popped off the stack; the result is pushed on the stack in their place. Any fractional part of an exponent is ignored.

`sx` The top of the stack is popped and stored into a register named `x`, where `x` may be any character. If the `s` is capitalized, `x` is treated as a stack and the value is pushed on it.

`lx` The value in register `x` is pushed on the stack. The register `x` is not altered. All registers start with zero value. If the `l` is capitalized, register `x` is treated as a stack and its top value is popped onto the main stack.

`d` The top value on the stack is duplicated.

`p` The top value on the stack is printed. The top value remains unchanged. `P` interprets the top of the stack as an ascii string, removes it, and prints it.

`f` All values on the stack and in registers are printed.

`q` exits the program. If executing a string, the recursion level is popped by two. If `q` is capitalized, the top value on the stack is popped and the string execution level is popped by that value.

- `x` treats the top element of the stack as a character string and executes it as a string of dc commands.
- `X` replaces the number on the top of the stack with its scale factor.
- `[...]` puts the bracketed ascii string onto the top of the stack.
- `<x >x =x` The top two elements of the stack are popped and compared. Register `x` is executed if they obey the stated relation.
- `v` replaces the top element on the stack by its square root. Any existing fractional part of the argument is taken into account, but otherwise the scale factor is ignored.
- `!` interprets the rest of the line as a UNIX command.
- `c` All values on the stack are popped.
- `i` The top value on the stack is popped and used as the number radix for further input. `I` pushes the input base on the top of the stack.
- `o` The top value on the stack is popped and used as the number radix for further output.
- `O` pushes the output base on the top of the stack.
- `k` the top of the stack is popped, and that value is used as a non-negative scale factor: the appropriate number of places are printed on output, and maintained during multiplication, division, and exponentiation. The interaction of scale factor, input base, and output base will be reasonable if all are changed together.
- `z` The stack level is pushed onto the stack.
- `Z` replaces the number on the top of the stack with its length.
- `?` A line of input is taken from the input source (usually the terminal) and executed.
- `;` `:` are used by bc for array operations.

An example which prints the first ten values of `n!` is

```
[1a1+dsa*pla10>y]sy  
0sa1  
lyx
```

SEE ALSO

bc(1), which is a preprocessor for dc providing infix notation and a C-like syntax which implements functions and reasonable control structures for programs.

DIAGNOSTICS

`x is unimplemented' where x is an octal number.
`stack empty' for not enough elements on the stack to do what was asked.
`Out of space' when the free list is exhausted (too many digits).
`Out of headers' for too many numbers being kept around.
`Out of pushdown' for too many items on the stack.
`Nesting Depth' for too many levels of nested execution.

NAME

dd - convert and copy a file

SYNOPSIS

dd [option=value] ...

DESCRIPTION

Dd copies the specified input file to the specified output with possible conversions. The standard input and output are used by default. The input and output block size may be specified to take advantage of raw physical I/O.

option	values
if=	input file name; standard input is default
of=	output file name; standard output is default
ibs=n	input block size n bytes (default 512)
obs=n	output block size (default 512)
bs=n	set both input and output block size, superseding ibs and obs; also, if no conversion is specified, it is particularly efficient since no copy need be done
cbs=n	conversion buffer size
skip=n	skip n input records before starting copy
files=n	copy n input files before terminating (makes sense only where input is a magtape or similar device).
seek=n	seek n records from beginning of output file before copying
count=n	copy only n input records
conv=ascii	convert EBCDIC to ASCII
ebcdic	convert ASCII to EBCDIC
ibm	slightly different map of ASCII to EBCDIC
block	convert variable length records to fixed length
unblock	convert fixed length records to variable length
lcase	map alphabetic to lower case
ucase	map alphabetic to upper case
swab	swap every pair of bytes
noerror	do not stop processing on an error
sync	pad every input record to ibs
...	, ... several comma-separated conversions

Where sizes are specified, a number of bytes is expected. A number may end with k, b or w to specify multiplication by 1024, 512, or 2 respectively; a pair of numbers may be separated by x to indicate a product.

Cbs is used only if ascii, unblock, ebcdic, ibm, or block conversion is specified. In the first two cases, cbs characters are placed into the conversion buffer, any specified character mapping is done, trailing blanks trimmed and new-

line added before sending the line to the output. In the latter three cases, characters are read into the conversion buffer, and blanks added to make up an output record of size cbs.

After completion, dd reports the number of whole and partial input and output blocks.

For example, to read an EBCDIC tape blocked ten 80-byte EBCDIC card images per record into the ASCII file x:

```
dd if=/dev/rmt0 of=x ibs=800 cbs=80 conv=ascii,lcase
```

Note the use of raw magtape. Dd is especially suited to I/O on the raw physical devices because it allows reading and writing in arbitrary record sizes.

SEE ALSO

cp(1), tr(1)

DIAGNOSTICS

f+p records in(out): numbers of full and partial records read(written)

BUGS

The ASCII/EBCDIC conversion tables are taken from the 256 character standard in the CACM Nov, 1968. The 'ibm' conversion, while less blessed as a standard, corresponds better to certain IBM print train conventions. There is no universal solution.

One must specify ``conv=noerror,sync'' when copying raw disks with bad sectors to insure dd stays synchronized.

Certain combinations of arguments to conv= are permitted. However, the block or unblock option cannot be combined with ascii, ebcidic or ibm. Invalid combinations silently ignore all but the last mutually-exclusive keyword.

NAME

deroff - remove nroff, troff, tbl and eqn constructs

SYNOPSIS

deroff [-w] file ...

DESCRIPTION

Deroff reads each file in sequence and removes all nroff and troff command lines, backslash constructions, macro definitions, eqn constructs (between ``.EQ'` and ``.EN'` lines or between delimiters), and table descriptions and writes the remainder on the standard output. Deroff follows chains of included files (``.so'` and ``.nx'` commands); if a file has already been included, a ``.so'` is ignored and a ``.nx'` terminates execution. If no input file is given, deroff reads from the standard input file.

If the `-w` flag is given, the output is a word list, one `'word'` (string of letters, digits, and apostrophes, beginning with a letter; apostrophes are removed) per line, and all other characters ignored. Otherwise, the output follows the original, with the deletions mentioned above.

SEE ALSO

troff(1), eqn(1), tbl(1)

BUGS

Deroff is not a complete troff interpreter, so it can be confused by subtle constructs. Most errors result in too much rather than too little output.

NAME

df - disk free

SYNOPSIS

df [-i] [filesystem ...] [file ...]

DESCRIPTION

Df prints out the amount of free disk space available on the specified filesystem, e.g. ```/dev/rp0a''`, or on the filesystem in which the specified file, e.g. ```$HOME''`, is contained. If no file system is specified, the free space on all of the normally mounted file systems is printed. The reported numbers are in kilobytes.

Other options are:

-i Report also the number of inodes which are used and free.

FILES

/etc/fstab list of normally mounted filesystems

SEE ALSO

fstab(5), icheck(8), quot(8)

NAME

diction, explain - print wordy sentences; thesaurus for diction

SYNOPSIS

diction [-ml] [-mm] [-n] [-f pfile] file ...
explain

DESCRIPTION

Diction finds all sentences in a document that contain phrases from a data base of bad or wordy diction. Each phrase is bracketed with []. Because diction runs deroff before looking at the text, formatting header files should be included as part of the input. The default macro package -ms may be overridden with the flag -mm. The flag -ml which causes deroff to skip lists, should be used if the document contains many lists of non-sentences. The user may supply her/his own pattern file to be used in addition to the default file with -f pfile. If the flag -n is also supplied the default file will be suppressed.

Explain is an interactive thesaurus for the phrases found by diction.

SEE ALSO

deroff(1)

BUGS

Use of non-standard formatting macros may cause incorrect sentence breaks. In particular, diction doesn't grok -me.

NAME

diff - differential file and directory comparator

SYNOPSIS

```
diff [ -l ] [ -r ] [ -s ] [ -cefhn ] [ -biwt ] dir1 dir2
diff [ -cefhn ] [ -biwt ] file1 file2
diff [ -Dstring ] [ -biw ] file1 file2
```

DESCRIPTION

If both arguments are directories, diff sorts the contents of the directories by name, and then runs the regular file diff algorithm (described below) on text files which are different. Binary files which differ, common subdirectories, and files which appear in only one directory are listed. Options when comparing directories are:

- l long output format; each text file diff is piped through pr(1) to paginate it, other differences are remembered and summarized after all text file differences are reported.
- r causes application of diff recursively to common subdirectories encountered.
- s causes diff to report files which are the same, which are otherwise not mentioned.
- Sname starts a directory diff in the middle beginning with file name.

When run on regular files, and when comparing text files which differ during directory comparison, diff tells what lines must be changed in the files to bring them into agreement. Except in rare circumstances, diff finds a smallest sufficient set of file differences. If neither file1 nor file2 is a directory, then either may be given as '-', in which case the standard input is used. If file1 is a directory, then a file in that directory whose file-name is the same as the file-name of file2 is used (and vice versa).

There are several options for output format; the default output format contains lines of these forms:

```
n1 a n3,n4
n1,n2 d n3
n1,n2 c n3,n4
```

These lines resemble ed commands to convert file1 into file2. The numbers after the letters pertain to file2. In fact, by exchanging 'a' for 'd' and reading backward one may ascertain equally how to convert file2 into file1. As in

ed, identical pairs where $n1 = n2$ or $n3 = n4$ are abbreviated as a single number.

Following each of these lines come all the lines that are affected in the first file flagged by `<<', then all the lines that are affected in the second file flagged by `>'.

Except for -b, -w, -i or -t which may be given with any of the others, the following options are mutually exclusive:

-e produces a script of a, c and d commands for the editor ed, which will recreate file2 from file1. In connection with -e, the following shell program may help maintain multiple versions of a file. Only an ancestral file (\$1) and a chain of version-to-version ed scripts (\$2,\$3,...) made by diff need be on hand. A 'latest version' appears on the standard output.

```
(shift; cat $*; echo '1,$p') | ed - $1
```

Extra commands are added to the output when comparing directories with -e, so that the result is a sh(1) script for converting text files which are common to the two directories from their state in dir1 to their state in dir2.

-f produces a script similar to that of -e, not useful with ed, and in the opposite order.

-n produces a script similar to that of -e, but in the opposite order and with a count of changed lines on each insert or delete command. This is the form used by rcsdiff(1).

-c produces a diff with lines of context. The default is to present 3 lines of context and may be changed, e.g to 10, by -c10. With -c the output format is modified slightly: the output beginning with identification of the files involved and their creation dates and then each change is separated by a line with a dozen *'s. The lines removed from file1 are marked with `-' ; those added to file2 are marked `+' . Lines which are changed from one file to the other are marked in both files with with `!' .

Changes which lie within <context> lines of each other are grouped together on output. (This is a change from the previous ``diff -c'' but the resulting output is usually much easier to interpret.)

- h does a fast, half-hearted job. It works only when changed stretches are short and well separated, but does work on files of unlimited length.

- Dstring causes diff to create a merged version of file1 and file2 on the standard output, with C preprocessor controls included so that a compilation of the result without defining string is equivalent to compiling file1, while defining string will yield file2.

- b causes trailing blanks (spaces and tabs) to be ignored, and other strings of blanks to compare equal.

- w is similar to -b but causes whitespace (blanks and tabs) to be totally ignored. E.g.,
 ``if (a == b)'`` will compare equal to
 ``if(a==b)'``.

- i ignores the case of letters. E.g., ``A'`` will compare equal to ``a'``.

- t will expand tabs in output lines. Normal or -c output adds character(s) to the front of each line which may screw up the indentation of the original source lines and make the output listing difficult to interpret. This option will preserve the original source's indentation.

FILES

/tmp/d?????
/usr/libexec/diffh for -h
/bin/diff for directory diffs
/bin/pr

SEE ALSO

cmp(1), cc(1), comm(1), ed(1), diff3(1)

DIAGNOSTICS

Exit status is 0 for no differences, 1 for some, 2 for trouble.

BUGS

Editing scripts produced under the -e or -f option are naive about creating lines consisting of a single `.'.

When comparing directories with the -b, -w or -i options specified, diff first compares the files ala cmp, and then decides to run the diff algorithm if they are not equal. This may cause a small amount of spurious output if the files then turn out to be identical because the only

differences are insignificant blank string or case differences.

NAME

diff3 - 3-way differential file comparison

SYNOPSIS

diff3 [-exEX3] file1 file2 file3

DESCRIPTION

Diff3 compares three versions of a file, and publishes disagreeing ranges of text flagged with these codes:

==== all three files differ

====1 file1 is different

====2 file2 is different

====3 file3 is different

The type of change suffered in converting a given range of a given file to some other is indicated in one of these ways:

f : n1 a Text is to be appended after line number n1
 in file f, where f = 1, 2, or 3.

f : n1 , n2 c Text is to be changed in the range line n1
 to line n2. If n1 = n2, the range may be
 abbreviated to n1.

The original contents of the range follows immediately after a c indication. When the contents of two files are identical, the contents of the lower-numbered file is suppressed.

Under the -e option, diff3 publishes a script for the editor that will incorporate into file1 all changes between file2 and file3, i.e. the changes that normally would be flagged with ==== and ====3. Option -x (-3) produces a script to incorporate only changes flagged with ==== (====3). The following command will apply the resulting script to 'file1'.

```
(cat script; echo '1,$p') | ed - file1
```

The -E and -X are similar to -e and -x, respectively, but treat overlapping changes (i.e., changes that would be flagged with ==== in the normal listing) differently. The overlapping lines from both files will be inserted by the edit script, bracketed by "<<<<<<" and ">>>>>>" lines.

For example, suppose lines 7-8 are changed in both file1 and file2. Applying the edit script generated by the command
 "diff3 -E file1 file2 file3"
to file1 results in the file:

```
lines 1-6
of file1
<<<<<< file1
lines 7-8
of file1
=====
lines 7-8
of file3
>>>>>> file3
rest of file1
```

The -E option is used by RCS merge(1) to insure that overlapping changes in the merged files are preserved and brought to someone's attention.

FILES

```
/tmp/d3?????
/usr/libexec/diff3
```

SEE ALSO

```
diff(1)
```

BUGS

Text lines that consist of a single `.' will defeat -e.

NAME

du - summarize disk usage

SYNOPSIS

du [-s] [-a] [name ...]

DESCRIPTION

Du gives the number of kilobytes contained in all files and, recursively, directories within each specified directory or file name. If name is missing, `.' is used.

The argument -s causes only the grand total to be given. The argument -a causes an entry to be generated for each file. Absence of either causes an entry to be generated for each directory only.

A file which has two links to it is only counted once.

SEE ALSO

df(1), quot(8)

BUGS

Non-directories given as arguments (not under -a option) are not listed.

If there are too many distinct linked files, du counts the excess files multiply.

NAME

echo - echo arguments

SYNOPSIS

echo [-n] [arg] ...

DESCRIPTION

Echo writes its arguments separated by blanks and terminated by a newline on the standard output. If the flag -n is used, no newline is added to the output.

Echo is useful for producing diagnostics in shell programs and for writing constant data on pipes. To send diagnostics to the standard error file, do `echo ... 1>&2'.

NAME

ed - text editor

SYNOPSIS

ed [-] [name]

DESCRIPTION

Ed is the standard text editor.

If a name argument is given, ed simulates an e command (see below) on the named file; that is to say, the file is read into ed's buffer so that it can be edited. The optional - suppresses the printing of explanatory output and should be used when the standard input is an editor script.

Ed operates on a copy of any file it is editing; changes made in the copy have no effect on the file until a w (write) command is given. The copy of the text being edited resides in a temporary file called the buffer.

Commands to ed have a simple and regular structure: zero or more addresses followed by a single character command, possibly followed by parameters to the command. These addresses specify one or more lines in the buffer. Missing addresses are supplied by default.

In general, only one command may appear on a line. Certain commands allow the addition of text to the buffer. While ed is accepting text, it is said to be in input mode. In this mode, no commands are recognized; all input is merely collected. Input mode is left by typing a period '.' alone at the beginning of a line.

Ed supports a limited form of regular expression notation. A regular expression specifies a set of strings of characters. A member of this set of strings is said to be matched by the regular expression. In the following specification for regular expressions the word 'character' means any character but newline.

1. Any character except a special character matches itself. Special characters are the regular expression delimiter plus '[.' and sometimes '^*\$.
2. A '.' matches any character.
3. A '\' followed by any character except a digit or '(' matches that character.
4. A nonempty string s bracketed [s] (or [^s]) matches any character in (or not in) s. In s, '\' has no special meaning, and] may only appear as the first letter. A

substring a-b, with a and b in ascending ASCII order, stands for the inclusive range of ASCII characters.

5. A regular expression of form 1-4 followed by * matches a sequence of 0 or more matches of the regular expression.
6. A regular expression, x, of form 1-8, bracketed \(\x\) matches what x matches.
7. A \ followed by a digit n matches a copy of the string that the bracketed regular expression beginning with the nth \(\) matched.
8. A regular expression of form 1-8, x, followed by a regular expression of form 1-7, y matches a match for x followed by a match for y, with the x match being as long as possible while still permitting a y match.
9. A regular expression of form 1-8 preceded by ^ (or followed by \$), is constrained to matches that begin at the left (or end at the right) end of a line.
10. A regular expression of form 1-9 picks out the longest among the leftmost matches in a line.
11. An empty regular expression stands for a copy of the last regular expression encountered.

Regular expressions are used in addresses to specify lines and in one command (see s below) to specify a portion of a line which is to be replaced. If it is desired to use one of the regular expression metacharacters as an ordinary character, that character may be preceded by \'. This also applies to the character bounding the regular expression (often '/') and to \' itself.

To understand addressing in ed it is necessary to know that at any time there is a current line. Generally speaking, the current line is the last line affected by a command; however, the exact effect on the current line is discussed under the description of the command. Addresses are constructed as follows.

1. The character '.' addresses the current line.
2. The character '\$' addresses the last line of the buffer.
3. A decimal number n addresses the n-th line of the buffer.

4. ``x'` addresses the line marked with the name `x`, which must be a lower-case letter. Lines are marked with the `k` command described below.
5. A regular expression enclosed in slashes ``/'` addresses the line found by searching forward from the current line and stopping at the first line containing a string that matches the regular expression. If necessary the search wraps around to the beginning of the buffer.
6. A regular expression enclosed in queries ``?'` addresses the line found by searching backward from the current line and stopping at the first line containing a string that matches the regular expression. If necessary the search wraps around to the end of the buffer.
7. An address followed by a plus sign ``+'` or a minus sign ``-'` followed by a decimal number specifies that address plus (resp. minus) the indicated number of lines. The plus sign may be omitted.
8. If an address begins with ``+'` or ``-'` the addition or subtraction is taken with respect to the current line; e.g. ``-5'` is understood to mean ``.-5'`.
9. If an address ends with ``+'` or ``-'`, then 1 is added (resp. subtracted). As a consequence of this rule and rule 8, the address ``-'` refers to the line before the current line. Moreover, trailing ``+'` and ``-'` characters have cumulative effect, so ``--'` refers to the current line less 2.
10. To maintain compatibility with earlier versions of the editor, the character ``^'` in addresses is equivalent to ``-'`.

Commands may require zero, one, or two addresses. Commands which require no addresses regard the presence of an address as an error. Commands which accept one or two addresses assume default addresses when insufficient are given. If more addresses are given than such a command requires, the last one or two (depending on what is accepted) are used.

Addresses are separated from each other typically by a comma ``,``. They may also be separated by a semicolon ``;'`. In this case the current line ``.`` is set to the previous address before the next address is interpreted. This feature can be used to determine the starting line for forward and backward searches (``/'`, ``?'`). The second address of any two-address sequence must correspond to a line following the line corresponding to the first address. The special form ``%'` is an abbreviation for the address pair

`1,\$'.

In the following list of ed commands, the default addresses are shown in parentheses. The parentheses are not part of the address, but are used to show that the given addresses are the default.

As mentioned, it is generally illegal for more than one command to appear on a line. However, most commands may be suffixed by `p' or by `l', in which case the current line is either printed or listed respectively in the way discussed below. Commands may also be suffixed by `n', meaning the output of the command is to be line numbered. These suffixes may be combined in any order.

(.)a
<text>

.

The append command reads the given text and appends it after the addressed line. `.' is left on the last line input, if there were any, otherwise at the addressed line. Address `0' is legal for this command; text is placed at the beginning of the buffer.

(.,.)c
<text>

.

The change command deletes the addressed lines, then accepts input text which replaces these lines. `.' is left at the last line input; if there were none, it is left at the line preceding the deleted lines.

(.,.)d

The delete command deletes the addressed lines from the buffer. The line originally after the last line deleted becomes the current line; if the lines deleted were originally at the end, the new last line becomes the current line.

e filename

The edit command causes the entire contents of the buffer to be deleted, and then the named file to be read in. `.' is set to the last line of the buffer. The number of characters read is typed. `filename' is remembered for possible use as a default file name in a subsequent r or w command. If `filename' is missing, the remembered name is used.

E filename

This command is the same as e, except that no diagnostic results when no w has been given since the last buffer alteration.

f filename

The filename command prints the currently remembered file name. If 'filename' is given, the currently remembered file name is changed to 'filename'.

(1,\$)g/regular expression/command list

In the global command, the first step is to mark every line which matches the given regular expression. Then for every such line, the given command list is executed with '.' initially set to that line. A single command or the first of multiple commands appears on the same line with the global command. All lines of a multi-line list except the last line must be ended with '\'. A, i, and c commands and associated input are permitted; the '.' terminating input mode may be omitted if it would be on the last line of the command list. The commands g and v are not permitted in the command list.

(.)i

<text>

.

This command inserts the given text before the addressed line. '.' is left at the last line input, or, if there were none, at the line before the addressed line. This command differs from the a command only in the placement of the text.

(.,.+1)j

This command joins the addressed lines into a single line; intermediate newlines simply disappear. '.' is left at the resulting line.

(.)kx

The mark command marks the addressed line with name x, which must be a lower-case letter. The address form 'x' then addresses this line.

(.,.)l

The list command prints the addressed lines in an unambiguous way: non-graphic characters are printed in two-digit octal, and long lines are folded. The l command may be placed on the same line after any non-i/o command.

(.,.)ma

The move command repositions the addressed lines after the line addressed by a. The last of the moved lines becomes the current line.

(.,.)p

The print command prints the addressed lines. '.' is

left at the last line printed. The p command may be placed on the same line after any non-i/o command.

(., .)P

This command is a synonym for p.

q The quit command causes ed to exit. No automatic write of a file is done.

Q This command is the same as q, except that no diagnostic results when no w has been given since the last buffer alteration.

(\$)r filename

The read command reads in the given file after the addressed line. If no file name is given, the remembered file name, if any, is used (see e and f commands). The file name is remembered if there was no remembered file name already. Address `0' is legal for r and causes the file to be read at the beginning of the buffer. If the read is successful, the number of characters read is typed. `.' is left at the last line read in from the file.

(., .)s/regular expression/replacement/ or,

(., .)s/regular expression/replacement/g

The substitute command searches each addressed line for an occurrence of the specified regular expression. On each line in which a match is found, all matched strings are replaced by the replacement specified, if the global replacement indicator `g' appears after the command. If the global indicator does not appear, only the first occurrence of the matched string is replaced. It is an error for the substitution to fail on all addressed lines. Any punctuation character may be used instead of `/` to delimit the regular expression and the replacement. `.' is left at the last line substituted.

An ampersand `&' appearing in the replacement is replaced by the string matching the regular expression. The special meaning of `&' in this context may be suppressed by preceding it by `\''. The characters `\'n' where n is a digit, are replaced by the text matched by the n-th regular subexpression enclosed between `\'(' and `\'')'. When nested, parenthesized subexpressions are present, n is determined by counting occurrences of `\'(' starting from the left.

Lines may be split by substituting new-line characters into them. The new-line in the replacement string must be escaped by preceding it by `\'\'.

One or two trailing delimiters may be omitted, implying the 'p' suffix. The special form 's' followed by no delimiters repeats the most recent substitute command on the addressed lines. The 's' may be followed by the letters r (use the most recent regular expression for the left hand side, instead of the most recent left hand side of a substitute command), p (complement the setting of the p suffix from the previous substitution), or g (complement the setting of the g suffix). These letters may be combined in any order.

(., .)ta

This command acts just like the m command, except that a copy of the addressed lines is placed after address a (which may be 0). '.' is left on the last line of the copy.

(., .)u

The undo command restores the buffer to it's state before the most recent buffer modifying command. The current line is also restored. Buffer modifying commands are a, c, d, g, i, k, and v. For purposes of undo, g and v are considered to be a single buffer modifying command. Undo is its own inverse.

When ed runs out of memory (at about 8000 lines on any 16 bit mini-computer such as the PDP-11) This full undo is not possible, and u can only undo the effect of the most recent substitute on the current line. This restricted undo also applies to editor scripts when ed is invoked with the - option.

(1, \$)v/regular expression/command list

This command is the same as the global command g except that the command list is executed g with '.' initially set to every line except those matching the regular expression.

(1, \$)w filename

The write command writes the addressed lines onto the given file. If the file does not exist, it is created. The file name is remembered if there was no remembered file name already. If no file name is given, the remembered file name, if any, is used (see e and f commands). '.' is unchanged. If the command is successful, the number of characters written is printed.

(1, \$)W filename

This command is the same as w, except that the addressed lines are appended to the file.

(1, \$)wq filename

This command is the same as `w` except that afterwards a `q` command is done, exiting the editor after the file is written.

`(.+1)z` or,
`(.+1)zn`

This command scrolls through the buffer starting at the addressed line. 22 (or `n`, if given) lines are printed. The last line printed becomes the current line. The value `n` is sticky, in that it becomes the default for future `z` commands.

`($)=` The line number of the addressed line is typed. ``.'` is unchanged by this command.

`!<shell command>`

The remainder of the line after the `!'` is sent to `sh(1)` to be interpreted as a command. ``.'` is unchanged.

`(.+1,.+1)<newline>`

An address alone on a line causes the addressed line to be printed. A blank line alone is equivalent to ``.+1p'`; it is useful for stepping through text. If two addresses are present with no intervening semicolon, `ed` prints the range of lines. If they are separated by a semicolon, the second line is printed.

If an interrupt signal (ASCII DEL) is sent, `ed` prints ``?interrupted'` and returns to its command level.

Some size limitations: 512 characters per line, 256 characters per global command list, 64 characters per file name, and, on mini computers, 128K characters in the temporary file. The limit on the number of lines depends on the amount of core: each line takes 2 words.

When reading a file, `ed` discards ASCII NUL characters and all characters after the last newline. It refuses to read files containing non-ASCII characters.

FILES

`/tmp/e*`

`edhup`: work is saved here if terminal hangs up

SEE ALSO

B. W. Kernighan, A Tutorial Introduction to the ED Text Editor

B. W. Kernighan, Advanced editing on UNIX
`ex(1)`, `sed(1)`, `crypt(1)`

DIAGNOSTICS

`?name' for inaccessible file; `?self-explanatory message'
for other errors.

To protect against throwing away valuable work, a q or e
command is considered to be in error, unless a w has
occurred since the last buffer change. A second q or e will
be obeyed regardless.

BUGS

The l command mishandles DEL.

The undo command causes marks to be lost on affected lines.

NAME

efl - Extended Fortran Language

SYNOPSIS

efl [option ...] [filename ...]

DESCRIPTION

Efl compiles a program written in the EFL language into clean Fortran. Efl provides the same control flow constructs as does ratfor(1), which are essentially identical to those in C:

statement grouping with braces;
decision-making with if, if-else, and switch-case;
while, for, Fortran do, repeat, and repeat...until loops; multi-level break and next. In addition, EFL has C-like data structures, and more uniform and convenient input/output syntax, generic functions. EFL also provides some syntactic sugar to make programs easier to read and write:

free form input:
multiple statements/line; automatic continuation statement label names (not just numbers),

comments:
this is a comment

translation of relationals:
>, >=, etc., become .GT., .GE., etc.

return (expression)
returns expression to caller from function

define:
define name replacement

include:
include filename

The Efl command option -w suppresses warning messages. The option -C causes comments to be copied through to the Fortran output (default); -# prevents comments from being copied through. If a command argument contains an embedded equal sign, that argument is treated as if it had appeared in an option statement at the beginning of the program. Efl is best used with f77(1).

SEE ALSO

f77(1), ratfor(1).

S. I. Feldman, The Programming Language EFL, Bell Labs Computing Science Technical Report #78.

NAME

error - analyze and disperse compiler error messages

SYNOPSIS

```
error [ -n ] [ -s ] [ -q ] [ -v ] [ -t suffixlist ] [ -I  
ignorefile ] [ name ]
```

DESCRIPTION

Error analyzes and optionally disperses the diagnostic error messages produced by a number of compilers and language processors to the source file and line where the errors occurred. It can replace the painful, traditional methods of scribbling abbreviations of errors on paper, and permits error messages and source code to be viewed simultaneously without machinations of multiple windows in a screen editor.

Error looks at the error messages, either from the specified file name or from the standard input, and attempts to determine which language processor produced each error message, determines the source file and line number to which the error message refers, determines if the error message is to be ignored or not, and inserts the (possibly slightly modified) error message into the source file as a comment on the line preceding to which the line the error message refers. Error messages which can't be categorized by language processor or content are not inserted into any file, but are sent to the standard output. Error touches source files only after all input has been read. By specifying the `-q` query option, the user is asked to confirm any potentially dangerous (such as touching a file) or verbose action. Otherwise error proceeds on its merry business. If the `-t` touch option and associated suffix list is given, error will restrict itself to touch only those files with suffices in the suffix list. Error also can be asked (by specifying `-v`) to invoke `vi(1)` on the files in which error messages were inserted; this obviates the need to remember the names of the files with errors.

Error is intended to be run with its standard input connected via a pipe to the error message source. Some language processors put error messages on their standard error file; others put their messages on the standard output. Hence, both error sources should be piped together into error. For example, when using the `cs`h syntax,

```
make -s lint |& error -q -v
```

will analyze all the error messages produced by whatever programs make runs when making lint.

Error knows about the error messages produced by: `make`, `cc`, `cpp`, `ccom`, `as`, `ld`, `lint`, `pi`, `pc`, `f77`, and DEC Western

Research Modula-2. Error knows a standard format for error messages produced by the language processors, so is sensitive to changes in these formats. For all languages except Pascal, error messages are restricted to be on one line. Some error messages refer to more than one line in more than one files; error will duplicate the error message and insert it at all of the places referenced.

Error will do one of six things with error messages.

synchronize

Some language processors produce short errors describing which file it is processing. Error uses these to determine the file name for languages that don't include the file name in each error message. These synchronization messages are consumed entirely by error.

discard Error messages from lint that refer to one of the two lint libraries, /usr/share/lint/l1ib-lc and /usr/share/lint/l1ib-port are discarded, to prevent accidentally touching these libraries. Again, these error messages are consumed entirely by error.

nullify Error messages from lint can be nullified if they refer to a specific function, which is known to generate diagnostics which are not interesting. Nullified error messages are not inserted into the source file, but are written to the standard output. The names of functions to ignore are taken from either the file named .errorrc in the users's home directory, or from the file named by the -I option. If the file does not exist, no error messages are nullified. If the file does exist, there must be one function name per line.

not file specific

Error messages that can't be intuited are grouped together, and written to the standard output before any files are touched. They will not be inserted into any source file.

file specific

Error message that refer to a specific file, but to no specific line, are written to the standard output when that file is touched.

true errors

Error messages that can be intuited are candidates for insertion into the file to which they refer.

Only true error messages are candidates for inserting into the file they refer to. Other error messages are consumed entirely by error or are written to the standard output. Error inserts the error messages into the source file on the line preceding the line the language processor found in error. Each error message is turned into a one line comment for the language, and is internally flagged with the string ``###'' at the beginning of the error, and ``%%%' at the end of the error. This makes pattern searching for errors easier with an editor, and allows the messages to be easily removed. In addition, each error message contains the source line number for the line the message refers to. A reasonably formatted source program can be recompiled with the error messages still in it, without having the error messages themselves cause future errors. For poorly formatted source programs in free format languages, such as C or Pascal, it is possible to insert a comment into another comment, which can wreak havoc with a future compilation. To avoid this, programs with comments and source on the same line should be formatted so that language statements appear before comments.

Options available with error are:

- n Do not touch any files; all error messages are sent to the standard output.
- q The user is queried whether s/he wants to touch the file. A ``y'' or ``n'' to the question is necessary to continue. Absence of the -q option implies that all referenced files (except those referring to discarded error messages) are to be touched.
- v After all files have been touched, overlay the visual editor vi with it set up to edit all files touched, and positioned in the first touched file at the first error. If vi can't be found, try ex or ed from standard places.
- t Take the following argument as a suffix list. Files whose suffixes do not appear in the suffix list are not touched. The suffix list is dot separated, and ``*' wildcards work. Thus the suffix list:


```
".c.y.foo*.h"
```

 allows error to touch files ending with ``.c'', ``.y'', ``.foo*' and ``.y'.'.
- s Print out statistics regarding the error categorization. Not too useful.

Error catches interrupt and terminate signals, and if in the insertion phase, will orderly terminate what it is doing.

AUTHOR

Robert Henry

FILES

~/.errorrc	function names to ignore for lint error messages
/dev/tty	user's teletype

BUGS

Opens the teletype directly to do user querying.

Source files with links make a new copy of the file with only one link to it.

Changing a language processor's format of error messages may cause error to not understand the error message.

Error, since it is purely mechanical, will not filter out subsequent errors caused by 'floodgating' initiated by one syntactically trivial error. Humans are still much better at discarding these related errors.

Pascal error messages belong after the lines affected (error puts them before). The alignment of the '|' marking the point of error is also disturbed by error.

Error was designed for work on CRT's at reasonably high speed. It is less pleasant on slow speed terminals, and has never been used on hardcopy terminals.

NAME

ex, edit - text editor

SYNOPSIS

```
ex [ - ] [ -v ] [ -t tag ] [ -r ] [ +command ] [ -l ] name
...
edit [ ex options ]
```

DESCRIPTION

Ex is the root of a family of editors: edit, ex and vi. Ex is a superset of ed, with the most notable extension being a display editing facility. Display based editing is the focus of vi.

If you have not used ed, or are a casual user, you will find that the editor edit is convenient for you. It avoids some of the complexities of ex used mostly by systems programmers and persons very familiar with ed.

If you have a CRT terminal, you may wish to use a display based editor; in this case see vi(1), which is a command which focuses on the display editing portion of ex.

DOCUMENTATION

The document Edit: A tutorial (USD:14) provides a comprehensive introduction to edit assuming no previous knowledge of computers or the UNIX system.

The Ex Reference Manual - Version 3.7 (USD:16) is a comprehensive and complete manual for the command mode features of ex, but you cannot learn to use the editor by reading it. For an introduction to more advanced forms of editing using the command mode of ex see the editing documents written by Brian Kernighan for the editor ed; the material in the introductory and advanced documents works also with ex.

An Introduction to Display Editing with Vi (USD:15) introduces the display editor vi and provides reference material on vi. In addition, the Vi Quick Reference card summarizes the commands of vi in a useful, functional way, and is useful with the Introduction.

FILES

/usr/share/misc/exstrings	error messages
/usr/libexec/exrecover	recover command
/usr/sbin/expreserve	preserve command
/etc/termcap	describes capabilities of terminals
~/.exrc	editor startup file
/tmp/Exnnnnn	editor temporary
/tmp/Rxnnnnn	named buffer temporary
/usr/preserve	preservation directory

SEE ALSO

awk(1), ed(1), grep(1), sed(1), grep(1), vi(1), termcap(5),
environ(7)

AUTHOR

Originally written by William Joy
Mark Horton has maintained the editor since version 2.7,
adding macros, support for many unusual terminals, and other
features such as word abbreviation mode.

BUGS

The undo command causes all marks to be lost on lines
changed and then restored if the marked lines were changed.

Undo never clears the buffer modified condition.

The z command prints a number of logical rather than physi-
cal lines. More than a screen full of output may result if
long lines are present.

File input/output errors don't print a name if the command
line '-' option is used.

There is no easy way to do a single scan ignoring case.

The editor does not warn if text is placed in named buffers
and not used before exiting the editor.

Null characters are discarded in input files, and cannot
appear in resultant files.

NAME

expand, unexpand - expand tabs to spaces, and vice versa

SYNOPSIS

```
expand [ -tabstop ] [ -tab1,tab2,...,tabn ] [ file ... ]  
unexpand [ -a ] [ file ... ]
```

DESCRIPTION

Expand processes the named files or the standard input writing the standard output with tabs changed into blanks. Backspace characters are preserved into the output and decrement the column count for tab calculations. Expand is useful for pre-processing character files (before sorting, looking at specific columns, etc.) that contain tabs.

If a single tabstop argument is given, then tabs are set tabstop spaces apart instead of the default 8. If multiple tabstops are given then the tabs are set at those specific columns.

Unexpand puts tabs back into the data from the standard input or the named files and writes the result on the standard output. By default, only leading blanks and tabs are reconverted to maximal strings of tabs. If the -a option is given, then tabs are inserted whenever they would compress the resultant file by replacing two or more characters.

NAME

expr - evaluate arguments as an expression

SYNOPSIS

expr arg ...

DESCRIPTION

The arguments are taken as an expression. After evaluation, the result is written on the standard output. Each token of the expression is a separate argument.

The operators and keywords are listed below. The list is in order of increasing precedence, with equal precedence operators grouped.

expr | expr

yields the first expr if it is neither null nor `0', otherwise yields the second expr.

expr & expr

yields the first expr if neither expr is null or `0', otherwise yields `0'.

expr relop expr

where relop is one of < <= = != >= >, yields `1' if the indicated comparison is true, `0' if false. The comparison is numeric if both expr are integers, otherwise lexicographic.

expr + expr

expr - expr

addition or subtraction of the arguments.

expr * expr

expr / expr

expr % expr

multiplication, division, or remainder of the arguments.

expr : expr

The matching operator compares the string first argument with the regular expression second argument; regular expression syntax is the same as that of ed(1). The \(...\) pattern symbols can be used to select a portion of the first argument. Otherwise, the matching operator yields the number of characters matched (`0' on failure).

(expr)

parentheses for grouping.

Examples:

To add 1 to the Shell variable a:

```
a=`expr $a + 1`
```

To find the filename part (least significant part) of the pathname stored in variable a, which may or may not contain '/':

```
expr $a : '.*\/(.*\)' '|' $a
```

Note the quoted Shell metacharacters.

SEE ALSO

sh(1), test(1)

DIAGNOSTICS

Expr returns the following exit codes:

- 0 if the expression is neither null nor `0',
- 1 if the expression is null or `0',
- 2 for invalid expressions.

NAME

f77 - Fortran77 compiler

SYNTAX

f77 [option] ... file ...

DESCRIPTION

F77 is the UNIX Fortran77 compiler. It accepts several types of arguments:

Arguments whose names end with ``.f'` are taken to be Fortran77 source programs; they are compiled, and each object program is left on the file in the current directory whose name is that of the source with ``.o'` substituted for ``.f'`.

Arguments whose names end with ``.r'` or ``.e'` are taken to be Ratfor or EFL source programs, respectively; these are first transformed by the appropriate preprocessor, then compiled by f77.

In the same way, arguments whose names end with ``.c'` or ``.s'` are taken to be C or assembly source programs and are compiled or assembled, producing a ``.o'` file.

The following options have the same meaning as in `cc(1)`. See `ld(1)` for load-time options.

- c Suppress loading and produce ``.o'` files for each source file.
- p Prepare object files for profiling, see `prof(1)`
- O Invoke an object-code optimizer.
- S Compile the named programs, and leave the assembler-language output on corresponding files suffixed ``.s'`. (No ``.o'` is created.)
- f Use a floating point interpreter (for PDP11's that lack 11/70-style floating point).
- o output
Name the final output file output instead of ``.a.out'`.

The following options are peculiar to f77:

- onetrip
Compile DO loops that are performed at least once if reached. (Fortran77 DO loops are not performed at all if the upper limit is smaller than the lower limit.)
- u Make the default type of a variable ``.undefined'` rather

than using the default Fortran rules.

- C Compile code to check that subscripts are within declared array bounds.
- w Suppress all warning messages. If the option is '-w66', only Fortran 66 compatibility warnings are suppressed.
- F Apply EFL and Ratfor preprocessor to relevant files, put the result in the file with the suffix changed to '.f', but do not compile.
- m Apply the M4 preprocessor to each '.r' or '.e' file before transforming it with the Ratfor or EFL preprocessor.
- Ex Use the string x as an EFL option in processing '.e' files.
- Rx Use the string x as a Ratfor option in processing '.r' files.
- U Do not convert upper case letters to lower case.
- I2 Make default integer size 16 bit.
- I4 Make default integer size 32 bit (default).
- v Verbose. Print information showing what compiler is doing.
- d Debug prints out intermediate information, leaves temporary files in /tmp and often produces a core file.

Other arguments are taken to be either loader option arguments, or F77-compatible object programs, typically produced by an earlier run, or perhaps libraries of F77-compatible routines. These programs, together with the results of any compilations specified, are loaded (in the order given) to produce an executable program with name 'a.out'.

FILES

file.[fresc]	input file
file.o	object file
a.out	loaded output
/usr/libexec/f77pass1compiler	pass 1
/lib/c1	compiler pass 2
/lib/c2	optional optimizer
/usr/lib/libF77.a	intrinsic function library
/usr/lib/libI77.a	Fortran I/O library
/usr/lib/libU77.a	Fortran system call library

/lib/libc.a C library, see section 3
/temp/fortPID.[xsad SopzA]temporary files

Different versions of the compiler tools may be used with the following flags followed immediately (no space) by the path name of the desired module:

```
-T1  pass1      /lib/f77pass1
-T2  pass2      /lib/c1
-Ta  assembler  /bin/as
-Tl  loader      /bin/ld
-TF  footname   /lib/crt0.o
-TM  macro pack m4
```

SEE ALSO

S. I. Feldman, P. J. Weinberger, A Portable Fortran77 Compiler
cc(1), ld(1), prof(1)

DIAGNOSTICS

The diagnostics produced by f77 itself are intended to be self-explanatory. Occasional messages may be produced by the loader.

-d causes the intermediate files to be saved in /tmp and causes the compiler to print out what it is doing.

Run-time diagnostics for the input/output library are as follows:

```
/* 100 */ "error in format"
          See error message output for the location
          of the error in the format. Can be caused
          by more than 10 levels of nested (), or
          an extremely long format statement.
```

```
/* 101 */ "illegal unit number"
        It is illegal to close logical unit 0.
        Negative unit numbers are not allowed.
        The upper limit is system dependent.

/* 102 */ "formatted io not allowed"
        The logical unit was opened for
        unformatted I/O.

/* 103 */ "unformatted io not allowed"
        The logical unit was opened for
        formatted I/O.

/* 104 */ "direct io not allowed"
        The logical unit was opened for sequential
        access, or the logical record length was
        specified as 0.

/* 105 */ "sequential io not allowed"
        The logical unit was opened for direct
        access I/O.

/* 106 */ "can't backspace file"
        The file associated with the logical unit
        can't seek. May be a device or a pipe.

/* 107 */ "off beginning of record"
        The format specified a left tab off the
        beginning of the record.

/* 108 */ "can't stat file"
        The system can't return status information
        about the file. Perhaps the directory is
        unreadable.

/* 109 */ "no * after repeat count"
        Repeat counts in list-directed I/O must be
        followed by an * with no blank spaces.

/* 110 */ "off end of record"
        A formatted write tried to go beyond the
        logical end-of-record. An unformatted read
        or write will also cause this.

/* 111 */ "truncation failed"
        The truncation of external sequential files
        on 'close', 'backspace', or 'rewind' tries
        to do a copy. It failed. Perhaps the temp
        file couldn't be created.

/* 112 */ "incomprehensible list input"
```

List input has to be just right.

```
/* 113 */  "out of free space"
           The library dynamically creates buffers for
           internal use. You ran out of memory for this.
           Your program is too big!

/* 114 */  "unit not connected"
           The logical unit was not open.

/* 115 */  "read unexpected character"
           Certain format conversions can't tolerate
           non-numeric data. Logical data must be
           T or F.

/* 116 */  "blank logical input field"

/* 117 */  "'new' file exists"
           You tried to open an existing file with
           "status='new'".

/* 118 */  "can't find 'old' file"
           You tried to open a nonexistent file
           with "status='old'".

/* 119 */  "unknown system error"
           Shouldn't happen, but .....
           (Send me a documented example.)

/* 120 */  "requires seek ability"
           Direct access requires seek ability.
           Sequential unformatted I/O requires seek
           ability on the file due to the special
           data structure required. Tabbing left
           also requires seek ability.

/* 121 */  "illegal argument"
           Certain arguments to 'open', etc. will be
           checked for legitimacy. Often only non-
           default forms are looked for.

/* 122 */  "negative repeat count"

/* 123 */  "illegal operation for channel or device"
```

BUGS

The Fortran66 subset of the language has been exercised extensively; the newer features have not.

Fortran style read/write routines take up 23 Kbytes of addressing space.

The compiler is not intelligent enough to know when to split up assemblies and loads. Occasionally this causes the loader ld(1) to produce the informative local symbol botch error message when local symbols like argument names are defined with different types. Thus one must split up such offensive modules into separate compilations.

All mathematics for reals is done in double precision. Integer*4 byte alignment is unlike DEC and everyone else's. There is no symbolic debugger.

The optimizer should be used with caution. It is known to occasionally produce incorrect code.

EXAMPLES

```
f77 -O -c myprog.f
    creates myprog.o using C optimizer
```

```
f77 -i -O myprog.f another.f anon.o -lplot
    compiles .f files, loads all files using separate i&d
    space and linking in routines in the plot library.
```

```
f77 myprog.f mine.c >&errors
    Compiles and loads both files putting error output into
    file called errors. This is the C shell (csh) version.
    The Bourne shell (sh) version is:
```

```
f77 myprog.f mine.c 2>errors 1>errors
```

NAME

false, true - provide truth values

SYNOPSIS

true

false

DESCRIPTION

True and false are usually used in a Bourne shell script. They test for the appropriate status "true" or "false" before running (or failing to run) a list of commands.

EXAMPLE

```
while false
do
    command list
done
```

SEE ALSO

csh(1), sh(1), true(1)

DIAGNOSTICS

False has exit status nonzero.

NAME

file - determine file type

SYNOPSIS

file file ...

DESCRIPTION

File performs a series of tests on each argument in an attempt to classify it. If an argument appears to be ascii, file examines the first 512 bytes and tries to guess its language.

BUGS

It often makes mistakes. In particular it often suggests that command files are C programs.

Does not recognize Pascal or LISP.

NAME

find - find files

SYNOPSIS

find pathname-list expression
find pattern

DESCRIPTION

In the first form above, find recursively descends the directory hierarchy for each pathname in the pathname-list (i.e., one or more pathnames) seeking files that match a boolean expression written in the primaries given below. In the descriptions, the argument n is used as a decimal integer where +n means more than n, -n means less than n and n means exactly n.

The second form rapidly searches a database for all pathnames which match pattern. Usually the database is recomputed weekly and contains the pathnames of all files which are publicly accessible. If escaped, normal shell "globbing" characters (`*', `?', `[', and `']) may be used in pattern, but the matching differs in that no characters (e.g. `/') have to be matched explicitly. As a special case, a simple pattern containing no globbing characters is matched as though it were *pattern*; if any globbing character appears there are no implicit globbing characters.

-name filename

True if the filename argument matches the current file name. Normal shell argument syntax may be used if escaped (watch out for `[`, `?` and `*').

-perm onum

True if the file permission flags exactly match the octal number onum (see chmod(1)). If onum is prefixed by a minus sign, more flag bits (017777, see stat(2)) become significant and the flags are compared: (flags&onum)==onum.

-type c True if the type of the file is c, where c is b, c, d, f, l or s for block special file, character special file, directory, plain file, symbolic link, or socket.

-links n True if the file has n links.

-user uname

True if the file belongs to the user uname (login name or numeric user ID).

-nouser True if the file belongs to a user not in the /etc/passwd database.

- group gname
True if the file belongs to group gname (group name or numeric group ID).
- nogroup True if the file belongs to a group not in the /etc/group database.
- size n True if the file is n blocks long (512 bytes per block).
- inum n True if the file has inode number n.
- atime n True if the file has been accessed in n days.
- mtime n True if the file has been modified in n days.
- exec command
True if the executed command returns a zero value as exit status. The end of the command must be punctuated by an escaped semicolon. A command argument `{}` is replaced by the current pathname.
- ok command
Like -exec except that the generated command is written on the standard output, then the standard input is read and the command executed only upon response y.
- print Always true; causes the current pathname to be printed.
- ls Always true; causes current pathname to be printed together with its associated statistics. These include (respectively) inode number, size in kilobytes (1024 bytes), protection mode, number of hard links, user, group, size in bytes, and modification time. If the file is a special file the size field will instead contain the major and minor device numbers. If the file is a symbolic link the pathname of the linked-to file is printed preceded by `->'. The format is identical to that of `ls -gilds' (note however that formatting is done internally, without executing the ls program).
- newer file
True if the current file has been modified more recently than the argument file.
- cpio file
Write the current file on the argument file in cpio format.

`-xdev` Always true; causes find not to traverse down into a file system different from the one on which current argument pathname resides.

The primaries may be combined using the following operators (in order of decreasing precedence):

- 1) A parenthesized group of primaries and operators (parentheses are special to the Shell and must be escaped).
- 2) The negation of a primary (`!` is the unary not operator).
- 3) Concatenation of primaries (the and operation is implied by the juxtaposition of two primaries).
- 4) Alternation of primaries (`-o` is the or operator).

EXAMPLES

To find all accessible files whose pathname contains ``find'`:

```
find find
```

To typeset all variants of manual pages for ``ls'`:

```
vtroff -man `find '*man*/ls.?'`
```

To remove all files named ``a.out'` or ``*.o'` that have not been accessed for a week:

```
find / \( -name a.out -o -name '*.o' \) -atime +7 -exec rm  
{ } \;
```

FILES

```
/etc/passwd  
/etc/group  
/var/db/find.codes      coded pathnames database
```

SEE ALSO

```
sh(1), test(1), fs(5)  
Relevant paper in February, 1983 issue of ;login:.
```

BUGS

The first form's syntax is painful, and the second form's exact semantics is confusing and can vary from site to site.

More than one ``-newer'` option does not work properly.

NAME

finger - user information lookup program

SYNOPSIS

finger [-lmsp] [user [@host] ...]

DESCRIPTION

The finger command has two basic output formats providing essentially the same information.

The -s option of finger displays the user's login name, real name, terminal name and write status (as a ``*'' after the terminal name if write permission is denied), idle time, login time, office location and office phone number.

Idle time is in minutes if it is a single integer, hours and minutes if a ``:' is present, or days if a ``d' is present. Login time is displayed as month, day, hours and minutes, unless more than six months ago, in which case the year is displayed rather than the hours and minutes.

Unknown devices as well as nonexistent idle and login times are displayed as single asterisks.

The -l option produces a multi-line format displaying all of the information described for the -s option as well as the user's home directory, home phone number, login shell, and the contents of the files ``.plan' and ``.project' from the user's home directory.

If idle time is at least a minute and less than a day, it is presented in the form ``hh:mm''. Idle times greater than a day are presented as ``d day[s] hh:mm''.

Phone numbers specified as eleven digits are printed as ``+N-NNN-NNN-NNNN''. Numbers specified as ten or seven digits are printed as the appropriate subset of that string. Numbers specified as five digits are printed as ``xN-NNNN''.

If write permission is denied to the device, the phrase ``(messages off)'' is appended to the line containing the device name. One entry per user is displayed with the -l option; if a user is logged on multiple times, terminal information is repeated once per login.

The -p option prevents the -l option of finger from displaying the contents of the ``.plan' and ``.project' files.

Note that some fields may be missing, in either format, if information is not available for them.

If no operands are specified, finger will print an entry for each user currently logged into the system. If no options are specified, finger defaults to the -l style output if operands are provided, otherwise to the -s style.

User is usually a login name; however, matching will also be done on the users' real names, unless the -m option is supplied. All name matching performed by finger is case insensitive.

Finger may be used to look up users on a remote machine. The format is to specify a user as ``user@host'', or ``@host'', where the default output format for the former is the -l style, and the default output format for the latter is the -s style. The -l option is the only option that may be passed to a remote machine.

SEE ALSO

chpass(1), w(1), who(1), getpwent(3)

NAME

fmt - simple text formatter

SYNOPSIS

fmt [name ...]

DESCRIPTION

Fmt is a simple text formatter which reads the concatenation of input files (or standard input if none are given) and produces on standard output a version of its input with lines as close to 72 characters long as possible. The spacing at the beginning of the input lines is preserved in the output, as are blank lines and interword spacing.

Fmt is meant to format mail messages prior to sending, but may also be useful for other simple tasks. For instance, within visual mode of the ex editor (e.g. vi) the command
!}fmt
will reformat a paragraph, evening the lines.

SEE ALSO

nroff(1), mail(1)

AUTHOR

Kurt Shoens

BUGS

The program was designed to be simple and fast - for more complex operations, the standard text processors are likely to be more appropriate.

NAME

fold - fold long lines for finite width output device

SYNOPSIS

fold [-width] [file ...]

DESCRIPTION

Fold is a filter which will fold the contents of the specified files, or the standard input if no files are specified, breaking the lines to have maximum width width. The default for width is 80. Width should be a multiple of 8 if tabs are present, or the tabs should be expanded using `expand(1)` before coming to fold.

SEE ALSO

`expand(1)`

BUGS

If underlining is present it may be messed up by folding.

NAME

fpr - print Fortran file

SYNOPSIS

fpr

DESCRIPTION

Fpr is a filter that transforms files formatted according to Fortran's carriage control conventions into files formatted according to UNIX line printer conventions.

Fpr copies its input onto its output, replacing the carriage control characters with characters that will produce the intended effects when printed using lpr(1). The first character of each line determines the vertical spacing as follows:

+-----+-----+	
Character	Vertical Space Before Printing
+-----+-----+	
Blank	One line
0	Two lines
1	To first line of next page
+	No advance
+-----+-----+	

A blank line is treated as if its first character is a blank. A blank that appears as a carriage control character is deleted. A zero is changed to a newline. A one is changed to a form feed. The effects of a "+" are simulated using backspaces.

EXAMPLES

a.out | fpr | lpr

fpr < f77.output | lpr

BUGS

Results are undefined for input lines longer than 170 characters.

NAME

from - who is my mail from?

SYNOPSIS

from [-s sender] [user]

DESCRIPTION

From prints out the mail header lines in your mailbox file to show you who your mail is from. If user is specified, then user's mailbox is examined instead of your own. If the -s option is given, then only headers for mail sent by sender are printed.

FILES

/usr/spool/mail/*

SEE ALSO

biff(1), mail(1)

NAME

fsplit - split a multi-routine Fortran file into individual files

SYNOPSIS

fsplit [-e efile] ... [file]

DESCRIPTION

Fsplit takes as input either a file or standard input containing Fortran source code. It attempts to split the input into separate routine files of the form name.f, where name is the name of the program unit (e.g. function, subroutine, block data or program). The name for unnamed block data subprograms has the form blkdtanNN.f where NNN is three digits and a file of this name does not already exist. For unnamed main programs the name has the form mainNNN.f. If there is an error in classifying a program unit, or if name.f already exists, the program unit will be put in a file of the form zzzNNN.f where zzzNNN.f does not already exist.

Normally each subprogram unit is split into a separate file. When the -e option is used, only the specified subprogram units are split into separate files. E.g.:

```
fsplit -e readit -e doit prog.f
```

will split readit and doit into separate files.

DIAGNOSTICS

If names specified via the -e option are not found, a diagnostic is written to standard error.

AUTHOR

Asa Romberger and Jerry Berkman

BUGS

Fsplit assumes the subprogram name is on the first noncomment line of the subprogram unit. Nonstandard source formats may confuse fsplit.

It is hard to use -e for unnamed main programs and block data subprograms since you must predict the created file name.

NAME

ftp - ARPANET file transfer program

SYNOPSIS

ftp [-v] [-d] [-i] [-n] [-g] [host]

DESCRIPTION

Ftp is the user interface to the ARPANET standard File Transfer Protocol. The program allows a user to transfer files to and from a remote network site.

The client host with which ftp is to communicate may be specified on the command line. If this is done, ftp will immediately attempt to establish a connection to an FTP server on that host; otherwise, ftp will enter its command interpreter and await instructions from the user. When ftp is awaiting commands from the user the prompt "ftp>" is provided to the user. The following commands are recognized by ftp:

! [command [args]]

Invoke an interactive shell on the local machine. If there are arguments, the first is taken to be a command to execute directly, with the rest of the arguments as its arguments.

\$ macro-name [args]

Execute the macro macro-name that was defined with the macdef command. Arguments are passed to the macro unglobbed.

account [passwd]

Supply a supplemental password required by a remote system for access to resources once a login has been successfully completed. If no argument is included, the user will be prompted for an account password in a non-echoing input mode.

append local-file [remote-file]

Append a local file to a file on the remote machine. If remote-file is left unspecified, the local file name is used in naming the remote file after being altered by any ntrans or nmap setting. File transfer uses the current settings for type, format, mode, and structure.

ascii

Set the file transfer type to network ASCII. This is the default type.

bell Arrange that a bell be sounded after each file transfer command is completed.

binary
Set the file transfer type to support binary image transfer.

bye Terminate the FTP session with the remote server and exit ftp. An end of file will also terminate the session and exit.

case Toggle remote computer file name case mapping during mget commands. When case is on (default is off), remote computer file names with all letters in upper case are written in the local directory with the letters mapped to lower case.

cd remote-directory
Change the working directory on the remote machine to remote-directory.

cdup Change the remote machine working directory to the parent of the current remote machine working directory.

chmod mode file-name
Change the permission modes the file file-name on the remote sytem to mode.

close
Terminate the FTP session with the remote server, and return to the command interpreter. Any defined macros are erased.

cr Toggle carriage return stripping during ascii type file retrieval. Records are denoted by a carriage return/linefeed sequence during ascii type file transfer. When cr is on (the default), carriage returns are stripped from this sequence to conform with the UNIX single linefeed record delimiter. Records on non-UNIX remote systems may contain single linefeeds; when an ascii type transfer is made, these linefeeds may be distinguished from a record delimiter only when cr is off.

delete remote-file
Delete the file remote-file on the remote machine.

debug [debug-value]
Toggle debugging mode. If an optional debug-value is specified it is used to set the debugging level. When debugging is on, ftp prints each command sent to the remote machine, preceded by the string "-->".

dir [remote-directory] [local-file]
Print a listing of the directory contents in the

directory, remote-directory, and, optionally, placing the output in local-file. If interactive prompting is on, ftp will prompt the user to verify that the last argument is indeed the target local file for receiving dir output. If no directory is specified, the current working directory on the remote machine is used. If no local file is specified, or local-file is -, output comes to the terminal.

disconnect

A synonym for close.

form format

Set the file transfer form to format. The default format is "file".

get remote-file [local-file]

Retrieve the remote-file and store it on the local machine. If the local file name is not specified, it is given the same name it has on the remote machine, subject to alteration by the current case, ntrans, and nmap settings. The current settings for type, form, mode, and structure are used while transferring the file.

glob Toggle filename expansion for mdelete, mget and mput.

If globbing is turned off with glob, the file name arguments are taken literally and not expanded. Globbing for mput is done as in csh(1). For mdelete and mget, each remote file name is expanded separately on the remote machine and the lists are not merged. Expansion of a directory name is likely to be different from expansion of the name of an ordinary file: the exact result depends on the foreign operating system and ftp server, and can be previewed by doing `m!ls remote-files -'. Note: mget and mput are not meant to transfer entire directory subtrees of files. That can be done by transferring a tar(1) archive of the subtree (in binary mode).

hash Toggle hash-sign (`#') printing for each data block transferred. The size of a data block is 1024 bytes.

help [command]

Print an informative message about the meaning of command. If no argument is given, ftp prints a list of the known commands.

idle [seconds]

Set the inactivity timer on the remote server to seconds seconds. If seconds is omitted, the current inactivity timer is printed.

`lcd [directory]`

Change the working directory on the local machine. If no directory is specified, the user's home directory is used.

`ls [remote-directory] [local-file]`

Print a listing of the contents of a directory on the remote machine. The listing includes any system-dependent information that the server chooses to include; for example, most UNIX systems will produce output from the command "`ls -l`". (See also `nlist`.) If `remote-directory` is left unspecified, the current working directory is used. If interactive prompting is on, `ftp` will prompt the user to verify that the last argument is indeed the target local file for receiving `ls` output. If no local file is specified, or if `local-file` is `-`, the output is sent to the terminal.

`macdef macro-name`

Define a macro. Subsequent lines are stored as the macro `macro-name`; a null line (consecutive newline characters in a file or carriage returns from the terminal) terminates macro input mode. There is a limit of 16 macros and 4096 total characters in all defined macros. Macros remain defined until a close command is executed. The macro processor interprets `'$'` and `'\'` as special characters. A `'$'` followed by a number (or numbers) is replaced by the corresponding argument on the macro invocation command line. A `'$'` followed by an `'i'` signals that macro processor that the executing macro is to be looped. On the first pass `'$i'` is replaced by the first argument on the macro invocation command line, on the second pass it is replaced by the second argument, and so on. A `'\'` followed by any character is replaced by that character. Use the `'\'` to prevent special treatment of the `'$'`.

`mdelete [remote-files]`

Delete the remote-files on the remote machine.

`mdir remote-files local-file`

Like `dir`, except multiple remote files may be specified. If interactive prompting is on, `ftp` will prompt the user to verify that the last argument is indeed the target local file for receiving `mdir` output.

`mget remote-files`

Expand the remote-files on the remote machine and do a get for each file name thus produced. See `glob` for details on the filename expansion. Resulting file names will then be processed according to case, `ntrans`, and `nmap` settings. Files are transferred into the

local working directory, which can be changed with ``lcd directory'`; new local directories can be created with ``! mkdir directory'`.

`mkdir directory-name`

Make a directory on the remote machine.

`mls remote-files local-file`

Like `nlist`, except multiple remote files may be specified, and the local-file must be specified. If interactive prompting is on, ftp will prompt the user to verify that the last argument is indeed the target local file for receiving mls output.

`mode [mode-name]`

Set the file transfer mode to mode-name. The default mode is "stream" mode.

`modtime file-name`

Show the last modification time of the file on the remote machine.

`mput local-files`

Expand wild cards in the list of local files given as arguments and do a put for each file in the resulting list. See `glob` for details of filename expansion. Resulting file names will then be processed according to `ntrans` and `nmap` settings.

`newer file-name`

Get the file only if the modification time of the remote file is more recent than the file on the current system. If the file does not exist on the current system, the remote file is considered newer. Otherwise, this command is identical to `get`.

`nlist [remote-directory] [local-file]`

Print a list of the files of a directory on the remote machine. If remote-directory is left unspecified, the current working directory is used. If interactive prompting is on, ftp will prompt the user to verify that the last argument is indeed the target local file for receiving nlist output. If no local file is specified, or if local-file is `-`, the output is sent to the terminal.

`nmap [inpattern outpattern]`

Set or unset the filename mapping mechanism. If no arguments are specified, the filename mapping mechanism is unset. If arguments are specified, remote filenames are mapped during `mput` commands and `put` commands issued without a specified remote target filename. If

arguments are specified, local filenames are mapped during mget commands and get commands issued without a specified local target filename. This command is useful when connecting to a non-UNIX remote computer with different file naming conventions or practices. The mapping follows the pattern set by inpattern and outpattern. Inpattern is a template for incoming filenames (which may have already been processed according to the ntrans and case settings). Variable templating is accomplished by including the sequences '\$1', '\$2', ..., '\$9' in inpattern. Use '\\' to prevent this special treatment of the '\$' character. All other characters are treated literally, and are used to determine the nmap inpattern variable values. For example, given inpattern \$1.\$2 and the remote file name "mydata.data", \$1 would have the value "mydata", and \$2 would have the value "data". The outpattern determines the resulting mapped filename. The sequences '\$1', '\$2', ..., '\$9' are replaced by any value resulting from the inpattern template. The sequence '\$0' is replaced by the original filename. Additionally, the sequence '[seq1,seq2]' is replaced by seq1 if seq1 is not a null string; otherwise it is replaced by seq2. For example, the command "nmap \$1.\$2.\$3 [\$1,\$2].[\$2,file]" would yield the output filename "myfile.data" for input filenames "myfile.data" and "myfile.data.old", "myfile.file" for the input filename "myfile", and "myfile.myfile" for the input filename ".myfile". Spaces may be included in outpattern, as in the example: nmap \$1 |sed "s/ *\$//" > \$1 . Use the '\\' character to prevent special treatment of the '\$', '[', ']', and ',' characters.

ntrans [inchars [outchars]]

Set or unset the filename character translation mechanism. If no arguments are specified, the filename character translation mechanism is unset. If arguments are specified, characters in remote filenames are translated during mput commands and put commands issued without a specified remote target filename. If arguments are specified, characters in local filenames are translated during mget commands and get commands issued without a specified local target filename. This command is useful when connecting to a non-UNIX remote computer with different file naming conventions or practices. Characters in a filename matching a character in inchars are replaced with the corresponding character in outchars. If the character's position in inchars is longer than the length of outchars, the character is deleted from the file name.

open host [port]

Establish a connection to the specified host FTP server. An optional port number may be supplied, in which case, ftp will attempt to contact an FTP server at that port. If the auto-login option is on (default), ftp will also attempt to automatically log the user in to the FTP server (see below).

prompt

Toggle interactive prompting. Interactive prompting occurs during multiple file transfers to allow the user to selectively retrieve or store files. If prompting is turned off (default is on), any mget or mput will transfer all files, and any mdelete will delete all files.

proxy ftp-command

Execute an ftp command on a secondary control connection. This command allows simultaneous connection to two remote ftp servers for transferring files between the two servers. The first proxy command should be an open, to establish the secondary control connection. Enter the command "proxy ?" to see other ftp commands executable on the secondary connection. The following commands behave differently when prefaced by proxy: open will not define new macros during the auto-login process, close will not erase existing macro definitions, get and mget transfer files from the host on the primary control connection to the host on the secondary control connection, and put, mput, and append transfer files from the host on the secondary control connection to the host on the primary control connection. Third party file transfers depend upon support of the ftp protocol PASV command by the server on the secondary control connection.

put local-file [remote-file]

Store a local file on the remote machine. If remote-file is left unspecified, the local file name is used after processing according to any ntrans or nmap settings in naming the remote file. File transfer uses the current settings for type, format, mode, and structure.

pwd Print the name of the current working directory on the remote machine.

quit A synonym for bye.

quote arg1 arg2 ...

The arguments specified are sent, verbatim, to the remote FTP server.

`recv remote-file [local-file]`
A synonym for `get`.

`reget remote-file [local-file]`
Reget acts like `get`, except that if `local-file` exists and is smaller than `remote-file`, `local-file` is presumed to be a partially transferred copy of `remote-file` and the transfer is continued from the apparent point of failure. This command is useful when transferring very large files over networks that are prone to dropping connections.

`remotehelp [command-name]`
Request help from the remote FTP server. If a `command-name` is specified it is supplied to the server as well.

`remotestatus [file-name]`
With no arguments, show status of remote machine. If `file-name` is specified, show status of `file-name` on remote machine.

`rename [from] [to]`
Rename the file `from` on the remote machine, to the file `to`.

`reset`
Clear reply queue. This command re-synchronizes command/reply sequencing with the remote ftp server. Resynchronization may be necessary following a violation of the ftp protocol by the remote server.

`restart marker`
Restart the immediately following `get` or `put` at the indicated marker. On UNIX systems, marker is usually a byte offset into the file.

`rmdir directory-name`
Delete a directory on the remote machine.

`runique`
Toggle storing of files on the local system with unique filenames. If a file already exists with a name equal to the target local filename for a `get` or `mget` command, a ".1" is appended to the name. If the resulting name matches another existing file, a ".2" is appended to the original name. If this process continues up to ".99", an error message is printed, and the transfer does not take place. The generated unique filename will be reported. Note that `runique` will not affect local files generated from a shell command (see below). The default value is off.

send local-file [remote-file]
A synonym for put.

sendport

Toggle the use of PORT commands. By default, ftp will attempt to use a PORT command when establishing a connection for each data transfer. The use of PORT commands can prevent delays when performing multiple file transfers. If the PORT command fails, ftp will use the default data port. When the use of PORT commands is disabled, no attempt will be made to use PORT commands for each data transfer. This is useful for certain FTP implementations which do ignore PORT commands but, incorrectly, indicate they've been accepted.

site arg1 arg2 ...

The arguments specified are sent, verbatim, to the remote FTP server as a SITE command.

size file-name

Return size of file-name on remote machine.

status

Show the current status of ftp.

struct [struct-name]

Set the file transfer structure to struct-name. By default "stream" structure is used.

sunique

Toggle storing of files on remote machine under unique file names. Remote ftp server must support ftp protocol STOU command for successful completion. The remote server will report unique name. Default value is off.

system

Show the type of operating system running on the remote machine.

tenex

Set the file transfer type to that needed to talk to TENEX machines.

trace

Toggle packet tracing.

type [type-name]

Set the file transfer type to type-name. If no type is specified, the current type is printed. The default type is network ASCII.

umask [newmask]

Set the default umask on the remote server to newmask.
If newmask is omitted, the current umask is printed.

user user-name [password] [account]

Identify yourself to the remote FTP server. If the password is not specified and the server requires it, ftp will prompt the user for it (after disabling local echo). If an account field is not specified, and the FTP server requires it, the user will be prompted for it. If an account field is specified, an account command will be relayed to the remote server after the login sequence is completed if the remote server did not require it for logging in. Unless ftp is invoked with "auto-login" disabled, this process is done automatically on initial connection to the FTP server.

verbose

Toggle verbose mode. In verbose mode, all responses from the FTP server are displayed to the user. In addition, if verbose is on, when a file transfer completes, statistics regarding the efficiency of the transfer are reported. By default, verbose is on.

? [command]

A synonym for help.

Command arguments which have embedded spaces may be quoted with quote (") marks.

ABORTING A FILE TRANSFER

To abort a file transfer, use the terminal interrupt key (usually Ctrl-C). Sending transfers will be immediately halted. Receiving transfers will be halted by sending a ftp protocol ABOR command to the remote server, and discarding any further data received. The speed at which this is accomplished depends upon the remote server's support for ABOR processing. If the remote server does not support the ABOR command, an "ftp>" prompt will not appear until the remote server has completed sending the requested file.

The terminal interrupt key sequence will be ignored when ftp has completed any local processing and is awaiting a reply from the remote server. A long delay in this mode may result from the ABOR processing described above, or from unexpected behavior by the remote server, including violations of the ftp protocol. If the delay results from unexpected remote server behavior, the local ftp program must be killed by hand.

FILE NAMING CONVENTIONS

Files specified as arguments to ftp commands are processed according to the following rules.

- 1) If the file name "-" is specified, the stdin (for reading) or stdout (for writing) is used.
- 2) If the first character of the file name is "|", the remainder of the argument is interpreted as a shell command. Ftp then forks a shell, using popen(3) with the argument supplied, and reads (writes) from the stdout (stdin). If the shell command includes spaces, the argument must be quoted; e.g. "| ls -lt". A particularly useful example of this mechanism is: "dir |more".
- 3) Failing the above checks, if ``globbing'' is enabled, local file names are expanded according to the rules used in the csh(1); c.f. the glob command. If the ftp command expects a single local file (.e.g. put), only the first filename generated by the "globbing" operation is used.
- 4) For mget commands and get commands with unspecified local file names, the local filename is the remote filename, which may be altered by a case, ntrans, or nmap setting. The resulting filename may then be altered if runique is on.
- 5) For mput commands and put commands with unspecified remote file names, the remote filename is the local filename, which may be altered by a ntrans or nmap setting. The resulting filename may then be altered by the remote server if sunique is on.

FILE TRANSFER PARAMETERS

The FTP specification specifies many parameters which may affect a file transfer. The type may be one of "ascii", "image" (binary), "ebcdic", and "local byte size" (for PDP-10's and PDP-20's mostly). Ftp supports the ascii and image types of file transfer, plus local byte size 8 for tenex mode transfers.

Ftp supports only the default values for the remaining file transfer parameters: mode, form, and struct.

OPTIONS

Options may be specified at the command line, or to the command interpreter.

The -v (verbose on) option forces ftp to show all responses from the remote server, as well as report on data transfer statistics.

The -n option restrains ftp from attempting "auto-login" upon initial connection. If auto-login is enabled, ftp will

check the `.netrc` (see below) file in the user's home directory for an entry describing an account on the remote machine. If no entry exists, `ftp` will prompt for the remote machine login name (default is the user identity on the local machine), and, if necessary, prompt for a password and an account with which to login.

The `-i` option turns off interactive prompting during multiple file transfers.

The `-d` option enables debugging.

The `-g` option disables file name globbing.

THE `.netrc` FILE

The `.netrc` file contains login and initialization information used by the auto-login process. It resides in the user's home directory. The following tokens are recognized; they may be separated by spaces, tabs, or new-lines:

machine name

Identify a remote machine name. The auto-login process searches the `.netrc` file for a machine token that matches the remote machine specified on the `ftp` command line or as an open command argument. Once a match is made, the subsequent `.netrc` tokens are processed, stopping when the end of file is reached or another machine or a default token is encountered.

default

This is the same as machine name except that default matches any name. There can be only one default token, and it must be after all machine tokens. This is normally used as:

```
default login anonymous password user@site
```

thereby giving the user automatic anonymous `ftp` login to machines not specified in `.netrc`. This can be overridden by using the `-n` flag to disable auto-login.

login name

Identify a user on the remote machine. If this token is present, the auto-login process will initiate a login using the specified name.

password string

Supply a password. If this token is present, the auto-login process will supply the specified string if the remote server requires a password as part of the login process. Note that if this token is present in the `.netrc` file for any user other than anonymous, `ftp` will abort the auto-login process if the `.netrc` is readable by anyone besides the user.

account string

Supply an additional account password. If this token is present, the auto-login process will supply the specified string if the remote server requires an additional account password, or the auto-login process will initiate an ACCT command if it does not.

macdef name

Define a macro. This token functions like the ftp macdef command functions. A macro is defined with the specified name; its contents begin with the next .netrc line and continue until a null line (consecutive new-line characters) is encountered. If a macro named init is defined, it is automatically executed as the last step in the auto-login process.

SEE ALSO

ftpd(8)

BUGS

Correct execution of many commands depends upon proper behavior by the remote server.

An error in the treatment of carriage returns in the 4.2BSD UNIX ascii-mode transfer code has been corrected. This correction may result in incorrect transfers of binary files to and from 4.2BSD servers using the ascii type. Avoid this problem by using the binary image type.

NAME

gcore - get core image of running process

SYNOPSIS

gcore [-s][-c core] pid

DESCRIPTION

gcore creates a core image of each specified process, suitable for use with adb(1). By default the core image is written to the file <pid>.core.

The options are:

- c Write the core file to the specified file instead of <pid>.core.
- s Stop the process while creating the core image and resume it when done. This makes sure that the core dump will be in a consistent state. The process is resumed even if it was already stopped. Of course, you can obtain the same result by manually stopping the process with kill(1).

The core image name was changed from core.<pid> to <pid>.core to prevent matching names like core.h and core.c when using programs such as find(1).

FILES

<process-id>.core The core image.

BUGS

If gcore encounters an error while creating the core image and the -s option was used the process will remain stopped.

Swapped out processes and system processes (the swapper) may not be gcore'd.

NAME

graph - draw a graph

SYNOPSIS

graph [option] ...

DESCRIPTION

Graph with no options takes pairs of numbers from the standard input as abscissas and ordinates of a graph. Successive points are connected by straight lines. The graph is encoded on the standard output for display by the plot(1G) filters.

If the coordinates of a point are followed by a nonnumeric string, that string is printed as a label beginning on the point. Labels may be surrounded with quotes "...", in which case they may be empty or contain blanks and numbers; labels never contain newlines.

The following options are recognized, each as a separate argument.

- a Supply abscissas automatically (they are missing from the input); spacing is given by the next argument (default 1). A second optional argument is the starting point for automatic abscissas (default 0 or lower limit given by -x).
- b Break (disconnect) the graph after each label in the input.
- c Character string given by next argument is default label for each point.
- g Next argument is grid style, 0 no grid, 1 frame with ticks, 2 full grid (default).
- l Next argument is label for graph.
- m Next argument is mode (style) of connecting lines: 0 disconnected, 1 connected (default). Some devices give distinguishable line styles for other small integers.
- s Save screen, don't erase before plotting.
- x [1]
If 1 is present, x axis is logarithmic. Next 1 (or 2) arguments are lower (and upper) x limits. Third argument, if present, is grid spacing on x axis. Normally these quantities are determined automatically.
- y [1]

Similarly for y.

- h Next argument is fraction of space for height.
- w Similarly for width.
- r Next argument is fraction of space to move right before plotting.
- u Similarly to move up before plotting.
- t Transpose horizontal and vertical axes. (Option -x now applies to the vertical axis.)

A legend indicating grid range is produced with a grid unless the -s option is present.

If a specified lower limit exceeds the upper limit, the axis is reversed.

SEE ALSO

spline(1G), plot(1G)

BUGS

Graph stores all points internally and drops those for which there isn't room.

Segments that run out of bounds are dropped, not windowed.

Logarithmic axes may not be reversed.

NAME

grep, egrep, fgrep - search a file for a pattern

SYNOPSIS

grep [option] ... expression [file] ...

egrep [option] ... [expression] [file] ...

fgrep [option] ... [strings] [file]

DESCRIPTION

Commands of the grep family search the input files (standard input default) for lines matching a pattern. Normally, each line found is copied to the standard output. Grep patterns are limited regular expressions in the style of `ex(1)`; it uses a compact nondeterministic algorithm. Egrep patterns are full regular expressions; it uses a fast deterministic algorithm that sometimes needs exponential space. Fgrep patterns are fixed strings; it is fast and compact. The following options are recognized.

- v All lines but those matching are printed.
- x (Exact) only lines matched in their entirety are printed (fgrep only).
- c Only a count of matching lines is printed.
- l The names of files with matching lines are listed (once) separated by newlines.
- n Each line is preceded by its relative line number in the file.
- b Each line is preceded by the block number on which it was found. This is sometimes useful in locating disk block numbers by context.
- i The case of letters is ignored in making comparisons - that is, upper and lower case are considered identical. This applies to grep and fgrep only.
- s Silent mode. Nothing is printed (except error messages). This is useful for checking the error status.
- w The expression is searched for as a word (as if surrounded by ``\<'` and ``\>'`, see `ex(1)`.) (grep only)
- e expression
Same as a simple expression argument, but useful when the expression begins with a `-`.

`-f file`

The regular expression (egrep) or string list (fgrep) is taken from the file.

In all cases the file name is shown if there is more than one input file. Care should be taken when using the characters `$ * [^ | ()` and `\` in the expression as they are also meaningful to the Shell. It is safest to enclose the entire expression argument in single quotes `' '`.

Fgrep searches for lines that contain one of the (newline-separated) strings.

Egrep accepts extended regular expressions. In the following description ``character'` excludes newline:

A `\` followed by a single character other than newline matches that character.

The character `^` matches the beginning of a line.

The character `$` matches the end of a line.

A `.` (period) matches any character.

A single character not otherwise endowed with special meaning matches that character.

A string enclosed in brackets `[]` matches any single character from the string. Ranges of ASCII character codes may be abbreviated as in ``a-z0-9'`. A `]` may occur only as the first character of the string. A literal `-` must be placed where it can't be mistaken as a range indicator.

A regular expression followed by an `*` (asterisk) matches a sequence of 0 or more matches of the regular expression. A regular expression followed by a `+` (plus) matches a sequence of 1 or more matches of the regular expression. A regular expression followed by a `?` (question mark) matches a sequence of 0 or 1 matches of the regular expression.

Two regular expressions concatenated match a match of the first followed by a match of the second.

Two regular expressions separated by `|` or newline match either a match for the first or a match for the second.

A regular expression enclosed in parentheses matches a match for the regular expression.

The order of precedence of operators at the same parenthesis level is [] then *+? then concatenation then | and newline.

Ideally there should be only one grep, but we don't know a single algorithm that spans a wide enough range of space-time tradeoffs.

SEE ALSO

ex(1), sed(1), sh(1)

DIAGNOSTICS

Exit status is 0 if any matches are found, 1 if none, 2 for syntax errors or inaccessible files.

BUGS

Lines are limited to 256 characters; longer lines are truncated.

NAME

groups - show group memberships

SYNOPSIS

groups [user]

DESCRIPTION

The groups utility has been obsoleted by the id(1) utility, and is equivalent to ``id -Gn [user]''. The command ``id -p'' is suggested for normal interactive use.

The groups utility displays the groups to which you (or the optionally specified user) belong.

The groups utility exits 0 on success, and >0 if an error occurs.

SEE ALSO

id(1)

NAME

head - give first few lines

SYNOPSIS

head [-count] [file ...]

DESCRIPTION

This filter gives the first count lines of each of the specified files, or of the standard input. If count is omitted it defaults to 10.

SEE ALSO

tail(1)

NAME

hostid - set or print identifier of current host system

SYNOPSIS

hostid [identifier]

DESCRIPTION

The hostid command prints the identifier of the current host in hexadecimal. This numeric value is expected to be unique across all hosts and is commonly set to the host's Internet address. The super-user can set the hostid by giving a hexadecimal argument or the hostname; this is usually done in the startup script /etc/rc.local.

SEE ALSO

gethostid(2), sethostid(2)

NAME

hostname - set or print name of current host system

SYNOPSIS

hostname [-s] [nameofhost]

DESCRIPTION

Hostname prints the name of the current host. The super-user can set the hostname by supplying an argument; this is usually done in the network initialization script /etc/netstart , normally run at boot time.

Options:

-s Trims off any domain information from the printed name.

SEE ALSO

gethostname(2), sethostname(2)

HISTORY

The hostname command appeared in BSD 4.2 .

NAME

indent - indent and format C program source

SYNOPSIS

```
indent [ input-file [ output-file ] ] [ -bad | -nbad ]
      [ -bap | -nbap ] [ -bbb | -nbbb ] [ -bc | -nbc ]
      [ -bl | -br ] [ -cn ] [ -cdn ] [ -cdb | -ncdb ]
      [ -ce | -nce ] [ -cin ] [ -clin ] [ -dn ] [ -din ]
      [ -dj | -ndj ] [ -ei | -nei ] [ -fc1 | -nfc1 ]
      [ -in ] [ -ip | -nip ] [ -ln ] [ -lcn ]
      [ -lp | -nlp ] [ -npro ] [ -pcs | -npcs ]
      [ -ps | -nps ] [ -psl | -npsl ] [ -sc | -nsc ]
      [ -sob | -nsob ] [ -st ] [ -troff ] [ -v | -nv ]
```

DESCRIPTION

Indent is a C program formatter. It reformats the C program in the input-file according to the switches. The switches which can be specified are described below. They may appear before or after the file names.

NOTE: If you only specify an input-file, the formatting is done 'in-place', that is, the formatted file is written back into input-file and a backup copy of input-file is written in the current directory. If input-file is named '/blah/blah/file', the backup file is named file.BAK.

If output-file is specified, indent checks to make sure it is different from input-file.

OPTIONS

The options listed below control the formatting style imposed by indent.

-bad,-nbad If -bad is specified, a blank line is forced after every block of declarations. Default: -nbad.

-bap,-nbap If -bap is specified, a blank line is forced after every procedure body. Default: -nbap.

-bbb,-nbbb If -bbb is specified, a blank line is forced before every block comment. Default: -nbbb.

-bc,-nbc If -bc is specified, then a newline is forced after each comma in a declaration. -nbc turns off this option. The default is -nbc.

-br,-bl Specifying -bl lines up compound statements like this:
 if (...)
 {
 code

- ```
 }
 Specifying -br (the default) makes them look
 like this:
 if (...) {
 code
 }
```
- cn        The column in which comments on code start.  
The default is 33.
- cdn       The column in which comments on declarations  
start. The default is for these comments to  
start in the same column as those on code.
- cdb,-ncdb    Enables (disables) the placement of comment  
delimiters on blank lines. With this option  
enabled, comments look like this:
- ```
    /*
     * this is a comment
     */
```
- Rather than like this:
- ```
 /* this is a comment */
```
- This only affects block comments, not com-  
ments to the right of code. The default is  
-cdb.
- ce,-nce       Enables (disables) forcing `else's to cuddle  
up to the immediately preceding `}'. The  
default is -ce.
- cin        Sets the continuation indent to be n. Con-  
tinuation lines will be indented that far  
from the beginning of the first line of the  
statement. Parenthesized expressions have  
extra indentation added to indicate the nest-  
ing, unless -lp is in effect. -ci defaults  
to the same value as -i.
- clin       Causes case labels to be indented n tab stops  
to the right of the containing switch state-  
ment. -cli0.5 causes case labels to be  
indented half a tab stop. The default is  
-cli0. (This is the only option that takes a  
fractional argument.)
- dn        Controls the placement of comments which are  
not to the right of code. Specifying -d1  
means that such comments are placed one  
indentation level to the left of code. The  
default -d0 lines up these comments with the  
code. See the section on comment indentation  
below.

- `-din` Specifies the indentation, in character positions, from a declaration keyword to the following identifier. The default is `-dil6`.
- `-dj,-ndj` `-dj` left justifies declarations. `-ndj` indents declarations the same as code. The default is `-ndj`.
- `-ei,-nei` Enables (disables) special else-if processing. If enabled, ifs following elses will have the same indentation as the preceding if statement. The default is `-ei`.
- `-fc1,-nfc1` Enables (disables) the formatting of comments that start in column 1. Often, comments whose leading ``/'` is in column 1 have been carefully hand formatted by the programmer. In such cases, `-nfc1` should be used. The default is `-fc1`.
- `-in` The number of spaces for one indentation level. The default is 8.
- `-ip,-nip` Enables (disables) the indentation of parameter declarations from the left margin. The default is `-ip`.
- `-ln` Maximum length of an output line. The default is 78.
- `-lp,-nlp` Lines up code surrounded by parenthesis in continuation lines. If a line has a left paren which is not closed on that line, then continuation lines will be lined up to start at the character position just after the left paren. For example, here is how a piece of continued code looks with `-nlp` in effect:
- ```
p1 = first_procedure(second_procedure(p2, p3),
    third_procedure(p4, p5));
```
- With `-lp` in effect (the default) the code looks somewhat clearer:
- ```
p1 = first_procedure(second_procedure(p2, p3),
 third_procedure(p4, p5));
```
- Inserting two more newlines we get:
- ```
p1 = first_procedure(second_procedure(p2,
    p3),
    third_procedure(p4,
    p5));
```
- `-npro` Causes the profile files, ``./.indent.pro'` and ``~/.indent.pro'`, to be ignored.

- `-pcs,-npcs` If true (`-pcs`) all procedure calls will have a space inserted between the name and the ``('`. The default is `-npcs`.
- `-ps,-nps` If true (`-ps`) the pointer following operator ``->'` will be surrounded by spaces on either side. The default is `-nps`.
- `-psl,-npsl` If true (`-psl`) the names of procedures being defined are placed in column 1 - their types, if any, will be left on the previous lines. The default is `-psl`.
- `-sc,-nsc` Enables (disables) the placement of asterisks (``*'s`) at the left edge of all comments. The default is `-sc`.
- `-sob,-nsob` If `-sob` is specified, indent will swallow optional blank lines. You can use this to get rid of blank lines after declarations. Default: `-nsob`.
- `-st` Causes indent to take its input from stdin, and put its output to stdout.
- `-Ttypename` Adds typename to the list of type keywords. Names accumulate: `-T` can be specified more than once. You need to specify all the typenames that appear in your program that are defined by typedefs - nothing will be harmed if you miss a few, but the program won't be formatted as nicely as it should. This sounds like a painful thing to have to do, but it's really a symptom of a problem in C: typedef causes a syntactic change in the language and indent can't find all typedefs.
- `-troff` Causes indent to format the program for processing by troff. It will produce a fancy listing in much the same spirit as vgrind. If the output file is not specified, the default is standard output, rather than formatting in place.
- `-v,-nv` `-v` turns on ``verbose'` mode; `-nv` turns it off. When in verbose mode, indent reports when it splits one line of input into two or more lines of output, and gives some size statistics at completion. The default is `-nv`.

FURTHER DESCRIPTION

You may set up your own 'profile' of defaults to indent by creating a file called `.indent.pro` in either your login directory and/or the current directory and including whatever switches you like. Switches in `.indent.pro` in the current directory override those in your login directory (with the exception of `-T` type definitions, which just accumulate). If `indent` is run and a profile file exists, then it is read to set up the program's defaults. The switches should be separated by spaces, tabs or newlines. Switches on the command line, however, override profile switches.

Comments

'Box' comments. Indent assumes that any comment with a dash or star immediately after the start of comment (that is, `/*-` or `/**`) is a comment surrounded by a box of stars. Each line of such a comment is left unchanged, except that its indentation may be adjusted to account for the change in indentation of the first line of the comment.

Straight text. All other comments are treated as straight text. Indent fits as many words (separated by blanks, tabs, or newlines) on a line as possible. Blank lines break paragraphs.

Comment indentation

If a comment is on a line with code it is started in the 'comment column', which is set by the `-cn` command line parameter. Otherwise, the comment is started at `n` indentation levels less than where code is currently being placed, where `n` is specified by the `-dn` command line parameter. If the code on a line extends past the comment column, the comment starts further to the right, and the right margin may be automatically extended in extreme cases.

Preprocessor lines

In general, `indent` leaves preprocessor lines alone. The only reformatting that it will do is to straighten up trailing comments. It leaves embedded comments alone. Conditional compilation (`#ifdef...#endif`) is recognized and `indent` attempts to correctly compensate for the syntactic peculiarities introduced.

C syntax

Indent understands a substantial amount about the syntax of C, but it has a 'forgiving' parser. It attempts to cope with the usual sorts of incomplete and misformed syntax. In particular, the use of macros like:

```
#define forever for(;;)
```

is handled properly.

FILES

./indent.pro profile file
~/.indent.pro profile file

BUGS

Indent has even more switches than ls.

A common mistake that often causes grief is typing:

```
indent *.c
```

to the shell in an attempt to indent all the C programs in a directory. This is probably a bug, not a feature.

NAME

install - install binaries

SYNOPSIS

install [-cs] [-f flags] [-g group] [-m mode] [-o owner]
file1 file2

install [-cs] [-f flags] [-g group] [-m mode] [-o owner]
file1 ... fileN directory

DESCRIPTION

The file(s) are moved (or copied if the -c option is specified) to the target file or directory. If the destination is a directory, then the file is moved into directory with its original filename. If the target file already exists, it is overwritten if permissions allow.

- c Copy the file. This flag turns off the default behavior of install where it deletes the original file after creating the target.
- f Specify the target's file flags. (See chflags(1) for a list of possible flags and their meanings.)
- g Specify a group.
- m Specify an alternate mode. The default mode is set to rwxr-xr-x (0755). The specified mode may be either an octal or symbolic value; see chmod(1) for a description of possible mode values.
- o Specify an owner.
- s Install exec's the command strip(1) to strip binaries so that install can be portable over a large number of systems and binary types.

By default, install preserves all file flags, with the exception of the ``nodump'' flag.

The install utility attempts to prevent moving a file onto itself.

Installing /dev/null creates an empty file.

Upon successful completion a value of 0 is returned. Otherwise, a value of 1 is returned.

SEE ALSO

chflags(1), chgrp(1), chmod(1), cp(1), mv(1), strip(1), chown(8)

HISTORY

The install utility appeared in 4.2BSD.

NAME

iostat - report I/O statistics

SYNOPSIS

iostat [drives] [interval [count]]

DESCRIPTION

Iostat iteratively reports the number of characters read and written to terminals per second, and, for each disk, the number of transfers per second, kilobytes transferred per second, and the milliseconds per average seek. It also gives the percentage of time the system has spent in user mode, in user mode running low priority (niced) processes, in system mode, and idling.

To compute this information, for each disk, seeks and data transfer completions and number of words transferred are counted; for terminals collectively, the number of input and output characters are counted. Also, each sixtieth of a second, the state of each disk is examined and a tally is made if the disk is active. From these numbers and given the transfer rates of the devices it is possible to determine average seek times for each device.

The optional interval argument causes iostat to report once each interval seconds. The first report is for all time since a reboot and each subsequent report is for the last interval only.

The optional count argument restricts the number of reports.

If more than 4 disk drives are configured in the system, iostat displays only the first 4 drives, with priority given to Massbus disk drives (i.e. if both Unibus and Massbus drives are present and the total number of drives exceeds 4, then some number of Unibus drives will not be displayed in favor of the Massbus drives). To force iostat to display specific drives, their names may be supplied on the command line.

FILES

/dev/kmem
/vmunix

SEE ALSO

vmstat(1)

NAME

join - relational database operator

SYNOPSIS

join [options] file1 file2

DESCRIPTION

Join forms, on the standard output, a join of the two relations specified by the lines of file1 and file2. If file1 is '-', the standard input is used.

File1 and file2 must be sorted in increasing ASCII collating sequence on the fields on which they are to be joined, normally the first in each line.

There is one line in the output for each pair of lines in file1 and file2 that have identical join fields. The output line normally consists of the common field, then the rest of the line from file1, then the rest of the line from file2.

Fields are normally separated by blank, tab or newline. In this case, multiple separators count as one, and leading separators are discarded.

These options are recognized:

-an In addition to the normal output, produce a line for each unpairable line in file n, where n is 1 or 2.

-e s Replace empty output fields by string s.

-jn m
Join on the mth field of file n. If n is missing, use the mth field in each file.

-o list
Each output line comprises the fields specified in list, each element of which has the form n.m, where n is a file number and m is a field number.

-tc Use character c as a separator (tab character). Every appearance of c in a line is significant.

SEE ALSO

sort(1), comm(1), awk(1)

BUGS

With default field separation, the collating sequence is that of sort -b; with -t, the sequence is that of a plain sort.

The conventions of join, sort, comm, uniq, look and awk(1) are wildly incongruous.

NAME

jove - an interactive display-oriented text editor

SYNOPSIS

```
jove [-d directory] [-w] [-t tag] [+[n] file] [-p file] [files]
jove -r
```

DESCRIPTION

JOVE is Jonathan's Own Version of Emacs. It is based on the original EMACS editor written at MIT by Richard Stallman. Although JOVE is meant to be compatible with EMACS, there are some major differences between the two editors and you shouldn't rely on their behaving identically.

JOVE works on any reasonable display terminal that is described in the termcap file (see TERMCAP(5) for more details). When you start up JOVE, it checks to see whether you have your TERM environment variable set. On most systems that will automatically be set up for you, but if it's not JOVE will ask you what kind of terminal you are using. To avoid having to type this every time you run JOVE you can set your TERM environment variable yourself. How you do this depends on which shell you are running. If you are running the C Shell, as most of you are, you type

```
% setenv TERM type
```

and with the Bourne Shell, you type

```
$ TERM= type ; export TERM
```

where type is the name of the kind of terminal you are using (e.g., vt100). If neither of these works get somebody to help you.

INVOKING JOVE

If you run JOVE with no arguments you will be placed in an empty buffer, called Main. Otherwise, any arguments you supply are considered file names and each is "given" its own buffer. Only the first file is actually read in--reading other files is deferred until you actually try to use the buffers they are attached to. This is for efficiency's sake: most of the time, when you run JOVE on a big list of files, you end up editing only a few of them.

The names of all of the files specified on the command line are saved in a buffer, called *minibuf*. The mini-buffer is a special JOVE buffer that is used when JOVE is prompting for some input to many commands (for example, when JOVE is prompting for a file name). When you are being prompted for a file name, you can type C-N (that's Control-N) and C-P to cycle through the list of files that were specified on the

command line. The file name will be inserted where you are typing and then you can edit it as if you typed it in yourself.

JOVE recognizes the following switches:

-d The following argument is taken to be the name of the current directory. This is for systems that don't have a version of C shell that automatically maintains the CWD environment variable. If -d is not specified on a system without a modified C shell, JOVE will have to figure out the current directory itself, and that can be VERY slow. You can simulate the modified C shell by putting the following lines in your C shell initialization file (.cshrc):

```
alias cd      'cd \!*; setenv CWD $cwd'
alias popd    'popd \!*; setenv CWD $cwd'
alias pushd   'pushd \!*; setenv CWD $cwd'
```

+n Reads the file, designated by the following argument, and positions point at the n'th line instead of the (default) 1'st line. This can be specified more than once but it doesn't make sense to use it twice on the same file; in that case the second one wins. If no numeric argument is given after the +, the point is positioned at the end of the file.

-p Parses the error messages in the file designated by the following argument. The error messages are assumed to be in a format similar to the C compiler, LINT, or GREP output.

-t Runs the find-tag command on the string of characters immediately following the -t if there is one (as in -tTagname), or on the following argument (as in -t Tagname) otherwise (see ctags(1)).

-w Divides the window in two. When this happens, either the same file is displayed in both windows, or the second file in the list is read in and displayed in its window.

RECOVERING BUFFERS AFTER A CRASH

The -r option of jove runs the JOVE recover program. Use this when the system crashes, or JOVE crashes, or you accidentally get logged out while in JOVE. If there are any buffers to be recovered, this will find them.

Recover looks for JOVE buffers that are left around and are owned by you. (You cannot recover other peoples' buffers, obviously.) If there were no buffers that were modified at

the time of the crash or there were but recover can't get its hands on them, you will be informed with the message, "There is nothing to recover." Otherwise, recover prints the date and time of the version of the buffers it has, and then waits for you type a command.

To get a list of the buffers recover knows about, use the list command. This will list all the buffers and the files and the number of lines associated with them. Next to each buffer is a number. When you want to recover a buffer, use the get command. The syntax is get buffer filename where buffer is either the buffer's name or the number at the beginning of the line. If you don't type the buffer name or the filename, recover will prompt you for them.

If there are a lot of buffers and you want to recover all of them, use the recover command. This will recover each buffer to the name of the buffer with ".#" prepended to the name (so that the original isn't over-written). It asks for each file and if you want to restore that buffer to that name you type "yes". If you want to recover the file but to a different name, just type the new name in. If you type "no" recover will skip that file and go on to the next one.

If you want to look at a buffer before deciding to recover it, use the print command. The syntax for this is print buffer where buffer again is either its name or the number. You can type ^C if you want to abort printing the file to the terminal, and recover will respond with an appropriate message.

When you're done and have all the buffers you want, type the quit command to leave. You will then be asked whether it's okay to delete the tmp files. Most of the time that's okay and you should type "yes". When you say that, JOVE removes all traces of those buffers and you won't be able to look at them again. (If you recovered some buffers they will still be around, so don't worry.) So, if you're not sure whether you've gotten all the buffers, you should answer "no" so that you'll be able to run recover again at a later time (presumably after you've figured out which ones you want to save).

If you type ^C at any time other than when you're printing a file to the terminal, recover will exit without a word. If you do this but wish you hadn't, just type "jove -r" to the shell again, and you will be put back with no loss.

GETTING HELP

Once in JOVE, there are several commands available to get help. To execute any JOVE command, you type "<ESC> X command-name" followed by <Return>. To get a list of all

the JOVE commands you type "<ESC> X" followed by "?". The describe-bindings command can be used to get a list containing each key, and its associated command (that is, the command that gets executed when you type that key). If you want to save the list of bindings, you can set the jove variable send-typeout-to-buffer to ON (using the set command), and then execute the describe-bindings command. This will create a buffer and put in it the bindings list it normally would have printed on the screen. Then you can save that buffer to a file and print it to use as a quick reference card. (See VARIABLES below.)

Once you know the name of a command, you can find out what it does with the describe-command command, which you can invoke quickly by typing "ESC ?". The apropos command will give you a list of all the command with a specific string in their names. For example, if you want to know the names of all the commands that are concerned with windows, you can run "apropos" with the keyword window.

If you're not familiar with the EMACS command set, it would be worth your while to use run TEACHJOVE. Do do that, just type "teachjove" to your shell and you will be placed in JOVE in a file which contains directions. I highly recommend this for beginners; you may save yourself a lot of time and headaches.

KEY BINDINGS and VARIABLES

You can alter the key bindings in JOVE to fit your personal tastes. That is, you can change what a key does every time you strike it. For example, by default the C-N key is bound to the command next-line and so when you type it you move down a line. If you want to change a binding or add a new one, you use the bind-to-key command. The syntax is "bind-to-key <command> key".

You can also change the way JOVE behaves in little ways by changing the value of some variables with the set command. The syntax is "set <variable> value", where value is a number or a string, or "on" or "off", depending on the context. For example, if you want JOVE to make backup files, you set the "make-backup-files" variable to "on". To see the value of a variable, use the "print <variable>" command.

INITIALIZATION

JOVE automatically reads commands from an initialization file in your HOME directory, called ".joverc". In this file you can place commands that you would normally type in JOVE. If you like to rearrange the key bindings and set some variables every time you get into JOVE, you should put them in your initialization file. Here are a few lines from mine:

```
set match-regular-expressions on
```

```

auto-execute-command auto-fill /tmp/Re\|.*drft
bind-to-key i-search-forward ^\
bind-to-key i-search-reverse ^R
bind-to-key find-tag-at-point ^[ ^T
bind-to-key scroll-down ^C
bind-to-key grow-window ^Xg
bind-to-key shrink-window ^Xs

```

(Note that the Control Characters can be either two character sequences (e.g. ^ and C together as ^C) or the actual control character. If you want to use an ^ by itself you must BackSlash it (e.g., bind-to-key grow-window ^X\^ binds grow-window to "^X^").

SOME MINOR DETAILS

You should type C-\ instead of C-S in many instances. For example, the way to search for a string is documented as being "C-S" but in reality you should type "C-\". This is because C-S is the XOFF character (what gets sent when you type the NO SCROLL key), and clearly that won't work. The XON character is "C-Q" (what gets sent when you type NO SCROLL again) which is documented as the way to do a quoted-insert. The alternate key for this is "C-^" (typed as "C-`" on vt100's and its look-alikes). If you want to enable C-S and C-Q and you know what you are doing, you can put the line:

```

set allow-^S-and-^Q on
in your initialization file.

```

If your terminal has a metakey, JOVE will use it if you turn on the "meta-key" variable. JOVE will automatically turn on "meta-key" if the METAKEY environment variable exists. This is useful for if you have different terminals (e.g., one at home and one at work) and one has a metakey and the other doesn't.

FILES

```

/usr/new/lib/jove/.joverc - system wide initialization file
~/.joverc - personal initialization file
/tmp - where temporary files are stored
/usr/new/lib/jove/teach-jove - the interactive tutorial
/usr/new/lib/jove/portsrv - for running shells in windows
(pdp11 only)

```

SEE ALSO

```

ed(1) - for a description of regular expressions
teachjove(1) - for an interactive JOVE tutorial.

```

DIAGNOSTICS

JOVE diagnostics are meant to be self-explanatory, but you are advised to seek help whenever you are confused. You can easily lose a lot of work if you don't know EXACTLY what you are doing.

BUGS

Lines can't be more than 1024 characters long.

Searches can't cross line boundaries.

AUTHOR

Jonathan Payne

NAME

kill - terminate a process with extreme prejudice

SYNOPSIS

```
kill [ -sig ] processid ...  
kill -l
```

DESCRIPTION

Kill sends the TERM (terminate, 15) signal to the specified processes. If a signal name or number preceded by '-' is given as first argument, that signal is sent instead of terminate (see sigvec(2)). The signal names are listed by 'kill -l', and are as given in /usr/include/signal.h, stripped of the common SIG prefix.

The terminate signal will kill processes that do not catch the signal; 'kill -9 ...' is a sure kill, as the KILL (9) signal cannot be caught. By convention, if process number 0 is specified, all members in the process group (i.e. processes resulting from the current login) are signaled (but beware: this works only if you use sh(1); not if you use csh(1).) Negative process numbers also have special meanings; see kill(2) for details.

The killed processes must belong to the current user unless he is the super-user.

The process number of an asynchronous process started with '&' is reported by the shell. Process numbers can also be found by using ps(1). Kill is a built-in to csh(1); it allows job specifiers of the form '%...' as arguments so process id's are not as often used as kill arguments. See csh(1) for details.

SEE ALSO

csh(1), ps(1), kill(2), sigvec(2)

BUGS

A replacement for 'kill 0' for csh(1) users should be provided.

NAME

last - indicate last logins of users and teletypes

SYNOPSIS

last [-f filename] [-N] [name ...] [tty ...]

DESCRIPTION

Last will look back in the wtmp file which records all logins and logouts for information about a user, a teletype or any group of users and teletypes. Arguments specify names of users or teletypes of interest. Names of teletypes may be given fully or abbreviated. For example 'last 0' is the same as 'last tty0'. If multiple arguments are given, the information which applies to any of the arguments is printed. For example 'last root console' would list all of "root's" sessions as well as all sessions on the console terminal. Last will print the sessions of the specified users and teletypes, most recent first, indicating the times at which the session began, the duration of the session, and the teletype which the session took place on. If the session is still continuing or was cut short by a reboot, last so indicates.

The pseudo-user reboot logs in at reboots of the system, thus

last reboot

will give an indication of mean time between reboot.

Last with no arguments prints a record of all logins and logouts, in reverse order. The -f filename option allows the user to examine an alternate wtmp file. The -N option limits the report to N lines.

If last is interrupted, it indicates how far the search has progressed in wtmp. If interrupted with a quit signal (generated by a control-\) last indicates how far the search has progressed so far, and the search continues.

FILES

/usr/adm/wtmp login data base
/usr/adm/shutdownlog which records shutdowns and reasons
for same

SEE ALSO

wtmp(5), ac(8), lastcomm(1)

AUTHOR

Howard Katseff

NAME

lastcomm - show last commands executed in reverse order

SYNOPSIS

lastcomm [-f file] [command name] ... [user name] ...
[terminal name] ...

DESCRIPTION

Lastcomm gives information on previously executed commands.

Option:

-f file Read from file rather than the default accounting
 file.

With no arguments, lastcomm prints information about all the commands recorded during the current accounting file's lifetime. If called with arguments, only accounting entries with a matching command name, user name, or terminal name are printed. So, for example,

```
lastcomm a.out root ttyd0
```

would produce a listing of all the executions of commands named a.out by user root on the terminal ttyd0.

For each process entry, the following are printed.

- The name of the user who ran the process.
- Flags, as accumulated by the accounting facilities in the system.
- The command name under which the process was called.
- The amount of cpu time used by the process (in seconds).
- The time the process exited.

The flags are encoded as follows: ``S'' indicates the command was executed by the super-user, ``F'' indicates the command ran after a fork, but without a following exec, ``C'' indicates the command was run in PDP-11 compatibility mode (VAX only), ``D'' indicates the command terminated with the generation of a core file, and ``X'' indicates the command was terminated with a signal.

FILES

/usr/adm/acct Default accounting file.

SEE ALSO

last(1), sigvec(2), acct(8), core(5)

NAME

ld - link editor (2BSD)

SYNOPSIS

ld [option] ... file ...

DESCRIPTION

Ld combines several object programs into one, resolves external references, and searches libraries. In the simplest case several object files are given, and ld combines them, producing an object module which can be either executed or become the input for a further ld run. (In the latter case, the -r option must be given to preserve the relocation bits.) The output of ld is left on a.out. This file is made executable only if no errors occurred during the load.

The argument routines are concatenated in the order specified. The entry point of the output is the beginning of the first routine (unless the -e option is specified).

If any argument is a library, it is searched exactly once at the point it is encountered in the argument list. Only those routines defining an unresolved external reference are loaded. If a routine from a library references another routine in the library, and the library has not been processed by ranlib(1), the referenced routine must appear after the referencing routine in the library. Thus the order of programs within libraries may be important. The first member of a library should be a file named `__SYMDEF', which is understood to be a dictionary for the library as produced by ranlib(1); the dictionary is searched iteratively to satisfy as many references as possible.

The symbols `_etext', `_edata' and `_end' (`etext', `edata' and `end' in C) are reserved, and if referred to, are set to the first location above the program, the first location above initialized data, and the first location above all data respectively. It is erroneous to define these symbols.

Ld understands several options. Except for -l, they should appear before the file names.

- D Take the next argument as a decimal number and pad the data segment with zero bytes to the indicated length.
- d Force definition of common storage even if the -r flag is present.
- e The following argument is taken to be the name of the entry point of the loaded program; location 0 is the default.

-Ldir

Add dir to the list of directories in which libraries are searched for. Directories specified with -L are searched before the standard directories.

-lx This option is an abbreviation for the library name ``libx.a'`, where x is a string. Ld searches for libraries first in any directories specified with -L options, then in the standard directories ``/lib'`, ``/usr/lib'`, and ``/usr/local/lib'`. A library is searched when its name is encountered, so the placement of a -l is significant.

-M produce a primitive load map, listing the names of the files which will be loaded.

-n Arrange (by giving the output file a 0410 "magic number") that when the output file is executed, the text portion will be read-only and shared among all users executing the file. This involves moving the data areas up to the first possible 8K byte boundary following the end of the text. This option creates a ``pure executable'` format.

-i When the output file is executed, the program text and data areas will live in separate address spaces. The only difference between this option and -n is that here the text and data segments are in separate address spaces and both start at location 0. This option creates a ``separate executable'` format.

-z This option is a synonym for the -i option. On other systems (4.3BSD for example) the -z option causes a demand paged executable to be built. This option was added to 2.11BSD because some systems (those which use gcc) do not safely ignore (with a warning) the -i option. Adding the -z option to 2.11BSD allows makefiles to be copied freely between multiple platforms once again.

-O This is a text replacement overlay file; only the text segment will be replaced by `execve(2)`. Shared data must have the same size as the program overlaid, otherwise the `execve(2)` will fail. The entry point to the overlay may be defined with the -e option. This option allows the creation of a ``replacement executable'` format.

-o The name argument after -o is used as the name of the ld output file, instead of a.out.

-r Generate relocation bits in the output file so that it

can be the subject of another ld run. This flag also prevents final definitions from being given to common symbols, and suppresses the 'undefined symbol' diagnostics. (Note that this option cannot be used with overlays (-Z) since they cannot be reloaded.)

- s 'Strip' the output, that is, remove the symbol table and relocation bits to save space (but impair the usefulness of the debuggers). This information can also be removed by strip(1).
- q ("quiet") Suppress the reporting of undefined symbols. Normally only used when building networked kernels - the large number of undefined symbols is normal (due to the three phase link procedure) but can be distracting none the less.
- t ("trace") Print the name of each file as it is processed.
- u Take the following argument as a symbol and enter it as undefined in the symbol table. This is useful for loading wholly from a library, since initially the symbol table is empty and an unresolved reference is needed to force the loading of the first routine.
- v ("verbose") Print the VM statistics. Printing out the number of pages swapped to and from the VM tmp file is now optional and only used when a problem is suspected (or if you are voyeuristic).
- X Save local symbols except for those whose names begin with 'L'. This option is used by cc(1) to discard internally-generated labels while retaining symbols local to routines.
- x Do not preserve local (non-.globl) symbols in the output symbol table; only enter external symbols. This option saves some space in the output file. It also allows temporary labels to be discarded to prevent redefinition in successive ld's. Warning: adb uses these local symbols, especially when debugging overlaid programs, so some debugging information is necessarily lost if this option is used.
- Z Indicate the creation of an automatic-overlay format. In addition a -i or -n must be present as overlays only work with shared text objects. Repeated instances of -Z bracket the modules that will be loaded into a given overlay. Modules before the first -Z or after the concluding -Y will be loaded into the non-overlaid text (base) area. Note that there may be a maximum of NOVL

(currently 15) overlays. This option produces the 'overlaid pure executable' and the 'overlaid separate executable' formats. The loader creates a small entry interface in the base segment for each subroutine in an overlay. This interface ("thunk") arranges for the correct overlay to be present before the actual routine is entered.

-Y Terminate text overlays. This allows any remaining modules or libraries to be loaded into the base area. Note that the -Y option used to be -L, but had to be changed when the loader was brought up to date with the 4.3BSD loader which uses -L to indicate a directory to be searched for library references.

AUTOMATIC TEXT OVERLAY EXAMPLE

To set up an automatic text overlay object with the loader, use a command of the form:

```
ld -n -X /lib/crt0.o base.o base2.o
-Z ov1a.o ov1b.o ...
-Z ov2a.o ov2b.o ...
-Y base3.o ... -lc
```

Assembly source code must be compiled using the assembler overlay flags: "as -V prog.s" which causes the assembler to leave certain symbols unresolved so that ld may rearrange them. The various system compilers automatically use this option.

When arranging modules into overlays, the following rules control the maximum sizes for an executable file. The magic numbers are due to the granularity of PDP-11 segmentation registers (there are 8 registers, each controlling an 8192-byte segment). The program is made up of four areas: base text, overlay text, data + bss, and stack sections. The size of the overlay section is controlled by the size of the largest of the overlays. Each section starts at an address that is a multiple of 8Kb, thus the size of each section is rounded up to a multiple of 8Kb.

In the case of separate overlaid executable files, the text and overlays share one 64Kb byte address space; and the data + bss and stack share the other. Thus, the total of the base text size (rounded up to an 8Kb boundary) plus the maximum overlay size (similarly rounded) must be less than or equal to 64Kb. Or, put another way, since there are only 8 segmentation registers available, the number of segmentation registers needed for an overlaid object must be less than or equal to 8. As an example, if the base text segment has 36800 bytes and the largest overlay takes 14144, the base

will fit in 5 segments and the overlays in 2 segments; leaving one to spare. The data and bss together must fit in 7 segments (56K bytes), leaving one 8Kb segment for the stack. All of the limits can be checked by using `checkobj(1)`.

For pure overlaid programs, the rules are similar except that all four sections share one 64K-byte address space. The number of segments required by the text, overlay, data + bss and stack are calculated in the same way. The sum of the segments required, including one for the stack, must be less than or equal to 8. Example: a program has 8128 bytes of base text, the largest overlay is 16248 bytes, and the data and bss total 19500. The text requires 1 8Kb segment, the overlays 2, and data and bss use 4, leaving one for the stack.

FILES

```
/lib/lib*.a      libraries
/usr/lib/lib*.a   more libraries
/usr/local/lib/lib*.a still more libraries
a.out            output file
```

SEE ALSO

```
adb(1), ar(1), as(1), cc(1), checkobj(1), f77(1), file(1),
ranlib(1), size(1), a.out(5)
```

BUGS

The text overlay scheme presented is unique to the PDP-11 and 2BSD. The `-i`, `-P`, `-Z`, `-Y` options are specific to 2BSD. The `-q` and `-v` options are new with 2.11BSD.

NAME

learn - computer aided instruction about UNIX

SYNOPSIS

learn [-directory] [subject [lesson]]

DESCRIPTION

Learn gives Computer Aided Instruction courses and practice in the use of UNIX, the C Shell, and the Berkeley text editors. To get started simply type learn. If you had used learn before and left your last session without completing a subject, the program will use information in \$HOME/.learnrc to start you up in the same place you left off. Your first time through, learn will ask questions to find out what you want to do. Some questions may be bypassed by naming a subject, and more yet by naming a lesson. You may enter the lesson as a number that learn gave you in a previous session. If you do not know the lesson number, you may enter the lesson as a word, and learn will look for the first lesson containing it. If the lesson is '-', learn prompts for each lesson; this is useful for debugging.

The subject's presently handled are

- files
- editor
- vi
- morefiles
- macros
- eqn
- C

There are a few special commands. The command `bye' terminates a learn session and `where' tells you of your progress, with `where m' telling you more. The command `again' re-displays the text of the lesson and `again lesson' lets you review lesson. There is no way for learn to tell you the answers it expects in English, however, the command `hint' prints the last part of the lesson script used to evaluate a response, while `hint m' prints the whole lesson script. This is useful for debugging lessons and might possibly give you an idea about what it expects.

The -directory option allows one to exercise a script in a nonstandard place.

FILES

- /usr/share/learn subtree for all dependent directories and files
- /usr/tmp/pl* playpen directories
- \$HOME/.learnrc startup information

SEE ALSO

cs(1), ex(1)

B. W. Kernighan and M. E. Lesk, LEARN - Computer-Aided Instruction on UNIX

BUGS

The main strength of learn, that it asks the student to use the real UNIX, also makes possible baffling mistakes. It is helpful, especially for nonprogrammers, to have a UNIX initiate near at hand during the first sessions.

Occasionally lessons are incorrect, sometimes because the local version of a command operates in a non-standard way. Occasionally a lesson script does not recognize all the different correct responses, in which case the 'hint' command may be useful. Such lessons may be skipped with the 'skip' command, but it takes some sophistication to recognize the situation.

To find a lesson given as a word, learn does a simple fgrep(1) through the lessons. It is unclear whether this sort of subject indexing is better than none.

Spawning a new shell is required for each of many user and internal functions.

The 'vi' lessons are provided separately from the others. To use them see your system administrator.

NAME

leave - remind you when you have to leave

SYNOPSIS

leave [[+]hhmm]

DESCRIPTION

Leave waits until the specified time, then reminds you that you have to leave. You are reminded 5 minutes and 1 minute before the actual time, at the time, and every minute thereafter. When you log off, leave exits just before it would have printed the next message.

The time of day is in the form hhmm where hh is a time in hours (on a 12 or 24 hour clock). All times are converted to a 12 hour clock, and assumed to be in the next 12 hours.

If the time is preceeded by '+', the alarm will go off in hours and minutes from the current time.

If no argument is given, leave prompts with "When do you have to leave?". A reply of newline causes leave to exit, otherwise the reply is assumed to be a time. This form is suitable for inclusion in a .login or .profile.

Leave ignores interrupts, quits, and terminates. To get rid of it you should either log off or use ``kill -9'' giving its process id.

SEE ALSO

calendar(1)

NAME

lex - generator of lexical analysis programs

SYNOPSIS

lex [-tvfn] [file] ...

DESCRIPTION

Lex generates programs to be used in simple lexical analysis of text. The input files (standard input default) contain regular expressions to be searched for, and actions written in C to be executed when expressions are found.

A C source program, 'lex.yy.c' is generated, to be compiled thus:

```
cc lex.yy.c -ll
```

This program, when run, copies unrecognized portions of the input to the output, and executes the associated C action for each regular expression that is recognized.

The options have the following meanings.

- t Place the result on the standard output instead of in file "lex.yy.c".
- v Print a one-line summary of statistics of the generated analyzer.
- n Opposite of -v; -n is default.
- f "Faster" compilation: don't bother to pack the resulting tables; limited to small programs.

EXAMPLE

```
lex lexcommands
```

would draw lex instructions from the file lexcommands, and place the output in lex.yy.c

```
%%  
[A-Z] putchar(yytext[0]+'a'-'A');  
[ ]+$ ;  
[ ]+   putchar(' ');
```

is an example of a lex program that would be put into a lex command file. This program converts upper case to lower, removes blanks at the end of lines, and replaces multiple blanks by single blanks.

SEE ALSO

yacc(1), sed(1)

M. E. Lesk and E. Schmidt, LEX - Lexical Analyzer Generator

NAME

lint - a C program verifier

SYNOPSIS

lint [-abchnpuvx] file ...

DESCRIPTION

Lint attempts to detect features of the C program files which are likely to be bugs, or non-portable, or wasteful. It also checks the type usage of the program more strictly than the compilers. Among the things which are currently found are unreachable statements, loops not entered at the top, automatic variables declared and not used, and logical expressions whose value is constant. Moreover, the usage of functions is checked to find functions which return values in some places and not in others, functions called with varying numbers of arguments, and functions whose values are not used.

By default, it is assumed that all the files are to be loaded together; they are checked for mutual compatibility. Function definitions for certain libraries are available to lint; these libraries are referred to by a conventional name, such as `~-lm'`, in the style of `ld(1)`. Arguments ending in `.ln` are also treated as library files. To create lint libraries, use the `-C` option:

```
lint -Cfoo files . . .
```

where files are the C sources of library foo. The result is a file `llib-lfoo.ln` in the correct library format suitable for linting programs using foo.

Any number of the options in the following list may be used. The `-D`, `-U`, and `-I` options of `cc(1)` are also recognized as separate arguments.

- p Attempt to check portability to the IBM and GCOS dialects of C.
- h Apply a number of heuristic tests to attempt to intuit bugs, improve style, and reduce waste.
- b Report break statements that cannot be reached. (This is not the default because, unfortunately, most lex and many yacc outputs produce dozens of such comments.)
- v Suppress complaints about unused arguments in functions.
- x Report variables referred to by extern declarations, but never used.

- a Report assignments of long values to int variables.
- c Complain about casts which have questionable portability.
- u Do not complain about functions and variables used and not defined, or defined and not used (this is suitable for running lint on a subset of files out of a larger program).
- n Do not check compatibility against the standard library.
- z Do not complain about structures that are never defined (e.g. using a structure pointer without knowing its contents.).

Exit(2) and other functions which do not return are not understood; this causes various lies.

Certain conventional comments in the C source will change the behavior of lint:

```
/*NOTREACHED*/
    at appropriate points stops comments about unreachable
    code.

/*VARARGSn*/
    suppresses the usual checking for variable numbers of
    arguments in the following function declaration. The
    data types of the first n arguments are checked; a
    missing n is taken to be 0.

/*NOSTRICT*/
    shuts off strict type checking in the next expression.

/*ARGSUSED*/
    turns on the -v option for the next function.

/*LINTLIBRARY*/
    at the beginning of a file shuts off complaints about
    unused functions in this file.
```

AUTHOR

S.C. Johnson. Lint library construction implemented by Edward Wang.

FILES

```
/usr/libexec/lint/lint[12]    programs
/usr/share/lint/l1ib-lc.ln    declarations for standard
functions
/usr/share/lint/l1ib-lc      human readable version of
```

above
/usr/share/lint/llib-port.ln declarations for portable
functions
/usr/share/lint/llib-port human readable . . .
llib-1*.ln library created with -C

SEE ALSO

cc(1)
S. C. Johnson, Lint, a C Program Checker

BUGS

There are some things you just can't get lint to shut up
about.

/*NOSTRICT*/ is not implemented in the current version
(alas).

NAME

lisp - lisp interpreter

SYNOPSIS

lisp

DESCRIPTION

Lisp is a provisional lisp interpreter. It only runs in eval mode. Built in functions are named in lower case, and case is distinguished. It is being transmuted from a subset of lisp as provided by the Harvard UNIX lisp in use at UCB, to a subset of MIT's MACLISP.

The following functions are provided as machine code:

Lambda functions:

atom	dptr	load	putd	rplacd
bcdp	drain	null	putprop	set
car	eq	numberp	ratom	terpr
cdr	equal	outfile	read	
close	eval	patom	readc	
concat	get	pntlen	rethrk	
cons	getd	portp	return	
cont	infile	print	rplaca	

Nlambda functions (possibly simulating ones which are normally lambdas):

add1	difference	onep	quotient	zerop
and	exit	or	reset	
break	go	plus	setq	
cond	minus	product	sub1	
cond	mod	prog	sum	
def	not	quote	times	

The following functions are provided as lisp code (and at the moment must be read in by saying (load 'auxfns):

add	copy	length	numbp
append	defevq	linelength	pp_etc
apply*	defprop	member	reverse
charcnt	defprop	memcar	terpri
chrct	diff	memcdr	
conc	last	nconc	

All of the above functions are documented in the ``Harvard Lisp Manual.''

The following functions are provided as in MIT's MACLISP.

alphalessp	do	mapc	setsyntax
------------	----	------	-----------

apply	explodec	mapcar	throw
ascii	exploden	prog2	tyi
catch	funcall	progn	tyipeek
defun	implode	progv	tyo

``Hairy control structure'' is provided by the Nlambda (process command inport outport) where command is an atom whose print name is some command that you would wish typed at the terminal, e.g. ``neqn | nroff -ms''; where inport and outport are atoms which will be bound to port descriptors for use in communication with the subprocess. Inport is a port to a pipe which will be read by the subprocess as its standard input. If Inport is nil (or not present), the subprocess inherits the standard input, and lisp waits for the subprocess to die. If Inport is the atom t lisp continues without waiting.

AUTHORS

Originally written by Jeff Levinsky, Mike Curry, and John Breedlove. Keith Sklower made it work and is maintaining the current version. The garbage collector was implemented by Bill Rowan.

SEE ALSO

Harvard UNIX Lisp Manual
MACLISP Manual
UCB Franz Lisp Manual

BUGS

The status bits for setsyntax are not the same as for MACLISP.

Closing down a pipe doesn't always seem to work correctly.

Arrays are not implemented in version 1.

NAME

ln - make links

SYNOPSIS

```
ln [ -s ] sourcename [ targetname ]  
ln [ -s ] sourcename1 sourcename2 [ sourcename3 ... ] tar-  
getdirectory
```

DESCRIPTION

A link is a directory entry referring to a file; the same file (together with its size, all its protection information, etc.) may have several links to it. There are two kinds of links: hard links and symbolic links.

By default ln makes hard links. A hard link to a file is indistinguishable from the original directory entry; any changes to a file are effective independent of the name used to reference the file. Hard links may not span file systems and may not refer to directories.

The -s option causes ln to create symbolic links. A symbolic link contains the name of the file to which it is linked. The referenced file is used when an open(2) operation is performed on the link. A stat(2) on a symbolic link will return the linked-to file; an lstat(2) must be done to obtain information about the link. The readlink(2) call may be used to read the contents of a symbolic link. Symbolic links may span file systems and may refer to directories.

Given one or two arguments, ln creates a link to an existing file sourcename. If targetname is given, the link has that name; targetname may also be a directory in which to place the link; otherwise it is placed in the current directory. If only the directory is specified, the link will be made to the last component of sourcename.

Given more than two arguments, ln makes links in target-directory to all the named source files. The links made will have the same name as the files being linked to.

SEE ALSO

rm(1), cp(1), mv(1), link(2), readlink(2), stat(2), symlink(2)

NAME

lock - reserve a terminal

SYNOPSIS

lock [-p] [-t timeout]

DESCRIPTION

Lock requests a password from the user, reads it again for verification and then will normally not relinquish the terminal until the password is repeated. There are two other conditions under which it will terminate: it will timeout after some interval of time and it may be killed by someone with the appropriate permission. The default time limit is 15 minutes, which may be changed with the -t option where timeout is the time limit in minutes. The -p option causes lock to use the user's current password instead of requesting another one.

NAME

logger - make entries in the system log

SYNOPSIS

```
logger [ -t tag ] [ -p pri ] [ -i ] [ -f file ] [ message  
... ]
```

ARGUMENTS

-t tag Mark every line in the log with the specified tag.

-p pri Enter the message with the specified priority. The priority may be specified numerically or as a ``facility.level'' pair. For example, ``-p local3.info'' logs the message(s) as informational level in the local3 facility. The default is ``user.notice.''

-i Log the process id of the logger process with each line.

-f file Log the specified file.

message The message to log; if not specified, the -f file or standard input is logged.

DESCRIPTION

Logger provides a program interface to the syslog(3) system log module.

A message can be given on the command line, which is logged immediately, or a file is read and each line is logged.

EXAMPLES

```
logger System rebooted
```

```
logger -p local0.notice -t HOSTIDM -f /dev/idmc
```

SEE ALSO

syslog(3), syslogd(8)

NAME

login - sign on

SYNOPSIS

login [-p] [username]

DESCRIPTION

The login command is used when a user initially signs on, or it may be used at any time to change from one user to another. The latter case is the one summarized above and described here. See "How to Get Started" for how to dial up initially.

If login is invoked without an argument, it asks for a user name, and, if appropriate, a password. Echoing is turned off (if possible) during the typing of the password, so it will not appear on the written record of the session.

After a successful login, accounting files are updated and the user is informed of the existence of mail. The message of the day is printed, as is the time of his last login. Both are suppressed if he has a ".hushlogin" file in his home directory; this is mostly used to make life easier for non-human users, such as uucp.

Login initializes the user and group IDs and the working directory, then executes a command interpreter (usually csh(1)) according to specifications found in a password file. Argument 0 of the command interpreter is the name of the command interpreter with a leading dash ("-").

Login also modifies the environment environ(7) with information specifying home directory, command interpreter, terminal type (if available) and user name. The '-p' argument causes the remainder of the environment to be preserved, otherwise any previous environment is discarded.

If the file /etc/nologin exists, login prints its contents on the user's terminal and exits. This is used by shutdown(8) to stop users logging in when the system is about to go down.

Login is recognized by sh(1) and csh(1) and executed directly (without forking).

FILES

/var/run/utmp	accounting
/usr/adm/wtmp	accounting
/usr/spool/mail/*	mail
/etc/motd	message-of-the-day
/etc/passwd	password file
/etc/nologinstops	logins

.hushlogin makes login quieter

SEE ALSO

init(8), getty(8), mail(1), passwd(1), passwd(5),
environ(7), shutdown(8), rlogin(1c)

DIAGNOSTICS

"Login incorrect," if the name or the password is bad.
"No Shell", "cannot open password file", "no directory":
consult a programming counselor.

BUGS

An undocumented option, -r is used by the remote login
server, rlogind(8C) to force login to enter into an initial
connection protocol. -h is used by telnetd(8C) and other
servers to list the host from which the connection was
received.

NAME

look - find lines in a sorted list

SYNOPSIS

look [-df] string [file]

DESCRIPTION

Look consults a sorted file and prints all lines that begin with string. It uses binary search.

The options d and f affect comparisons as in sort(1):

d `Dictionary' order: only letters, digits, tabs and blanks participate in comparisons.

f Fold. Upper case letters compare equal to lower case.

If no file is specified, /usr/dict/words is assumed with collating sequence -df.

FILES

/usr/dict/words

SEE ALSO

sort(1), grep(1)

NAME

indxbib, lookbib - build inverted index for a bibliography,
find references in a bibliography

SYNOPSIS

indxbib database ...
lookbib [-n] database

DESCRIPTION

Indxbib makes an inverted index to the named databases (or files) for use by lookbib(1) and refer(1). These files contain bibliographic references (or other kinds of information) separated by blank lines.

A bibliographic reference is a set of lines, constituting fields of bibliographic information. Each field starts on a line beginning with a ``%'', followed by a key-letter, then a blank, and finally the contents of the field, which may continue until the next line starting with ``%''.

Indxbib is a shell script that calls /usr/libexec/refer/mkey and /usr/libexec/refer/inv. The first program, mkey, truncates words to 6 characters, and maps upper case to lower case. It also discards words shorter than 3 characters, words among the 100 most common English words, and numbers (dates) < 1900 or > 2000. These parameters can be changed; see page 4 of the Refer document by Mike Lesk. The second program, inv, creates an entry file (.ia), a posting file (.ib), and a tag file (.ic), all in the working directory.

Lookbib uses an inverted index made by indxbib to find sets of bibliographic references. It reads keywords typed after the ``>'' prompt on the terminal, and retrieves records containing all these keywords. If nothing matches, nothing is returned except another ``>'' prompt.

Lookbib will ask if you need instructions, and will print some brief information if you reply ``y''. The ``-n'' flag turns off the prompt for instructions.

It is possible to search multiple databases, as long as they have a common index made by indxbib. In that case, only the first argument given to indxbib is specified to lookbib.

If lookbib does not find the index files (the .i[abc] files), it looks for a reference file with the same name as the argument, without the suffixes. It creates a file with a '.ig' suffix, suitable for use with fgrep. It then uses this fgrep file to find references. This method is simpler to use, but the .ig file is slower to use than the .i[abc] files, and does not allow the use of multiple reference files.

FILES

x.ia, x.ib, x.ic, where x is the first argument, or if these are not present, then x.ig, x

SEE ALSO

refer(1), addbib(1), sortbib(1), roffbib(1), lookbib(1)

BUGS

Probably all dates should be indexed, since many disciplines refer to literature written in the 1800s or earlier.

NAME

lorder - find ordering relation for an object library

SYNOPSIS

lorder file ...

DESCRIPTION

The input is one or more object or library archive (see `ar(1)`) files. The standard output is a list of pairs of object file names, meaning that the first file of the pair refers to external identifiers defined in the second. The output may be processed by `tsort(1)` to find an ordering of a library suitable for one-pass access by `ld(1)`.

This brash one-liner intends to build a new library from existing `.o` files.

```
ar cr library `lorder *.o | tsort`
```

The need for `lorder` may be vitiated by use of `ranlib(1)`, which converts an ordered archive into a randomly accessed library.

FILES

`*symref`, `*symdef`
`nm(1)`, `sed(1)`, `sort(1)`, `join(1)`

SEE ALSO

`tsort(1)`, `ld(1)`, `ar(1)`, `ranlib(1)`

BUGS

The names of object files, in and out of libraries, must end with `.o`; nonsense results otherwise.

NAME

lpq - spool queue examination program

SYNOPSIS

```
lpq [ +[ n ] ] [ -l ] [ -Pprinter ] [ job # ... ] [ user ... ]
```

DESCRIPTION

lpq examines the spooling area used by lpd(8) for printing files on the line printer, and reports the status of the specified jobs or all jobs associated with a user. lpq invoked without any arguments reports on any jobs currently in the queue. A -P flag may be used to specify a particular printer, otherwise the default line printer is used (or the value of the PRINTER variable in the environment). If a + argument is supplied, lpq displays the spool queue until it empties. Supplying a number immediately after the + sign indicates that lpq should sleep n seconds in between scans of the queue. All other arguments supplied are interpreted as user names or job numbers to filter out only those jobs of interest.

For each job submitted (i.e. invocation of lpr(1)) lpq reports the user's name, current rank in the queue, the names of files comprising the job, the job identifier (a number which may be supplied to lprm(1) for removing a specific job), and the total size in bytes. The -l option causes information about each of the files comprising the job to be printed. Normally, only as much information as will fit on one line is displayed. Job ordering is dependent on the algorithm used to scan the spooling directory and is supposed to be FIFO (First in First Out). File names comprising a job may be unavailable (when lpr(1) is used as a sink in a pipeline) in which case the file is indicated as ``(standard input)".

If lpq warns that there is no daemon present (i.e. due to some malfunction), the lpc(8) command can be used to restart the printer daemon.

FILES

/etc/termcap	for manipulating the screen for repeated display
/etc/printcap	to determine printer characteristics
/usr/spool/*	the spooling directory, as determined from printcap
/usr/spool/*/cf*	control files specifying jobs
/usr/spool/*/lock	the lock file to obtain the currently active job

SEE ALSO

lpr(1), lprm(1), lpc(8), lpd(8)

BUGS

Due to the dynamic nature of the information in the spooling

directory lpq may report unreliably. Output formatting is sensitive to the line length of the terminal; this can results in widely spaced columns.

DIAGNOSTICS

Unable to open various files. The lock file being malformed. Garbage files when there is no daemon active, but files in the spooling directory.

NAME

lpr - off line print

SYNOPSIS

```
lpr [ -Pprinter ] [ -#num ] [ -C class ] [ -J job ] [ -T
title ] [ -i [ numcols ] ] [ -1234 font ] [ -wnum ] [
-pltndgvcfrmhs ] [ name ... ]
```

DESCRIPTION

Lpr uses a spooling daemon to print the named files when facilities become available. If no names appear, the standard input is assumed. The -P option may be used to force output to a specific printer. Normally, the default printer is used (site dependent), or the value of the environment variable PRINTER is used.

The following single letter options are used to notify the line printer spooler that the files are not standard text files. The spooling daemon will use the appropriate filters to print the data accordingly.

- p Use pr(1) to format the files (equivalent to print).
- l Use a filter which allows control characters to be printed and suppresses page breaks.
- t The files are assumed to contain data from troff(1) (cat phototypesetter commands).
- n The files are assumed to contain data from ditroff (device independent troff).
- d The files are assumed to contain data from tex(1) (DVI format from Stanford).
- g The files are assumed to contain standard plot data as produced by the plot(3X) routines (see also plot(1G) for the filters used by the printer spooler).
- v The files are assumed to contain a raster image for devices like the Benson Varian.
- c The files are assumed to contain data produced by cifplot(1).
- f Use a filter which interprets the first character of each line as a standard FORTRAN carriage control character.

The remaining single letter options have the following meaning.

- r Remove the file upon completion of spooling or upon completion of printing (with the -s option).
- m Send mail upon completion.
- h Suppress the printing of the burst page.
- s Use symbolic links. Usually files are copied to the spool directory.

The -C option takes the following argument as a job classification for use on the burst page. For example,

```
lpr -C EECS foo.c
```

causes the system name (the name returned by `hostname(1)`) to be replaced on the burst page by EECS, and the file `foo.c` to be printed.

The -J option takes the following argument as the job name to print on the burst page. Normally, the first file's name is used.

The -T option uses the next argument as the title used by `pr(1)` instead of the file name.

To get multiple copies of output, use the -#num option, where num is the number of copies desired of each file named. For example,

```
lpr -#3 foo.c bar.c more.c
```

would result in 3 copies of the file `foo.c`, followed by 3 copies of the file `bar.c`, etc. On the other hand,

```
cat foo.c bar.c more.c | lpr -#3
```

will give three copies of the concatenation of the files.

The -i option causes the output to be indented. If the next argument is numeric, it is used as the number of blanks to be printed before each line; otherwise, 8 characters are printed.

The -w option takes the immediately following number to be the page width for `pr`.

The -s option will use `symlink(2)` to link data files rather than trying to copy them so large files can be printed. This means the files should not be modified or removed until they have been printed.

The option `-1234` Specifies a font to be mounted on font position `i`. The daemon will construct a `.railmag` file referencing `/usr/share/vfont/name.size`.

FILES

<code>/etc/passwd</code>	personal identification
<code>/etc/printcap</code>	printer capabilities data base
<code>/usr/sbin/lpd</code>	line printer daemon
<code>/usr/spool/*</code>	directories used for spooling
<code>/usr/spool/*/cf*</code>	daemon control files
<code>/usr/spool/*/df*</code>	data files specified in "cf" files
<code>/usr/spool/*/tf*</code>	temporary copies of "cf" files

SEE ALSO

`lpq(1)`, `lprm(1)`, `pr(1)`, `symlink(2)`, `printcap(5)`, `lpc(8)`,
`lpd(8)`

DIAGNOSTICS

If you try to spool too large a file, it will be truncated. Lpr will object to printing binary files. If a user other than root prints a file and spooling is disabled, lpr will print a message saying so and will not put jobs in the queue. If a connection to lpd on the local machine cannot be made, lpr will say that the daemon cannot be started. Diagnostics may be printed in the daemon's log file regarding missing spool files by lpd.

BUGS

Fonts for troff and tex reside on the host with the printer. It is currently not possible to use local font libraries.

NAME

lprm - remove jobs from the line printer spooling queue

SYNOPSIS

lprm [-Pprinter] [-] [job # ...] [user ...]

DESCRIPTION

Lprm will remove a job, or jobs, from a printer's spool queue. Since the spooling directory is protected from users, using lprm is normally the only method by which a user may remove a job.

Lprm without any arguments will delete the currently active job if it is owned by the user who invoked lprm.

If the - flag is specified, lprm will remove all jobs which a user owns. If the super-user employs this flag, the spool queue will be emptied entirely. The owner is determined by the user's login name and host name on the machine where the lpr command was invoked.

Specifying a user's name, or list of user names, will cause lprm to attempt to remove any jobs queued belonging to that user (or users). This form of invoking lprm is useful only to the super-user.

A user may dequeue an individual job by specifying its job number. This number may be obtained from the lpq(1) program, e.g.

```
% lpq -l
```

```
1st: ken                      [job #013ucbarpa]
      (standard input)        100 bytes
```

```
% lprm 13
```

Lprm will announce the names of any files it removes and is silent if there are no jobs in the queue which match the request list.

Lprm will kill off an active daemon, if necessary, before removing any spooling files. If a daemon is killed, a new one is automatically restarted upon completion of file removals.

The -P option may be used to specify the queue associated with a specific printer (otherwise the default printer, or the value of the PRINTER variable in the environment is used).

FILES

/etc/printcap printer characteristics file

```
/usr/spool/* spooling directories
/usr/spool/*/lock lock file used to obtain the pid of the current
                  daemon and the job number of the currently active job
```

SEE ALSO

lpr(1), lpq(1), lpd(8)

DIAGNOSTICS

``Permission denied" if the user tries to remove files other than his own.

BUGS

Since there are race conditions possible in the update of the lock file, the currently active job may be incorrectly identified.

NAME

lpctest - generate lineprinter ripple pattern

SYNOPSIS

lpctest [length [count]]

DESCRIPTION

Lptest writes the traditional "ripple test" pattern on standard output. In 96 lines, this pattern will print all 96 printable ASCII characters in each position. While originally created to test printers, it is quite useful for testing terminals, driving terminal ports for debugging purposes, or any other task where a quick supply of random data is needed.

The length argument specifies the output line length if the the default length of 79 is inappropriate.

The count argument specifies the number of output lines to be generated if the default count of 200 is inappropriate. Note that if count is to be specified, length must be also be specified.

SEE ALSO

BUGS

NAME

ls - list contents of directory

SYNOPSIS

ls [-acdfgiloqrstulACLFR] name ...

DESCRIPTION

For each directory argument, ls lists the contents of the directory; for each file argument, ls repeats its name and any other information requested. By default, the output is sorted alphabetically. When no argument is given, the current directory is listed. When several arguments are given, the arguments are first sorted appropriately, but file arguments are processed before directories and their contents.

There are a large number of options:

- l List in long format, giving mode, number of links, owner, size in bytes, and time of last modification for each file. (See below.) If the file is a special file the size field will instead contain the major and minor device numbers. If the file is a symbolic link the pathname of the linked-to file is printed preceded by ```->''`.
- o Include the file flags in a long (-l) output.
- g Include the group ownership of the file in a long output.
- t Sort by time modified (latest first) instead of by name.
- a List all entries; in the absence of this option, entries whose names begin with a period (.) are not listed.
- s Give size in kilobytes of each file.
- d If argument is a directory, list only its name; often used with -l to get the status of a directory.
- L If argument is a symbolic link, list the file or directory the link references rather than the link itself.
- r Reverse the order of sort to get reverse alphabetic or oldest first as appropriate.
- u Use time of last access instead of last modification for sorting (with the -t option) and/or printing (with the -l option).

- c Use time of file creation for sorting or printing.
- i For each file, print the i-number in the first column of the report.
- f Output is not sorted.
- F cause directories to be marked with a trailing `/', sockets with a trailing `=', symbolic links with a trailing `@', and executable files with a trailing `*'.
-R recursively list subdirectories encountered.
- l force one entry per line output format; this is the default when output is not to a terminal.
- C force multi-column output; this is the default when output is to a terminal.
- q force printing of non-graphic characters in file names as the character `?'; this is the default when output is to a terminal.

The mode printed under the -l option contains 11 characters which are interpreted as follows: the first character is

d	if the entry is a directory;
b	if the entry is a block-type special file;
c	if the entry is a character-type special file;
l	if the entry is a symbolic link;
s	if the entry is a socket, or
-	if the entry is a plain file.

The next 9 characters are interpreted as three sets of three bits each. The first set refers to owner permissions; the next refers to permissions to others in the same user-group; and the last to all others. Within each set the three characters indicate permission respectively to read, to write, or to execute the file as a program. For a directory, 'execute' permission is interpreted to mean permission to search the directory. The permissions are indicated as follows:

r	if the file is readable;
w	if the file is writable;
x	if the file is executable;
-	if the indicated permission is not granted.

The group-execute permission character is given as s if the file has the set-group-id bit set; likewise the user-execute permission character is given as s if the file has the set-user-id bit set.

The last character of the mode (normally `x' or `-') is t if the 1000 bit of the mode is on. See `chmod(1)` for the meaning of this mode.

When the sizes of the files in a directory are listed, a total count of blocks, including indirect blocks is printed.

FILES

/etc/passwd to get user id's for ``ls -l'`.
/etc/group to get group id's for ``ls -g'`.

BUGS

Newline and tab are considered printing characters in file names.

The output device is assumed to be 80 columns wide.

The option setting based on whether the output is a teletype is undesirable as ``ls -s'` is much different than ``ls -s | lpr'`. On the other hand, not doing this setting would make old shell scripts which used `ls` almost certain losers.

NAME

lxref - lisp cross reference program

SYNOPSIS

lxref [-N] xref-file ... [-a source-file ...]

DESCRIPTION

Lxref reads cross reference file(s) written by the lisp compiler liszt and prints a cross reference listing on the standard output. Liszt will create a cross reference file during compilation when it is given the -x switch. Cross reference files usually end in '.x' and consequently lxref will append a '.x' to the file names given if necessary. The first option to lxref is a decimal integer, N, which sets the ignorelevel. If a function is called more than ignorelevel times, the cross reference listing will just print the number of calls instead of listing each one of them. The default for ignorelevel is 50.

The -a option causes lxref to put limited cross reference information in the sources named. lxref will scan the source and when it comes across a definition of a function (that is a line beginning with '(def' it will precede that line with a list of the functions which call this function, written as a comment preceded by ';.. '. All existing lines beginning with ';.. ' will be removed from the file. If the source file contains a line beginning ';.-' then this will disable this annotation process from this point on until a ';.+' is seen (however, lines beginning with ';.. ' will continue to be deleted). After the annotation is done, the original file 'foo.l' is renamed to " '#.foo.l'" and the new file with annotation is named 'foo.l'

AUTHOR

John Foderaro

SEE ALSO

lisp(1), liszt(1)

BUGS

NAME

m4 - macro processor

SYNOPSIS

m4 [-Dname[=value]] [-Uname] [files ...]

DESCRIPTION

m4 is a macro processor intended as a front end for any language (e.g., C, ratfor, fortran, lex, and yacc). m4 reads from the standard input and writes the processed text to the standard output. Each of the optional argument files is processed in order.

Macro calls have the form

name(arg1,arg2, . . . , argn)

The '(' must immediately follow the name of the macro. If a defined macro name is not followed by a '(', it is deemed to have no arguments. Leading unquoted blanks, tabs, and newlines are ignored while collecting arguments. Potential macro names consist of alphabetic letters, digits, and underscore '_', where the first character is not a digit.

Left and right single quotes (``') are used to quote strings. The value of a quoted string is the string stripped of the quotes.

When a macro name is recognized, its arguments are collected by searching for a matching right parenthesis. Macro evaluation proceeds normally during the collection of the arguments, and any commas or right parentheses which happen to turn up within the value of a nested call are as effective as those in the original input text. After argument collection, the value of the macro is pushed back onto the input stream and rescanned.

-Dname[=value]

Define the symbol name to have some value (or NULL).

-Uname Undefine the symbol name.

m4 makes available the following built-in macros. They may be redefined, but once this is done the original meaning is lost. Their values are null unless otherwise stated.

define The second argument is installed as the value of the macro whose name is the first argument. Each occurrence of \$n in the replacement text, where n is a digit, is replaced by the n-th argument. Argument 0 is the name of the macro; missing

arguments are replaced by the null string.

undefine removes the definition of the macro named in its argument.

ifdef If the first argument is defined, the value is the second argument, otherwise the third. If there is no third argument, the value is null.

changeocom Change the start and end comment sequences. The default is the pound sign '#' and the newline character. With no arguments comments are turned off. The maximum length for a comment marker is five characters.

changequote
Change quote characters to the first and second arguments. Changequote without arguments restores the original values (i.e., `').

decr Decrements the argument by 1. The argument must be a valid numeric string.

divert m4 maintains 10 output streams, numbered 0-9. The final output is the concatenation of the streams in numerical order; initially stream 0 is the current stream. The divert macro changes the current output stream to its (digit-string) argument. Output diverted to a stream other than 0 through 9 is discarded.

undivert causes immediate output of text from diversions named as arguments, or all diversions if no argument. Text may be undiverted into another diversion. Undiverting discards the diverted text.

defn Returns the quoted definition for each argument. This can be used to rename macro definitions (even for builtin macros).

divnum returns the value of the current output stream.

dnl reads and discards characters up to and including the next newline.

expr This is an alias for eval.

ifelse has three or more arguments. If the first argument is the same string as the second, then the value is the third argument. If not, and if there are more than four arguments, the process is repeated with arguments 4, 5, 6 and 7. Otherwise,

the value is either the fourth string, or, if it is not present, null.

- incr** returns the value of its argument incremented by 1. The value of the argument is calculated by interpreting an initial digit-string as a decimal number.
- eval** evaluates its argument as an arithmetic expression, using 32-bit arithmetic. Operators include +, -, *, /, %, ^ (exponentiation); relationals; parentheses.
- len** returns the number of characters in its argument.
- m4exit** Immediately exits with the return value specified by the first argument, 0 if none.
- m4wrap** Allows you to define what happens at the final EOF, usually for cleanup purposes. (e.g., `m4wrap("cleanup(tempfile)")` causes the macro cleanup to be invoked after all processing is done.)
- index** returns the position in its first argument where the second argument begins (zero origin), or -1 if the second argument does not occur.
- substr** returns a substring of its first argument. The second argument is a zero origin number selecting the first character; the third argument indicates the length of the substring. A missing third argument is taken to be large enough to extend to the end of the first string.
- translit** transliterates the characters in its first argument from the set given by the second argument to the set given by the third. No abbreviations are permitted.
- include** returns the contents of the file named in the argument.
- sinclude** is identical to include, except that it says nothing if the file is inaccessible.
- syscmd** executes the UNIX command given in the first argument. No value is returned.
- maketemp** fills in a string of XXXXX in its argument with the current process id.

paste Includes the contents of the file specified by the first argument without any macro processing. Aborts with an error message if the file cannot be included.

popdef Restores the pushdef'd definition for each argument.

pushdef Takes the same arguments as define, but it saves the definition on a stack for later retrieval by popdef.

shift Returns all but the first argument, the remaining arguments are quoted and pushed back with commas in between. The quoting nullifies the effect of the extra scan that will subsequently be performed.

spaste Similar to paste, except it ignores any errors.

syscal Returns the return value from the last syscmd.

errprint prints its argument on the diagnostic output file.

dumpdef prints current names and definitions, for the named items, or for all if no arguments are given.

unix A pre-defined macro for testing the OS platform.

SEE ALSO

B. W. Kernighan and D. M. Ritchie, The m4 Macro Processor

HISTORY

An m4 command appeared in Version 6 AT&T UNIX.

AUTHOR

Ozan Yigit <oz@sis.yorku.ca>

NAME

mail - send and receive mail

SYNOPSIS

```
mail [ -v ] [ -i ] [ -n ] [ -s subject ] [ user ... ]  
mail [ -v ] [ -i ] [ -n ] -f [ name ]  
mail [ -v ] [ -i ] [ -n ] -u user
```

INTRODUCTION

Mail is a intelligent mail processing system, which has a command syntax reminiscent of ed with lines replaced by messages.

The -v flag puts mail into verbose mode; the details of delivery are displayed on the users terminal. The -i flag causes tty interrupt signals to be ignored. This is particularly useful when using mail on noisy phone lines. The -n flag inhibits the reading of /etc/Mail.rc.

Sending mail. To send a message to one or more people, mail can be invoked with arguments which are the names of people to whom the mail will be sent. You are then expected to type in your message, followed by an EOT (control-D) at the beginning of a line. A subject may be specified on the command line by using the -s flag. (Only the first argument after the -s flag is used as a subject; be careful to quote subjects containing spaces.) The section below, labeled Replying to or originating mail, describes some features of mail available to help you compose your letter.

Reading mail. In normal usage mail is given no arguments and checks your mail out of the post office, then prints out a one line header of each message there. The current message is initially the first message (numbered 1) and can be printed using the print command (which can be abbreviated p). You can move among the messages much as you move between lines in ed, with the commands '+' and '-' moving backwards and forwards, and simple numbers.

Disposing of mail. After examining a message you can delete (d) the message or reply (r) to it. Deletion causes the mail program to forget about the message. This is not irreversible; the message can be undeleted (u) by giving its number, or the mail session can be aborted by giving the exit (x) command. Deleted messages will, however, usually disappear never to be seen again.

Specifying messages. Commands such as print and delete can be given a list of message numbers as arguments to apply to a number of messages at once. Thus ``delete 1 2'' deletes messages 1 and 2, while ``delete 1-5'' deletes messages 1 through 5. The special name ``*' addresses all messages,

and ``\$' addresses the last message; thus the command `top` which prints the first few lines of a message could be used in `top *` to print the first few lines of all messages.

Replying to or originating mail. You can use the reply command to set up a response to a message, sending it back to the person who it was from. Text you then type in, up to an end-of-file, defines the contents of the message. While you are composing a message, mail treats lines beginning with the character `~` specially. For instance, typing `~m` (alone on a line) will place a copy of the current message into the response right shifting it by a tabstop. Other escapes will set up subject fields, add and delete recipients to the message and allow you to escape to an editor to revise the message or to a shell to run some commands. (These options are given in the summary below.)

Ending a mail processing session. You can end a mail session with the quit (`q`) command. Messages which have been examined go to your mbox file unless they have been deleted in which case they are discarded. Unexamined messages go back to the post office. The `-f` option causes mail to read in the contents of your mbox (or the specified file) for processing; when you quit, mail writes undeleted messages back to this file. The `-u` flag is a short way of doing `"mail -f /usr/spool/mail/user"`.

Personal and systemwide distribution lists. It is also possible to create a personal distribution lists so that, for instance, you can send mail to `cohorts` and have it go to a group of people. Such lists can be defined by placing a line like

```
alias cohorts bill ozalp jkf mark kridle@ucbcory
```

in the file `.mailrc` in your home directory. The current list of such aliases can be displayed with the `alias` (`a`) command in mail. System wide distribution lists can be created by editing `/etc/aliases`, see `aliases(5)` and `sendmail(8)`; these are kept in a different syntax. In mail you send, personal aliases will be expanded in mail sent to others so that they will be able to reply to the recipients. System wide aliases are not expanded when the mail is sent, but any reply returned to the machine will have the system wide alias expanded as all mail goes through `sendmail`.

Network mail (ARPA, UUCP, Berknet) See `mailaddr(7)` for a description of network addresses.

Mail has a number of options which can be set in the `.mailrc` file to alter its behavior; thus `set askcc` enables the `askcc` feature. (These options are summarized below.)

SUMMARY

(Adapted from the 'Mail Reference Manual')

Each command is typed on a line by itself, and may take arguments following the command word. The command need not be typed in its entirety - the first command which matches the typed prefix is used. For commands which take message lists as arguments, if no message list is given, then the next message forward which satisfies the command's requirements is used. If there are no messages forward of the current message, the search proceeds backwards, and if there are no good messages at all, mail types ``No applicable messages'' and aborts the command.

- Goes to the previous message and prints it out.
If given a numeric argument n, goes to the n-th
previous message and prints it.

? Prints a brief summary of commands.

! Executes the UNIX shell command which follows.

Print (P) Like print but also prints out ignored
header fields. See also print , ignore and
retain.

Reply (R) Reply to originator. Does not reply to other
recipients of the original message.

Type (T) Identical to the Print command.

alias (a) With no arguments, prints out all
currently-defined aliases. With one argument,
prints out that alias. With more than one argu-
ment, creates an new or changes an on old alias.

alternates (alt) The alternates command is useful if you
have accounts on several machines. It can be
used to inform mail that the listed addresses
are really you. When you reply to messages, mail
will not send a copy of the message to any of
the addresses listed on the alternates list. If
the alternates command is given with no argu-
ment, the current set of alternate names is
displayed.

chdir (c) Changes the user's working directory to that
specified, if given. If no directory is given,
then changes to the user's login directory.

copy (co) The copy command does the same thing that
save does, except that it does not mark the

messages it is used on for deletion when you quit.

- delete** (d) Takes a list of messages as argument and marks them all as deleted. Deleted messages will not be saved in mbox, nor will they be available for most other commands.
- dp** (also dt) Deletes the current message and prints the next message. If there is no next message, mail says ``at EOF.''
- edit** (e) Takes a list of messages and points the text editor at each one in turn. On return from the editor, the message is read back in.
- exit** (ex or x) Effects an immediate return to the Shell without modifying the user's system mailbox, his mbox file, or his edit file in -f.
- file** (fi) The same as folder.
- folders** List the names of the folders in your folder directory.
- folder** (fo) The folder command switches to a new mail file or folder. With no arguments, it tells you which file you are currently reading. If you give it an argument, it will write out changes (such as deletions) you have made in the current file and read in the new file. Some special conventions are recognized for the name. # means the previous file, % means your system mailbox, %user means user's system mailbox, & means your ~/mbox file, and +folder means a file in your folder directory.
- from** (f) Takes a list of messages and prints their message headers.
- headers** (h) Lists the current range of headers, which is an 18 message group. If a ``+' argument is given, then the next 18 message group is printed, and if a ``-' argument is given, the previous 18 message group is printed.
- help** A synonym for ?
- hold** (ho, also preserve) Takes a message list and marks each message therein to be saved in the user's system mailbox instead of in mbox. Does not override the delete command.

- ignore N.B.: Ignore has been superseded by retain.
Add the list of header fields named to the ignored list. Header fields in the ignore list are not printed on your terminal when you print a message. This command is very handy for suppression of certain machine-generated header fields. The Type and Print commands can be used to print a message in its entirety, including ignored fields. If ignore is executed with no arguments, it lists the current set of ignored fields.
- mail (m) Takes as argument login names and distribution group names and sends mail to those people.
- mbx Indicate that a list of messages be sent to mbox in your home directory when you quit. This is the default action for messages if you do not have the hold option set.
- next (n like + or CR) Goes to the next message in sequence and types it. With an argument list, types the next matching message.
- preserve (pre) A synonym for hold.
- print (p) Takes a message list and types out each message on the user's terminal.
- quit (q) Terminates the session, saving all undeleted, unsaved messages in the user's mbox file in his login directory, preserving all messages marked with hold or preserve or never referenced in his system mailbox, and removing all other messages from his system mailbox. If new mail has arrived during the session, the message ``You have new mail'' is given. If given while editing a mailbox file with the -f flag, then the edit file is rewritten. A return to the Shell is effected, unless the rewrite of edit file fails, in which case the user can escape with the exit command.
- reply (r) Takes a message list and sends mail to the sender and all recipients of the specified message. The default message must not be deleted.
- respond A synonym for reply.
- retain Add the list of header fields named to the retained list. Only the header fields in the retain list are shown on your terminal when you

print a message. All other header fields are suppressed. The Type and Print commands can be used to print a message in its entirety. If retain is executed with no arguments, it lists the current set of retained fields.

- save (s) Takes a message list and a filename and appends each message in turn to the end of the file. The filename in quotes, followed by the line count and character count is echoed on the user's terminal.
- set (se) With no arguments, prints all variable values. Otherwise, sets option. Arguments are of the form ``option=value'' (no space before or after =) or ``option.''
- shell (sh) Invokes an interactive version of the shell.
- size Takes a message list and prints out the size in characters of each message.
- source (so) The source command reads mail commands from a file.
- top Takes a message list and prints the top few lines of each. The number of lines printed is controlled by the variable toplines and defaults to five.
- type (t) A synonym for print.
- unalias Takes a list of names defined by alias commands and discards the remembered groups of users. The group names no longer have any significance.
- undelete (u) Takes a message list and marks each message as not being deleted.
- unread (U) Takes a message list and marks each message as not having been read.
- unset Takes a list of option names and discards their remembered values; the inverse of set.
- visual (v) Takes a message list and invokes the display editor on each message.
- write (w) Similar to save, except that only the message body (without the header) is saved. Extremely useful for such tasks as sending and

receiving source program text over the message system.

xit (x) A synonym for exit.

z Mail presents message headers in windowfuls as described under the headers command. You can move mail's attention forward to the next window with the z command. Also, you can move to the previous window by using z-.

Here is a summary of the tilde escapes, which are used when composing messages to perform special functions. Tilde escapes are only recognized at the beginning of lines. The name ``tilde escape'' is somewhat of a misnomer since the actual escape character can be set by the option escape.

~!command Execute the indicated shell command, then return to the message.

~b name ... Add the given names to the list of carbon copy recipients but do not make the names visible in the Cc: line ("blind" carbon copy).

~c name ... Add the given names to the list of carbon copy recipients.

~d Read the file ``dead.letter'' from your home directory into the message.

~e Invoke the text editor on the message collected so far. After the editing session is finished, you may continue appending text to the message.

~f messages Read the named messages into the message being sent. If no messages are specified, read in the current message.

~h Edit the message header fields by typing each one in turn and allowing the user to append text to the end or modify the field by using the current terminal erase and kill characters.

~m messages Read the named messages into the message being sent, shifted right one tab. If no messages are specified, read the current message.

~p Print out the message collected so far, prefaced by the message header fields.

~q Abort the message being sent, copying the message to ``dead.letter'' in your home directory

if save is set.

~r filename Read the named file into the message.

~s string Cause the named string to become the current subject field.

~t name ... Add the given names to the direct recipient list.

~v Invoke an alternate editor (defined by the VISUAL option) on the message collected so far. Usually, the alternate editor will be a screen editor. After you quit the editor, you may resume appending text to the end of your message.

~w filename Write the message onto the named file.

~|command Pipe the message through the command as a filter. If the command gives no output or terminates abnormally, retain the original text of the message. The command `fmt(1)` is often used as command to rejustify the message.

~~string Insert the string of text in the message prefaced by a single ~. If you have changed the escape character, then you should double that character in order to send it.

Options are controlled via the `set` and `unset` commands. Options may be either binary, in which case it is only significant to see whether they are set or not; or string, in which case the actual value is of interest. The binary options include the following:

`append` Causes messages saved in `mbox` to be appended to the end rather than prepended. (This is set in `/etc/Mail.rc` on 2.11BSD systems.)

`ask` Causes mail to prompt you for the subject of each message you send. If you respond with simply a newline, no subject field will be sent.

`askcc` Causes you to be prompted for additional carbon copy recipients at the end of each message. Responding with a newline indicates your satisfaction with the current list.

`autoprint` Causes the `delete` command to behave like `dp` - thus, after deleting a message, the next one

will be typed automatically.

- debug** Setting the binary option debug is the same as specifying -d on the command line and causes mail to output all sorts of information useful for debugging mail.
- dot** The binary option dot causes mail to interpret a period alone on a line as the terminator of a message you are sending.
- hold** This option is used to hold messages in the system mailbox by default.
- ignore** Causes interrupt signals from your terminal to be ignored and echoed as @'s.
- ignoreeof** An option related to dot is ignoreeof which makes mail refuse to accept a control-d as the end of a message. Ignoreeof also applies to mail command mode.
- metoo** Usually, when a group is expanded that contains the sender, the sender is removed from the expansion. Setting this option causes the sender to be included in the group.
- nosave** Normally, when you abort a message with two RUBOUT, mail copies the partial letter to the file ``dead.letter'' in your home directory. Setting the binary option nosave prevents this.
- Replyall** Reverses the sense of reply and Reply commands.
- quiet** Suppresses the printing of the version when first invoked.
- verbose** Setting the option verbose is the same as using the -v flag on the command line. When mail runs in verbose mode, the actual delivery of messages is displayed on the users terminal.

The following options have string values:

- EDITOR** Pathname of the text editor to use in the edit command and ~e escape. If not defined, then a default editor is used.
- PAGER** Pathname of the program to use in the more

command or when crt variable is set. A default paginator is used if this option is not defined.

SHELL	Pathname of the shell to use in the ! command and the ~! escape. A default shell is used if this option is not defined.
VISUAL	Pathname of the text editor to use in the visual command and ~v escape.
crt	The valued option crt is used as a threshold to determine how long a message must be before PAGER is used to read it.
escape	If defined, the first character of this option gives the character to use in the place of ~ to denote escapes.
folder	The name of the directory to use for storing folders of messages. If this name begins with a `/', mail considers it to be an absolute pathname; otherwise, the folder directory is found relative to your home directory.
record	If defined, gives the pathname of the file used to record all outgoing mail. If not defined, then outgoing mail is not so saved.
toplines	If defined, gives the number of lines of a message to be printed out with the top command; normally, the first five lines are printed.

FILES

/usr/spool/mail/*	post office
~/mbox	your old mail
~/.mailrc	file giving initial mail commands
/tmp/R#	temporary for editor escape
/usr/share/misc/Mail.help*help files	
/etc/Mail.rc	system initialization file
Message*	temporary for editing messages

SEE ALSO

binmail(1), fmt(1), newaliases(1), aliases(5),
mailaddr(7), sendmail(8)
'The Mail Reference Manual'

BUGS

There are many flags that are not documented here. Most are not useful to the general user.
Usually, mail is just a link to Mail, which can be

confusing.

AUTHOR

Kurt Shoens

NAME

make - maintain program groups

SYNOPSIS

make [-f makefile] [option] ... file ...

DESCRIPTION

Make executes commands in makefile to update one or more target names. Name is typically a program. If no -f option is present, 'makefile' and 'Makefile' are tried in order. If makefile is '-', the standard input is taken. More than one -f option may appear.

Make updates a target if it depends on prerequisite files that have been modified since the target was last modified, or if the target does not exist.

Makefile contains a sequence of entries that specify dependencies. The first line of an entry is a blank-separated list of targets, then a colon, then a list of prerequisite files. Text following a semicolon, and all following lines that begin with a tab, are shell commands to be executed to update the target. If a name appears on the left of more than one 'colon' line, then it depends on all of the names on the right of the colon on those lines, but only one command sequence may be specified for it. If a name appears on a line with a double colon :: then the command sequence following that line is performed only if the name is out of date with respect to the names to the right of the double colon, and is not affected by other double colon lines on which that name may appear.

The special form of the name: a(b) means the file named b stored in the archive named a.

Sharp and newline surround comments.

The following makefile says that 'pgm' depends on two files 'a.o' and 'b.o', and that they in turn depend on '.c' files and a common file 'incl'.

```
pgm: a.o b.o
    cc a.o b.o -lm -o pgm
a.o: incl a.c
    cc -c a.c
b.o: incl b.c
    cc -c b.c
```

Makefile entries of the form

```
string1 = string2
```

are macro definitions. Subsequent appearances of `$(string1)` or `${string1}` are replaced by `string2`. If `string1` is a single character, the parentheses or braces are optional.

All environment variables are assumed to be macro definitions and processed as such. The environment variables are processed before any makefile macro definitions; thus, macro assignments in a makefile override environmental variables. The `-e` option causes the environment to override the macro assignments in a makefile. Finally, command line options of the form `string1=string2` override both environment and makefile macro definitions.

Make infers prerequisites for files for which makefile gives no construction commands. For example, a `.c` file may be inferred as prerequisite for a `.o` file and be compiled to produce the `.o` file. Thus the preceding example can be done more briefly:

```
pgm: a.o b.o
    cc a.o b.o -lm -o pgm
a.o b.o: incl
```

Prerequisites are inferred according to selected suffixes listed as the `'prerequisites'` for the special name `'.SUFFIXES'`; multiple lists accumulate; an empty list clears what came before. Order is significant; the first possible name for which both a file and a rule as described in the next paragraph exist is inferred. The default list is

```
.SUFFIXES: .out .o .c .e .r .f .y .l .s .p
```

The rule to create a file with suffix `s2` that depends on a similarly named file with suffix `s1` is specified as an entry for the `'target'` `s1s2`. In such an entry, the special macro `$(*)` stands for the target name with suffix deleted, `$(@)` for the full target name, `$(<)` for the complete list of prerequisites, and `$(?)` for the list of prerequisites that are out of date. For example, a rule for making optimized `.o` files from `.c` files is

```
.c.o: ; cc -c -O -o $(@) $(*)
```

Certain macros are used by the default inference rules to communicate optional arguments to any resulting compilations. In particular, `'CFLAGS'` is used for `cc(1)` options, `'FFLAGS'` for `f77(1)` options, `'PFLAGS'` for `pc(1)` options, and `'LFLAGS'` and `'YFLAGS'` for `lex` and `yacc(1)` options. In addition, the macro `'MFLAGS'` is filled in with the initial command line options supplied to make. This simplifies maintaining a hierarchy of makefiles as one may then invoke make on makefiles in subdirectories and pass along useful options

such as `-k`.

Another special macro is ``VPATH'`. The ``VPATH'` macro should be set to a list of directories separated by colons. When make searches for a file as a result of a dependency relation, it will first search the current directory and then each of the directories on the ``VPATH'` list. If the file is found, the actual path to the file will be used, rather than just the filename. If ``VPATH'` is not defined, then only the current directory is searched.

One use for ``VPATH'` is when one has several programs that compile from the same source. The source can be kept in one directory and each set of object files (along with a separate makefile) would be in a separate subdirectory. The ``VPATH'` macro would point to the source directory in this case.

Command lines are executed one at a time, each by its own shell. A line is printed when it is executed unless the special target ``.SILENT'` is in makefile, or the first character of the command is ``@'`.

Commands returning nonzero status (see `intro(1)`) cause make to terminate unless the special target ``.IGNORE'` is in makefile or the command begins with `<tab><hyphen>`.

Interrupt and quit cause the target to be deleted unless the target is a directory or depends on the special name ``.PRECIOUS'`.

Other options:

- `-e` Environmental variables override assignments within makefiles.
- `-i` Equivalent to the special entry ``.IGNORE:'`.
- `-k` When a command returns nonzero status, abandon work on the current entry, but continue on branches that do not depend on the current entry.
- `-n` Trace and print, but do not execute the commands needed to update the targets.
- `-t` Touch, i.e. update the modified date of targets, without executing any commands.
- `-r` Equivalent to an initial special entry ``.SUFFIXES:'` with no list.
- `-s` Equivalent to the special entry ``.SILENT:'`.

FILES

makefile, Makefile

SEE ALSO

sh(1), touch(1), f77(1), pc(1), getenv(3)
S. I. Feldman Make - A Program for Maintaining Computer Programs

BUGS

Some commands return nonzero status inappropriately. Use -i to overcome the difficulty.

Commands that are directly executed by the shell, notably cd(1), are ineffectual across newlines in make.

`VPATH' is intended to act like the System V `VPATH' support, but there is no guarantee that it functions identically.

NAME

man - print out the manual

SYNOPSIS

man [-] [-a] [-M path] [section] title ...

DESCRIPTION

Man is the program which provides on-line access to the UNIX manual. If a section specifier is given, man looks in that section of the manual for the given title(s). Section is either an Arabic section number ('`3' for example), or one of the words ``local'', ``new'', or ``old''. (The abbreviations ``l'', ``n'', and ``o'' are also allowed.) If section is omitted, man searches all sections of the manual, giving preference to commands over library subroutines, and displays the first manual page it finds, if any. If the -a option is supplied, man displays all applicable manual pages.

Normally man checks in standard locations (/usr/man and /usr/local/man) for manual information. This can be changed by supplying a search path (a la the Bourne shell) with the -M flag. The search path is a colon (``:') separated list of directories in which man expects to find the standard manual subdirectories. This search path can also be set with the environmental variable MANPATH.

Since some manual pages are intended for use only on certain machines, man only searches those directories applicable to the current machine. Man's determination of the current machine type can be overridden by setting the environmental variable MACHINE.

If the standard output is a teletype, and the - flag is not provided, man uses more(1), or the pager provided by the environmental variable PAGER, to display the manual page.

The FORTRAN version of section 3 of the manual may be specified by supplying man with the section ``3f''. Also, a specific section of the local manual may be specified by appending a number to the section, i.e. ``l5'' would indicate section 5 of the local manual.

FILES

/usr/man	standard manual area
/usr/man/cat?/*	directories containing standard manual pages
/usr/local/man/cat?/*	directories containing local manual pages
/usr/src/man	directories containing unformatted manual pages

SEE ALSO

apropos(1), more(1), whatis(1), whereis(1)

BUGS

The manual is supposed to be reproducible either on the phototypesetter or on a typewriter, however, on a typewriter, some information is necessarily lost.

NAME

msg - permit or deny messages

SYNOPSIS

msg [n] [y]

DESCRIPTION

Msg with argument n forbids messages via write and talk(1) by revoking non-user write permission on the user's terminal. Msg with argument y reinstates permission. All by itself, msg reports the current state without changing it.

FILES

/dev/tty*

SEE ALSO

write(1), talk(1)

DIAGNOSTICS

Exit status is 0 if messages are receivable, 1 if not, 2 on error.

NAME

mkdep - construct Makefile dependency list

SYNOPSIS

mkdep [-f makefile] [-p] flags file ...

DESCRIPTION

Mkdep takes a set of flags for the C compiler and a list of C source files as arguments and constructs a set of include file dependencies. It attaches this dependency list to the end of the file ``Makefile''. An example of its use in a makefile might be:

```
CFLAGS= -O -I../include -I.  
SRCS= file1.c file2.c
```

depend:

```
mkdep ${CFLAGS} ${SRCS}
```

where the macro SRCS is the list of C source files and the macro CFLAGS is the list of flags for the C compiler. The -f option provides mkdep with a name other than ``Makefile'' to be edited. If the -p option is provided, mkdep produces dependencies of the form ``program: program.c'' so that subsequent makes will produce program directly from its C module rather than using an intermediate .o module. This is useful in directories that contain many programs, each of whose source is contained in a single C module.

SEE ALSO

cc(1), make(1)

NAME

mkdir - make directories

SYNOPSIS

mkdir [-p directory_name ...]

DESCRIPTION

Mkdir creates the directories named as operands, in the order specified, using mode 0777 modified by the current `umask(2)`.

The options are as follows:

-p Create intermediate directories as required. If this option is not specified, the full path prefix of each operand must already exist.

The user must have write permission in the parent directory.

Mkdir exits 0 if successful, and >0 if an error occurred.

SEE ALSO

rmdir(1)

STANDARDS

Mkdir is POSIX 1003.2 compliant. This manual page is derived from the POSIX 1003.2 manual page.

NAME

mkstr - create an error message file by massaging C source

SYNOPSIS

mkstr [-] messagefile prefix file ...

DESCRIPTION

Mkstr is used to create files of error messages. Its use can make programs with large numbers of error diagnostics much smaller, and reduce system overhead in running the program as the error messages do not have to be constantly swapped in and out.

Mkstr will process each of the specified files, placing a massaged version of the input file in a file whose name consists of the specified prefix and the original name. A typical usage of mkstr would be

```
mkstr pistrings xx *.c
```

This command would cause all the error messages from the C source files in the current directory to be placed in the file pistrings and processed copies of the source for these files to be placed in files whose names are prefixed with xx.

To process the error messages in the source to the message file mkstr keys on the string ``error("`` in the input stream. Each time it occurs, the C string starting at the ``` is placed in the message file followed by a null character and a new-line character; the null character terminates the message so it can be easily used when retrieved, the new-line character makes it possible to sensibly cat the error message file to see its contents. The massaged copy of the input file then contains a lseek pointer into the file which can be used to retrieve the message, i.e.:

```
char efilename[] = "/usr/share/pascal/pi_strings";
int  efil = -1;

error(a1, a2, a3, a4)
{
    char buf[256];

    if (efil < 0) {
        efil = open(efilename, 0);
        if (efil < 0) {
oops:
            perror(efilename);
            exit(1);
        }
    }
}
```

```
    }  
    if (lseek(efil, (long) a1, 0) <= 0 ||  
        read(efil, buf, 256) <= 0)  
        goto oops;  
    printf(buf, a2, a3, a4);  
}
```

The optional - causes the error messages to be placed at the end of the specified message file for recompiling part of a large mkstred program.

SEE ALSO

lseek(2), xstr(1)

NAME

more, page - file perusal filter for crt viewing

SYNOPSIS

```
more [ -cdflsu ] [ -n ] [ +linenumber ] [ +/pattern ] [ name  
... ]
```

page more options

DESCRIPTION

More is a filter which allows examination of a continuous text one screenful at a time on a soft-copy terminal. It normally pauses after each screenful, printing --More-- at the bottom of the screen. If the user then types a carriage return, one more line is displayed. If the user hits a space, another screenful is displayed. Other possibilities are enumerated later.

The command line options are:

- n An integer which is the size (in lines) of the window which more will use instead of the default.
- c More will draw each page by beginning at the top of the screen and erasing each line just before it draws on it. This avoids scrolling the screen, making it easier to read while more is writing. This option will be ignored if the terminal does not have the ability to clear to the end of a line.
- d More will prompt the user with the message "Press space to continue, 'q' to quit." at the end of each screenful, and will respond to subsequent illegal user input by printing "Press 'h' for instructions." instead of ringing the bell. This is useful if more is being used as a filter in some setting, such as a class, where many users may be unsophisticated.
- f This causes more to count logical, rather than screen lines. That is, long lines are not folded. This option is recommended if nroff output is being piped through ul, since the latter may generate escape sequences. These escape sequences contain characters which would ordinarily occupy screen positions, but which do not print when they are sent to the terminal as part of an escape sequence. Thus more may think that lines are longer than they actually are, and fold lines erroneously.
- l Do not treat ^L (form feed) specially. If this option is not given, more will pause after any line that contains a ^L, as if the end of a screenful had been

reached. Also, if a file begins with a form feed, the screen will be cleared before the file is printed.

- s Squeeze multiple blank lines from the output, producing only one blank line. Especially helpful when viewing nroff output, this option maximizes the useful information present on the screen.
- u Normally, more will handle underlining such as produced by nroff in a manner appropriate to the particular terminal: if the terminal can perform underlining or has a stand-out mode, more will output appropriate escape sequences to enable underlining or stand-out mode for underlined information in the source file. The -u option suppresses this processing.

+linenumber
Start up at linenumber.

+/*pattern*
Start up two lines before the line containing the regular expression *pattern*.

If the program is invoked as *page*, then the screen is cleared before each screenful is printed (but only if a full screenful is being printed), and $k - 1$ rather than $k - 2$ lines are printed in each screenful, where k is the number of lines the terminal can display.

More looks in the file */etc/termcap* to determine terminal characteristics, and to determine the default window size. On a terminal capable of displaying 24 lines, the default window size is 22 lines.

More looks in the environment variable *MORE* to pre-set any flags desired. For example, if you prefer to view files using the -c mode of operation, the *cs*h command *setenv MORE -c* or the *sh* command sequence *MORE='-c' ; export MORE* would cause all invocations of *more*, including invocations by programs such as *man* and *msgs*, to use this mode. Normally, the user will place the command sequence which sets up the *MORE* environment variable in the *.cshrc* or *.profile* file.

If *more* is reading from a file, rather than a pipe, then a percentage is displayed along with the --More-- prompt. This gives the fraction of the file (in characters, not lines) that has been read so far.

Other sequences which may be typed when *more* pauses, and their effects, are as follows (*i* is an optional integer argument, defaulting to 1) :

`i<space>`
display `i` more lines, (or another screenful if no argument is given)

`^D` display 11 more lines (a ``scroll``). If `i` is given, then the scroll size is set to `i`.

`d` same as `^D` (control-D)

`iz` same as typing a space except that `i`, if present, becomes the new window size.

`is` skip `i` lines and print a screenful of lines

`if` skip `i` screenfuls and print a screenful of lines

`ib` skip back `i` screenfuls and print a screenful of lines

`i^B` same as `b`

`q` or `Q`
Exit from more.

`=` Display the current line number.

`v` Start up the editor `vi` at the current line.

`h` Help command; give a description of all the more commands.

`i/expr`
search for the `i`-th occurrence of the regular expression `expr`. If there are less than `i` occurrences of `expr`, and the input is a file (rather than a pipe), then the position in the file remains unchanged. Otherwise, a screenful is displayed, starting two lines before the place where the expression was found. The user's erase and kill characters may be used to edit the regular expression. Erasing back past the first column cancels the search command.

`in` search for the `i`-th occurrence of the last regular expression entered.

`'` (single quote) Go to the point from which the last search started. If no search has been performed in the current file, this command goes back to the beginning of the file.

`!command`
invoke a shell with `command`. The characters ``%'` and ``!'` in "`command`" are replaced with the current file name

and the previous shell command respectively. If there is no current file name, '%' is not expanded. The sequences "%\" and "\!" are replaced by "%" and "!" respectively.

i:n skip to the i-th next file given in the command line (skips to last file if n doesn't make sense)

i:p skip to the i-th previous file given in the command line. If this command is given in the middle of printing out a file, then more goes back to the beginning of the file. If i doesn't make sense, more skips back to the first file. If more is not reading from a file, the bell is rung and nothing else happens.

:f display the current file name and line number.

:q or :Q
exit from more (same as q or Q).

. (dot) repeat the previous command.

The commands take effect immediately, i.e., it is not necessary to type a carriage return. Up to the time when the command character itself is given, the user may hit the line kill character to cancel the numerical argument being formed. In addition, the user may hit the erase character to redisplay the --More--(xx%) message.

At any time when output is being sent to the terminal, the user can hit the quit key (normally control-\\). More will stop sending output, and will display the usual --More-- prompt. The user may then enter one of the above commands in the normal manner. Unfortunately, some output is lost when this is done, due to the fact that any characters waiting in the terminal's output queue are flushed when the quit signal occurs.

The terminal is set to noecho mode by this program so that the output can be continuous. What you type will thus not show on your terminal, except for the / and ! commands.

If the standard output is not a teletype, then more acts just like cat, except that a header is printed before each file (if there is more than one).

A sample usage of more in previewing nroff output would be

```
nroff -ms +2 doc.n | more -s
```

FILES

/etc/termcap Terminal data base

/usr/share/misc/more.help Help file

SEE ALSO

csh(1), man(1), msgs(1), script(1), sh(1), environ(7)

BUGS

Skipping backwards is too slow on large files.

NAME

mset - retrieve ASCII to IBM 3270 keyboard map

SYNOPSIS

mset

DESCRIPTION

Mset retrieves mapping information for the ASCII keyboard to IBM 3270 terminal special functions. Normally, these mappings are found in /usr/share/misc/map3270 (see map3270(5)). This information is used by the tn3270 command (see tn3270(1)).

Mset can be used store the mapping information in the process environment in order to avoid scanning /usr/share/misc/map3270 each time tn3270 is invoked. To do this, place the following command in your .login file:

```
set noglob; setenv MAP3270 "`mset`; unset noglob
```

Mset first determines the user's terminal type from the environment variable TERM. Normally mset then uses the file /usr/share/misc/map3270 to find the keyboard mapping for that terminal. However, if the environment variable MAP3270 exists and contains the entry for the specified terminal, then that definition is used. If the value of MAP3270 begins with a slash ('/') then it is assumed to be the full pathname of an alternate mapping file and that file is searched first. In any case, if the mapping for the terminal is not found in the environment, nor in an alternate map file, nor in the standard map file, then the same search is performed for an entry for a terminal type of unknown. If that search also fails, then a default mapping is used.

FILES

/usr/share/misc/map3270 keyboard mapping for known terminals

SEE ALSO

tn3270(1), map3270(5)

BUGS

If the entry for the specific terminal exceeds 1024 bytes, csh(1) will fail to set the environment variable. Mset should probably detect this case and output the path to the map3270 file instead of the terminal entry.

NAME

msgs - system messages and junk mail program

SYNOPSIS

msgs [-fhlpq] [number] [-number]

msgs -s

msgs -c [-days]

DESCRIPTION

Msgs is used to read system messages. These messages are sent by mailing to the login `msgs' and should be short pieces of information which are suitable to be read once by most users of the system.

Msgs is normally invoked each time you login, by placing it in the file .login (.profile if you use /bin/sh). It will then prompt you with the source and subject of each new message. If there is no subject line, the first few non-blank lines of the message will be displayed. If there is more to the message, you will be told how long it is and asked whether you wish to see the rest of the message. The possible responses are:

y type the rest of the message.

RETURN synonym for y.

n skip this message and go on to the next message.

- redisplay the last message.

q drops you out of msgs; the next time you run the program it will pick up where you left off.

s append the current message to the file ``Messages'' in the current directory; `s-' will save the previously displayed message. A `s' or `s-' may be followed by a space and a file name to receive the message replacing the default ``Messages''.

m or `m-' causes a copy of the specified message to be placed in a temporary mailbox and mail(1) to be invoked on that mailbox. Both `m' and `s' accept a numeric argument in place of the `-'.

Msgs keeps track of the next message you will see by a number in the file .msgsrc in your home directory. In the directory /usr/msgs it keeps a set of files whose names are the (sequential) numbers of the messages they represent. The file /usr/msgs/bounds shows the low and high number of

the messages in the directory so that msgs can quickly determine if there are no messages for you. If the contents of bounds is incorrect it can be fixed by removing it; msgs will make a new bounds file the next time it is run.

The -s option is used for setting up the posting of messages. The line

```
msgs: "| /usr/ucb/msgs -s"
```

should be include in /etc/aliases to enable posting of messages.

The -c option is used for performing cleanup on /usr/msgs. An entry with the -c option should be placed in /etc/crontab to run every night. This will remove all messages over 21 days old. A different expiration may be specified on the command line to override the default.

Options when reading messages include:

- f which causes it not to say ``No new messages.''. This is useful in your .login file since this is often the case here.
- q Queries whether there are messages, printing ``There are new messages.' ' if there are. The command ``msgs -q' ' is often used in login scripts.
- h causes msgs to print the first part of messages only.
- l option causes only locally originated messages to be reported.
- num A message number can be given on the command line, causing msgs to start at the specified message rather than at the next message indicated by your .msgsrc file. Thus

```
msgs -h 1
```

prints the first part of all messages.
- number will cause msgs to start number messages back from the one indicated by your .msgsrc file, useful for reviews of recent messages.
- p causes long messages to be piped through more(1).

Within msgs you can also go to any specific message by typing its number when msgs requests input as to what to do.

FILES

/usr/msgsg/* database
~/.msgsrc number of next message to be presented

AUTHORS

William Joy
David Wasley

SEE ALSO

aliases(5), crontab(5), mail(1), more(1)

BUGS

NAME

mt - magnetic tape manipulating program

SYNOPSIS

mt [-f tapename] command [count]

DESCRIPTION

Mt is used to give commands to a magnetic tape drive. If a tape name is not specified, the environment variable TAPE is used; if TAPE does not exist, mt uses the device /dev/rmt12. Note that tapename must reference a raw (not block) tape device. By default mt performs the requested operation once. Operations may be performed multiple times by specifying count.

The available commands are listed below. Only as many characters as are required to uniquely identify a command need be specified.

eof, weof

Write count end-of-file marks at the current position on the tape.

fsf Forward space count files.

fsr Forward space count records.

bsf Back space count files.

bsr Back space count records.

rewind

Rewind the tape (Count is ignored).

offline, rewoffl

Rewind the tape and place the tape unit off-line (Count is ignored).

status

Print status information about the tape unit.

cacheon

Enable the readahead/writebehind cache on the tape unit.

cacheoff

Turn off the readahead/writebehind cache on the tape unit.

Mt returns a 0 exit status when the operation(s) were successful, 1 if the command was unrecognized, and 2 if an operation failed.

The `cacheon` and `cacheoff` commands currently only apply to the `TMSCP` driver and then only for certain drives such as the `TU81+`. No error is produced by the `TMSCP` driver if the cache commands are used. Other drivers will return an error code since they do not recognize the `MTCACHE` and `MTNOCACHE` functions codes. See `mtio(4)`.

FILES

`/dev/rmt*` Raw magnetic tape interface

SEE ALSO

`mtio(4)`, `tmSCP(4)`, `dd(1)`, `ioctl(2)`, `environ(7)`

NAME

mv - move or rename files

SYNOPSIS

mv [-i] [-f] [-] file1 file2

mv [-i] [-f] [-] file ... directory

DESCRIPTION

Mv moves (changes the name of) file1 to file2.

If file2 already exists, it is removed before file1 is moved. If file2 has a mode which forbids writing, mv prints the mode (see chmod(2)) and reads the standard input to obtain a line; if the line begins with y, the move takes place; if not, mv exits.

In the second form, one or more files (plain files or directories) are moved to the directory with their original file-names.

Mv refuses to move a file onto itself.

Options:

- i stands for interactive mode. Whenever a move is to supercede an existing file, the user is prompted by the name of the file followed by a question mark. If he answers with a line starting with 'y', the move continues. Any other reply prevents the move from occurring.
- f stands for force. This option overrides any mode restrictions or the -i switch.
- means interpret all the following arguments to mv as file names. This allows file names starting with minus.

SEE ALSO

cp(1), ln(1)

BUGS

If file1 and file2 lie on different file systems, mv must copy the file and delete the original. In this case the owner name becomes that of the copying process and any linking relationship with other files is lost.

NAME

netstat - show network status

SYNOPSIS

```
netstat [ -Aan ] [ -f address_family ] [ system ] [ core ]
netstat [ -himnrs ] [ -f address_family ] [ system ] [ core ]
netstat [ -n ] [ -I interface ] interval [ system ] [ core ]
```

DESCRIPTION

The netstat command symbolically displays the contents of various network-related data structures. There are a number of output formats, depending on the options for the information presented. The first form of the command displays a list of active sockets for each protocol. The second form presents the contents of one of the other network data structures according to the option selected. Using the third form, with an interval specified, netstat will continuously display the information regarding packet traffic on the configured network interfaces.

The options have the following meaning:

- A With the default display, show the address of any protocol control blocks associated with sockets; used for debugging.
- a With the default display, show the state of all sockets; normally sockets used by server processes are not shown.
- h Show the state of the IMP host table.
- i Show the state of interfaces which have been auto-configured (interfaces statically configured into a system, but not located at boot time are not shown).
- I interface
Show information only about this interface; used with an interval as described below.
- m Show statistics recorded by the memory management routines (the network manages a private pool of memory buffers).
- n Show network addresses as numbers (normally netstat interprets addresses and attempts to display them symbolically). This option may be used with any of the display formats.
- s Show per-protocol statistics.

`-r` Show the routing tables. When `-s` is also present, show routing statistics instead.

`-f address_family`

Limit statistics or address control block reports to those of the specified address family. The following address families are recognized: `inet`, for `AF_INET`, `ns`, for `AF_NS`, and `unix`, for `AF_UNIX`.

The arguments, `system` and `core` allow substitutes for the defaults ```/vmunix''` and ```/dev/kmem''`.

The default display, for active sockets, shows the local and remote addresses, send and receive queue sizes (in bytes), protocol, and the internal state of the protocol. Address formats are of the form ```host.port''` or ```network.port''` if a socket's address specifies a network but no specific host address. When known the host and network addresses are displayed symbolically according to the data bases `/etc/hosts` and `/etc/networks`, respectively. If a symbolic name for an address is unknown, or if the `-n` option is specified, the address is printed numerically, according to the address family. For more information regarding the Internet ```dot format''` refer to `inet(3N)`. Unspecified, or ```wildcard''`, addresses and ports appear as ```*''`.

The interface display provides a table of cumulative statistics regarding packets transferred, errors, and collisions. The network addresses of the interface and the maximum transmission unit (```mtu''`) are also displayed.

The routing table display indicates the available routes and their status. Each route consists of a destination host or network and a gateway to use in forwarding packets. The flags field shows the state of the route (```U''` if ```up''`), whether the route is to a gateway (```G''`), and whether the route was created dynamically by a redirect (```D''`). Direct routes are created for each interface attached to the local host; the gateway field for such entries shows the address of the outgoing interface. The `refcnt` field gives the current number of active uses of the route. Connection oriented protocols normally hold on to a single route for the duration of a connection while connectionless protocols obtain a route while sending to the same destination. The `use` field provides a count of the number of packets sent using that route. The interface entry indicates the network interface utilized for the route.

When `netstat` is invoked with an interval argument, it displays a running count of statistics related to network interfaces. This display consists of a column for the primary interface (the first interface found during

autoconfiguration) and a column summarizing information for all interfaces. The primary interface may be replaced with another interface with the `-I` option. The first line of each screen of information contains a summary since the system was last rebooted. Subsequent lines of output show values accumulated over the preceding interval.

SEE ALSO

`iostat(1)`, `vmstat(1)`, `hosts(5)`, `networks(5)`, `protocols(5)`, `services(5)`, `trpt(8C)`

BUGS

The notion of errors is ill-defined. Collisions mean something else for the IMP.

NAME

newaliases - rebuild the data base for the mail aliases file

SYNOPSIS

newaliases

DESCRIPTION

Newaliases rebuilds the random access data base for the mail aliases file /etc/aliases. It must be run each time /etc/aliases is changed in order for the change to take effect.

SEE ALSO

aliases(5), sendmail(8)

NAME

nice, nohup - run a command at low priority (sh only)

SYNOPSIS

nice [-number] command [arguments]

nohup command [arguments]

DESCRIPTION

Nice executes command with low scheduling priority. If the number argument is present, the priority is incremented (higher numbers mean lower priorities) by that amount up to a limit of 20. The default number is 10.

The super-user may run commands with priority higher than normal by using a negative priority, e.g. '--10'.

Nohup executes command immune to hangup and terminate signals from the controlling terminal. The priority is incremented by 5. Nohup should be invoked from the shell with '&' in order to prevent it from responding to interrupts by or stealing the input from the next person who logs in on the same terminal.

FILES

nohup.out standard output and standard error file under nohup

SEE ALSO

csh(1), setpriority(2), renice(8)

DIAGNOSTICS

Nice returns the exit status of the subject command.

BUGS

Nice and nohup are particular to sh(1). If you use csh(1), then commands executed with '&' are automatically immune to hangup signals while in the background. There is a builtin command nohup which provides immunity from terminate, but it does not redirect output to nohup.out.

Nice is built into csh(1) with a slightly different syntax than described here. The form ``nice +10'' nices to positive nice, and ``nice -10'' can be used by the super-user to give a process more of the processor.

NAME

nm - print name list (2BSD)

SYNOPSIS

nm [-gnopru] [file ...]

DESCRIPTION

Nm prints the name list (symbol table) of each object file in the argument list. If an argument is an archive, a listing for each object file in the archive will be produced. If no file is given, the symbols in "a.out" are listed.

Each symbol name is preceded by its value (blanks if undefined) and one of the letters U (undefined), A (absolute), T (text segment symbol), D (data segment symbol), B (bss segment symbol), C (common symbol), or f file name. If the symbol is local (non-external) the type letter is in lower case. If the file is an overlaid executable, the overlay number is printed after the name. The number is printed if the symbol is in an overlay or if it is the entry point (in the base segment) for a subroutine in an overlay. The output is sorted alphabetically.

Options are:

- g Print only global (external) symbols.
- n Sort numerically rather than alphabetically.
- o Prepend file or archive element name to each output line rather than only once.
- p Don't sort; print in symbol-table order.
- r Sort in reverse order.
- u Print only undefined symbols.

SEE ALSO

ar(1), ar(5), a.out(5), stab(5)

NAME

nroff - text formatting

SYNOPSIS

nroff [option] ... [file] ...

DESCRIPTION

Nroff formats text in the named files for typewriter-like devices. See also troff(1). The full capabilities of nroff are described in the Nroff/Troff User's Manual.

If no file argument is present, the standard input is read. An argument consisting of a single minus (-) is taken to be a file name corresponding to the standard input.

The options, which may appear in any order so long as they appear before the files, are:

- olist Print only pages whose page numbers appear in the comma-separated list of numbers and ranges. A range N-M means pages N through M; an initial -N means from the beginning to page N; and a final N- means from N to the end.
- nN Number first generated page N.
- sN Stop every N pages. Nroff will halt prior to every N pages (default N=1) to allow paper loading or changing, and will resume upon receipt of a newline.
- mname Prepend the macro file /usr/share/tmac/tmac.name to the input files.
- raN Set register a (one-character) to N.
- i Read standard input after the input files are exhausted.
- q Invoke the simultaneous input-output mode of the rd request.
- Tname Prepare output for specified terminal. Known names are 37 for the (default) Teletype Corporation Model 37 terminal, tn300 for the GE TermiNet 300 (or any terminal without half-line capability), 300S for the DASI-300S, 300 for the DASI-300, and 450 for the DASI-450 (Diablo Hyterm).
- e Produce equally-spaced words in adjusted lines, using full terminal resolution.
- h Use output tabs during horizontal spacing to speed

output and reduce output character count. Tab settings are assumed to be every 8 nominal character widths.

FILES

/tmp/ta*	temporary file
/usr/share/tmac/tmac.*	standard macro files
/usr/share/term/*	terminal driving tables for nroff

SEE ALSO

J. F. Ossanna, Nroff/Troff user's manual
B. W. Kernighan, A TROFF Tutorial
troff(1), eqn(1), tbl(1), ms(7), me(7), man(7), col(1)

NAME

nslookup - query name servers interactively

SYNOPSIS

```
nslookup [ host-to-find | - [ server address | server name ] ]
```

DESCRIPTION

Nslookup is a program to query DARPA Internet domain name servers. Nslookup has two modes: interactive and non-interactive. Interactive mode allows the user to query the name server for information about various hosts and domains or print a list of hosts in the domain. Non-interactive mode is used to print just the name and Internet address of a host or domain.

ARGUMENTS

Interactive mode is entered in the following cases:

- a) when no arguments are given (the default name server will be used), and
- b) when the first argument is a hyphen (-) and the second argument is the host name of a name server.

Non-interactive mode is used when the name of the host to be looked up is given as the first argument. The optional second argument specifies a name server.

INTERACTIVE COMMANDS

Commands may be interrupted at any time by typing a control-C. To exit, type a control-D (EOF). The command line length must be less than 80 characters. N.B. an unrecognized command will be interpreted as a host name.

host [server]

Look up information for host using the current default server or using server if it is specified.

server domain

lserver domain

Change the default server to domain. Lserver uses the initial server to look up information about domain while server uses the current default server. If an authoritative answer can't be found, the names of servers that might have the answer are returned.

root Changes the default server to the server for the root of the domain name space. Currently, the host sri-nic.arpa is used. (This command is a synonym for the lserver sri-nic.arpa.) The name of the root server can be changed with the set root command.

finger [name] [> filename]

finger [name] [>> filename]

Connects with the finger server on the current host.

The current host is defined when a previous lookup for a host was successful and returned address information (see the set querytype=A command). Name is optional. > and >> can be used to redirect output in the usual manner.

ls domain [> filename]

ls domain [>> filename]

ls -a domain [> filename]

ls -a domain [>> filename]

ls -h domain [> filename]

ls -h domain [>> filename]

ls -d domain [> filename]

List the information available for domain. The default output contains host names and their Internet addresses. The -a option lists aliases of hosts in the domain. The -h option lists CPU and operating system information for the domain. The -d option lists all contents of a zone transfer. When output is directed to a file, hash marks are printed for every 50 records received from the server.

view filename

Sorts and lists the output of previous ls command(s) with more(1).

help

? Prints a brief summary of commands.

set keyword[=value]

This command is used to change state information that affects the lookups. Valid keywords are:

all Prints the current values of the various options to set. Information about the current default server and host is also printed.

[no]debug
Turn debugging mode on. A lot more information is printed about the packet sent to the server and the resulting answer.
(Default = nodebug, abbreviation = [no]deb)

[no]d2
Turn exhaustive debugging mode on. Essentially all fields of every packet are printed.
(Default = nod2)

[no]defname
Append the default domain name to every lookup.
(Default = defname, abbreviation = [no]def)

[no]search
With defname, search for each name in parent domains of the current domain.
(Default = search)

domain=name
Change the default domain name to name. The default domain name is appended to all lookup requests if the defname option has been set. The search list is set to parents of the domain with at least two components in their names.
(Default = value in hostname or /etc/resolv.conf, abbreviation = do)

querytype=value

type=value
Change the type of information returned from a query to one of:

A	the host's Internet address (the default).
CNAME	the canonical name for an alias.
HINFO	the host CPU and operating system type.
MD	the mail destination.
MX	the mail exchanger.
MG	the mail group member.
MINFO	the mailbox or mail list information.
MR	the mail rename domain name.

NS nameserver for the named zone.
Other types specified in the RFC883 document are valid but aren't very useful.
(Abbreviation = q)

[no]recurse
Tell the name server to query other servers if it does not have the information.
(Default = recurse, abbreviation = [no]rec)

retry=number
Set the number of retries to number. When a reply to a request is not received within a certain amount of time (changed with set timeout), the request is resent. The retry value controls how many times a request is resent before giving up.
(Default = 2, abbreviation = ret)

root=host
Change the name of the root server to host. This affects the root command.
(Default = sri-nic.arpa, abbreviation = ro)

timeout=number
Change the time-out interval for waiting for a reply to number seconds.
(Default = 10 seconds, abbreviation = t)

[no]vc
Always use a virtual circuit when sending requests to the server.
(Default = novc, abbreviation = [no]v)

DIAGNOSTICS

If the lookup request was not successful, an error message is printed. Possible errors are:

Time-out

The server did not respond to a request after a certain amount of time (changed with set timeout=value) and a certain number of retries (changed with set retry=value).

No information

Depending on the query type set with the set querytype command, no information about the host was available, though the host name is valid.

Non-existent domain

The host or domain name does not exist.

Connection refused

Network is unreachable

The connection to the name or finger server could not be made at the current time. This error commonly occurs with finger requests.

Server failure

The name server found an internal inconsistency in its database and could not return a valid answer.

Refused

The name server refused to service the request.

The following error should not occur and it indicates a bug in the program.

Format error

The name server found that the request packet was not in the proper format.

FILES

/etc/resolv.conf initial domain name and name server addresses.

SEE ALSO

resolver(3), resolver(5), named(8), RFC882, RFC883

AUTHOR

Andrew Cherenon

NAME

od - octal, decimal, hex, ascii dump

SYNOPSIS

od [-format] [file] [[+]*offset*[.][*b*] [*label*]]

DESCRIPTION

Od displays file, or it's standard input, in one or more dump formats as selected by the first argument. If the first argument is missing, -o is the default. Dumping continues until end-of-file.

The meanings of the format argument characters are:

- a Interpret bytes as characters and display them with their ACSII names. If the p character is given also, then bytes with even parity are underlined. The P character causes bytes with odd parity to be underlined. Otherwise the parity bit is ignored.
- b Interpret bytes as unsigned octal.
- c Interpret bytes as ASCII characters. Certain non-graphic characters appear as C escapes: null=\0, backspace=\b, formfeed=\f, newline=\n, return=\r, tab=\t; others appear as 3-digit octal numbers. Bytes with the parity bit set are displayed in octal.
- d Interpret (short) words as unsigned decimal.
- f Interpret long words as floating point.
- h Interpret (short) words as unsigned hexadecimal.
- i Interpret (short) words as signed decimal.
- l Interpret long words as signed decimal.
- o Interpret (short) words as unsigned octal.
- s[n] Look for strings of ascii graphic characters, terminated with a null byte. N specifies the minimum length string to be recognized. By default, the minimum length is 3 characters.
- v Show all data. By default, display lines that are identical to the last line shown are not output, but are indicated with an ``*'' in column 1.
- w[n] Specifies the number of input bytes to be interpreted and displayed on each output line. If w is not specified, 16 bytes are read for each display line. If n is

not specified, it defaults to 32.

x Interpret (short) words as hexadecimal.

An upper case format character implies the long or double precision form of the object.

The offset argument specifies the byte offset into the file where dumping is to commence. By default this argument is interpreted in octal. A different radix can be specified; If ``.`' is appended to the argument, then offset is interpreted in decimal. If offset begins with ``x' or ``0x'', it is interpreted in hexadecimal. If ``b' ('`B'') is appended, the offset is interpreted as a block count, where a block is 512 (1024) bytes. If the file argument is omitted, an offset argument must be preceded by ``+'.

The radix of the displayed address will be the same as the radix of the offset, if specified; otherwise it will be octal.

Label will be interpreted as a pseudo-address for the first byte displayed. It will be shown in ``()' following the file offset. It is intended to be used with core images to indicate the real memory address. The syntax for label is identical to that for offset.

SEE ALSO

adb(1)

BUGS

A file name argument can't start with ``+'. A hexadecimal offset can't be a block count. Only one file name argument can be given.

It is an historical botch to require specification of object, radix, and sign representation in a single character argument.

NAME

pagesize - print system page size

SYNOPSIS

pagesize

DESCRIPTION

Pagesize prints the size of a page of memory in bytes, as returned by `getpagesize(2)`. This program is useful in constructing portable shell scripts.

SEE ALSO

`getpagesize(2)`

NAME

passwd - change password file information

SYNOPSIS

passwd [user]

DESCRIPTION

Passwd changes the user's password. First, the user is prompted for their old password, and then, if that is correct, for the new password. The new password must be entered twice to forestall any typing errors. The super-user is not prompted for the old password.

Once the password has been verified, passwd uses mkpasswd(8) to update the user database. This is run in the background, and, at very large sites could take several minutes. Until this update is completed, the password file is unavailable for other updates and the new password will not be available to programs.

FILES

/etc/master.passwd the user database

SEE ALSO

chpass(1), login(1), passwd(5), mkpasswd(8), vipw(8)
Robert Morris and Ken Thompson, UNIX password security

NAME

patch - a program for applying a diff file to an original

SYNOPSIS

patch [options] orig patchfile [+ [options] orig]

but usually just

patch <patchfile

DESCRIPTION

Patch will take a patch file containing any of the three forms of difference listing produced by the diff program and apply those differences to an original file, producing a patched version. By default, the patched version is put in place of the original, with the original file backed up to the same name with the extension ".orig" or "~", or as specified by the -b switch. You may also specify where you want the output to go with a -o switch. If patchfile is omitted, or is a hyphen, the patch will be read from standard input.

Upon startup, patch will attempt to determine the type of the diff listing, unless over-ruled by a -c, -e, or -n switch. Context diffs and normal diffs are applied by the patch program itself, while ed diffs are simply fed to the ed editor via a pipe.

Patch will try to skip any leading garbage, apply the diff, and then skip any trailing garbage. Thus you could feed an article or message containing a diff listing to patch, and it should work. If the entire diff is indented by a consistent amount, this will be taken into account.

With context diffs, and to a lesser extent with normal diffs, patch can detect when the line numbers mentioned in the patch are incorrect, and will attempt to find the correct place to apply each hunk of the patch. As a first guess, it takes the line number mentioned for the hunk, plus or minus any offset used in applying the previous hunk. If that is not the correct place, patch will scan both forwards and backwards for a set of lines matching the context given in the hunk. First patch looks for a place where all lines of the context match. If no such place is found, and it's a context diff, and the maximum fuzz factor is set to 1 or more, then another scan takes place ignoring the first and last line of context. If that fails, and the maximum fuzz factor is set to 2 or more, the first two and last two lines of context are ignored, and another scan is made. (The default maximum fuzz factor is 2.) If patch cannot find a place to install that hunk of the patch, it will put the hunk out to a reject file, which normally is the name of the

output file plus ".rej" or "#" . (Note that the rejected hunk will come out in context diff form whether the input patch was a context diff or a normal diff. If the input was a normal diff, many of the contexts will simply be null.) The line numbers on the hunks in the reject file may be different than in the patch file: they reflect the approximate location patch thinks the failed hunks belong in the new file rather than the old one.

As each hunk is completed, you will be told whether the hunk succeeded or failed, and which line (in the new file) patch thought the hunk should go on. If this is different from the line number specified in the diff you will be told the offset. A single large offset MAY be an indication that a hunk was installed in the wrong place. You will also be told if a fuzz factor was used to make the match, in which case you should also be slightly suspicious.

If no original file is specified on the command line, patch will try to figure out from the leading garbage what the name of the file to edit is. In the header of a context diff, the filename is found from lines beginning with "****" or "---", with the shortest name of an existing file winning. Only context diffs have lines like that, but if there is an "Index:" line in the leading garbage, patch will try to use the filename from that line. The context diff header takes precedence over an Index line. If no filename can be intuited from the leading garbage, you will be asked for the name of the file to patch.

(If the original file cannot be found, but a suitable SCCS or RCS file is handy, patch will attempt to get or check out the file.)

Additionally, if the leading garbage contains a "Prereq: " line, patch will take the first word from the prerequisites line (normally a version number) and check the input file to see if that word can be found. If not, patch will ask for confirmation before proceeding.

The upshot of all this is that you should be able to say, while in a news interface, the following:

```
| patch -d /usr/src/local/blurfl
```

and patch a file in the blurfl directory directly from the article containing the patch.

If the patch file contains more than one patch, patch will try to apply each of them as if they came from separate patch files. This means, among other things, that it is assumed that the name of the file to patch must be

determined for each diff listing, and that the garbage before each diff listing will be examined for interesting things such as filenames and revision level, as mentioned previously. You can give switches (and another original file name) for the second and subsequent patches by separating the corresponding argument lists by a '+'. (The argument list for a second or subsequent patch may not specify a new patch file, however.)

Patch recognizes the following switches:

- b causes the next argument to be interpreted as the backup extension, to be used in place of ".orig" or "~".
- B causes the next argument to be interpreted as a prefix to the backup file name. If this argument is specified any argument from -b will be ignored. This argument is an extension to Larry Wall's patch v2.0.1.4, patchlevel 8, made by M. Greim (greim@sbsvax.uucp).
- c forces patch to interpret the patch file as a context diff.
- d causes patch to interpret the next argument as a directory, and cd to it before doing anything else.
- D causes patch to use the "#ifdef...#endif" construct to mark changes. The argument following will be used as the differentiating symbol. Note that, unlike the C compiler, there must be a space between the -D and the argument.
- e forces patch to interpret the patch file as an ed script.
- f forces patch to assume that the user knows exactly what he or she is doing, and to not ask any questions. It does not suppress commentary, however. Use -s for that.
- F<number>
sets the maximum fuzz factor. This switch only applies to context diffs, and causes patch to ignore up to that many lines in looking for places to install a hunk. Note that a larger fuzz factor increases the odds of a faulty patch. The default fuzz factor is 2, and it may not be set to more than the number of lines of context in the context diff, ordinarily 3.
- l causes the pattern matching to be done loosely, in case the tabs and spaces have been munged in your input

file. Any sequence of whitespace in the pattern line will match any sequence in the input file. Normal characters must still match exactly. Each line of the context must still match a line in the input file.

- n forces patch to interpret the patch file as a normal diff.
- N causes patch to ignore patches that it thinks are reversed or already applied. See also -R .
- o causes the next argument to be interpreted as the output file name.
- p<number>
sets the pathname strip count, which controls how pathnames found in the patch file are treated, in case the you keep your files in a different directory than the person who sent out the patch. The strip count specifies how many slashes are to be stripped from the front of the pathname. (Any intervening directory names also go away.) For example, supposing the filename in the patch file was

 /u/howard/src/blurfl/blurfl.c

setting -p or -p0 gives the entire pathname unmodified, -p1 gives

 u/howard/src/blurfl/blurfl.c

without the leading slash, -p4 gives

 blurfl/blurfl.c

and not specifying -p at all just gives you "blurfl.c". Whatever you end up with is looked for either in the current directory, or the directory specified by the -d switch.
- r causes the next argument to be interpreted as the reject file name.
- R tells patch that this patch was created with the old and new files swapped. (Yes, I'm afraid that does happen occasionally, human nature being what it is.) Patch will attempt to swap each hunk around before applying it. Rejects will come out in the swapped format. The -R switch will not work with ed diff scripts because there is too little information to reconstruct the reverse operation.

If the first hunk of a patch fails, patch will reverse the hunk to see if it can be applied that way. If it can, you will be asked if you want to have the -R switch set. If it can't, the patch will continue to be applied normally. (Note: this method cannot detect a reversed patch if it is a normal diff and if the first command is an append (i.e. it should have been a delete) since appends always succeed, due to the fact that a null context will match anywhere. Luckily, most patches add or change lines rather than delete them, so most reversed normal diffs will begin with a delete, which will fail, triggering the heuristic.)

-s makes patch do its work silently, unless an error occurs.

-S causes patch to ignore this patch from the patch file, but continue on looking for the next patch in the file. Thus

```
patch -S + -S + <patchfile
```

will ignore the first and second of three patches.

-v causes patch to print out it's revision header and patch level.

-x<number>
sets internal debugging flags, and is of interest only to patch patchers.

ENVIRONMENT

No environment variables are used by patch.

FILES

/tmp/patch*

SEE ALSO

diff(1)

NOTES FOR PATCH SENDERS

There are several things you should bear in mind if you are going to be sending out patches. First, you can save people a lot of grief by keeping a patchlevel.h file which is patched to increment the patch level as the first diff in the patch file you send out. If you put a Prereq: line in with the patch, it won't let them apply patches out of order without some warning. Second, make sure you've specified the filenames right, either in a context diff header, or with an Index: line. If you are patching something in a subdirectory, be sure to tell the patch user to specify a -p switch as needed. Third, you can create a file by sending

out a diff that compares a null file to the file you want to create. This will only work if the file you want to create doesn't exist already in the target directory. Fourth, take care not to send out reversed patches, since it makes people wonder whether they already applied the patch. Fifth, while you may be able to get away with putting 582 diff listings into one file, it is probably wiser to group related patches into separate files in case something goes haywire.

DIAGNOSTICS

Too many to list here, but generally indicative that patch couldn't parse your patch file.

The message "Hmm..." indicates that there is unprocessed text in the patch file and that patch is attempting to intuit whether there is a patch in that text and, if so, what kind of patch it is.

Patch will exit with a non-zero status if any reject files were created. When applying a set of patches in a loop it behooves you to check this exit status so you don't apply a later patch to a partially patched file.

CAVEATS

Patch cannot tell if the line numbers are off in an ed script, and can only detect bad line numbers in a normal diff when it finds a "change" or a "delete" command. A context diff using fuzz factor 3 may have the same problem. Until a suitable interactive interface is added, you should probably do a context diff in these cases to see if the changes made sense. Of course, compiling without errors is a pretty good indication that the patch worked, but not always.

Patch usually produces the correct results, even when it has to do a lot of guessing. However, the results are guaranteed to be correct only when the patch is applied to exactly the same version of the file that the patch was generated from.

BUGS

Could be smarter about partial matches, excessively deviant offsets and swapped code, but that would take an extra pass.

If code has been duplicated (for instance with `#ifdef OLD-CODE ... #else ... #endif`), patch is incapable of patching both versions, and, if it works at all, will likely patch the wrong one, and tell you that it succeeded to boot.

If you apply a patch you've already applied, patch will think it is a reversed patch, and offer to un-apply the patch. This could be construed as a feature.

NAME

pdx - pascal debugger

SYNOPSIS

pdx [-r] [objfile]

DESCRIPTION

Pdx is a tool for source level debugging and execution of Pascal programs. The objfile is an object file produced by the Pascal translator pi(1). If no objfile is specified, pdx looks for a file named ``obj'' in the current directory. The object file contains a symbol table which includes the name of the all the source files translated by pi to create it. These files are available for perusal while using the debugger.

If the file ``.pdxinit'' exists in the current directory, then the debugger commands in it are executed.

The -r option causes the objfile to be executed immediately; if it terminates successfully pdx exits. Otherwise it reports the reason for termination and offers the user the option of entering the debugger or simply letting px continue with a traceback. If -r is not specified, pdx just prompts and waits for a command.

The commands are:

run [args] [< filename] [> filename]

Start executing objfile, passing args as command line arguments; < or > can be used to redirect input or output in the usual manner.

trace [in procedure/function] [if condition]

trace source-line-number [if condition]

trace procedure/function [in procedure/function] [if condition]

trace expression at source-line-number [if condition]

trace variable [in procedure/function] [if condition]

Have tracing information printed when the program is executed. A number is associated with the command that is used to turn the tracing off (see the delete command).

The first argument describes what is to be traced. If it is a source-line-number, then the line is printed immediately prior to being executed. Source line numbers in a file other than the current one must be preceded by the name of the file and a colon, e.g. ``mumble.p:17''.

If the argument is a procedure or function name then every time it is called, information is printed telling

what routine called it, from what source line it was called, and what parameters were passed to it. In addition, its return is noted, and if it's a function then the value it is returning is also printed.

If the argument is an expression with an at clause then the value of the expression is printed whenever the identified source line is reached.

If the argument is a variable then the name and value of the variable is printed whenever it changes. Execution is substantially slower during this form of tracing.

If no argument is specified then all source lines are printed before they are executed. Execution is substantially slower during this form of tracing.

The clause ``in procedure/function'' restricts tracing information to be printed only while executing inside the given procedure or function.

Condition is a Pascal boolean expression and is evaluated prior to printing the tracing information; if it is false then the information is not printed.

There is no restriction on the amount of information that can be traced.

stop if condition

stop at source-line-number [if condition]

stop in procedure/function [if condition]

stop variable [if condition]

Stop execution when the given line is reached, procedure or function called, variable changed, or condition true.

delete command-number

The trace or stop corresponding to the given number is removed. The numbers associated with traces and stops are printed by the status command.

status [> filename]

Print out the currently active trace and stop commands.

cont Continue execution from where it stopped. This can only be done when the program was stopped by an interrupt or through use of the stop command.

step Execute one source line.

next Execute up to the next source line. The difference

between this and step is that if the line contains a call to a procedure or function the step command will stop at the beginning of that block, while the next command will not.

`print expression [, expression ...]`

Print out the values of the Pascal expressions. Variables declared in an outer block but having the same identifier as one in the current block may be referenced as ```block-name . variable''`.

`whatis identifier`

Print the declaration of the given identifier.

`which identifier`

Print the full qualification of the given identifier, i.e. the outer blocks that the identifier is associated with.

`assign variable expression`

Assign the value of the expression to the variable.

`call procedure(parameters)`

Execute the object code associated with the named procedure or function.

`help` Print out a synopsis of pdx commands.

`gripe`

Invokes a mail program to send a message to the person in charge of pdx.

`where`

Print out a list of the active procedures and functions and the respective source line where they are called.

`source filename`

Read pdx commands from the given filename. Especially useful when the filename has been created by redirecting a status command from an earlier debugging session.

`dump [> filename]`

Print the names and values of all active data.

`list [source-line-number [, source-line-number]]`

`list procedure/function`

List the lines in the current source file from the first line number to the second inclusive. As in the editor ```$''` can be used to refer to the last line. If no lines are specified, the entire file is listed. If the name of a procedure or function is given lines `n-k` to `n+k` are listed where `n` is the first statement in the

procedure or function and k is small.

file [filename]

Change the current source file name to filename. If none is specified then the current source file name is printed.

edit [filename]

edit procedure/function-name

Invoke an editor on filename or the current source file if none is specified. If a procedure or function name is specified, the editor is invoked on the file that contains it. Which editor is invoked by default depends on the installation. The default can be overridden by setting the environment variable EDITOR to the name of the desired editor.

pi Recompile the program and read in the new symbol table information.

sh command-line

Pass the command line to the shell for execution. The SHELL environment variable determines which shell is used.

alias new-command-name old-command-name

This command makes pdx respond to new-command-name the way it used to respond to old-command-name.

quit Exit pdx.

The following commands deal with the program at the px instruction level rather than source level. They are not intended for general use.

tracei [address] [if cond]

tracei [variable] [at address] [if cond]

stopi [address] [if cond]

stopi [at] [address] [if cond]

Turn on tracing or set a stop using a px machine instruction addresses.

xi address [, address]

Print the instructions starting at the first address. Instructions up to the second address are printed.

xd address [, address]

Print in octal the specified data location(s).

FILES

obj	Pascal object file
.pdxinit	Pdx initialization file

SEE ALSO

pi(1), px(1)
An Introduction to Pdx

BUGS

Pdx does not understand sets, and provides no information about files.

The whatis command doesn't quite work for variant records.

Bad things will happen if a procedure invoked with the call command does a non-local goto.

The commands step and next should be able to take a count that specifies how many lines to execute.

There should be commands stepi and nexti that correspond to step and next but work at the instruction level.

There should be a way to get an address associated with a line number, procedure or function, and variable.

Most of the command names are too long.

The alias facility is quite weak.

A csh-like history capability would improve the situation.

NAME

pi - Pascal interpreter code translator

SYNOPSIS

pi [-blnpstuwz] [-i name ...] name.p

DESCRIPTION

Pi translates the program in the file name.p leaving interpreter code in the file obj in the current directory. The interpreter code can be executed using px. Pix performs the functions of pi and px for 'load and go' Pascal.

The following flags are interpreted by pi; the associated options can also be controlled in comments within the program as described in the Berkeley Pascal User's Manual.

- b Block buffer the file output.
- i Enable the listing for any specified procedures and functions and while processing any specified include files.
- l Make a program listing during translation.
- n Begin each listed include file on a new page with a banner line.
- p Suppress the post-mortem control flow backtrace if an error occurs; suppress statement limit counting.
- s Accept standard Pascal only; non-standard constructs cause warning diagnostics.
- t Suppress runtime tests of subrange variables and treat assert statements as comments.
- u Card image mode; only the first 72 characters of input lines are used.
- w Suppress warning diagnostics.
- z Allow execution profiling with pxp by generating statement counters, and arranging for the creation of the profile data file pmon.out when the resulting object is executed.

FILES

file.p	input file
file.i	include file(s)
/usr/share/pascal/pi_stringtext	of the error messages
/usr/share/pascal/how_pi*basic	usage explanation
obj	interpreter code output

SEE ALSO

Berkeley Pascal User's Manual
pcc(1), pix(1), px(1), pxp(1), pxref(1)

DIAGNOSTICS

For a basic explanation do

pi

In the diagnostic output of the translator, lines containing syntax errors are listed with a flag indicating the point of error. Diagnostic messages indicate the action which the recovery mechanism took in order to be able to continue parsing. Some diagnostics indicate only that the input is 'malformed.' This occurs if the recovery can find no simple correction to make the input syntactically valid.

Semantic error diagnostics indicate a line in the source text near the point of error. Some errors evoke more than one diagnostic to help pinpoint the error; the follow-up messages begin with an ellipsis '...'.

The first character of each error message indicates its class:

E	Fatal error; no code will be generated.
e	Non-fatal error.
w	Warning - a potential problem.
s	Non-standard Pascal construct warning.

If a severe error occurs which inhibits further processing, the translator will give a diagnostic and then 'QUIT'.

AUTHORS

Charles B. Haley, William N. Joy, and Ken Thompson

BUGS

Formal parameters which are procedures and functions are not supported.

The keyword packed and the function dispose are recognized but have no effect.

For clarity, semantic errors should be flagged at an appropriate place in the source text, and multiple instances of the 'same' semantic error should be summarized at the end of a procedure or function rather than evoking many diagnostics.

When include files are present, diagnostics relating to the last procedure in one file may appear after the beginning of the listing of the next.

NAME

pix - Pascal interpreter and executor

SYNOPSIS

pix [-blnpstuwz] [-i name ...] name.p [argument ...]

DESCRIPTION

Pix is a 'load and go' version of Pascal which combines the functions of the interpreter code translator pi and the executor px. It uses pi to translate the program in the file name.p and, if there were no fatal errors during translation, causes the resulting interpreter code to be executed by px with the specified arguments. A temporary file is used for the object code; the file obj is neither created nor destroyed.

FILES

/usr/bin/pi Pascal translator
/usr/bin/px Pascal executor
/tmp/pix???? temporary files
/usr/share/pascal/how_pixbasic explanation

SEE ALSO

Berkeley Pascal User's Manual
pi(1), px(1)

DIAGNOSTICS

For a basic explanation do

pix

AUTHOR

Susan L. Graham and William N. Joy

NAME

plot - graphics filters

SYNOPSIS

plot [-Tterminal] [-rresolution] [files...]

DESCRIPTION

These commands read plotting instructions (see plot(5)) from the standard input or the specified files, and in general produce plotting instructions suitable for a particular terminal on the standard output. The -r flag may be used to specify the device's output resolution (currently only the Imagen laser printer understands this option).

If no terminal type is specified, the environment parameter \$TERM (see environ(7)) is used. Known terminals are:

4013 Tektronix 4013 storage scope.

4014 or tek

Tektronix 4014 or 4015 storage scope with Enhanced Graphics Module. (Use 4013 for Tektronix 4014 or 4015 without the Enhanced Graphics Module).

450 DASI Hyterm 450 terminal (Diablo mechanism).

300 DASI 300 or GSI terminal (Diablo mechanism).

300S DASI 300S terminal (Diablo mechanism).

aed AED 512 color graphics terminal.

bitgraph or bg

BBN bitgraph graphics terminal.

imagen or ip

Imagen laser printer (default 240 dots-per-inch resolution).

crt Any crt terminal capable of running vi(1).

dumb Dumb terminals without cursor addressing or line printers.

vt125

DEC vt125 terminal.

hp2648 or hp or hp8

Hewlett Packard 2648 graphics terminal.

ver Versatec D1200A printer-plotter.

var Benson Varian printer-plotter.

These versions of plot use the -g option of lpr(1) to send the result directly to the plotter device rather than to the standard output.

FILES

- /usr/bin/t4013
- /usr/bin/tek
- /usr/bin/t450
- /usr/bin/t300
- /usr/bin/t300s
- /usr/bin/aedplot
- /usr/bin/bgplot
- /usr/bin/crtplot
- /usr/bin/dumbplot
- /usr/bin/gigipplot
- /usr/bin/hpplot
- /usr/bin/implot
- /usr/ucb/lpr

SEE ALSO

plot(3X), plot(3F), plot(5), lpr(1)

NAME

pmerge - pascal file merger

SYNOPSIS

pmerge name.p ...

DESCRIPTION

Pmerge assembles the named Pascal files into a single standard Pascal program. The resulting program is listed on the standard output. It is intended to be used to merge a collection of separately compiled modules so that they can be run through pi , or exported to other sites.

FILES

/usr/tmp/MG* default temporary files

SEE ALSO

pc(1), pi(1),
Auxiliary documentation Berkeley Pascal User's Manual.

AUTHOR

M. Kirk McKusick

BUGS

Very minimal error checking is done, so incorrect programs will produce unpredictable results. Block comments should be placed after the keyword to which they refer or they are likely to end up in bizarre places.

NAME

pr - print file

SYNOPSIS

pr [option] ... [file] ...

DESCRIPTION

Pr produces a printed listing of one or more files. The output is separated into pages headed by a date, the name of the file or a specified header, and the page number. If there are no file arguments, pr prints its standard input.

Options apply to all following files but may be reset between files:

- n Produce n-column output.
- +n Begin printing with page n.
- h Take the next argument as a page header.
- wn For purposes of multi-column output, take the width of the page to be n characters instead of the default 72.
- f Use formfeeds instead of newlines to separate pages. A formfeed is assumed to use up two blank lines at the top of a page. (Thus this option does not affect the effective page length.)
- ln Take the length of the page to be n lines instead of the default 66.
- t Do not print the 5-line header or the 5-line trailer normally supplied for each page.
- sc Separate columns by the single character c instead of by the appropriate amount of white space. A missing c is taken to be a tab.
- m Print all files simultaneously, each in one column,

Inter-terminal messages via write(1) are forbidden during a pr.

FILES

/dev/tty? to suspend messages.

SEE ALSO

cat(1)

DIAGNOSTICS

There are no diagnostics when pr is printing on a terminal.

NAME

printenv - print out the environment

SYNOPSIS

printenv [name]

DESCRIPTION

Printenv prints out the values of the variables in the environment. If a name is specified, only its value is printed.

If a name is specified and it is not defined in the environment, printenv returns exit status 1, else it returns status 0.

SEE ALSO

sh(1), environ(7), csh(1)

NAME

prof - display profile data

SYNOPSIS

```
prof [ -a ] [ -l ] [ -n ] [ -z ] [ -s ] [ -v [ -low [ -high  
] ] ] [ a.out [ mon.out ... ] ]
```

DESCRIPTION

Prof interprets the file produced by the monitor subroutine. Under default modes, the symbol table in the named object file (a.out default) is read and correlated with the profile file (mon.out default). For each external symbol, the percentage of time spent executing between that symbol and the next is printed (in decreasing order), together with the number of times that routine was called and the number of milliseconds per call. If more than one profile file is specified, the output represents the sum of the profiles.

In order for the number of calls to a routine to be tallied, the -p option of cc, f77 or pc must have been given when the file containing the routine was compiled. This option also arranges for the profile file to be produced automatically.

Options are:

- a all symbols are reported rather than just external symbols.
- l the output is sorted by symbol value.
- n the output is sorted by number of calls
- s a summary profile file is produced in mon.sum. This is really only useful when more than one profile file is specified.
- v all printing is suppressed and a graphic version of the profile is produced on the standard output for display by the plot(1) filters. When plotting, the numbers low and high, by default 0 and 100, may be given to cause a selected percentage of the profile to be plotted with accordingly higher resolution.
- z routines which have zero usage (as indicated by call counts and accumulated time) are nevertheless printed in the output.

FILES

mon.out for profile
a.out for namelist
mon.sum for summary profile

SEE ALSO

monitor(3), profil(2), cc(1), plot(1G)

BUGS

Beware of quantization errors.

Is confused by f77 which puts the entry points at the bottom of subroutines and functions.

NAME

ps - process status

SYNOPSIS

ps [acgklnrtuwXU# [core [swap [system]]]]

DESCRIPTION

Ps prints certain indicia about active processes. To get a complete printout on the console or lpr, use ``ps axlw''. For a quick snapshot of system activity, ``ps au'' is recommended. A hyphen may precede options with no effect. The following options may be specified.

- a asks for information about all processes with terminals (ordinarily only one's own processes are displayed).
- c causes only the comm field to be displayed instead of the arguments. (The comm field is the tail of the path name of the file the process last exec'ed.) This option speeds up ps somewhat and reduces the amount of output. It is also more reliable since the process can't scribble on top of it.
- g asks for all processes. Without this option, ps only prints ``interesting'' processes. Processes are deemed to be uninteresting if they are process group leaders, or if their arguments begin with a `-' . This normally eliminates shells and getty processes.
- k causes the file /usr/sys/core is used in place of /dev/kmem and /dev/mem. This is used for postmortem system debugging.
- l asks for a long listing. The short listing contains the user name, process ID, tty, the cumulative execution time of the process and an approximation to the command line.
- n asks for numeric rather than symbolic wchans. This flag implies the ``l'' flag.
- r asks for ``raw'' output. A non-human readable sequence of structures is output on the standard output. There is one structure for each process, the format is defined by <psout.h>

tttyname

restricts output to processes whose controlling tty is the specified ttyname (which should be specified as printed by ps, including t? for processes with no tty). This option must be the last one given.

u A user oriented output is produced. This includes the name of the owner of the process, process id, nice value, size, tty, cpu time used, and the command.

w tells ps you are on a wide terminal (132 columns). Ps normally assumes you are on an 80 column terminal. This information is used to decide how much of long commands to print. The w option may be repeated, e.g. ww, and the entire command, up to 128 characters, will be printed without regard to terminal width.

x asks even about processes with no terminal.

U causes ps to update a private database where it keeps system information. Thus ``ps -U'` should be included in the `/etc/rc` file.

A process number may be given, (indicated here by #), in which case the output is restricted to that process. This option must also be last.

A second argument tells ps where to look for core if the **k** option is given, instead of `/usr/sys/core`. A third argument is the name of a swap file to use instead of the default `/dev/swap`. If a fourth argument is given, it is taken to be the file containing the system's namelist. Otherwise, ```/unix''` is used.

The output is sorted by tty, then by process ID.

The long listing is columnar and contains

F Flags associated with the process. These are defined by `#define` lines in `/usr/include/sys/proc.h`.

S The state of the process. 0: nonexistent; S: sleeping; W: waiting; R: running; I: intermediate; Z: terminated; T: stopped.

UID The user id of the process owner.

PID The process ID of the process; as in certain cults it is possible to kill a process if you know its true name.

PPID The process ID of the parent process.

CPU Processor utilization for scheduling.

PRI The priority of the process; high numbers mean low priority.

NICE Used in priority computation.

ADDR The memory address of the process if resident, otherwise the disk address.

SZ The size in blocks (512 bytes) of the memory image of the process.

WCHAN

The event for which the process is waiting or sleeping; if blank, the process is running.

TTY The controlling tty for the process.

TIME The cumulative execution time for the process.

COMMAND

The command and its arguments.

A process that has exited and has a parent, but has not yet been waited for by the parent is marked <defunct>. Ps makes an educated guess as to the file name and arguments given when the process was created by examining memory or the swap area. The method is inherently somewhat unreliable and in any event a process is entitled to destroy this information, so the names cannot be counted on too much.

FILES

/unix system namelist
/dev/kmem kernel memory
/dev/swap swap device
/usr/sys/core core file
/dev searched to find swap device and tty
names
/var/run/psdatabase system namelist and device information

SEE ALSO

kill(1), w(1), pstat(8)

BUGS

Things can change while ps is running; the picture it gives is only a close approximation to reality.

Some processes, typically those in the background, are printed with null or garbaged arguments, even though the process has not swapped. (Sometimes ps even loses on its own arguments!) In these cases, the name of the command is printed in parentheses.

When automatic crash dumps are enabled, /usr/sys/core is not a sensible default core file name.

NAME

ptx - permuted index

SYNOPSIS

ptx [option] ... [input [output]]

DESCRIPTION

Ptx generates a permuted index to file input on file output (standard input and output default). It has three phases: the first does the permutation, generating one line for each keyword in an input line. The keyword is rotated to the front. The permuted file is then sorted. Finally, the sorted lines are rotated so the keyword comes at the middle of the page. Ptx produces output in the form:

```
.xx "tail" "before keyword" "keyword and after" "head"
```

where .xx may be an nroff or troff(1) macro for user-defined formatting. The before keyword and keyword and after fields incorporate as much of the line as will fit around the keyword when it is printed at the middle of the page. Tail and head, at least one of which is an empty string "", are wrapped-around pieces small enough to fit in the unused space at the opposite end of the line. When original text must be discarded, `/' marks the spot.

The following options can be applied:

- f Fold upper and lower case letters for sorting.
- t Prepare the output for the phototypesetter; the default line length is 100 characters.
- w n Use the next argument, n, as the width of the output line. The default line length is 72 characters.
- g n Use the next argument, n, as the number of characters to allow for each gap among the four parts of the line as finally printed. The default gap is 3 characters.
- o only
Use as keywords only the words given in the only file.
- i ignore
Do not use as keywords any words given in the ignore file. If the -i and -o options are missing, use /usr/share/misc/eign as the ignore file.
- b break
Use the characters in the break file to separate words. In any case, tab, newline, and space characters are always used as break characters.

-r Take any leading nonblank characters of each input line to be a reference identifier (as to a page or chapter) separate from the text of the line. Attach that identifier as a 5th field on each output line.

The index for this manual was generated using ptx.

FILES

/usr/bin/sort
/usr/share/misc/eign

BUGS

Line length counts do not account for overstriking or proportional spacing.

NAME

pwd - working directory name

SYNOPSIS

pwd

DESCRIPTION

Pwd prints the pathname of the working (current) directory.

SEE ALSO

cd(1), csh(1), getwd(3)

BUGS

In csh(1) the command dirs is always faster (although it can give a different answer in the rare case that the current directory or a containing directory was moved after the shell descended into it).

NAME

px - Pascal interpreter

SYNOPSIS

px [obj [argument ...]]

DESCRIPTION

Px interprets the abstract machine code generated by pi. The first argument is the file to be interpreted, and defaults to obj; remaining arguments are available to the Pascal program using the built-ins argv and argc. Px is also invoked by pix when running 'load and go'.

If the program terminates abnormally an error message and a control flow backtrace are printed. The number of statements executed and total execution time are printed after normal termination. The p option of pi suppresses all of this except the message indicating the cause of abnormal termination.

FILES

obj	default object file
pmon.out	profile data file

SEE ALSO

Berkeley Pascal User's Manual
pi(1), pix(1)

DIAGNOSTICS

Most run-time error messages are self-explanatory. Some of the more unusual ones are:

Reference to an inactive file

A file other than input or output was used before a call to reset or rewrite.

Statement count limit exceeded

The limit of 500,000 executed statements (which prevents excessive looping or recursion) has been exceeded.

Bad data found on integer read

Bad data found on real read

Usually, non-numeric input was found for a number. For reals, Pascal requires digits before and after the decimal point so that numbers like '.1' or '21.' evoke the second diagnostic.

panic: Some message

Indicates a internal inconsistency detected in px probably due to a Pascal system bug. Charles B. Haley, William N. Joy, and Ken Thompson

BUGS

Calls to the procedures `dispose` and `linelimit` are ignored.

Post-mortem traceback is not limited; infinite recursion leads to almost infinite traceback.

Because interrupts sometimes find the system in the middle of a procedure or function entry or exit, the error backtrace on an interrupt is occasionally meaningless. The current line is, however, always correct; only the call backtrace and the name of the current routine may be lost.

NAME

pxp - Pascal execution profiler

SYNOPSIS

```
pxp [ -acdefjnstuw_ ] [ -23456789 ] [ -z [ name ... ] ]  
name.p
```

DESCRIPTION

Pxp can be used to obtain execution profiles of Pascal programs or as a pretty-printer. To produce an execution profile all that is necessary is to translate the program specifying the z option to pi or pix, to execute the program, and to then issue the command

```
pxp -z name.p
```

A reformatted listing is output if none of the c, t, or z options are specified; thus

```
pxp old.p > new.p
```

places a pretty-printed version of the program in 'old.p' in the file 'new.p'.

The use of the following options of pxp is discussed in sections 2.6, 5.4, 5.5 and 5.10 of the Berkeley Pascal User's Manual.

- a Print the bodies of all procedures and functions in the profile; even those which were never executed.
- c Extract profile data from the file core.
- d Include declaration parts in a profile.
- e Eliminate include directives when reformatting a file; the include is replaced by the reformatted contents of the specified file.
- f Fully parenthesize expressions.
- j Left justify all procedures and functions.
- n Eject a new page as each file is included; in profiles, print a blank line at the top of the page.
- s Strip comments from the input text.
- t Print a table summarizing procedure and function call counts.
- u Card image mode; only the first 72 characters of input

lines are used.

- w Suppress warning diagnostics.
- z Generate an execution profile. If no names, are given the profile is of the entire program. If a list of names is given, then only any specified procedures or functions and the contents of any specified include files will appear in the profile.
- _ Underline keywords.
- d With d a digit, $2 < d < 9$, causes pxp to use d spaces as the basic indenting unit. The default is 4.

FILES

name.p	input file
name.i	include file(s)
pmon.out	profile data
core	profile data source with -c
/usr/share/pascal/how_pxpinformation on basic usage	

SEE ALSO

Berkeley Pascal User's Manual
pi(1), px(1)

DIAGNOSTICS

For a basic explanation do

pxp

Error diagnostics include 'No profile data in file' with the c option if the z option was not enabled to pi; 'Not a Pascal system core file' if the core is not from a px execution; 'Program and count data do not correspond' if the program was changed after compilation, before profiling; or if the wrong program is specified.

AUTHOR

William N. Joy

BUGS

Does not place multiple statements per line.

NAME

pxref - Pascal cross-reference program

SYNOPSIS

pxref [-] name

DESCRIPTION

Pxref makes a line numbered listing and a cross-reference of identifier usage for the program in name. The optional '-' argument suppresses the listing. The keywords goto and label are treated as identifiers for the purpose of the cross-reference. Include directives are not processed, but cause the placement of an entry indexed by '#include' in the cross-reference.

SEE ALSO

Berkeley Pascal User's Manual

AUTHOR

Niklaus Wirth

BUGS

Identifiers are trimmed to 10 characters.

NAME

quota - display disc usage and limits

SYNOPSIS

quota [-qv] [user]

DESCRIPTION

Quota displays users' disc usage and limits. Only the super-user may use the optional user argument to view the limits of users other than himself.

The -q flag prints a more terse message, containing only information on file systems where usage is over quota.

If a -v flag is supplied, quota will also display user's quotas on file systems where no storage is allocated.

Quota reports only on file systems which have disc quotas. If quota exits with a non-zero status, one or more file systems are over quota.

SEE ALSO

quota(2), quotaon(8)

NAME

quota - manipulate disk quotas

SYNOPSIS

```
#include <sys/quota.h>

quota(cmd, uid, arg, addr)
int cmd, uid, arg;
char *addr;
```

DESCRIPTION

The quota call manipulates disk quotas for file systems that have had quotas enabled with `setquota(2)`. The `cmd` parameter indicates a command to be applied to the user ID `uid`. `Arg` is a command specific argument and `addr` is the address of an optional, command specific, data structure that is copied in or out of the system. The interpretation of `arg` and `addr` is given with each command below.

Q_SETDLIM

Set disc quota limits and current usage for the user with ID `uid`. `Arg` is a major-minor device indicating a particular file system. `Addr` is a pointer to a struct `dqblk` structure (defined in `<sys/quota.h>`). This call is restricted to the super-user.

Q_GETDLIM

Get disc quota limits and current usage for the user with ID `uid`. The remaining parameters are as for **Q_SETDLIM**.

Q_SETDUSE

Set disc usage limits for the user with ID `uid`. `Arg` is a major-minor device indicating a particular file system. `Addr` is a pointer to a struct `dqusage` structure (defined in `<sys/quota.h>`). This call is restricted to the super-user.

Q_SYNC

Update the on-disc copy of quota usages. `Arg` is a major-minor device indicating the file system to be sync'ed. If the `arg` parameter is specified as `NODEV`, all file systems that have disc quotas will be sync'ed. The `uid` and `addr` parameters are ignored.

Q_SETUID

Change the calling process's quota limits to those of the user with ID `uid`. The `arg` and `addr` parameters are ignored. This call is restricted to the super-user.

Q_SETWARN

Alter the disc usage warning limits for the user with

ID uid. Arg is a major-minor device indicating a particular file system. Addr is a pointer to a struct dqwarn structure (defined in <sys/quota.h>). This call is restricted to the super-user.

Q_DOWARN

Warn the user with user ID uid about excessive disc usage. This call causes the system to check its current disc usage information and print a message on the terminal of the caller for each file system on which the user is over quota. If the user is under quota, his warning count is reset to MAX_*_WARN (defined in <sys/quota.h>). If the arg parameter is specified as NODEV, all file systems that have disc quotas will be checked. Otherwise, arg indicates a specific major-minor device to be checked. This call is restricted to the super-user.

RETURN VALUE

A successful call returns 0, otherwise the value -1 is returned and the global variable errno indicates the reason for the failure.

ERRORS

A quota call will fail when one of the following occurs:

- | | |
|----------|---|
| [EINVAL] | The kernel has not been compiled with the QUOTA option. |
| [EINVAL] | Cmd is invalid. |
| [ESRCH] | No disc quota is found for the indicated user. |
| [EPERM] | The call is privileged and the caller was not the super-user. |
| [ENODEV] | The arg parameter is being interpreted as a major-minor device and it indicates an unmounted file system. |
| [EFAULT] | An invalid addr is supplied; the associated structure could not be copied in or out of the kernel. |
| [EUSERS] | The quota table is full. |

SEE ALSO

setquota(2), quotaon(8), quotacheck(8)

BUGS

There should be some way to integrate this call with the

NAME

ranlib - table-of-contents for archive libraries

SYNOPSIS

ranlib [-t] file ...

DESCRIPTION

Ranlib creates a table of external references for archive libraries, normally used by the loader, ld(1). This table is named ``__SYMDEF'' and is prepended to the archive. Files in the archive which are not executable and symbols which are uninteresting to the loader are ignored.

The options are as follows:

-t Set the modification time of the __SYMDEF file. This time is compared by the loader with the modification time of the archive to verify that the table is up-to-date with respect to the archive. If the modification time has been changed without any change to the archive (for example, by a cp(1)), the -t option can be used to ``touch'' the modification time so that it appears that the table is up-to-date.

FILES

/tmp default temporary file directory

ranlib.XXXXXX temporary file names

SEE ALSO

ar(1), ld(1), lorder(1), nm(1), ranlib(5)

HISTORY

A ranlib command appeared in Version 7 AT&T UNIX.

NAME

ratfor - rational Fortran dialect

SYNOPSIS

ratfor [option ...] [filename ...]

DESCRIPTION

Ratfor converts a rational dialect of Fortran into ordinary irrational Fortran. Ratfor provides control flow constructs essentially identical to those in C:

statement grouping:

```
{ statement; statement; statement }
```

decision-making:

```
if (condition) statement [ else statement ]
switch (integer value) {
    case integer: statement
    ...
    [ default: ] statement
}
```

loops:

```
while (condition) statement
for (expression; condition; expression) statement
do limits statement
repeat statement [ until (condition) ]
break
next
```

and some syntactic sugar to make programs easier to read and write:

free form input:

multiple statements/line; automatic continuation

comments:

```
# this is a comment
```

translation of relationals:

>, >=, etc., become .GT., .GE., etc.

return (expression)

returns expression to caller from function

define:

define name replacement

include:

include filename

Ratfor is best used with f77(1).

SEE ALSO

f77(1)

B. W. Kernighan and P. J. Plauger, Software Tools, Addison-Wesley, 1976.

NAME

rcp - remote file copy

SYNOPSIS

```
rcp [ -p ] file1 file2
rcp [ -p ] [ -r ] file ... directory
```

DESCRIPTION

Rcp copies files between machines. Each file or directory argument is either a remote file name of the form ``rhost:path'``, or a local file name (containing no ``:`` characters, or a ``/'` before any ``:'`s).

If the `-r` option is specified and any of the source files are directories, rcp copies each subtree rooted at that name; in this case the destination must be a directory.

By default, the mode and owner of file2 are preserved if it already existed; otherwise the mode of the source file modified by the `umask(2)` on the destination host is used. The `-p` option causes rcp to attempt to preserve (duplicate) in its copies the modification times and modes of the source files, ignoring the `umask`.

If path is not a full path name, it is interpreted relative to your login directory on rhost. A path on a remote host may be quoted (using `\`, `"`, or `'`) so that the metacharacters are interpreted remotely.

Rcp does not prompt for passwords; your current local user name must exist on rhost and allow remote command execution via `rsh(1C)`.

Rcp handles third party copies, where neither source nor target files are on the current machine. Hostnames may also take the form ``rname@rhost'`` to use rname rather than the current user name on the remote host. The destination hostname may also take the form ``rhost.rname'`` to support destination machines that are running 4.2BSD versions of rcp.

SEE ALSO

`cp(1)`, `ftp(1C)`, `rsh(1C)`, `rlogin(1C)`

BUGS

Doesn't detect all cases where the target of a copy might be a file in cases where only a directory should be legal. Is confused by any output generated by commands in a `.login`, `.profile`, or `.cshrc` file on the remote host.

NAME

rcsintro - introduction to RCS commands

DESCRIPTION

The Revision Control System (RCS) manages multiple revisions of text files. RCS automates the storing, retrieval, logging, identification, and merging of revisions. RCS is useful for text that is revised frequently, for example programs, documentation, graphics, papers, form letters, etc.

The basic user interface is extremely simple. The novice only needs to learn two commands: ci(1L) and co(1L). Ci, short for "check in", deposits the contents of a text file into an archival file called an RCS file. An RCS file contains all revisions of a particular text file. Co, short for "check out", retrieves revisions from an RCS file.

Functions of RCS

- + Storage and retrieval of multiple revisions of text. RCS saves all old revisions in a space efficient way. Changes no longer destroy the original, because the previous revisions remain accessible. Revisions can be retrieved according to ranges of revision numbers, symbolic names, dates, authors, and states.
- + Maintenance of a complete history of changes. RCS logs all changes automatically. Besides the text of each revision, RCS stores the author, the date and time of check-in, and a log message summarizing the change. The logging makes it easy to find out what happened to a module, without having to compare source listings or having to track down colleagues.
- + Resolution of access conflicts. When two or more programmers wish to modify the same revision, RCS alerts the programmers and prevents one modification from corrupting the other.
- + Maintenance of a tree of Revisions. RCS can maintain separate lines of development for each module. It stores a tree structure that represents the ancestral relationships among revisions.
- + Merging of revisions and resolution of conflicts. Two separate lines of development of a module can be coalesced by merging. If the revisions to be merged affect the same sections of code, RCS alerts the user about the overlapping changes.
- + Release and configuration control. Revisions can be assigned symbolic names and marked as released, stable,

experimental, etc. With these facilities, configurations of modules can be described simply and directly.

- + Automatic identification of each revision with name, revision number, creation time, author, etc. The identification is like a stamp that can be embedded at an appropriate place in the text of a revision. The identification makes it simple to determine which revisions of which modules make up a given configuration.
- + Minimization of secondary storage. RCS needs little extra space for the revisions (only the differences). If intermediate revisions are deleted, the corresponding deltas are compressed accordingly.

Getting Started with RCS

Suppose you have a file `f.c` that you wish to put under control of RCS. Invoke the check-in command

```
ci f.c
```

This command creates the RCS file `f.c,v`, stores `f.c` into it as revision 1.1, and deletes `f.c`. It also asks you for a description. The description should be a synopsis of the contents of the file. All later check-in commands will ask you for a log entry, which should summarize the changes that you made.

Files ending in `,v` are called RCS files (`'v'` stands for `'versions'`), the others are called working files. To get back the working file `f.c` in the previous example, use the check-out command

```
co f.c
```

This command extracts the latest revision from `f.c,v` and writes it into `f.c`. You can now edit `f.c` and check it back in by invoking

```
ci f.c
```

Ci increments the revision number properly. If ci complains with the message

```
ci error: no lock set by <your login>
```

then your system administrator has decided to create all RCS files with the locking attribute set to `'strict'`. In this case, you should have locked the revision during the previous check-out. Your last check-out should have been

```
co -l f.c
```

Of course, it is too late now to do the check-out with locking, because you probably modified `f.c` already, and a second check-out would overwrite your modifications. Instead, invoke

```
rcs -l f.c
```

This command will lock the latest revision for you, unless somebody else got ahead of you already. In this case, you'll have to negotiate with that person.

Locking assures that you, and only you, can check in the next update, and avoids nasty problems if several people work on the same file. Even if a revision is locked, it can still be checked out for reading, compiling, etc. All that locking prevents is a CHECK-IN by anybody but the locker.

If your RCS file is private, i.e., if you are the only person who is going to deposit revisions into it, strict locking is not needed and you can turn it off. If strict locking is turned off, the owner of the RCS file need not have a lock for check-in; all others still do. Turning strict locking off and on is done with the commands

```
rcs -U f.c and rcs -L f.c
```

If you don't want to clutter your working directory with RCS files, create a subdirectory called `RCS` in your working directory, and move all your RCS files there. RCS commands will look first into that directory to find needed files. All the commands discussed above will still work, without any modification. (Actually, pairs of RCS and working files can be specified in 3 ways: (a) both are given, (b) only the working file is given, (c) only the RCS file is given. Both RCS and working files may have arbitrary path prefixes; RCS commands pair them up intelligently).

To avoid the deletion of the working file during check-in (in case you want to continue editing), invoke

```
ci -l f.c or ci -u f.c
```

These commands check in `f.c` as usual, but perform an implicit check-out. The first form also locks the checked in revision, the second one doesn't. Thus, these options save you one check-out operation. The first form is useful if locking is strict, the second one if not strict. Both update the identification markers in your working file (see below).

You can give `ci` the number you want assigned to a checked in revision. Assume all your revisions were numbered 1.1, 1.2, 1.3, etc., and you would like to start release 2. The command

```
ci -r2 f.c or ci -r2.1 f.c
```

assigns the number 2.1 to the new revision. From then on, `ci` will number the subsequent revisions with 2.2, 2.3, etc. The corresponding `co` commands

```
co -r2 f.c and co -r2.1 f.c
```

retrieve the latest revision numbered 2.x and the revision 2.1, respectively. `Co` without a revision number selects the latest revision on the "trunk", i.e., the highest revision with a number consisting of 2 fields. Numbers with more than 2 fields are needed for branches. For example, to start a branch at revision 1.3, invoke

```
ci -r1.3.1 f.c
```

This command starts a branch numbered 1 at revision 1.3, and assigns the number 1.3.1.1 to the new revision. For more information about branches, see `rcsfile(5L)`.

Automatic Identification

RCS can put special strings for identification into your source and object code. To obtain such identification, place the marker

```
$Header$
```

into your text, for instance inside a comment. RCS will replace this marker with a string of the form

```
$Header: filename revision_number date time
author state $
```

With such a marker on the first page of each module, you can always see with which revision you are working. RCS keeps the markers up to date automatically. To propagate the markers into your object code, simply put them into literal character strings. In C, this is done as follows:

```
static char rcsid[] = "$Header$";
```

The command `ident` extracts such markers from any file, even object code and dumps. Thus, `ident` lets you find out which revisions of which modules were used in a given program.

You may also find it useful to put the marker \$Log\$ into your text, inside a comment. This marker accumulates the log messages that are requested during check-in. Thus, you can maintain the complete history of your file directly inside it. There are several additional identification markers; see co(1L) for details.

IDENTIFICATION

Author: Walter F. Tichy, Purdue University, West Lafayette, IN, 47907.

Revision Number: 1.2 ; Release Date: 87/02/27 .

Copyright (C) 1982 by Walter F. Tichy.

SEE ALSO

ci(1L), co(1L), ident(1L), merge(1L), rcs(1L), rcsdiff(1L), rcsmerge(1L), rlog(1L)

Walter F. Tichy, "Design, Implementation, and Evaluation of a Revision Control System," in Proceedings of the 6th International Conference on Software Engineering, IEEE, Tokyo, Sept. 1982.

NAME

rdist - remote file distribution program

SYNOPSIS

```
rdist [ -nqbRhivwy ] [ -f distfile ] [ -d var=value ] [ -m
host ] [ name ... ]
```

```
rdist [ -nqbRhivwy ] -c name ... [login@]host[:dest]
```

DESCRIPTION

Rdist is a program to maintain identical copies of files over multiple hosts. It preserves the owner, group, mode, and mtime of files if possible and can update programs that are executing. Rdist reads commands from distfile to direct the updating of files and/or directories. If distfile is '-', the standard input is used. If no -f option is present, the program looks first for 'distfile', then 'Distfile' to use as the input. If no names are specified on the command line, rdist will update all of the files and directories listed in distfile. Otherwise, the argument is taken to be the name of a file to be updated or the label of a command to execute. If label and file names conflict, it is assumed to be a label. These may be used together to update specific files using specific commands.

The -c option forces rdist to interpret the remaining arguments as a small distfile. The equivalent distfile is as follows.

```
( name ... ) -> [login@]host
install      [dest] ;
```

Other options:

- d Define var to have value. The -d option is used to define or override variable definitions in the distfile. Value can be the empty string, one name, or a list of names surrounded by parentheses and separated by tabs and/or spaces.
- m Limit which machines are to be updated. Multiple -m arguments can be given to limit updates to a subset of the hosts listed the distfile.
- n Print the commands without executing them. This option is useful for debugging distfile.
- q Quiet mode. Files that are being modified are normally printed on standard output. The -q option suppresses this.

- R Remove extraneous files. If a directory is being updated, any files that exist on the remote host that do not exist in the master directory are removed. This is useful for maintaining truly identical copies of directories.
- h Follow symbolic links. Copy the file that the link points to rather than the link itself.
- i Ignore unresolved links. Rdist will normally try to maintain the link structure of files being transferred and warn the user if all the links cannot be found.
- v Verify that the files are up to date on all the hosts. Any files that are out of date will be displayed but no files will be changed nor any mail sent.
- w Whole mode. The whole file name is appended to the destination directory name. Normally, only the last component of a name is used when renaming files. This will preserve the directory structure of the files being copied instead of flattening the directory structure. For example, renaming a list of files such as (dir1/f1 dir2/f2) to dir3 would create files dir3/dir1/f1 and dir3/dir2/f2 instead of dir3/f1 and dir3/f2.
- y Younger mode. Files are normally updated if their mtime and size (see stat(2)) disagree. The -y option causes rdist not to update files that are younger than the master copy. This can be used to prevent newer copies on other hosts from being replaced. A warning message is printed for files which are newer than the master copy.
- b Binary comparison. Perform a binary comparison and update files if they differ rather than comparing dates and sizes.

Distfile contains a sequence of entries that specify the files to be copied, the destination hosts, and what operations to perform to do the updating. Each entry has one of the following formats.

```
<variable name> '=' <name list>
[ label: ] <source list> '->' <destination list> <command list>
[ label: ] <source list> '::' <time_stamp file> <command list>
```

The first format is used for defining variables. The second format is used for distributing files to other hosts. The third format is used for making lists of files that have been changed since some given date. The source list

specifies a list of files and/or directories on the local host which are to be used as the master copy for distribution. The destination list is the list of hosts to which these files are to be copied. Each file in the source list is added to a list of changes if the file is out of date on the host which is being updated (second format) or the file is newer than the time stamp file (third format).

Labels are optional. They are used to identify a command for partial updates.

Newlines, tabs, and blanks are only used as separators and are otherwise ignored. Comments begin with '#' and end with a newline.

Variables to be expanded begin with '\$' followed by one character or a name enclosed in curly braces (see the examples at the end).

The source and destination lists have the following format:

```
<name>
or
`(' <zero or more names separated by white-space> `')
```

The shell meta-characters '['', ']', '{', '}', '*', and '?' are recognized and expanded (on the local host only) in the same way as csh(1). They can be escaped with a backslash. The '~' character is also expanded in the same way as csh but is expanded separately on the local and destination hosts. When the -w option is used with a file name that begins with '~', everything except the home directory is appended to the destination name. File names which do not begin with '/' or '~' use the destination user's home directory as the root directory for the rest of the file name.

The command list consists of zero or more commands of the following format.

```
`install' <options> opt_dest_name `;'
`notify' <name list> `;'
`except' <name list> `;'
`except_pat' <pattern list> `;'
`special' <name list> string `;'
```

The install command is used to copy out of date files and/or directories. Each source file is copied to each host in the destination list. Directories are recursively copied in the same way. Opt_dest_name is an optional parameter to rename files. If no install command appears in the command list or the destination name is not specified, the source file name

is used. Directories in the path name will be created if they do not exist on the remote host. To help prevent disasters, a non-empty directory on a target host will never be replaced with a regular file or a symbolic link. However, under the `-R` option a non-empty directory will be removed if the corresponding filename is completely absent on the master host. The options are `-R`, `-h`, `-i`, `-v`, `-w`, `-y`, and `-b` and have the same semantics as options on the command line except they only apply to the files in the source list. The login name used on the destination host is the same as the local host unless the destination name is of the format `login@host`.

The `notify` command is used to mail the list of files updated (and any errors that may have occurred) to the listed names. If no `@` appears in the name, the destination host is appended to the name (e.g., `name1@host`, `name2@host`, ...).

The `except` command is used to update all of the files in the source list except for the files listed in name list. This is usually used to copy everything in a directory except certain files.

The `except_pat` command is like the `except` command except that pattern list is a list of regular expressions (see `ed(1)` for details). If one of the patterns matches some string within a file name, that file will be ignored. Note that since `\` is a quote character, it must be doubled to become part of the regular expression. Variables are expanded in pattern list but not shell file pattern matching characters. To include a `$`, it must be escaped with `\`.

The `special` command is used to specify `sh(1)` commands that are to be executed on the remote host after the file in name list is updated or installed. If the name list is omitted then the shell commands will be executed for every file updated or installed. The shell variable `FILE` is set to the current filename before executing the commands in string. String starts and ends with `"` and can cross multiple lines in distfile. Multiple commands to the shell should be separated by `;`. Commands are executed in the user's home directory on the host being updated. The special command can be used to rebuild private databases, etc. after a program has been updated.

The following is a small example.

```
HOSTS = ( matisse root@arpa)

FILES = ( /bin /lib /usr/bin /usr/games
          /usr/include/{*.h,{stand,sys,vax*,pascal,machine}/*.h}
          /usr/lib /usr/man/man? /usr/ucb /usr/local/rdist )
```

```

EXLIB = ( Mail.rc aliases aliases.dir aliases.pag crontab dshrc
          sendmail.cf sendmail.fc sendmail.hf sendmail.st uucp vfont )

${FILES} -> ${HOSTS}
        install -R ;
        except /usr/lib/${EXLIB} ;
        except /usr/games/lib ;
        special /usr/sbin/sendmail "/usr/sbin/sendmail -bz" ;

srcs:
/usr/src/bin -> arpa
        except_pat ( \\..o\ $ /SCCS\ $ ) ;

IMAGEN = (ips dviimp catdvi)

imagen:
/usr/local/${IMAGEN} -> arpa
        install /usr/local/lib ;
        notify ralph ;

${FILES} :: stamp.cory
        notify root@cory ;

```

FILES

```

distfile      input command file
/tmp/rdist*   temporary file for update lists

```

SEE ALSO

```
sh(1), csh(1), stat(2)
```

DIAGNOSTICS

A complaint about mismatch of rdist version numbers may really stem from some problem with starting your shell, e.g., you are in too many groups.

BUGS

Source files must reside on the local host where rdist is executed.

There is no easy way to have a special command executed after all files in a directory have been updated.

Variable expansion only works for name lists; there should be a general macro facility.

Rdist aborts on files which have a negative mtime (before Jan 1, 1970).

There should be a 'force' option to allow replacement of non-empty directories by regular files or symlinks. A means of updating file modes and owners of otherwise identical

files is also needed.

NAME

refer - find and insert literature references in documents

SYNOPSIS

```
refer [ -a ] [ -b ] [ -c ] [ -e ] [ -fn ] [ -kx ] [ -lm,n ]  
[ -n ] [ -p bib ] [ -skeys ] [ -Bl.m ] [ -P ] [ -S ] [ file  
... ]
```

DESCRIPTION

Refer is a preprocessor for nroff or troff(1) that finds and formats references for footnotes or endnotes. It is also the base for a series of programs designed to index, search, sort, and print stand-alone bibliographies, or other data entered in the appropriate form.

Given an incomplete citation with sufficiently precise keywords, refer will search a bibliographic database for references containing these keywords anywhere in the title, author, journal, etc. The input file (or standard input) is copied to standard output, except for lines between .[and .] delimiters, which are assumed to contain keywords, and are replaced by information from the bibliographic database. The user may also search different databases, override particular fields, or add new fields. The reference data, from whatever source, are assigned to a set of troff strings. Macro packages such as ms(7) print the finished reference text from these strings. By default references are flagged by footnote numbers.

The following options are available:

- an Reverse the first n author names (Jones, J. A. instead of J. A. Jones). If n is omitted all author names are reversed.
- b Bare mode: do not put any flags in text (neither numbers nor labels).
- ckeys Capitalize (with CAPS SMALL CAPS) the fields whose key-letters are in keys.
- e Instead of leaving the references where encountered, accumulate them until a sequence of the form
.[
\$LIST\$
.]
is encountered, and then write out all references collected so far. Collapse references to same source.
- fn Set the footnote number to n instead of the default of 1 (one). With labels rather than numbers, this flag

is a no-op.

- kx Instead of numbering references, use labels as specified in a reference data line beginning %x; by default x is L.
- lm,n Instead of numbering references, use labels made from the senior author's last name and the year of publication. Only the first m letters of the last name and the last n digits of the date are used. If either m or n is omitted the entire name or date respectively is used.
- n Do not search the default file /usr/dict/papers/Ind. If there is a REFER environment variable, the specified file will be searched instead of the default file; in this case the -n flag has no effect.
- p bib Take the next argument bib as a file of references to be searched. The default file is searched last.
- skeys Sort references by fields whose key-letters are in the keys string; permute reference numbers in text accordingly. Implies -e. The key-letters in keys may be followed by a number to indicate how many such fields are used, with + taken as a very large number. The default is AD which sorts on the senior author and then date; to sort, for example, on all authors and then title, use -sA+T.
- Bl.m Bibliography mode. Take a file composed of records separated by blank lines, and turn them into troff input. Label l will be turned into the macro .m with l defaulting to %X and .m defaulting to .AP (annotation paragraph).
- P Place punctuation marks .,:;?! after the reference signal, rather than before. (Periods and commas used to be done with strings.)
- S Produce references in the Natural or Social Science format.

To use your own references, put them in the format described below. They can be searched more rapidly by running `indx-bib(1)` on them before using `refer`; failure to index results in a linear search. When `refer` is used with the `eqn`, `neqn` or `tbl` preprocessors `refer` should be first, to minimize the volume of data passed through pipes.

The refer preprocessor and associated programs expect input from a file of references composed of records separated by blank lines. A record is a set of lines (fields), each containing one kind of information. Fields start on a line beginning with a ``%'', followed by a key-letter, then a blank, and finally the contents of the field, and continue until the next line starting with ``%''. The output ordering and formatting of fields is controlled by the macros specified for nroff/troff (for footnotes and endnotes) or roffbib (for stand-alone bibliographies). For a list of the most common key-letters and their corresponding fields, see addbib(1). An example of a refer entry is given below.

EXAMPLE

```
%A   M. E. Lesk
%T   Some Applications of Inverted Indexes on the UNIX System
%B   UNIX Programmer's Manual
%V   2b
%I   Bell Laboratories
%C   Murray Hill, NJ
%D   1978
```

FILES

```
/usr/dict/papers  directory of default publication lists
/usr/libexec/refer  directory of companion programs
```

SEE ALSO

addbib(1), sortbib(1), roffbib(1), indxbib(1), lookbib(1)

AUTHOR

Mike Lesk

BUGS

Blank spaces at the end of lines in bibliography fields will cause the records to sort and reverse incorrectly. Sorting large numbers of references causes a core dump.

NAME

rev - reverse lines of a file

SYNOPSIS

rev [file] ...

DESCRIPTION

Rev copies the named files to the standard output, reversing the order of characters in every line. If no file is specified, the standard input is copied.

NAME

rlogin - remote login

SYNOPSIS

```
rlogin rhost [ -ec ] [ -8 ] [ -L ] [ -l username ]  
rhost [ -ec ] [ -8 ] [ -L ] [ -l username ]
```

DESCRIPTION

Rlogin connects your terminal on the current local host system lhost to the remote host system rhost.

Each host has a file /etc/hosts.equiv which contains a list of rhost's with which it shares account names. (The host names must be the standard names as described in rsh(1C).) When you rlogin as the same user on an equivalent host, you don't need to give a password. Each user may also have a private equivalence list in a file .rhosts in his login directory. Each line in this file should contain an rhost and a username separated by a space, giving additional cases where logins without passwords are to be permitted. If the originating user is not equivalent to the remote user, then a login and password will be prompted for on the remote machine as in login(1). To avoid some security problems, the .rhosts file must be owned by either the remote user or root.

The remote terminal type is the same as your local terminal type (as given in your environment TERM variable). The terminal or window size is also copied to the remote system if the server supports the option, and changes in size are reflected as well. All echoing takes place at the remote site, so that (except for delays) the rlogin is transparent. Flow control via ^S and ^Q and flushing of input and output on interrupts are handled properly. The optional argument -8 allows an eight-bit input data path at all times; otherwise parity bits are stripped except when the remote side's stop and start characters are other than ^S/^Q. The argument -L allows the rlogin session to be run in litout mode. A line of the form ``~.''' disconnects from the remote host, where ``~'' is the escape character. Similarly, the line ``~^Z'' (where ^Z, control-Z, is the suspend character) will suspend the rlogin session. Substitution of the delayed-suspend character (normally ^Y) for the suspend character suspends the send portion of the rlogin, but allows output from the remote system. A different escape character may be specified by the -e option. There is no space separating this option flag and the argument character.

SEE ALSO

rsh(1C)

FILES

/usr/hosts/* for rhost version of the command

BUGS

More of the environment should be propagated.

NAME

rm, rmdir - remove (unlink) files or directories

SYNOPSIS

rm [-f] [-r] [-i] [-] file ...

rmdir dir ...

DESCRIPTION

Rm removes the entries for one or more files from a directory. If an entry was the last link to the file, the file is destroyed. Removal of a file requires write permission in its directory, but neither read nor write permission on the file itself.

If a file has no write permission and the standard input is a terminal, its permissions are printed and a line is read from the standard input. If that line begins with `y' the file is deleted, otherwise the file remains. No questions are asked and no errors are reported when the -f (force) option is given.

If a designated file is a directory, an error comment is printed unless the optional argument -r has been used. In that case, rm recursively deletes the entire contents of the specified directory, and the directory itself.

If the -i (interactive) option is in effect, rm asks whether to delete each file, and, under -r, whether to examine each directory.

The null option - indicates that all the arguments following it are to be treated as file names. This allows the specification of file names starting with a minus.

Rmdir removes entries for the named directories, which must be empty.

SEE ALSO

rm(1), unlink(2), rmdir(2)

NAME

rmail - handle remote mail received via uucp

SYNOPSIS

rmail user ...

DESCRIPTION

Rmail interprets incoming mail received via uucp(1C), collapsing ``From'' lines in the form generated by binmail(1) into a single line of the form ``return-path!sender'', and passing the processed mail on to sendmail(8).

Rmail is explicitly designed for use with uucp and sendmail.

SEE ALSO

binmail(1), uucp(1C), sendmail(8)

BUGS

Rmail should not reside in /bin.

NAME

rmkdir, rm - remove (unlink) directories or files

SYNOPSIS

rmkdir dir ...

rm [-f] [-r] [-i] [-] file ...

DESCRIPTION

Rmkdir removes entries for the named directories, which must be empty.

Rm removes the entries for one or more files from a directory. If an entry was the last link to the file, the file is destroyed. Removal of a file requires write permission in its directory, but neither read nor write permission on the file itself.

If a file has no write permission and the standard input is a terminal, its permissions are printed and a line is read from the standard input. If that line begins with 'y' the file is deleted, otherwise the file remains. No questions are asked and no errors are reported when the -f (force) option is given.

If a designated file is a directory, an error comment is printed unless the optional argument -r has been used. In that case, rm recursively deletes the entire contents of the specified directory, and the directory itself.

If the -i (interactive) option is in effect, rm asks whether to delete each file, and, under -r, whether to examine each directory.

The null option - indicates that all the arguments following it are to be treated as file names. This allows the specification of file names starting with a minus.

SEE ALSO

rm(1), unlink(2), rmkdir(2)

NAME

roffbib - run off bibliographic database

SYNOPSIS

```
roffbib [ -e ] [ -h ] [ -n ] [ -o ] [ -r ] [ -s ] [ -Tterm ]  
[ -x ] [ -m mac ] [ -V ] [ -Q ] [ file ... ]
```

DESCRIPTION

Roffbib prints out all records in a bibliographic database, in bibliography format rather than as footnotes or endnotes. Generally it is used in conjunction with sortbib:

```
sortbib database | roffbib
```

Roffbib accepts most of the options understood by nroff(1), most importantly the -T flag to specify terminal type.

If abstracts or comments are entered following the %X field key, roffbib will format them into paragraphs for an annotated bibliography. Several %X fields may be given if several annotation paragraphs are desired. The -x flag will suppress the printing of these abstracts.

A user-defined set of macros may be specified after the -m option. There should be a space between the -m and the macro filename. This set of macros will replace the ones defined in /usr/share/tmac/tmac.bib. The -V flag will send output to the Versatec; the -Q flag will queue output for the phototypesetter.

Four command-line registers control formatting style of the bibliography, much like the number registers of ms(7). The command-line argument -rN1 will number the references starting at one (1). The flag -rV2 will double space the bibliography, while -rV1 will double space references but single space annotation paragraphs. The line length can be changed from the default 6.5 inches to 6 inches with the -rL6i argument, and the page offset can be set from the default of 0 to one inch by specifying -rO1i (capital O, not zero). Note: with the -V and -Q flags the default page offset is already one inch.

FILES

/usr/share/tmac/tmac.bib file of macros used by nroff/troff

SEE ALSO

refer(1), addbib(1), sortbib(1), indxbib(1), lookbib(1)

BUGS

Users have to rewrite macros to create customized formats.

NAME

rsh - remote shell

SYNOPSIS

```
rsh host [ -l username ] [ -n ] command
host [ -l username ] [ -n ] command
```

DESCRIPTION

Rsh connects to the specified host, and executes the specified command. Rsh copies its standard input to the remote command, the standard output of the remote command to its standard output, and the standard error of the remote command to its standard error. Interrupt, quit and terminate signals are propagated to the remote command; rsh normally terminates when the remote command does.

The remote username used is the same as your local username, unless you specify a different remote name with the `-l` option. This remote name must be equivalent (in the sense of `rlogin(1C)`) to the originating account; no provision is made for specifying a password with a command.

If you omit command, then instead of executing a single command, you will be logged in on the remote host using `rlogin(1C)`.

Shell metacharacters which are not quoted are interpreted on local machine, while quoted metacharacters are interpreted on the remote machine. Thus the command

```
rsh otherhost cat remotefile >> localfile
```

appends the remote file remotefile to the localfile localfile, while

```
rsh otherhost cat remotefile ">>" otherremotefile
```

appends remotefile to otherremotefile.

Host names are given in the file `/etc/hosts`. Each host has one standard name (the first name given in the file), which is rather long and unambiguous, and optionally one or more nicknames. The host names for local machines are also commands in the directory `/usr/hosts`; if you put this directory in your search path then the rsh can be omitted.

FILES

```
/etc/hosts
/usr/hosts/*
```

SEE ALSO

`rlogin(1C)`

BUGS

If you are using `cs(1)` and put a `rsh(1C)` in the background without redirecting its input away from the terminal, it will block even if no reads are posted by the remote command. If no input is desired you should redirect the input of `rsh` to `/dev/null` using the `-n` option.

You cannot run an interactive command (like `rogue(6)` or `vi(1)`); use `rlogin(1C)`.

Stop signals stop the local `rsh` process only; this is arguably wrong, but currently hard to fix for reasons too complicated to explain here.

NAME

ruptime - show host status of local machines

SYNOPSIS

ruptime [-a] [-r] [-l] [-t] [-u]

DESCRIPTION

Ruptime gives a status line like uptime for each machine on the local network; these are formed from packets broadcast by each host on the network once a minute.

Machines for which no status report has been received for 11 minutes are shown as being down.

Users idle an hour or more are not counted unless the -a flag is given.

Normally, the listing is sorted by host name. The -l , -t , and -u flags specify sorting by load average, uptime, and number of users, respectively. The -r flag reverses the sort order.

FILES

/usr/spool/rwho/whod.* data files

SEE ALSO

rwho(1C)

NAME

rwho - who's logged in on local machines

SYNOPSIS

rwho [-a]

DESCRIPTION

The rwho command produces output similar to who, but for all machines on the local network. If no report has been received from a machine for 5 minutes then rwho assumes the machine is down, and does not report users last known to be logged into that machine.

If a users hasn't typed to the system for a minute or more, then rwho reports this idle time. If a user hasn't typed to the system for an hour or more, then the user will be omitted from the output of rwho unless the -a flag is given.

FILES

/usr/spool/rwho/whod.* information about other machines

SEE ALSO

ruptime(1C), rwhod(8C)

BUGS

This is unwieldy when the number of machines on the local net is large.

NAME

sccs - front end for the SCCS subsystem

SYNOPSIS

sccs [-r] [-dpath] [-ppath] command [flags] [args]

DESCRIPTION

Sccs is a front end to the SCCS programs that helps them mesh more cleanly with the rest of UNIX. It also includes the capability to run ``set user id'' to another user to provide additional protection.

Basically, sccs runs the command with the specified flags and args. Each argument is normally modified to be prepended with ``SCCS/s.''.

Flags to be interpreted by the sccs program must be before the command argument. Flags to be passed to the actual SCCS program must come after the command argument. These flags are specific to the command and are discussed in the documentation for that command.

Besides the usual SCCS commands, several ``pseudo-commands'' can be issued. These are:

edit Equivalent to ``get -e''.

delget Perform a delta on the named files and then get new versions. The new versions will have id keywords expanded, and will not be editable. The -m, -p, -r, -s, and -y flags will be passed to delta, and the -b, -c, -e, -i, -k, -l, -s, and -x flags will be passed to get.

deledit Equivalent to ``delget'' except that the ``get'' phase includes the ``-e'' flag. This option is useful for making a ``checkpoint'' of your current editing phase. The same flags will be passed to delta as described above, and all the flags listed for ``get'' above except -e and -k are passed to ``edit''.

create Creates an SCCS file, taking the initial contents from the file of the same name. Any flags to ``admin'' are accepted. If the creation is successful, the files are renamed with a comma on the front. These should be removed when you are convinced that the SCCS files have been created successfully.

fix Must be followed by a -r flag. This command essentially removes the named delta, but leaves

you with a copy of the delta with the changes that were in it. It is useful for fixing small compiler bugs, etc. Since it doesn't leave audit trails, it should be used carefully.

- clean** This routine removes everything from the current directory that can be recreated from SCCS files. It will not remove any files being edited. If the **-b** flag is given, branches are ignored in the determination of whether they are being edited; this is dangerous if you are keeping the branches in the same directory.
- unedit** This is the opposite of an ```edit''` or a ```get -e''`. It should be used with extreme caution, since any changes you made since the `get` will be irretrievably lost.
- info** Gives a listing of all files being edited. If the **-b** flag is given, branches (i.e., SID's with two or fewer components) are ignored. If the **-u** flag is given (with an optional argument) then only files being edited by you (or the named user) are listed.
- check** Like ```info''` except that nothing is printed if nothing is being edited, and a non-zero exit status is returned if anything is being edited. The intent is to have this included in an ```install''` entry in a makefile to insure that everything is included into the SCCS file before a version is installed.
- tell** Gives a newline-separated list of the files being edited on the standard output. Takes the **-b** and **-u** flags like ```info''` and ```check''`.
- diffs** Gives a ```diff''` listing between the current version of the program(s) you have out for editing and the versions in SCCS format. The **-r**, **-c**, **-i**, **-x**, and **-t** flags are passed to `get`; the **-l**, **-s**, **-e**, **-f**, **-h**, and **-b** options are passed to `diff`. The **-C** flag is passed to `diff` as **-c**.
- print** This command prints out verbose information about the named files.

The **-r** flag runs `sccs` as the real user rather than as whatever effective user `sccs` is ```set user id''` to. The **-d** flag gives a root directory for the SCCS files. The default is the current directory. The **-p** flag defines the pathname of the directory in which the SCCS files will be found;

```SCCS''` is the default. The `-p` flag differs from the `-d` flag in that the `-d` argument is prepended to the entire pathname and the `-p` argument is inserted before the final component of the pathname. For example, ```sccs -d/x -py get a/b''` will convert to ```get /x/a/y/s.b''`. The intent here is to create aliases such as ```alias syssccs sccs -d/usr/src''` which will be used as ```syssccs get cmd/who.c''`. Also, if the environment variable `PROJECT` is set, its value is used to determine the `-d` flag. If it begins with a slash, it is taken directly; otherwise, the home directory of a user of that name is examined for a sub-directory ```src''` or ```source''`. If such a directory is found, it is used.

Certain commands (such as `admin`) cannot be run ```set user id''` by all users, since this would allow anyone to change the authorizations. These commands are always run as the real user.

#### EXAMPLES

To get a file for editing, edit it, and produce a new delta:

```
sccs get -e file.c
ex file.c
sccs delta file.c
```

To get a file from another directory:

```
sccs -p/usr/src/sccs/s. get cc.c
```

or

```
sccs get /usr/src/sccs/s.cc.c
```

To make a delta of a large number of files in the current directory:

```
sccs delta *.c
```

To get a list of files being edited that are not on branches:

```
sccs info -b
```

To delta everything being edited by you:

```
sccs delta `sccs tell -u`
```

In a makefile, to get source files from an SCCS file if it does not already exist:

```
SRCS = <list of source files>
$(SRCS):
 sccs get $(REL) $@
```

#### SEE ALSO

admin(SCCS), chghist(SCCS), comb(SCCS), delta(SCCS),  
get(SCCS), help(SCCS), prt(SCCS), rmdel(SCCS),  
sccsdiff(SCCS), what(SCCS)  
Eric Allman, An Introduction to the Source Code Control Sys-  
tem

#### BUGS

It should be able to take directory arguments on pseudo-  
commands like the SCCS commands do.



## NAME

script - make typescript of terminal session

## SYNOPSIS

script [ -a ] [ file ]

## DESCRIPTION

Script makes a typescript of everything printed on your terminal. The typescript is written to file, or appended to file if the -a option is given. It can be sent to the line printer later with lpr. If no file name is given, the typescript is saved in the file typescript.

The script ends when the forked shell exits.

This program is useful when using a crt and a hard-copy record of the dialog is desired, as for a student handing in a program that was developed on a crt when hard-copy terminals are in short supply.

## BUGS

Script places everything in the log file. This is not what the naive user expects.

## NAME

sed - stream editor

## SYNOPSIS

sed [ -n ] [ -e script ] [ -f sfile ] [ file ] ...

## DESCRIPTION

Sed copies the named files (standard input default) to the standard output, edited according to a script of commands. The -f option causes the script to be taken from file sfile; these options accumulate. If there is just one -e option and no -f's, the flag -e may be omitted. The -n option suppresses the default output.

A script consists of editing commands, one per line, of the following form:

[address [, address] ] function [arguments]

In normal operation sed cyclically copies a line of input into a pattern space (unless there is something left after a 'D' command), applies in sequence all commands whose addresses select that pattern space, and at the end of the script copies the pattern space to the standard output (except under -n) and deletes the pattern space.

An address is either a decimal number that counts input lines cumulatively across files, a '\$' that addresses the last line of input, or a context address, '/regular expression/', in the style of ed(1) modified thus:

The escape sequence '\n' matches a newline embedded in the pattern space.

A command line with no addresses selects every pattern space.

A command line with one address selects each pattern space that matches the address.

A command line with two addresses selects the inclusive range from the first pattern space that matches the first address through the next pattern space that matches the second. (If the second address is a number less than or equal to the line number first selected, only one line is selected.) Thereafter the process is repeated, looking again for the first address.

Editing commands can be applied only to non-selected pattern spaces by use of the negation function '!' (below).

In the following list of functions the maximum number of permissible addresses for each function is indicated in parentheses.

An argument denoted text consists of one or more lines, all but the last of which end with '\\' to hide the newline. Backslashes in text are treated like backslashes in the replacement string of an 's' command, and may be used to protect initial blanks and tabs against the stripping that is done on every script line.

An argument denoted rfile or wfile must terminate the command line and must be preceded by exactly one blank. Each wfile is created before processing begins. There can be at most 10 distinct wfile arguments.

(1)a\  
text

Append. Place text on the output before reading the next input line.

(2)b label

Branch to the ':' command bearing the label. If label is empty, branch to the end of the script.

(2)c\  
text

Change. Delete the pattern space. With 0 or 1 address or at the end of a 2-address range, place text on the output. Start the next cycle.

(2)d Delete the pattern space. Start the next cycle.

(2)D Delete the initial segment of the pattern space through the first newline. Start the next cycle.

(2)g Replace the contents of the pattern space by the contents of the hold space.

(2)G Append the contents of the hold space to the pattern space.

(2)h Replace the contents of the hold space by the contents of the pattern space.

(2)H Append the contents of the pattern space to the hold space.

(1)i\  
text

Insert. Place text on the standard output.

- (2)n Copy the pattern space to the standard output. Replace the pattern space with the next line of input.
- (2)N Append the next line of input to the pattern space with an embedded newline. (The current line number changes.)
- (2)p Print. Copy the pattern space to the standard output.
- (2)P Copy the initial segment of the pattern space through the first newline to the standard output.
- (1)q Quit. Branch to the end of the script. Do not start a new cycle.
- (2)r rfile  
Read the contents of rfile. Place them on the output before reading the next input line.
- (2)s/regular expression/replacement/flags  
Substitute the replacement string for instances of the regular expression in the pattern space. Any character may be used instead of `/' . For a fuller description see ed(1). Flags is zero or more of
  - g Global. Substitute for all nonoverlapping instances of the regular expression rather than just the first one.
  - p Print the pattern space if a replacement was made.
  - w wfile  
Write. Append the pattern space to wfile if a replacement was made.
- (2)t label  
Test. Branch to the `:' command bearing the label if any substitutions have been made since the most recent reading of an input line or execution of a `t'. If label is empty, branch to the end of the script.
- (2)w wfile  
Write. Append the pattern space to wfile.
- (2)x Exchange the contents of the pattern and hold spaces.
- (2)y/string1/string2/  
Transform. Replace all occurrences of characters in string1 with the corresponding character in string2. The lengths of string1 and string2 must be equal.
- (2)! function

Don't. Apply the function (or group, if function is  
{') only to lines not selected by the address(es).

(0): label

This command does nothing; it bears a label for 'b' and  
't' commands to branch to.

(1)= Place the current line number on the standard output as  
a line.

(2){ Execute the following commands through a matching '}'  
only when the pattern space is selected.

(0) An empty command is ignored.

SEE ALSO

ed(1), grep(1), awk(1), lex(1)

## NAME

sendbug - mail a system bug report to 4bsd-bugs

## SYNOPSIS

sendbug [ address ]

## DESCRIPTION

Bug reports sent to `4bsd-bugs@Berkeley.EDU' are intercepted by a program which expects bug reports to conform to a standard format. Sendbug is a shell script to help the user compose and mail bug reports in the correct format. Sendbug works by invoking the editor specified by the environment variable EDITOR on a temporary copy of the bug report format outline. The user must fill in the appropriate fields and exit the editor. The default editor is vi(1). Sendbug then mails the completed report to `4bsd-bugs@Berkeley.EDU' or the address specified on the command line.

## FILES

/usr/ucb/bugformat            contains the bug report outline

## SEE ALSO

vi(1), environ(7), sendmail(8)

## NAME

sh, for, case, if, while, :, ., break, continue, cd, eval, exec, exit, export, login, read, readonly, set, shift, times, trap, umask, wait - command language

## SYNOPSIS

sh [ -ceiknrstuvx ] [ arg ] ...

## DESCRIPTION

Sh is a command programming language that executes commands read from a terminal or a file. See invocation for the meaning of arguments to the shell.

## Commands.

A simple-command is a sequence of non blank words separated by blanks (a blank is a tab or a space). The first word specifies the name of the command to be executed. Except as specified below the remaining words are passed as arguments to the invoked command. The command name is passed as argument 0 (see `execve(2)`). The value of a simple-command is its exit status if it terminates normally or 200+status if it terminates abnormally (see `sigvec(2)` for a list of status values).

A pipeline is a sequence of one or more commands separated by `|`. The standard output of each command but the last is connected by a pipe(2) to the standard input of the next command. Each command is run as a separate process; the shell waits for the last command to terminate.

A list is a sequence of one or more pipelines separated by `;`, `&`, `&&` or `||` and optionally terminated by `;` or `&.` `;` and `&` have equal precedence which is lower than that of `&&` and `||`, `&&` and `||` also have equal precedence. A semicolon causes sequential execution; an ampersand causes the preceding pipeline to be executed without waiting for it to finish. The symbol `&&` (`||`) causes the list following to be executed only if the preceding pipeline returns a zero (non zero) value. Newlines may appear in a list, instead of semicolons, to delimit commands.

A command is either a simple-command or one of the following. The value returned by a command is that of the last simple-command executed in the command.

for name [in word ...] do list done

Each time a for command is executed name is set to the next word in the for word list. If in word ... is omitted, in "\$@" is assumed. Execution ends when there are no more words in the list.

case word in [pattern [ | pattern ] ... ) list ;;] ... esac

A case command executes the list associated with the first pattern that matches word. The form of the patterns is the same as that used for file name generation.

```
if list then list [elif list then list] ... [else list] fi
```

The list following if is executed and if it returns zero the list following then is executed. Otherwise, the list following elif is executed and if its value is zero the list following then is executed. Failing that the else list is executed.

```
while list [do list] done
```

A while command repeatedly executes the while list and if its value is zero executes the do list; otherwise the loop terminates. The value returned by a while command is that of the last executed command in the do list. until may be used in place of while to negate the loop termination test.

```
(list)
```

Execute list in a subshell.

```
{ list }
```

list is simply executed.

The following words are only recognized as the first word of a command and when not quoted.

```
if then else elif fi case in esac for while until do
done { }
```

Command substitution.

The standard output from a command enclosed in a pair of back quotes (``) may be used as part or all of a word; trailing newlines are removed.

Parameter substitution.

The character \$ is used to introduce substitutable parameters. Positional parameters may be assigned values by set. Variables may be set by writing

```
name=value [name=value] ...
```

```
${parameter}
```

A parameter is a sequence of letters, digits or underscores (a name), a digit, or any of the characters \* @ # ? - \$ !. The value, if any, of the parameter is substituted. The braces are required only when parameter is followed by a letter, digit, or underscore that is not to be interpreted as part of its name. If parameter is a digit, it is a positional parameter. If



parameter is \* or @ then all the positional parameters, starting with \$1, are substituted separated by spaces. \$0 is set from argument zero when the shell is invoked.

`${parameter-word}`

If parameter is set, substitute its value; otherwise substitute word.

`${parameter=word}`

If parameter is not set, set it to word; the value of the parameter is then substituted. Positional parameters may not be assigned to in this way.

`${parameter?word}`

If parameter is set, substitute its value; otherwise, print word and exit from the shell. If word is omitted, a standard message is printed.

`${parameter+word}`

If parameter is set, substitute word; otherwise substitute nothing.

In the above word is not evaluated unless it is to be used as the substituted string. (So that, for example, `echo ${d-'pwd'}` will only execute `pwd` if `d` is unset.)

The following parameters are automatically set by the shell.

- # The number of positional parameters in decimal.
- Options supplied to the shell on invocation or by set.
- ? The value returned by the last executed command in decimal.
- \$ The process number of this shell.
- ! The process number of the last background command invoked.

The following parameters are used but not set by the shell.

- HOME The default argument (home directory) for the `cd` command.
- PATH The search path for commands (see execution).
- MAIL If this variable is set to the name of a mail file, the shell informs the user of the arrival of mail in the specified file.
- PS1 Primary prompt string, by default '\$ '.
- PS2 Secondary prompt string, by default '> '.
- IFS Internal field separators, normally space, tab, and newline. IFS is ignored if `sh` is running as root or if the effective user id differs from the real user id.

Blank interpretation.

After parameter and command substitution, any results of substitution are scanned for internal field separator characters (those found in \$IFS) and split into distinct arguments where such characters are found. Explicit null arguments (" or ') are retained. Implicit null arguments (those resulting from parameters that have no values) are removed.

File name generation.

Following substitution, each command word is scanned for the characters \*, ? and [. If one of these characters appears, the word is regarded as a pattern. The word is replaced with alphabetically sorted file names that match the pattern. If no file name is found that matches the pattern, the word is left unchanged. The character . at the start of a file name or immediately following a /, and the character /, must be matched explicitly.

\* Matches any string, including the null string.

? Matches any single character.

[...]

Matches any one of the characters enclosed. A pair of characters separated by - matches any character lexicographically between the pair.

Quoting.

The following characters have a special meaning to the shell and cause termination of a word unless quoted.

; & ( ) | < > newline space tab

A character may be quoted by preceding it with a \. \newline is ignored. All characters enclosed between a pair of quote marks ('), except a single quote, are quoted. Inside double quotes (") parameter and command substitution occurs and \ quotes the characters \ ' " and \$.

"\$\*" is equivalent to "\$1 \$2 ..." whereas

"\$@" is equivalent to "\$1" "\$2" ... .

Prompting.

When used interactively, the shell prompts with the value of PS1 before reading a command. If at any time a newline is typed and further input is needed to complete a command, the secondary prompt (\$PS2) is issued.

Input output.

Before a command is executed its input and output may be redirected using a special notation interpreted by the shell. The following may appear anywhere in a simple-command or may precede or follow a command and are not

passed on to the invoked command. Substitution occurs before word or digit is used.

<word

Use file word as standard input (file descriptor 0).

>word

Use file word as standard output (file descriptor 1).

If the file does not exist, it is created; otherwise it is truncated to zero length.

>>word

Use file word as standard output. If the file exists, output is appended (by seeking to the end); otherwise the file is created.

<<word

The shell input is read up to a line the same as word, or end of file. The resulting document becomes the standard input. If any character of word is quoted, no interpretation is placed upon the characters of the document; otherwise, parameter and command substitution occurs, \newline is ignored, and \ is used to quote the characters \ \$ ' and the first character of word.

<&digit

The standard input is duplicated from file descriptor digit; see dup(2). Similarly for the standard output using >.

<&- The standard input is closed. Similarly for the standard output using >.

If one of the above is preceded by a digit, the file descriptor created is that specified by the digit (instead of the default 0 or 1). For example,

... 2>&1

creates file descriptor 2 to be a duplicate of file descriptor 1.

If a command is followed by & then the default standard input for the command is the empty file (/dev/null). Otherwise, the environment for the execution of a command contains the file descriptors of the invoking shell as modified by input output specifications.

Environment.

The environment is a list of name-value pairs that is passed to an executed program in the same way as a normal argument list; see execve(2) and environ(7). The shell interacts

with the environment in several ways. On invocation, the shell scans the environment and creates a parameter for each name found, giving it the corresponding value. Executed commands inherit the same environment. If the user modifies the values of these parameters or creates new ones, none of these affects the environment unless the export command is used to bind the shell's parameter to the environment. The environment seen by any executed command is thus composed of any unmodified name-value pairs originally inherited by the shell, plus any modifications or additions, all of which must be noted in export commands.

The environment for any simple-command may be augmented by prefixing it with one or more assignments to parameters. Thus these two lines are equivalent

```
TERM=450 cmd args
(export TERM; TERM=450; cmd args)
```

If the -k flag is set, all keyword arguments are placed in the environment, even if they occur after the command name. The following prints 'a=b c' and 'c':

```
echo a=b c
set -k
echo a=b c
```

#### Signals.

The INTERRUPT and QUIT signals for an invoked command are ignored if the command is followed by &; otherwise signals have the values inherited by the shell from its parent. (But see also trap.)

#### Execution.

Each time a command is executed the above substitutions are carried out. Except for the 'special commands' listed below a new process is created and an attempt is made to execute the command via an `execve(2)`.

The shell parameter `$PATH` defines the search path for the directory containing the command. Each alternative directory name is separated by a colon (:). The default path is `:/bin:/usr/bin`. If the command name contains a /, the search path is not used. Otherwise, each directory in the path is searched for an executable file. If the file has execute permission but is not an `a.out` file, it is assumed to be a file containing shell commands. A subshell (i.e., a separate process) is spawned to read it. A parenthesized command is also executed in a subshell.

#### Special commands.

The following commands are executed in the shell process and except where specified no input output redirection is

permitted for such commands.

#        For non-interactive shells, everything following the #  
is treated as a comment, i.e. the rest of the line is  
ignored. For interactive shells, the # has no special  
effect.

:

No effect; the command does nothing.

. file

Read and execute commands from file and return. The  
search path \$PATH is used to find the directory con-  
taining file.

break [n]

Exit from the enclosing for or while loop, if any. If  
n is specified, break n levels.

continue [n]

Resume the next iteration of the enclosing for or while  
loop. If n is specified, resume at the n-th enclosing  
loop.

cd [arg]

Change the current directory to arg. The shell param-  
eter \$HOME is the default arg.

eval [arg ...]

The arguments are read as input to the shell and the  
resulting command(s) executed.

exec [arg ...]

The command specified by the arguments is executed in  
place of this shell without creating a new process.  
Input output arguments may appear and if no other argu-  
ments are given cause the shell input output to be  
modified.

exit [n]

Causes a non interactive shell to exit with the exit  
status specified by n. If n is omitted, the exit status  
is that of the last command executed. (An end of file  
will also exit from the shell.)

export [name ...]

The given names are marked for automatic export to the  
environment of subsequently-executed commands. If no  
arguments are given, a list of exportable names is  
printed.

login [arg ...]

Equivalent to 'exec login arg ...'.

read name ...

One line is read from the standard input; successive  
words of the input are assigned to the variables name  
in order, with leftover words to the last variable.  
The return code is 0 unless the end-of-file is encoun-  
tered.

readonly [name ...]

The given names are marked readonly and the values of  
the these names may not be changed by subsequent

assignment. If no arguments are given, a list of all readonly names is printed.

`set [-eknptuvx [arg ...]]`

- e If non interactive, exit immediately if a command fails.
- k All keyword arguments are placed in the environment for a command, not just those that precede the command name.
- n Read commands but do not execute them.
- t Exit after reading and executing one command.
- u Treat unset variables as an error when substituting.
- v Print shell input lines as they are read.
- x Print commands and their arguments as they are executed.
- Turn off the -x and -v options.

These flags can also be used upon invocation of the shell. The current set of flags may be found in `$-`.

Remaining arguments are positional parameters and are assigned, in order, to `$1`, `$2`, etc. If no arguments are given, the values of all names are printed.

`shift`

The positional parameters from `$2...` are renamed `$1...`

`times`

Print the accumulated user and system times for processes run from the shell.

`trap [arg] [n] ...`

Arg is a command to be read and executed when the shell receives signal(s) `n`. (Note that `arg` is scanned once when the trap is set and once when the trap is taken.) Trap commands are executed in order of signal number. If `arg` is absent, all `trap(s) n` are reset to their original values. If `arg` is the null string, this signal is ignored by the shell and by invoked commands. If `n` is 0, the command `arg` is executed on exit from the shell, otherwise upon receipt of signal `n` as numbered in `sigvec(2)`. Trap with no arguments prints a list of commands associated with each signal number.

`umask [ nnn ]`

The user file creation mask is set to the octal value `nnn` (see `umask(2)`). If `nnn` is omitted, the current value of the mask is printed.

`wait [n]`

Wait for the specified process and report its termination status. If `n` is not given, all currently active child processes are waited for. The return code from

this command is that of the process waited for.

#### Invocation.

If the first character of argument zero is `-`, commands are read from `$HOME/.profile`, if such a file exists. Commands are then read as described below. The following flags are interpreted by the shell when it is invoked.

- `-c string` If the `-c` flag is present, commands are read from `string`.
- `-s` If the `-s` flag is present or if no arguments remain then commands are read from the standard input. Shell output is written to file descriptor 2.
- `-i` If the `-i` flag is present or if the shell input and output are attached to a terminal (as told by `gtty`) then this shell is interactive. In this case the terminate signal `SIGTERM` (see `sigvec(2)`) is ignored (so that `'kill 0'` does not kill an interactive shell) and the interrupt signal `SIGINT` is caught and ignored (so that `wait` is interruptible). In all cases `SIGQUIT` is ignored by the shell.

The remaining flags and arguments are described under the `set` command.

#### FILES

`$HOME/.profile`  
`/tmp/sh*`  
`/dev/null`

#### SEE ALSO

`csh(1)`, `test(1)`, `execve(2)`, `environ(7)`

#### DIAGNOSTICS

Errors detected by the shell, such as syntax errors cause the shell to return a non zero exit status. If the shell is being used non interactively then execution of the shell file is abandoned. Otherwise, the shell returns the exit status of the last command executed (see also `exit`).

#### BUGS

If `<<` is used to provide standard input to an asynchronous process invoked by `&`, the shell gets mixed up about naming the input document. A garbage file `/tmp/sh*` is created, and the shell complains about not being able to find the file by another name.

## NAME

size - size of an object file

## SYNOPSIS

size [ object ... ]

## DESCRIPTION

Size prints the (decimal) number of bytes required by the text, data, and bss portions, and their sum in hex and decimal, of each object-file argument. If no file is specified, a.out is used.

## SEE ALSO

a.out(5)



## NAME

sleep - suspend execution for an interval

## SYNOPSIS

sleep time

## DESCRIPTION

Sleep suspends execution for time seconds. It is used to execute a command after a certain amount of time as in:

```
(sleep 105; command)&
```

or to execute a command every so often, as in:

```
while true
do
 command
 sleep 37
done
```

## SEE ALSO

setitimer(2), alarm(3C), sleep(3)

## BUGS

Time must be less than 2,147,483,647 seconds.

## NAME

soelim - eliminate .so's from nroff input

## SYNOPSIS

soelim [ file ... ]

## DESCRIPTION

Soelim reads the specified files or the standard input and performs the textual inclusion implied by the nroff directives of the form

```
.so somefile
```

when they appear at the beginning of input lines. This is useful since programs such as tbl do not normally do this; it allows the placement of individual tables in separate files to be run as a part of a large document.

An argument consisting of a single minus (-) is taken to be a file name corresponding to the standard input.

Note that inclusion can be suppressed by using `''` instead of `.'`, i.e.

```
'so /usr/share/tmac/tmac.s
```

A sample usage of soelim would be

```
soelim exum?.n | tbl | nroff -ms | col | lpr
```

## SEE ALSO

colcrt(1), more(1)

## BUGS

The format of the source commands must involve no strangeness - exactly one blank must precede and no blanks follow the file name.

## NAME

sort - sort or merge files

## SYNOPSIS

```
sort [-mubdfinrt_____x] [+pos1 [-pos2]] ... [-o name] [
-T directory] [name] ...
```

## DESCRIPTION

Sort sorts lines of all the named files together and writes the result on the standard output. The name '-' means the standard input. If no input files are named, the standard input is sorted.

The default sort key is an entire line. Default ordering is lexicographic by bytes in machine collating sequence. The ordering is affected globally by the following options, one or more of which may appear.

- b        Ignore leading blanks (spaces and tabs) in field comparisons.
- d        'Dictionary' order: only letters, digits and blanks are significant in comparisons.
- f        Fold upper case letters onto lower case.
- i        Ignore characters outside the ASCII range 040-0176 in nonnumeric comparisons.
- n        An initial numeric string, consisting of optional blanks, optional minus sign, and zero or more digits with optional decimal point, is sorted by arithmetic value. Option n implies option b.
- r        Reverse the sense of comparisons.
- tx      'Tab character' separating fields is x.

The notation +pos1 -pos2 restricts a sort key to a field beginning at pos1 and ending just before pos2. Pos1 and pos2 each have the form m.n, optionally followed by one or more of the flags bdfinr, where m tells a number of fields to skip from the beginning of the line and n tells a number of characters to skip further. If any flags are present they override all the global ordering options for this key. If the b option is in effect n is counted from the first nonblank in the field; b is attached independently to pos2. A missing .n means .0; a missing -pos2 means the end of the line. Under the -tx option, fields are strings separated by x; otherwise fields are nonempty nonblank strings separated by blanks.

When there are multiple sort keys, later keys are compared only after all earlier keys compare equal. Lines that otherwise compare equal are ordered with all bytes significant.

These option arguments are also understood:

- c        Check that the input file is sorted according to the ordering rules; give no output unless the file is out of sort.
- m        Merge only, the input files are already sorted.
- o        The next argument is the name of an output file to use instead of the standard output. This file may be the same as one of the inputs.
- T        The next argument is the name of a directory in which temporary files should be made.
- u        Suppress all but one in each set of equal lines. Ignored bytes and bytes outside keys do not participate in this comparison.

#### EXAMPLES

Print in alphabetical order all the unique spellings in a list of words. Capitalized words differ from uncapitalized.

```
sort -u +0f +0 list
```

Print the password file (passwd(5)) sorted by user id number (the 3rd colon-separated field).

```
sort -t: +2n /etc/passwd
```

Print the first instance of each month in an already sorted file of (month day) entries. The options -um with just one input file make the choice of a unique representative from a set of equal lines predictable.

```
sort -um +0 -1 dates
```

#### FILES

/usr/tmp/stm\*, /tmp/\*     first and second tries for temporary files

#### SEE ALSO

uniq(1), comm(1), rev(1), join(1)

#### DIAGNOSTICS

Comments and exits with nonzero status for various trouble conditions and for disorder discovered under option -c.

## BUGS

Very long lines are silently truncated.

## NAME

sortbib - sort bibliographic database

## SYNOPSIS

sortbib [ -sKEYS ] database ...

## DESCRIPTION

Sortbib sorts files of records containing refer key-letters by user-specified keys. Records may be separated by blank lines, or by `.[` and `.]` delimiters, but the two styles may not be mixed together. This program reads through each database and pulls out key fields, which are sorted separately. The sorted key fields contain the file pointer, byte offset, and length of corresponding records. These records are delivered using disk seeks and reads, so sortbib may not be used in a pipeline to read standard input.

By default, sortbib alphabetizes by the first `%A` and the `%D` fields, which contain the senior author and date. The `-s` option is used to specify new KEYS. For instance, `-sATD` will sort by author, title, and date, while `-sA+D` will sort by all authors, and date. Sort keys past the fourth are not meaningful. No more than 16 databases may be sorted together at one time. Records longer than 4096 characters will be truncated.

Sortbib sorts on the last word on the `%A` line, which is assumed to be the author's last name. A word in the final position, such as ``jr.'` or ``ed.'`, will be ignored if the name beforehand ends with a comma. Authors with two-word last names or unusual constructions can be sorted correctly by using the nroff convention ``\0'` in place of a blank. A `%Q` field is considered to be the same as `%A`, except sorting begins with the first, not the last, word. Sortbib sorts on the last word of the `%D` line, usually the year. It also ignores leading articles (like ``A'` or ``The'`) when sorting by titles in the `%T` or `%J` fields; it will ignore articles of any modern European language. If a sort-significant field is absent from a record, sortbib places that record before other records containing that field.

## SEE ALSO

refer(1), addbib(1), roffb(1), indxbib(1), lookbib(1)

## AUTHORS

Greg Shenaut, Bill Tuthill

## BUGS

Records with missing author fields should probably be sorted by title.

## NAME

spell, spellin, spellout - find spelling errors

## SYNOPSIS

spell [ -v ] [ -b ] [ -x ] [ -d hlist ] [ -s hstop ] [ -h  
spellhist ] [ file ] ...

spellin [ list ]

spellout [ -d ] list

## DESCRIPTION

Spell collects words from the named documents, and looks them up in a spelling list. Words that neither occur among nor are derivable (by applying certain inflections, prefixes or suffixes) from words in the spelling list are printed on the standard output. If no files are named, words are collected from the standard input.

Spell ignores most troff, tbl and eqn(1) constructions.

Under the -v option, all words not literally in the spelling list are printed, and plausible derivations from spelling list words are indicated.

Under the -b option, British spelling is checked. Besides preferring centre, colour, speciality, travelled, etc., this option insists upon -ise in words like standardise, Fowler and the OED to the contrary notwithstanding.

Under the -x option, every plausible stem is printed with '=' for each word.

The spelling list is based on many sources. While it is more haphazard than an ordinary dictionary, it is also more effective with proper names and popular technical words. Coverage of the specialized vocabularies of biology, medicine and chemistry is light.

The auxiliary files used for the spelling list, stop list, and history file may be specified by arguments following the -d, -s, and -h options. The default files are indicated below. Copies of all output may be accumulated in the history file. The stop list filters out misspellings (e.g. thier=thy-y+ier) that would otherwise pass.

Two routines help maintain the hash lists used by spell. Both expect a set of words, one per line, from the standard input. Spellin combines the words from the standard input and the preexisting list file and places a new list on the standard output. If no list file is specified, the new list is created from scratch. Spellout looks up each word from

the standard input and prints on the standard output those that are missing from (or present on, with option -d) the hashed list file. For example, to verify that hookey is not on the default spelling list, add it to your own private list, and then use it with spell,

```
echo hookey | spellout /usr/dict/hlista
echo hookey | spellin /usr/dict/hlista > myhlist
spell -d myhlist huckfinn
```

#### FILES

/usr/dict/hlist[ab] hashed spelling lists, American & British, default for -d  
/usr/dict/hstop hashed stop list, default for -s  
/dev/null history file, default for -h  
/tmp/spell.\$\$\* temporary files  
/usr/libexec/spell

#### SEE ALSO

deroff(1), sort(1), tee(1), sed(1)

#### BUGS

The spelling list's coverage is uneven; new installations will probably wish to monitor the output for several months to gather local additions.  
British spelling was done by an American.



## NAME

spline - interpolate smooth curve

## SYNOPSIS

spline [ option ] ...

## DESCRIPTION

Spline takes pairs of numbers from the standard input as abscissas and ordinates of a function. It produces a similar set, which is approximately equally spaced and includes the input set, on the standard output. The cubic spline output (R. W. Hamming, Numerical Methods for Scientists and Engineers, 2nd ed., 349ff) has two continuous derivatives, and sufficiently many points to look smooth when plotted, for example by graph(1G).

The following options are recognized, each as a separate argument.

-a Supply abscissas automatically (they are missing from the input); spacing is given by the next argument, or is assumed to be 1 if next argument is not a number.

-k The constant k used in the boundary value computation

$(2\text{nd deriv. at end}) = k \cdot (2\text{nd deriv. next to end})$

is set by the next argument. By default  $k = 0$ .

-n Space output points so that approximately n intervals occur between the lower and upper x limits. (Default n = 100.)

-p Make output periodic, i.e. match derivatives at ends. First and last input values should normally agree.

-x Next 1 (or 2) arguments are lower (and upper) x limits. Normally these limits are calculated from the data. Automatic abscissas start at lower limit (default 0).

## SEE ALSO

graph(1G), plot(1G)

## DIAGNOSTICS

When data is not strictly monotone in x, spline reproduces the input without interpolating extra points.

## BUGS

A limit of 1000 input points is enforced silently.

## NAME

split - split a file into pieces

## SYNOPSIS

split [ -n ] [ file [ name ] ]

## DESCRIPTION

Split reads file and writes it in n-line pieces (default 1000), as many as necessary, onto a set of output files. The name of the first output file is name with aa appended, and so on lexicographically. If no output name is given, x is default.

If no input file is given, or if - is given in its stead, then the standard input file is used.

## NAME

strcompact - string compaction for object files

## SYNOPSIS

strcompact [ object\_name ... ]

## DESCRIPTION

strcompact scans the symbol and string tables of an object file looking for multiple references in the symbol table to the same string. The string offset of symbol table entries is updated to preserve only one copy of the string.

strcompact cut the size of the kernel string table by about 25%.

The user must have write permission to the object/executable file.

strcompact writes to stderr the number of shared strings found.

strcompact exits 0 if successful, and >0 if an error occurred.

## SEE ALSO

sort(1), symcompact(1), symorder(1), uniq(1)

## BUGS

Execution speed leaves much to be desired - on a 11/73 it takes about 4 minutes to process the string table of the kernel. Fortunately this is only done once when the kernel is created.

Although strcompact may be run on .o files as well as executables but this is probably not worth the trouble since the linker will not create shared strings in the final executable.

## NAME

strings - find the printable strings in a object, or other binary, file

## SYNOPSIS

strings [ - ] [ -o ] [ -number ] file ...

## DESCRIPTION

Strings looks for ascii strings in a binary file. A string is any sequence of 4 or more printing characters ending with a newline or a null. Unless the - flag is given, strings only looks in the initialized data space of object files. If the -o flag is given, then each string is preceded by its offset in the file (in octal). If the -number flag is given then number is used as the minimum string length rather than 4.

Strings is useful for identifying random object files and many other things.

## SEE ALSO

od(1)

## BUGS

The algorithm for identifying strings is extremely primitive.

## NAME

strip - remove symbols and relocation bits

## SYNOPSIS

strip name ...

## DESCRIPTION

Strip removes the symbol table and relocation bits ordinarily attached to the output of the assembler and loader. This is useful to save space after a program has been debugged.

The effect of strip is the same as use of the -s option of ld.

## FILES

/tmp/stm? temporary file

## SEE ALSO

ld(1)

## NAME

struct - structure Fortran programs

## SYNOPSIS

struct [ option ] ... file

## DESCRIPTION

Struct translates the Fortran program specified by file (standard input default) into a Ratfor program. Wherever possible, Ratfor control constructs replace the original Fortran. Statement numbers appear only where still necessary. Cosmetic changes are made, including changing Hollerith strings into quoted strings and relational operators into symbols (e.g. ".GT." into ">"). The output is appropriately indented.

The following options may occur in any order.

- s Input is accepted in standard format, i.e. comments are specified by a c, C, or \* in column 1, and continuation lines are specified by a nonzero, nonblank character in column 6. Normally input is in the form accepted by f77(1)
- i Do not turn computed goto statements into switches. (Ratfor does not turn switches back into computed goto statements.)
- a Turn sequences of else ifs into a non-Ratfor switch of the form

```
switch
{ case pred1: code
 case pred2: code
 case pred3: code
 default: code
}
```

The case predicates are tested in order; the code appropriate to only one case is executed. This generalized form of switch statement does not occur in Ratfor.
- b Generate goto's instead of multilevel break statements.
- n Generate goto's instead of multilevel next statements.
- tn Make the nonzero integer n the lowest valued label in the output program (default 10).
- cn Increment successive labels in the output program by

the nonzero integer n (default 1).

-en If n is 0 (default), place code within a loop only if it can lead to an iteration of the loop. If n is nonzero, admit a small code segments to a loop if otherwise the loop would have exits to several places including the segment, and the segment can be reached only from the loop. 'Small' is close to, but not equal to, the number of statements in the code segment. Values of n under 10 are suggested.

#### FILES

/tmp/struct\*  
/usr/libexec/struct/\*

#### SEE ALSO

f77(1)

#### BUGS

Struct knows Fortran 66 syntax, but not full Fortran 77. If an input Fortran program contains identifiers which are reserved words in Ratfor, the structured version of the program will not be a valid Ratfor program. The labels generated cannot go above 32767. If you get a goto without a target, try -e .

## NAME

stty - set terminal options

## SYNOPSIS

stty [-a | -e] [-f file] [operands]

## DESCRIPTION

Stty sets certain I/O options on the current output terminal, placing its output on the diagnostic output. With no argument, it reports the speed of the terminal and the settings of the options which are different from their defaults.

The following options are available:

- a        Display everything stty knows. This has the same effect as using the operand all or everything. The distinction between all and everything has been removed.
- e        Same as -a above.
- f        Open and use the terminal named by file rather than using standard output. The file is opened using the O\_NONBLOCK flag of open(), making it possible to set or display settings on a terminal that might otherwise block on the open.

The following operands are special:

- all        Everything stty knows about is printed.
- everything    Same as all above.
- flushout     Flush the queues for the device. This is most useful when an exiting process is stuck waiting for terminal output to drain.
- speed        The terminal speed alone is printed on the standard output.
- size        The terminal (window) sizes are printed on the standard output, first rows and then columns.

Operands are selected from the following:

- even        allow even parity input
- even       disallow even parity input
- odd         allow odd parity input
- odd        disallow odd parity input



```

raw raw mode input (no input processing (erase, kill,
 interrupt, ...); parity bit passed back)
-raw negate raw mode

cooked same as '-raw'

cbreak make each character available to read(2) as
 received; no erase and kill processing, but all
 other processing (interrupt, suspend, ...) is per-
 formed
-cbreak make characters available to read only when new-
 line is received

-nl allow carriage return for new-line, and output
 CR-LF for carriage return or new-line
nl accept only new-line to end lines

echo echo back every character typed
-echo do not echo characters

tandem enable inbound software (xon/xoff) flow control,
 so that the system sends out the stop character
 when its internal queue is in danger of overflow-
 ing on input, and sends the start character when
 it is ready to accept further input
-tandem disable inbound software (xon/xoff) flow control

-tabs replace tabs by spaces when printing
tabs preserve tabs

```

For the following commands which take a character argument *c*, you may also specify *c* as ``undef'', to set the value to be undefined. A value of ``^x'', a 2 character sequence, is also interpreted as a control character, with ``^?' representing delete.

```

erase c set erase character to c (default `#', but often
 reset to ^H.)
kill c set kill character to c (default `@', but often
 reset to ^U.)
intr c set interrupt character to c (default DEL or ^?
 (delete), but often reset to ^C.)
quit c set quit character to c (default control \.)
start c set start character to c (default control Q.)
stop c set stop character to c (default control S.)
eof c set end of file character to c (default control
 D.)
brk c set break character to c (default undefined.) This
 character is an additional character causing
 wakeup.
dec set all modes suitable for Digital Equipment Corp.
 operating systems users; (erase, kill, and

```

interrupt characters to ^?, ^U, and ^C, decctlq and ``newcrt''.)

0 hang up phone line immediately  
 50 75 110 134 150 200 300 600 1200 1800 2400 4800 9600 exta extb  
 Set terminal baud rate to the number given, if possible. (These are the speeds supported by the DH-11 interface).

rows n The terminal size is recorded as having n rows.

columns n The terminal size is recorded as having n columns.

cols n is an alias for columns.

A teletype driver which supports the job control processing of csh(1) and more functionality than the basic driver is fully described in tty(4). The following options apply only to it.

new Use new driver (switching flushes typeahead).  
 crt Set options for a CRT (crtbs, ctlecho and, if >= 1200 baud, crterase and crtkill.)

crtbs Echo backspaces on erase characters.

prterase For printing terminal echo erased characters backwards within ``\'' and ``/''.  
 -prterase Leave erased characters visible; just backspace.

crterase Wipe out erased characters with ``backspace-space-backspace.''  
 -crterase Leave erased characters visible; just backspace.

crtkill Wipe out input on like kill ala crterase.  
 -crtkill Just echo line kill character and a newline on line kill.

ctlecho Echo control characters as ``^x'' (and delete as ``^?''.) Print two backspaces following the EOT character (control D).  
 -ctlecho Control characters echo as themselves; in cooked mode EOT (control-D) is not echoed.

decctlq After output is suspended (normally by ^S), only a start character (normally ^Q) will restart it. This is compatible with DEC's vendor supplied systems.  
 -decctlq After output is suspended, any character typed will restart it; the start character will restart output without providing any input. (This is the default.)

`tostop` Background jobs stop if they attempt terminal output.  
`-tostop` Output from background jobs to the terminal is allowed.

`flusho` Output is being discarded usually because user hit control O (internal state bit).  
`-flusho` Output is not being discarded.

`pendin` Input is pending after a switch from cbreak to cooked and will be re-input when a read becomes pending or more input arrives (internal state bit).  
`-pendin` Input is not pending.

`pass8` Passes all 8 bits through on input, in any mode.  
`-pass8` Strips the 0200 bit on input except in raw mode.

`mdmbuf` Start/stop output on carrier transitions (not implemented).  
`-mdmbuf` Return error if write attempted after carrier drops.

`litout` Send output characters without any processing.  
`-litout` Do normal output processing, inserting delays, etc.

`nohang` Don't send hangup signal if carrier drops.  
`-nohang` Send hangup signal to control process group when carrier drops.

The following special characters are applicable only to the new teletype driver and are not normally changed.

`susp c` set suspend process character to c (default control Z).  
`dsusp c` set delayed suspend process character to c (default control Y).  
`rprnt c` set reprint line character to c (default control R).  
`flush c` set flush output character to c (default control O).  
`werase c` set word erase character to c (default control W).  
`lnext c` set literal next character to c (default control V).

#### Modem Control Status:

These display the current state of modem control. They are only displayed for actual tty lines and not for pseudo tty lines (more precisely, it is only displayed for lines which support the TIOCMGET ioctl. See `tty(4)`).

While it is possible to change the state of the modem control lines, the hardware or other software may prevent the change from actually taking place, or may cause the state to immediately revert to the original state.

dcd (-dcd)        State of Data Carrier Detect.

dsr (-dsr)        State of Data Set Ready.

dtr (-dtr)        State of Data Terminal Ready.

cts (-cts)        State of Clear To Send.

rts (-rts)        State of Request To Send.

SEE ALSO

ioctl(2), tabs(1), tset(1), tty(4)

## NAME

style - analyze surface characteristics of a document

## SYNOPSIS

style [ -ml ] [ -mm ] [ -a ] [ -e ] [ -l num ] [ -r num ] [ -p ] [ -P ] file ...

## DESCRIPTION

Style analyzes the surface characteristics of the writing style of a document. It reports on readability, sentence length and structure, word length and usage, verb type, and sentence openers. Because style runs deroff before looking at the text, formatting header files should be included as part of the input. The default macro package -ms may be overridden with the flag -mm. The flag -ml, which causes deroff to skip lists, should be used if the document contains many lists of non-sentences. The other options are used to locate sentences with certain characteristics.

-a print all sentences with their length and readability index.

-e print all sentences that begin with an expletive.

-p print all sentences that contain a passive verb.

-l num  
print all sentences longer than num.

-r num  
print all sentences whose readability index is greater than num.

-P print parts of speech of the words in the document.

## SEE ALSO

deroff(1), diction(1)

## BUGS

Use of non-standard formatting macros may cause incorrect sentence breaks.

## NAME

su - substitute user id temporarily

## SYNOPSIS

su [ -f ] [ - ] [ userid ]

## DESCRIPTION

Su demands the password of the specified userid, and if it is given, changes to that userid and invokes the Shell sh(1) or csh(1) without changing the current directory. The user environment is unchanged except for HOME and SHELL, which are taken from the password file for the user being substituted (see environ(7)). The new user ID stays in force until the Shell exits.

If no userid is specified, ``root'' is assumed. Only users in the ``wheel'' group (group 0) can su to ``root'', even with the root password. To remind the super-user of his responsibilities, the Shell substitutes `#' for its usual prompt.

The -f option prevents csh(1) from executing the .cshrc file; thus making su start up faster.

The - option simulates a full login.

## SEE ALSO

sh(1), csh(1)

## NAME

sum - sum and count blocks in a file

## SYNOPSIS

sum file

## DESCRIPTION

Sum calculates and prints a 16-bit checksum for the named file, and also prints the number of blocks in the file. It is typically used to look for bad spots, or to validate a file communicated over some transmission line.

## SEE ALSO

wc(1)

## DIAGNOSTICS

`Read error' is indistinguishable from end of file on most devices; check the block count.

## NAME

symcompact - string compaction for object files

## SYNOPSIS

symcompact [ object\_name ... ]

## DESCRIPTION

symcompact reduces the symbol table size of an executable file. This is done by removing unnecessary overlay transfer vectors (text symbols beginning with a tilde). In a nonoverlaid program there is no need for both the underscore (`_foo`) and tilde (`~foo`) text symbol and only the underscore form is kept. For overlaid programs if the symbol is in the base segment the tilde form is not needed and again only the underscore form is preserved. Running symcompact typically reduces the kernel symbol table size by 250 or so symbols.

It is possible to run both symcompact and strcompact to achieve an even higher degree of symbol and string table compaction. The normal sequence is to run symcompact first followed by strcompact. If symcompact runs out of memory it will be necessary to reverse the order and run symcompact a second time - see the BUGS note below.

The user must have write permission to the object/executable file.

symcompact writes to stderr the count of symbols removed from the symbol table.

symcompact exits 0 if successful, and >0 if an error occurred.

## SEE ALSO

symcompact(1), symorder(1)

## BUGS

This program can partially negate the benefits of strcompact because multiple references to identical strings cause additional strings to be placed in the string table. Running strcompact again after running this program fixes this problem.

The register local symbol type is removed from the executable/object file. Since the debugger really doesn't know how to deal with those symbols this is not much of a loss and saves quite a bit of space both in the symbol table and the string table.

symcompact should not be run on .o files that will be passed to the linker. The linker will need the tilde form of the symbol if an overlaid executable is being created.



## NAME

symorder - rearrange name list

## SYNOPSIS

symorder symlist file

## DESCRIPTION

The file symlist contains a list of symbols to be found in file, one symbol per line.

The symbol table of file is updated in place; symbols read from symlist are relocated to the beginning of the table and in the order given.

This program was specifically designed to cut down on the overhead of getting symbols from the kernel name list.

## DIAGNOSTICS

The symorder(1) utility exits 0 on success, non zero if an error occurs.

## SEE ALSO

nm(1), nlist(3), strip(1)

## HISTORY

The symorder command appeared in 3.0BSD.

## NAME

tabs - set terminal tabs

## SYNOPSIS

tabs [ -n ] [ terminal ]

## DESCRIPTION

Tabs sets the tabs on a variety of terminals. Various terminal names given in term(7) are recognized; the default is, however, suitable for most 300 baud terminals. If the -n flag is present then the left margin is not indented as is normal.

## SEE ALSO

stty(1), term(7)

## BUGS

It's much better to use tset(1).

## NAME

tail - deliver the last part of a file

## SYNOPSIS

tail +number[lbc][rf] [ file ]

## DESCRIPTION

Tail copies the named file to the standard output beginning at a designated place. If no file is named, the standard input is used.

Copying begins at distance +number from the beginning, or -number from the end of the input. Number is counted in units of lines, blocks or characters, according to the appended option l, b or c. When no units are specified, counting is by lines.

Specifying r causes tail to print lines from the end of the file in reverse order. The default for r is to print the entire file this way. Specifying f causes tail to not quit at end of file, but rather wait and try to read repeatedly in hopes that the file will grow.

## SEE ALSO

dd(1)

## BUGS

Tails relative to the end of the file are treasured up in a buffer, and thus are limited in length.

Various kinds of anomalous behavior may happen with character special files.

## NAME

talk - talk to another user

## SYNOPSIS

talk person [ ttyname ]

## DESCRIPTION

Talk is a visual communication program which copies lines from your terminal to that of another user.

If you wish to talk to someone on you own machine, then person is just the person's login name. If you wish to talk to a user on another host, then person is of the form :

```
host!user or
host.user or
host:user or
user@host
```

though host@user is perhaps preferred.

If you want to talk to a user who is logged in more than once, the ttyname argument may be used to indicate the appropriate terminal name.

When first called, it sends the message

```
Message from TalkDaemon@his_machine...
talk: connection requested by your_name@your_machine.
talk: respond with: talk your_name@your_machine
```

to the user you wish to talk to. At this point, the recipient of the message should reply by typing

```
talk your_name@your_machine
```

It doesn't matter from which machine the recipient replies, as long as his login-name is the same. Once communication is established, the two parties may type simultaneously, with their output appearing in separate windows. Typing control L will cause the screen to be reprinted, while your erase, kill, and word kill characters will work in talk as normal. To exit, just type your interrupt character; talk then moves the cursor to the bottom of the screen and restores the terminal.

Permission to talk may be denied or granted by use of the mesg command. At the outset talking is allowed. Certain commands, in particular nroff and pr(1) disallow messages in order to prevent messy output.

## FILES

/etc/hosts       to find the recipient's machine  
/var/run/utmp    to find the recipient's tty

## SEE ALSO

mesg(1), who(1), mail(1), write(1)

## BUGS

The version of talk(1) released with 4.3BSD uses a protocol that is incompatible with the protocol used in the version released with 4.2BSD.

## NAME

tar - tape archiver

## SYNOPSIS

tar [ key ] [ name ... ]

## DESCRIPTION

Tar saves and restores multiple files on a single file (usually a magnetic tape, but it can be any file). Tar's actions are controlled by the key argument. The key is a string of characters containing at most one function letter and possibly one or more function modifiers. Other arguments to tar are file or directory names specifying which files to dump or restore. In all cases, appearance of a directory name refers to the files and (recursively) subdirectories of that directory.

The function portion of the key is specified by one of the following letters:

- r           The named files are written on the end of the tape. The c function implies this.
- x           The named files are extracted from the tape. If the named file matches a directory whose contents had been written onto the tape, this directory is (recursively) extracted. The owner, modification time, and mode are restored (if possible). If no file argument is given, the entire content of the tape is extracted. Note that if multiple entries specifying the same file are on the tape, the last one overwrites all earlier.
- t           The names of the specified files are listed each time they occur on the tape. If no file argument is given, all of the names on the tape are listed.
- u           The named files are added to the tape if either they are not already there or have been modified since last put on the tape.
- c           Create a new tape; writing begins on the beginning of the tape instead of after the last file. This command implies r.

The following characters may be used in addition to the letter which selects the function desired.

- o           On output, tar normally places information specifying owner and modes of directories in the archive. Former versions of tar, when encountering this information will give error message of

the form

"<name>/: cannot create".

This modifier will suppress the directory information.

- p            This modifier says to restore files to their original modes, ignoring the present umask(2). Setuid and sticky information will also be restored to the super-user.
- 0, ..., 9   This modifier selects an alternate drive on which the tape is mounted. The default is drive 0 at 1600 bpi, which is normally /dev/rmt8.
- v            Normally tar does its work silently. The v (verbose) option makes tar print the name of each file it treats preceded by the function letter. With the t function, the verbose option gives more information about the tape entries than just their names.
- w            Tar prints the action to be taken followed by file name, then wait for user confirmation. If a word beginning with 'y' is given, the action is done. Any other input means don't do it.
- f            Tar uses the next argument as the name of the archive instead of /dev/rmt?. If the name of the file is '-', tar writes to standard output or reads from standard input, whichever is appropriate. Thus, tar can be used as the head or tail of a filter chain. Tar can also be used to move hierarchies with the command  
cd fromdir; tar cf - . | (cd todir; tar xf -)
- b            Tar uses the next argument as the blocking factor for tape records. The default is 20 (the maximum). This option should only be used with raw magnetic tape archives (See f above). The block size is determined automatically when reading tapes (key letters 'x' and 't').
- l            tells tar to complain if it cannot resolve all of the links to the files dumped. If this is not specified, no error messages are printed.
- m            tells tar not to restore the modification times. The modification time will be the time of extraction.
- h            Force tar to follow symbolic links as if they were normal files or directories. Normally, tar does

not follow symbolic links.

B            Forces input and output blocking to 20 blocks per record. This option was added so that tar can work across a communications channel where the blocking may not be maintained.

C            If a file name is preceded by -C, then tar will perform a chdir(2) to that file name. This allows multiple directories not related by a close common parent to be archived using short relative path names. For example, to archive files from /usr/include and from /etc, one might use  
tar c -C /usr include -C / etc

Previous restrictions dealing with tar's inability to properly handle blocked archives have been lifted.

#### FILES

/dev/rmt?  
/tmp/tar\*

#### SEE ALSO

tar(5)

#### DIAGNOSTICS

Complaints about bad key characters and tape read/write errors.  
Complaints if enough memory is not available to hold the link tables.

#### BUGS

There is no way to ask for the n-th occurrence of a file.  
Tape errors are handled ungracefully.  
The u option can be slow.  
The current limit on file name length is 100 characters.  
There is no way selectively to follow symbolic links.  
When extracting tapes created with the r or u options, directory modification times may not be set correctly.



## NAME

tbl - format tables for nroff or troff

## SYNOPSIS

tbl [ files ] ...

## DESCRIPTION

Tbl is a preprocessor for formatting tables for nroff or troff(1). The input files are copied to the standard output, except for lines between and are reformatted. Details are given in the tbl(1) reference manual.

## EXAMPLE

As an example, letting \t represent a tab (which should be typed as a genuine tab) the input

```
.TS
c s s
c c s
c c c
l n n.
Household Population
Town\tHouseholds
\tNumber\tSize
Bedminster\t789\t3.26
Bernards Twp.\t3087\t3.74
Bernardsville\t2018\t3.30
Bound Brook\t3425\t3.04
Branchburg\t1644\t3.49
Bridgewater\t7897\t3.81
Far Hills\t240\t3.19
.TE
```

yields

Household Population			
Town	Households		
	Number	Size	
Bedminster	789	3.26	
Bernards Twp.	3087	3.74	
Bernardsville	2018	3.30	
Bound Brook	3425	3.04	
Branchburg	1644	3.49	
Bridgewater	7897	3.81	
Far Hills	240	3.19	

If no arguments are given, tbl reads the standard input, so it may be used as a filter. When tbl is used with eqn or neqn the tbl command should be first, to minimize the volume of data passed through pipes.

SEE ALSO

troff(1), eqn(1)

M. E. Lesk, TBL.

## NAME

tc - phototypesetter simulator

## SYNOPSIS

tc [ -t ] [ -sN ] [ -pL ] [ file ]

## DESCRIPTION

Tc interprets its input (standard input default) as device codes for a Graphic Systems phototypesetter (cat). The standard output of tc is intended for a Tektronix 4015 (a 4014 terminal with ASCII and APL character sets). The sixteen typesetter sizes are mapped into the 4014's four sizes; the entire TROFF character set is drawn using the 4014's character generator, using overstruck combinations where necessary. Typical usage:

```
troff -t file | tc
```

At the end of each page tc waits for a newline (empty line) from the keyboard before continuing on to the next page. In this wait state, the command e will suppress the screen erase before the next page; sN will cause the next N pages to be skipped; and !line will send line to the shell.

The command line options are:

-t Don't wait between pages; for directing output into a file.

-sN Skip the first N pages.

-pL Set page length to L. L may include the scale factors p (points), i (inches), c (centimeters), and P (picas); default is picas.

'-l w'  
Multiply the default aspect ratio, 1.5, of a displayed page by l/w.

## SEE ALSO

troff(1), plot(1G)

## BUGS

Font distinctions are lost.  
tc's character set is limited to ASCII in just one size.  
The aspect ratio option is unbelievable.

## NAME

tcopy - copy a mag tape

## SYNOPSIS

tcopy src [ dest ]

## DESCRIPTION

Tcopy is designed to copy magnetic tapes. The only assumption made about the tape is that there are two tape marks at the end. Tcopy with only a source tape specified will print information about the sizes of records and tape files. If a destination is specified, then, a copy will be made of the source tape. The blocking on the destination tape will be identical to that used on the source tape. Copying a tape will yield the same output as if just printing the sizes.

## SEE ALSO

mtio(4)

## NAME

tee - pipe fitting

## SYNOPSIS

tee [ -i ] [ -a ] [ file ] ...

## DESCRIPTION

Tee transcribes the standard input to the standard output and makes copies in the files. Option -i ignores interrupts; option -a causes the output to be appended to the files rather than overwriting them.

## NAME

telnet - user interface to the TELNET protocol

## SYNOPSIS

telnet [ host [ port ] ]

## DESCRIPTION

Telnet is used to communicate with another host using the TELNET protocol. If telnet is invoked without arguments, it enters command mode, indicated by its prompt ("telnet>"). In this mode, it accepts and executes the commands listed below. If it is invoked with arguments, it performs an open command (see below) with those arguments.

Once a connection has been opened, telnet enters an input mode. The input mode entered will be either "character at a time" or "line by line" depending on what the remote system supports.

In "character at a time" mode, most text typed is immediately sent to the remote host for processing.

In "line by line" mode, all text is echoed locally, and (normally) only completed lines are sent to the remote host. The "local echo character" (initially "^E") may be used to turn off and on the local echo (this would mostly be used to enter passwords without the password being echoed).

In either mode, if the localchars toggle is TRUE (the default in line mode; see below), the user's quit, intr, and flush characters are trapped locally, and sent as TELNET protocol sequences to the remote side. There are options (see toggle autoflush and toggle autosynch below) which cause this action to flush subsequent output to the terminal (until the remote host acknowledges the TELNET sequence) and flush previous terminal input (in the case of quit and intr).

While connected to a remote host, telnet command mode may be entered by typing the telnet "escape character" (initially "^]"). When in command mode, the normal terminal editing conventions are available.

## COMMANDS

The following commands are available. Only enough of each command to uniquely identify it need be typed (this is also true for arguments to the mode, set, toggle, and display commands).

open host [ port ]

Open a connection to the named host. If no port number

is specified, telnet will attempt to contact a TELNET server at the default port. The host specification may be either a host name (see hosts(5)) or an Internet address specified in the "dot notation" (see inet(3N)).

#### close

Close a TELNET session and return to command mode.

#### quit

Close any open TELNET session and exit telnet. An end of file (in command mode) will also close a session and exit.

#### z

Suspend telnet. This command only works when the user is using the csh(1).

#### mode type

Type is either line (for "line by line" mode) or character (for "character at a time" mode). The remote host is asked for permission to go into the requested mode. If the remote host is capable of entering that mode, the requested mode will be entered.

#### status

Show the current status of telnet. This includes the peer one is connected to, as well as the current mode.

#### display [ argument... ]

Displays all, or some, of the set and toggle values (see below).

#### ? [ command ]

Get help. With no arguments, telnet prints a help summary. If a command is specified, telnet will print the help information for just that command.

#### send arguments

Sends one or more special character sequences to the remote host. The following are the arguments which may be specified (more than one argument may be specified at a time):

#### escape

Sends the current telnet escape character (initially "^>").

#### synch

Sends the TELNET SYNCH sequence. This sequence causes the remote system to discard all previously typed (but not yet read) input. This sequence is sent as TCP urgent data (and may not work if the

remote system is a 4.2 BSD system -- if it doesn't work, a lower case "r" may be echoed on the terminal).

brk

Sends the TELNET BRK (Break) sequence, which may have significance to the remote system.

ip

Sends the TELNET IP (Interrupt Process) sequence, which should cause the remote system to abort the currently running process.

ao

Sends the TELNET AO (Abort Output) sequence, which should cause the remote system to flush all output from the remote system to the user's terminal.

ayt

Sends the TELNET AYT (Are You There) sequence, to which the remote system may or may not choose to respond.

ec

Sends the TELNET EC (Erase Character) sequence, which should cause the remote system to erase the last character entered.

el

Sends the TELNET EL (Erase Line) sequence, which should cause the remote system to erase the line currently being entered.

ga

Sends the TELNET GA (Go Ahead) sequence, which likely has no significance to the remote system.

nop

Sends the TELNET NOP (No Operation) sequence.

?

Prints out help information for the send command.

set argument value

Set any one of a number of telnet variables to a specific value. The special value "off" turns off the function associated with the variable. The values of variables may be interrogated with the display command. The variables which may be specified are:

echo

This is the value (initially "^E") which, when in



"line by line" mode, toggles between doing local echoing of entered characters (for normal processing), and suppressing echoing of entered characters (for entering, say, a password).

#### escape

This is the telnet escape character (initially "^[" ) which causes entry into telnet command mode (when connected to a remote system).

#### interrupt

If telnet is in localchars mode (see toggle localchars below) and the interrupt character is typed, a TELNET IP sequence (see send ip above) is sent to the remote host. The initial value for the interrupt character is taken to be the terminal's intr character.

#### quit

If telnet is in localchars mode (see toggle localchars below) and the quit character is typed, a TELNET BRK sequence (see send brk above) is sent to the remote host. The initial value for the quit character is taken to be the terminal's quit character.

#### flushoutput

If telnet is in localchars mode (see toggle localchars below) and the flushoutput character is typed, a TELNET AO sequence (see send ao above) is sent to the remote host. The initial value for the flush character is taken to be the terminal's flush character.

#### erase

If telnet is in localchars mode (see toggle localchars below), and if telnet is operating in "character at a time" mode, then when this character is typed, a TELNET EC sequence (see send ec above) is sent to the remote system. The initial value for the erase character is taken to be the terminal's erase character.

#### kill

If telnet is in localchars mode (see toggle localchars below), and if telnet is operating in "character at a time" mode, then when this character is typed, a TELNET EL sequence (see send el above) is sent to the remote system. The initial value for the kill character is taken to be the terminal's kill character.

#### eof

If telnet is operating in "line by line" mode, entering this character as the first character on a line will cause this character to be sent to the remote system. The initial value of the eof character is taken to be the terminal's eof character.

#### toggle arguments...

Toggle (between TRUE and FALSE) various flags that control how telnet responds to events. More than one argument may be specified. The state of these flags may be interrogated with the display command. Valid arguments are:

#### localchars

If this is TRUE, then the flush, interrupt, quit, erase, and kill characters (see set above) are recognized locally, and transformed into (hopefully) appropriate TELNET control sequences (respectively ao, ip, brk, ec, and el; see send above). The initial value for this toggle is TRUE in "line by line" mode, and FALSE in "character at a time" mode.

#### autoflush

If autoflush and localchars are both TRUE, then when the ao, intr, or quit characters are recognized (and transformed into TELNET sequences; see set above for details), telnet refuses to display any data on the user's terminal until the remote system acknowledges (via a TELNET Timing Mark option) that it has processed those TELNET sequences. The initial value for this toggle is TRUE if the terminal user had not done an "stty noflsh", otherwise FALSE (see stty(1)).

#### autosynch

If autosynch and localchars are both TRUE, then when either the intr or quit characters is typed (see set above for descriptions of the intr and quit characters), the resulting TELNET sequence sent is followed by the TELNET SYNCH sequence. This procedure should cause the remote system to begin throwing away all previously typed input until both of the TELNET sequences have been read and acted upon. The initial value of this toggle is FALSE.

#### crmod

Toggle carriage return mode. When this mode is enabled, most carriage return characters received from the remote host will be mapped into a

carriage return followed by a line feed. This mode does not affect those characters typed by the user, only those received from the remote host. This mode is not very useful unless the remote host only sends carriage return, but never line feed. The initial value for this toggle is FALSE.

#### debug

Toggles socket level debugging (useful only to the superuser). The initial value for this toggle is FALSE.

#### options

Toggles the display of some internal telnet protocol processing (having to do with TELNET options). The initial value for this toggle is FALSE.

#### netdata

Toggles the display of all network data (in hexadecimal format). The initial value for this toggle is FALSE.

#### ?

Displays the legal toggle commands.

#### BUGS

There is no adequate way for dealing with flow control.

On some remote systems, echo has to be turned off manually when in "line by line" mode.

There is enough settable state to justify a .telnetrc file.

No capability for a .telnetrc file is provided.

In "line by line" mode, the terminal's eof character is only recognized (and sent to the remote system) when it is the first character on a line.

## NAME

test - condition evaluation utility

## SYNOPSIS

test expression

## DESCRIPTION

The test utility evaluates the expression and, if it evaluates to true, returns a zero (true) exit status; otherwise it returns 1 (false). If there is no expression, test also returns 1 (false).

All operators and flags are separate arguments to the test utility.

The following primaries are used to construct expression:

- b file  
True if file exists and is a block special file.
- c file  
True if file exists and is a character special file.
- d file  
True if file exists and is a directory.
- e file  
True if file exists (regardless of type).
- f file  
True if file exists and is a regular file.
- g file  
True if file exists and its set group ID flag is set.
- h file  
True if file exists and is a symbolic link.
- n string  
True if the length of string is nonzero.
- p file  
True if file is a named pipe
- r file  
True if file exists and is readable.
- s file  
True if file exists and has a size greater than zero.
- t [file\_descriptor]  
True if the file whose file descriptor number is

file\_descriptor (default 1) is open and is associated with a terminal.

-u file

True if file exists and its set user ID flag is set.

-w file

True if file exists and is writable. True indicates only that the write flag is on. The file is not writable on a read-only file system even if this test indicates true.

-x file

True if file exists and is executable. True indicates only that the execute flag is on. If file is a directory, true indicates that file can be searched.

-z string

True if the length of string is zero.

string

True if string is not the null string.

s1 = s2

True if the strings s1 and s2 are identical.

s1 != s2

True if the strings s1 and s2 are not identical.

n1 -eq n2

True if the integers n1 and n2 are algebraically equal.

n1 -ne n2

True if the integers n1 and n2 are not algebraically equal.

n1 -gt n2

True if the integer n1 is algebraically greater than the integer n2 .

n1 -ge n2

True if the integer n1 is algebraically greater than or equal to the integer n2 .

n1 -lt n2

True if the integer n1 is algebraically less than the integer n2 .

n1 -le n2

True if the integer n1 is algebraically less than or equal to the integer n2 .

These primaries can be combined with the following operators:

`! expression`

True if expression is false.

`expression1 -a expression2`

True if both expression1 and expression2 are true.

`expression1 -o expression2`

True if either expression1 or expression2 are true.

`(expression)`

True if expression is true.

The

`-a` operator has higher precedence than the `-o` operator.

#### GRAMMAR AMBIGUITY

The test grammar is inherently ambiguous. In order to assure a degree of consistency, the cases described in the IEEE Std 1003.2 ("POSIX"), section D11.2/4.62.4, standard are evaluated consistently according to the rules specified in the standards document. All other cases are subject to the ambiguity in the command semantics.

#### RETURN VALUES

The test utility exits with one of the following values:

0        expression evaluated to true.

1        expression evaluated to false or expression was missing.

>1      An error occurred.

#### BUGS

Named pipes are not implemented in 2.11BSD.

#### STANDARDS

The test function is expected to be IEEE Std 1003.2 ("POSIX") compatible.

## NAME

tftp - trivial file transfer program

## SYNOPSIS

tftp [ host ]

## DESCRIPTION

Tftp is the user interface to the Internet TFTP (Trivial File Transfer Protocol), which allows users to transfer files to and from a remote machine. The remote host may be specified on the command line, in which case tftp uses host as the default host for future transfers (see the connect command below).

## COMMANDS

Once tftp is running, it issues the prompt tftp> and recognizes the following commands:

connect host-name [ port ]

Set the host (and optionally port) for transfers. Note that the TFTP protocol, unlike the FTP protocol, does not maintain connections between transfers; thus, the connect command does not actually create a connection, but merely remembers what host is to be used for transfers. You do not have to use the connect command; the remote host can be specified as part of the get or put commands.

mode transfer-mode

Set the mode for transfers; transfer-mode may be one of ascii or binary. The default is ascii.

put file

put localfile remotefile

put file1 file2 ... fileN remote-directory

Put a file or set of files to the specified remote file or directory. The destination can be in one of two forms: a filename on the remote host, if the host has already been specified, or a string of the form host:filename to specify both a host and filename at the same time. If the latter form is used, the host-name specified becomes the default for future transfers. If the remote-directory form is used, the remote host is assumed to be a UNIX machine.

get filename

get remotename localname

get file1 file2 ... fileN

Get a file or set of files from the specified sources. Source can be in one of two forms: a filename on the remote host, if the host has already been specified, or a string of the form host:filename to specify both a

host and filename at the same time. If the latter form is used, the last hostname specified becomes the default for future transfers.

quit Exit tftp. An end of file also exits.

verbose

Toggle verbose mode.

trace

Toggle packet tracing.

status

Show current status.

rexmt retransmission-timeout

Set the per-packet retransmission timeout, in seconds.

timeout total-transmission-timeout

Set the total transmission timeout, in seconds.

ascii

Shorthand for "mode ascii"

binary

Shorthand for "mode binary"

? [ command-name ... ]

Print help information.

## BUGS

Because there is no user-login or validation within the TFTP protocol, the remote site will probably have some sort of file-access restrictions in place. The exact methods are specific to each site and therefore difficult to document here.



## NAME

time - time a command

## SYNOPSIS

time command

## DESCRIPTION

The given command is executed; after it is complete, time prints the elapsed time during the command, the time spent in the system, and the time spent in execution of the command. Times are reported in seconds.

On a PDP-11, the execution time can depend on what kind of memory the program happens to land in; the user time in MOS is often half what it is in core.

The times are printed on the diagnostic output stream.

Time is built in to csh(1), using a different output format.

## BUGS

Elapsed time is accurate to the second, while the CPU times are measured to the 100th second. Thus the sum of the CPU times can be up to a second larger than the elapsed time.

Time is a built-in command to csh(1), with a much different syntax. This command is available as ``/bin/time'' to csh users.

## NAME

tip, cu - connect to a remote system

## SYNOPSIS

```
tip [-v] [-speed] system-name
tip [-v] [-speed] phone-number
cu phone-number [-t] [-s speed] [-a acu] [-l line]
[-#]
```

## DESCRIPTION

Tip and cu establish a full-duplex connection to another machine, giving the appearance of being logged in directly on the remote cpu. It goes without saying that you must have a login on the machine (or equivalent) to which you wish to connect. The preferred interface is tip. The cu interface is included for those people attached to the ``call UNIX'' command of version 7. This manual page describes only tip.

Typed characters are normally transmitted directly to the remote machine (which does the echoing as well). A tilde (~) appearing as the first character of a line is an escape signal; the following are recognized:

- ~^D ~.      Drop the connection and exit (you may still be logged in on the remote machine).
- ~c [name]      Change directory to name (no argument implies change to your home directory).
- ~!      Escape to a shell (exiting the shell will return you to tip).
- ~>      Copy file from local to remote.      Tip prompts for the name of a local file to transmit.
- ~<      Copy file from remote to local.      Tip prompts first for the name of the file to be sent, then for a command to be executed on the remote machine.
- ~p from [ to ]      Send a file to a remote UNIX host. The put command causes the remote UNIX system to run the command string ``cat > 'to'', while tip sends it the ``from'' file. If the ``to'' file isn't specified the ``from'' file name is used. This command is actually a UNIX specific version of the ``~>'' command.
- ~t from [ to ]      Take a file from a remote UNIX host. As in the put

command the ``to'' file defaults to the ``from'' file name if it isn't specified. The remote host executes the command string ``cat 'from';echo ^A'' to send the file to tip.

- ~|        Pipe the output from a remote command to a local UNIX process. The command string sent to the local UNIX system is processed by the shell.
- ~\$        Pipe the output from a local UNIX process to the remote host. The command string sent to the local UNIX system is processed by the shell.
- ~#        Send a BREAK to the remote system. For systems which don't support the necessary ioctl call the break is simulated by a sequence of line speed changes and DEL characters.
- ~s        Set a variable (see the discussion below).
- ~^Z       Stop tip (only available with job control).
- ~^Y       Stop only the ``local side'' of tip (only available with job control); the ``remote side'' of tip, the side that displays output from the remote host, is left running.
- ~?        Get a summary of the tilde escapes

Tip uses the file /etc/remote to find how to reach a particular system and to find out how it should operate while talking to the system; refer to remote(5) for a full description. Each system has a default baud rate with which to establish a connection. If this value is not suitable, the baud rate to be used may be specified on the command line, e.g. ``tip -300 mds''.

When tip establishes a connection it sends out a connection message to the remote system; the default value, if any, is defined in /etc/remote.

When tip prompts for an argument (e.g. during setup of a file transfer) the line typed may be edited with the standard erase and kill characters. A null line in response to a prompt, or an interrupt, will abort the dialogue and return you to the remote machine.

Tip guards against multiple users connecting to a remote system by opening modems and terminal lines with exclusive access, and by honoring the locking protocol used by uucp(1C).

During file transfers tip provides a running count of the number of lines transferred. When using the ~> and ~< commands, the ``eofread'' and ``eofwrite'' variables are used to recognize end-of-file when reading, and specify end-of-file when writing (see below). File transfers normally depend on tandem mode for flow control. If the remote system does not support tandem mode, ``echocheck'' may be set to indicate tip should synchronize with the remote system on the echo of each transmitted character.

When tip must dial a phone number to connect to a system it will print various messages indicating its actions. Tip supports the DEC DN-11 and Racal-Vadic 831 auto-call-units; the DEC DF02 and DF03, Ventel 212+, Racal-Vadic 3451, and Bizcomp 1031 and 1032 integral call unit/modems.

## VARIABLES

Tip maintains a set of variables which control its operation. Some of these variable are read-only to normal users (root is allowed to change anything of interest). Variables may be displayed and set through the ``s'' escape. The syntax for variables is patterned after vi(1) and Mail(1). Supplying ``all'' as an argument to the set command displays all variables readable by the user. Alternatively, the user may request display of a particular variable by attaching a `?' to the end. For example ``escape?'' displays the current escape character.

Variables are numeric, string, character, or boolean values. Boolean variables are set merely by specifying their name; they may be reset by prepending a `!' to the name. Other variable types are set by concatenating an `=' and the value. The entire assignment must not have any blanks in it. A single set command may be used to interrogate as well as set a number of variables. Variables may be initialized at run time by placing set commands (without the ``~s'' prefix in a file .tiprc in one's home directory). The -v option causes tip to display the sets as they are made. Certain common variables have abbreviations. The following is a list of common variables, their abbreviations, and their default values.

### beautify

(bool) Discard unprintable characters when a session is being scripted; abbreviated be.

### baudrate

(num) The baud rate at which the connection was established; abbreviated ba.

### dialtimeout

(num) When dialing a phone number, the time (in seconds) to wait for a connection to be established; abbreviated dial.

#### echocheck

(bool) Synchronize with the remote host during file transfer by waiting for the echo of the last character transmitted; default is off.

#### eofread

(str) The set of characters which signify an end-of-transmission during a ~< file transfer command; abbreviated eofr.

#### eofwrite

(str) The string sent to indicate end-of-transmission during a ~> file transfer command; abbreviated eofw.

#### eol

(str) The set of characters which indicate an end-of-line. Tip will recognize escape characters only after an end-of-line.

#### escape

(char) The command prefix (escape) character; abbreviated es; default value is '~'.

#### exceptions

(str) The set of characters which should not be discarded due to the beautification switch; abbreviated ex; default value is ``\t\n\f\b'`'.

#### force

(char) The character used to force literal data transmission; abbreviated fo; default value is '^P'.

#### framesize

(num) The amount of data (in bytes) to buffer between file system writes when receiving files; abbreviated fr.

#### host

(str) The name of the host to which you are connected; abbreviated ho.

#### prompt

(char) The character which indicates an end-of-line on the remote host; abbreviated pr; default value is '\n'. This value is used to synchronize during data transfers. The count of lines transferred during a file transfer command is based on receipt of this character.

**raise**

(bool) Upper case mapping mode; abbreviated ra; default value is off. When this mode is enabled, all lower case letters will be mapped to upper case by tip for transmission to the remote machine.

**raisechar**

(char) The input character used to toggle upper case mapping mode; abbreviated rc; default value is '^A'.

**record**

(str) The name of the file in which a session script is recorded; abbreviated rec; default value is ``tip.record''.

**script**

(bool) Session scripting mode; abbreviated sc; default is off. When script is true, tip will record everything transmitted by the remote machine in the script record file specified in record. If the beautify switch is on, only printable ASCII characters will be included in the script file (those characters between 040 and 0177). The variable exceptions is used to indicate characters which are an exception to the normal beautification rules.

**tabexpand**

(bool) Expand tabs to spaces during file transfers; abbreviated tab; default value is false. Each tab is expanded to 8 spaces.

**verbose**

(bool) Verbose mode; abbreviated verb; default is true. When verbose mode is enabled, tip prints messages while dialing, shows the current number of lines transferred during a file transfer operations, and more.

**SHELL**

(str) The name of the shell to use for the ~! command; default value is ``/bin/sh'', or taken from the environment.

**HOME**

(str) The home directory to use for the ~c command; default value is taken from the environment.

**FILES**

/etc/remote	global system descriptions
/etc/phones	global phone number data base
\${REMOTE}	private system descriptions
\${PHONES}	private phone numbers
~/.tiprc	initialization file.

/usr/spool/uucp/LCK..\*    lock file to avoid conflicts with uucp

#### DIAGNOSTICS

Diagnostics are, hopefully, self explanatory.

#### SEE ALSO

remote(5), phones(5)

#### BUGS

The full set of variables is undocumented and should, probably, be paired down.

## NAME

tk - paginator for the Tektronix 4014

## SYNOPSIS

tk [ -t ] [ -N ] [ -pL ] [ file ]

## DESCRIPTION

The output of tk is intended for a Tektronix 4014 terminal. Tk arranges for 66 lines to fit on the screen, divides the screen into N columns, and contributes an eight space page offset in the (default) single-column case. Tabs, spaces, and backspaces are collected and plotted when necessary. Teletype Model 37 half- and reverse-line sequences are interpreted and plotted. At the end of each page tk waits for a newline (empty line) from the keyboard before continuing on to the next page. In this wait state, the command !command will send the command to the shell.

The command line options are:

-t    Don't wait between pages; for directing output into a file.

-N    Divide the screen into N columns and wait after the last column.

-pL   Set page length to L lines.

## SEE ALSO

pr(1)



## NAME

tn3270 - full-screen remote login to IBM VM/CMS

## SYNOPSIS

tn3270 sysname

## DESCRIPTION

Tn3270 permits a full-screen, full-duplex connection from a VAX UNIX machine to an IBM machine running VM/CMS giving the appearance of being logged in directly to the remote machine on an IBM 3270 terminal. Of course you must have an account on the machine to which you wish to connect in order to log in. Tn3270 looks to the user in many respects like the Yale ASCII Terminal Communication System II. Tn3270 is actually a modification of the Arpanet TELNET user interface (see telnet(1)) that interprets and generates raw 3270 control streams.

Emulation of the 3270 terminal is done in the Unix process. This emulation involves mapping 3270-style commands from the host into appropriate sequences to control the user's terminal screen. Tn3270 uses curses(3x) and the /etc/termcap file to do this. The emulation also involves simulating the special 3270 keyboard keys (program function keys, etc.) by mapping sequences of keystrokes from the ASCII keyboard into appropriate 3270 control strings. This mapping is terminal dependent and is specified in a description file, /usr/share/misc/map3270, (see map3270(5)) or in an environment variable MAP3270 (see mset(1)). Any special function keys on the ASCII keyboard are used whenever possible. If an entry for the user's terminal is not found, tn3270 looks for an entry for the terminal type unknown. If this is not found, tn3270 uses a default keyboard mapping (see map3270(5)).

The first character of each special keyboard mapping sequence is either an ASCII escape (ESC), a control character, or an ASCII delete (DEL). If the user types an unrecognized function key sequence, tn3270 sends an ASCII bell (BEL), or a visual bell if defined in the user's termcap entry, to the user's terminal and nothing is sent to the IBM host.

If tn3270 is invoked without specifying a remote host system name, it enters local command mode, indicated by the prompt ``tn3270>''. In this mode, tn3270 accepts and executes the following commands:

open	connect to a remote host
close	close the current connection
quit	exit tn3270
z	suspend tn3270

```
status print connection status
? print help information
```

Other common telnet commands are not available in tn3270. Tn3270 command mode may also be entered, after connecting to a host, by typing a special escape character (typically control-C).

While in command mode, any host login session is still alive but temporarily suspended. The host login session may be resumed by entering an empty line (press the RETURN key) in response to the command prompt. A session may be terminated by logging off the foreign host, or by typing ``quit'' or ``close'' while in local command mode.

#### FILES

```
/etc/termcap
/usr/share/misc/map3270
```

#### AUTHOR

Greg Minshall

#### SEE ALSO

mset(1), telnet(1), termcap(3x), termcap(5), map3270(5),  
Yale ASCII Terminal Communication System II Program  
Description/Operator's Manual (IBM SB30-1911)

#### BUGS

Performance is slow and uses system resources prodigiously.

Not all 3270 functions are supported, nor all Yale enhancements.

## NAME

touch - update date last modified of a file

## SYNOPSIS

touch [ -c ] [ -f ] file ...

## DESCRIPTION

Touch attempts to set the modified date of each file. If a file exists, this is done by reading a character from the file and writing it back. If a file does not exist, an attempt will be made to create it unless the -c option is specified. The -f option will attempt to force the touch in spite of read and write permissions on a file.

## SEE ALSO

utimes(2)

## NAME

tp - manipulate tape archive

## SYNOPSIS

tp [ key ] [ name ... ]

## DESCRIPTION

Tp saves and restores files on DECTape or magtape. Its actions are controlled by the key argument. The key is a string of characters containing at most one function letter and possibly one or more function modifiers. Other arguments to the command are file or directory names specifying which files are to be dumped, restored, or listed. In all cases, appearance of a directory name refers to the files and (recursively) subdirectories of that directory.

The function portion of the key is specified by one of the following letters:

- r           The named files are written on the tape. If files with the same names already exist, they are replaced. 'Same' is determined by string comparison, so './abc' can never be the same as '/usr/dmr/abc' even if '/usr/dmr' is the current directory. If no file argument is given, '.' is the default.
- u           updates the tape. u is like r, but a file is replaced only if its modification date is later than the date stored on the tape; that is to say, if it has changed since it was dumped. u is the default command if none is given.
- d           deletes the named files from the tape. At least one name argument must be given. This function is not permitted on magtapes.
- x           extracts the named files from the tape to the file system. The owner and mode are restored. If no file argument is given, the entire contents of the tape are extracted.
- t           lists the names of the specified files. If no file argument is given, the entire contents of the tape is listed.

The following characters may be used in addition to the letter which selects the function desired.

m           Specifies magtape as opposed to DECTape.

0,...,7    This modifier selects the drive on which the tape

is mounted. For DEctape, x is default; for magtape `0' is the default.

- v                Normally tp does its work silently. The v (verbose) option causes it to type the name of each file it treats preceded by the function letter. With the t function, v gives more information about the tape entries than just the name.
- c                means a fresh dump is being created; the tape directory is cleared before beginning. Usable only with r and u. This option is assumed with magtape since it is impossible to selectively overwrite magtape.
- i                Errors reading and writing the tape are noted, but no action is taken. Normally, errors cause a return to the command level.
- f                Use the first named file, rather than a tape, as the archive. This option currently acts like m; i.e. r implies c, and neither d nor u are permitted.
- w                causes tp to pause before treating each file, type the indicative letter and the file name (as with v) and await the user's response. Response y means `yes', so the file is treated. Null response means `no', and the file does not take part in whatever is being done. Response x means `exit'; the tp command terminates immediately. In the x function, files previously asked about have been extracted already. With r, u, and d no change has been made to the tape.

#### FILES

/dev/tap?  
/dev/rmt?

#### SEE ALSO

ar(1), tar(1)

#### DIAGNOSTICS

Several; the non-obvious one is `Phase error', which means the file changed after it was selected for dumping but before it was dumped.

#### BUGS

A single file with several links to it is treated like several files.

Binary-coded control information makes magnetic tapes written by `tp` difficult to carry to other machines; `tar(1)` avoids the problem.

## NAME

tr - translate characters

## SYNOPSIS

tr [ -cds ] [ string1 [ string2 ] ]

## DESCRIPTION

Tr copies the standard input to the standard output with substitution or deletion of selected characters. Input characters found in string1 are mapped into the corresponding characters of string2. When string2 is short it is padded to the length of string1 by duplicating its last character. Any combination of the options -cds may be used: -c complements the set of characters in string1 with respect to the universe of characters whose ASCII codes are 01 through 0377 octal; -d deletes all input characters in string1; -s squeezes all strings of repeated output characters that are in string2 to single characters.

In either string the notation a-b means a range of characters from a to b in increasing ASCII order. The character '\\' followed by 1, 2 or 3 octal digits stands for the character whose ASCII code is given by those digits. A '\\' followed by any other character stands for that character.

The following example creates a list of all the words in 'file1' one per line in 'file2', where a word is taken to be a maximal string of alphabetic characters. The second string is quoted to protect '\\' from the Shell. 012 is the ASCII code for newline.

```
tr -cs A-Za-z '\\012' <file1 >file2
```

## SEE ALSO

ed(1), ascii(7), expand(1)

## BUGS

Won't handle ASCII NUL in string1 or string2; always deletes NUL from input.

## NAME

troff, nroff - text formatting and typesetting

## SYNOPSIS

troff [ option ] ... [ file ] ...

nroff [ option ] ... [ file ] ...

## DESCRIPTION

Troff formats text in the named files for printing on a Graphic Systems C/A/T phototypesetter; nroff is used for typewriter-like devices. Their capabilities are described in the Nroff/Troff user's manual.

If no file argument is present, the standard input is read. An argument consisting of a single minus (-) is taken to be a file name corresponding to the standard input. The options, which may appear in any order so long as they appear before the files, are:

- olist Print only pages whose page numbers appear in the comma-separated list of numbers and ranges. A range N-M means pages N through M; an initial -N means from the beginning to page N; and a final N- means from N to the end.
  - nN Number first generated page N.
  - sN Stop every N pages. Nroff will halt prior to every N pages (default N=1) to allow paper loading or changing, and will resume upon receipt of a newline. Troff will stop the phototypesetter every N pages, produce a trailer to allow changing cassettes, and resume when the typesetter's start button is pressed.
  - mname Prepend the macro file /usr/share/tmac/tmac.name to the input files.
  - raN Set register a (one-character) to N.
  - i Read standard input after the input files are exhausted.
  - q Invoke the simultaneous input-output mode of the rd request.
- Troff only
- t Direct output to the standard output instead of the phototypesetter.
  - f Refrain from feeding out paper and stopping



phototypesetter at the end of the run.

- w      Wait until phototypesetter is available, if currently busy.
- b      Report whether the phototypesetter is busy or available. No text processing is done.
- a      Send a printable ASCII approximation of the results to the standard output.
- pN     Print all characters in point size N while retaining all prescribed spacings and motions, to reduce phototypesetter elapsed time.
- Ffontdir  
      The directory fontdir contains the font width tables instead of the default directory /usr/share/fonts. This option can be used to produce output for devices besides the phototypesetter.

If the file /usr/adm/tracct is writable, troff keeps phototypesetter accounting records there. The integrity of that file may be secured by making troff a 'set user-id' program.

#### FILES

/tmp/ta*	temporary file
/usr/share/tmac/tmac.*	standard macro files
/usr/share/term/*	terminal driving tables for nroff
/usr/share/font/*	font width tables for troff
/dev/cat	phototypesetter
/usr/adm/tracct	accounting statistics for /dev/cat

#### SEE ALSO

J. F. Ossanna, Nroff/Troff user's manual  
B. W. Kernighan, A TROFF Tutorial  
eqn(1), tbl(1), ms(7), me(7), man(7), col(1)

## NAME

true, false - provide truth values

## SYNOPSIS

true

false

## DESCRIPTION

True and false are usually used in a Bourne shell script. They test for the appropriate status "true" or "false" before running (or failing to run) a list of commands.

## EXAMPLE

```
while true
do
 command list
done
```

## SEE ALSO

csh(1), sh(1), false(1)

## DIAGNOSTICS

True has exit status zero.

## NAME

tset - terminal dependent initialization

## SYNOPSIS

```
tset [options] [-m [ident][test baudrate]:type] ... [
type]

reset [options] [-m [ident][test baudrate]:type] ... [
type]
```

## DESCRIPTION

Tset sets up your terminal when you first log in to a UNIX system. It does terminal dependent processing such as setting erase and kill characters, setting or resetting delays, sending any sequences needed to properly initialize the terminal, and the like. It first determines the type of terminal involved, and then does necessary initializations and mode settings. The type of terminal attached to each UNIX port is specified in the /etc/ttys(5) database. Type names for terminals may be found in the termcap(5) database. If a port is not wired permanently to a specific terminal (not hardwired) it will be given an appropriate generic identifier such as dialup.

In the case where no arguments are specified, tset simply reads the terminal type out of the environment variable TERM and re-initializes the terminal. The rest of this manual concerns itself with mode and environment initialization, typically done once at login, and options used at initialization time to determine the terminal type and set up terminal modes.

When used in a startup script (.profile for sh(1) users or .login for csh(1) users) it is desirable to give information about the type of terminal you will usually use on ports which are not hardwired. These ports are identified in /etc/ttys as dialup or plugboard or arpanet, etc. To specify what terminal type you usually use on these ports, the -m (map) option flag is followed by the appropriate port type identifier, an optional baud rate specification, and the terminal type. (The effect is to ``map'' from some conditions to a terminal type, that is, to tell tset ``If I'm on this kind of port, guess that I'm on that kind of terminal''.) If more than one mapping is specified, the first applicable mapping prevails. A missing port type identifier matches all identifiers. Any of the alternate generic names given in termcap may be used for the identifier.

A baudrate is specified as with stty(1), and is compared with the speed of the diagnostic output (which should be the control terminal). The baud rate test may be any combination of: >, @, <, and !; @ means ``at'' and ! inverts the

sense of the test. To avoid problems with metacharacters, it is best to place the entire argument to `-m` within ````` characters; users of `csh(1)` must also put a ````` before any ````` used here.

Thus

```
tset -m 'dialup>300:adm3a' -m dialup:dw2 -m
'plugboard:?adm3a'
```

causes the terminal type to be set to an `adm3a` if the port in use is a dialup at a speed greater than 300 baud; to a `dw2` if the port is (otherwise) a dialup (i.e. at 300 baud or less). (NOTE: the examples given here appear to take up more than one line, for text processing reasons. When you type in real `tset` commands, you must enter them entirely on one line.) If the type finally determined by `tset` begins with a question mark, the user is asked if s/he really wants that type. A null response means to use that type; otherwise, another type can be entered which will be used instead. Thus, in the above case, the user will be queried on a plugboard port as to whether they are actually using an `adm3a`.

If no mapping applies and a final type option, not preceded by a `-m`, is given on the command line then that type is used; otherwise the type found in the `/etc/ttys` database will be taken to be the terminal type. This should always be the case for hardwired ports.

It is usually desirable to return the terminal type, as finally determined by `tset`, and information about the terminal's capabilities to a shell's environment. This can be done using the `-o` option; using the Bourne shell, `sh(1)`:

```
export TERM; TERM=`tset -o options...`
```

or using the C shell, `csh(1)`:

```
setenv TERM `tset -o options...`
```

With `csh` it is preferable to use the following command in your `.login` file to initialize the `TERM` and `TERMCAP` environment variables at the same time.

```
eval `tset -s options...`
```

It is also convenient to make an alias in your `.cshrc`:

```
alias tset 'eval `tset -s `!`*`
```

This allows the command:

```
tset 2621
```

to be invoked at any time to set the terminal and environment. Note to Bourne Shell users: It is not possible to get this aliasing effect with a shell script, because shell scripts cannot set the environment of their parent. (If a process could set its parent's environment, none of this nonsense would be necessary in the first place.)

These commands cause tset to place the name of your terminal in the variable TERM in the environment; see environ(7).

Once the terminal type is known, tset engages in terminal driver mode setting. This normally involves sending an initialization sequence to the terminal, setting the single character erase (and optionally the line-kill (full line erase)) characters, and setting special character delays. Tab and newline expansion are turned off during transmission of the terminal initialization sequence.

On terminals that can backspace but not overstrike (such as a CRT), and when the erase character is the default erase character (`#' on standard systems), the erase character is changed to BACKSPACE (Control-H).

The options are:

- ec set the erase character to be the named character c on all terminals, the default being the backspace character on the terminal, usually ^H. The character c can either be typed directly, or entered using the hat notation used here.
- kc is similar to -e but for the line kill character rather than the erase character; c defaults to ^X (for purely historical reasons). The kill characters is left alone if -k is not specified. The hat notation can also be used for this option.
- ic is similar to -e but for the interrupt character rather than the erase character; c defaults to ^C. The hat notation can also be used for this option.
- The name of the terminal finally decided upon is output on the standard output. This is intended to be captured by the shell and placed in the environment variable TERM.
- s Print the sequence of csh commands to initialize the environment variables TERM and TERMCAP based on the

name of the terminal finally decided upon.

-n On systems with the Berkeley 4BSD tty driver, specifies that the new tty driver modes should be initialized for this terminal. For a CRT, the CRTERASE and CRTKILL modes are set only if the baud rate is 1200 or greater. See tty(4) for more detail.

-I suppresses transmitting terminal initialization strings.

-Q suppresses printing the ``Erase set to'' and ``Kill set to'' messages.

If tset is invoked as reset, it will set cooked and echo modes, turn off cbreak and raw modes, turn on newline translation, and restore special characters to a sensible state before any terminal dependent processing is done. Any special character that is found to be NULL or ``-1'' is reset to its default value. All arguments to tset may be used with reset.

This is most useful after a program dies leaving a terminal in a funny state. You may have to type ``<LF>reset<LF>'' to get it to work since <CR> may not work in this state. Often none of this will echo.

#### EXAMPLES

These examples all assume the Bourne shell and use the - option. If you use csh, use one of the variations described above. Note that a typical use of tset in a .profile or .login will also use the -e and -k options, and often the -n or -Q options as well. These options have not been included here to keep the examples small. (NOTE: some of the examples given here appear to take up more than one line, for text processing reasons. When you type in real tset commands, you must enter them entirely on one line.)

At the moment, you are on a 2621. This is suitable for typing by hand but not for a .profile, unless you are always on a 2621.

```
export TERM; TERM=`tset - 2621`
```

You have an h19 at home which you dial up on, but your office terminal is hardwired and known in /etc/ttys.

```
export TERM; TERM=`tset - -m dialup:h19`
```

You have a switch which connects everything to everything, making it nearly impossible to key on what port you are coming in on. You use a vt100 in your office at 9600 baud, and

dial up to switch ports at 1200 baud from home on a 2621. Sometimes you use someone else's terminal at work, so you want it to ask you to make sure what terminal type you have at high speeds, but at 1200 baud you are always on a 2621. Note the placement of the question mark, and the quotes to protect the greater than and question mark from interpretation by the shell.

```
export TERM; TERM=`tset - -m 'switch>1200:?vt100' -m
'switch<=1200:2621'
```

All of the above entries will fall back on the terminal type specified in /etc/ttys if none of the conditions hold. The following entry is appropriate if you always dial up, always at the same baud rate, on many different kinds of terminals. Your most common terminal is an adm3a. It always asks you what kind of terminal you are on, defaulting to adm3a.

```
export TERM; TERM=`tset - ?adm3a`
```

If the file /etc/ttys is not properly installed and you want to key entirely on the baud rate, the following can be used:

```
export TERM; TERM=`tset - -m '>1200:vt100' 2621`
```

Here is a fancy example to illustrate the power of tset and to hopelessly confuse anyone who has made it this far. You dial up at 1200 baud or less on a concept100, sometimes over switch ports and sometimes over regular dialups. You use various terminals at speeds higher than 1200 over switch ports, most often the terminal in your office, which is a vt100. However, sometimes you log in from the university you used to go to, over the ARPANET; in this case you are on an ALTO emulating a dm2500. You also often log in on various hardwired ports, such as the console, all of which are properly entered in /etc/ttys. You want your erase character set to control H, your kill character set to control U, and don't want tset to print the ``Erase set to Backspace, Kill set to Control U'' message.

```
export TERM; TERM=`tset -e -k^U -Q - -m
'switch<=1200:concept100' -m 'switch:?vt100' -m
dialup:concept100 -m arpanet:dm2500`
```

#### FILES

```
/etc/ttys port name to terminal type mapping database
/etc/termcap terminal capability database
```

#### SEE ALSO

```
csh(1), sh(1), stty(1), ttys(5), termcap(5), environ(7)
```

## BUGS

The tset command is one of the first commands a user must master when getting started on a UNIX system. Unfortunately, it is one of the most complex, largely because of the extra effort the user must go through to get the environment of the login shell set. Something needs to be done to make all this simpler, either the login(1) program should do this stuff, or a default shell alias should be made, or a way to set the environment of the parent should exist.

This program can't intuit personal choices for erase, interrupt and line kill characters, so it leaves these set to the local system standards.

## NOTES

For compatibility with earlier versions of tset a number of flags are accepted whose use is discouraged:

- d type     equivalent to -m dialup:type
- p type     equivalent to -m plugboard:type
- a type     equivalent to -m arpanet:type
- prints the terminal type on the standard output
- r           prints the terminal type on the diagnostic output.



## NAME

tsort - topological sort

## SYNOPSIS

tsort [ file ]

## DESCRIPTION

Tsort produces on the standard output a totally ordered list of items consistent with a partial ordering of items mentioned in the input file. If no file is specified, the standard input is understood.

The input consists of pairs of items (nonempty strings) separated by blanks. Pairs of different items indicate ordering. Pairs of identical items indicate presence, but not ordering.

## SEE ALSO

lorder(1)

## DIAGNOSTICS

Odd data: there is an odd number of fields in the input file.

## BUGS

Uses a quadratic algorithm; not worth fixing for the typical use of ordering a library archive file.

## NAME

tty - get terminal name

## SYNOPSIS

tty [ -s ]

## DESCRIPTION

Tty prints the pathname of the user's terminal unless the -s (silent) is given. In either case, the exit value is zero if the standard input is a terminal and one if it is not.

## DIAGNOSTICS

`not a tty' if the standard input file is not a terminal.

## NAME

ul - do underlining

## SYNOPSIS

ul [ -i ] [ -t terminal ] [ name ... ]

## DESCRIPTION

Ul reads the named files (or standard input if none are given) and translates occurrences of underscores to the sequence which indicates underlining for the terminal in use, as specified by the environment variable TERM. The -t option overrides the terminal kind specified in the environment. The file /etc/termcap is read to determine the appropriate sequences for underlining. If the terminal is incapable of underlining, but is capable of a standout mode then that is used instead. If the terminal can overstrike, or handles underlining automatically, ul degenerates to cat(1). If the terminal cannot underline, underlining is ignored.

The -i option causes ul to indicate underlining onto by a separate line containing appropriate dashes `-' ; this is useful when you want to look at the underlining which is present in an nroff output stream on a crt-terminal.

## SEE ALSO

man(1), nroff(1), colcrt(1)

## BUGS

Nroff usually outputs a series of backspaces and underlines intermixed with the text to indicate underlining. No attempt is made to optimize the backward motion.

## NAME

unifdef - remove ifdef'ed lines

## SYNOPSIS

unifdef [ -t -l -c -Dsym -Usym -idsym -iusym ] ... [ file ]

## DESCRIPTION

Unifdef is useful for removing ifdef'ed lines from a file while otherwise leaving the file alone. Unifdef is like a stripped-down C preprocessor: it is smart enough to deal with the nested ifdefs, comments, single and double quotes of C syntax so that it can do its job, but it doesn't do any including or interpretation of macros. Neither does it strip out comments, though it recognizes and ignores them. You specify which symbols you want defined -Dsym or undefined -Usym and the lines inside those ifdefs will be copied to the output or removed as appropriate. The ifdef, ifndef, else, and endif lines associated with sym will also be removed. Ifdefs involving symbols you don't specify are untouched and copied out along with their associated ifdef, else, and endif lines. If an ifdef X occurs nested inside another ifdef X, then the inside ifdef is treated as if it were an unrecognized symbol. If the same symbol appears in more than one argument, only the first occurrence is significant.

The -l option causes unifdef to replace removed lines with blank lines instead of deleting them.

If you use ifdefs to delimit non-C lines, such as comments or code which is under construction, then you must tell unifdef which symbols are used for that purpose so that it won't try to parse for quotes and comments in those ifdef'ed lines. You specify that you want the lines inside certain ifdefs to be ignored but copied out with -idsym and -iusym similar to -Dsym and -Usym above.

If you want to use unifdef for plain text (not C code), use the -t option. This makes unifdef refrain from attempting to recognize comments and single and double quotes.

Unifdef copies its output to stdout and will take its input from stdin if no file argument is given. If the -c argument is specified, then the operation of unifdef is complemented, i.e. the lines that would have been removed or blanked are retained and vice versa.

## SEE ALSO

diff(1)

## DIAGNOSTICS

Premature EOF, inappropriate else or endif.

Exit status is 0 if output is exact copy of input, 1 if not,  
2 if trouble.

#### BUGS

Does not know how to deal with cpp constructs such as

```
#if defined(X) || defined(Y)
```

#### AUTHOR

Dave Yost

## NAME

uniq - report repeated lines in a file

## SYNOPSIS

uniq [ -udc [ +n ] [ -n ] ] [ input [ output ] ]

## DESCRIPTION

Uniq reads the input file comparing adjacent lines. In the normal case, the second and succeeding copies of repeated lines are removed; the remainder is written on the output file. Note that repeated lines must be adjacent in order to be found; see sort(1). If the -u flag is used, just the lines that are not repeated in the original file are output. The -d option specifies that one copy of just the repeated lines is to be written. The normal mode output is the union of the -u and -d mode outputs.

The -c option supersedes -u and -d and generates an output report in default style but with each line preceded by a count of the number of times it occurred.

The n arguments specify skipping an initial portion of each line in the comparison:

-n        The first n fields together with any blanks before each are ignored. A field is defined as a string of non-space, non-tab characters separated by tabs and spaces from its neighbors.

+n        The first n characters are ignored. Fields are skipped before characters.

## SEE ALSO

sort(1), comm(1)

## NAME

units - conversion program

## SYNOPSIS

units

## DESCRIPTION

Units converts quantities expressed in various standard scales to their equivalents in other scales. It works interactively in this fashion:

```
You have: inch
You want: cm
 * 2.54000e+00
 / 3.93701e-01
```

A quantity is specified as a multiplicative combination of units optionally preceded by a numeric multiplier. Powers are indicated by suffixed positive integers, division by the usual sign:

```
You have: 15 pounds force/in2
You want: atm
 * 1.02069e+00
 / 9.79730e-01
```

Units only does multiplicative scale changes. Thus it can convert Kelvin to Rankine, but not Centigrade to Fahrenheit. Most familiar units, abbreviations, and metric prefixes are recognized, together with a generous leavening of exotica and a few constants of nature including:

```
pi ratio of circumference to diameter
c speed of light
e charge on an electron
g acceleration of gravity
force same as g
mole Avogadro's number
water pressure head per unit height of water
au astronomical unit
```

`Pound' is a unit of mass. Compound names are run together, e.g. `lightyear'. British units that differ from their US counterparts are prefixed thus: `brgallon'. Currency is denoted `belgiumfranc', `britainpound', ...

For a complete list of units, `cat /usr/share/misc/units'.

## FILES

/usr/share/misc/units

## BUGS

Don't base your financial plans on the currency conversions.



## NAME

uptime - show how long system has been up

## SYNOPSIS

uptime

## DESCRIPTION

Uptime prints the current time, the length of time the system has been up, and the average number of jobs in the run queue over the last 1, 5 and 15 minutes. It is, essentially, the first line of a w(1) command.

## FILES

/vmunix    system name list

## SEE ALSO

w(1)

## NAME

users - compact list of users who are on the system

## SYNOPSIS

users

## DESCRIPTION

Users lists the login names of the users currently on the system in a compact, one-line format.

## FILES

/var/run/utmp

## SEE ALSO

who(1)

## NAME

uucp - unix to unix copy

## SYNOPSIS

uucp [ -acCdfmr ] [ -nuser ] [ -ggrade ] [ -sspool ] [ -xdebug ] source-file.... destination-file

## DESCRIPTION

Uucp copies files named by the source-file arguments to the destination-file argument. A file name may be a pathname on your machine, or may have the form

system-name!pathname

where 'system-name' is taken from a list of system names that uucp knows about. Shell metacharacters ?\*[] appearing in the pathname part will be expanded on the appropriate system.

Pathnames may be one of:

- (1) a full pathname;
- (2) a pathname preceded by ~user; where user is a userid on the specified system and is replaced by that user's login directory;
- (3) a pathname prefixed by ~, where ~ is expanded into the system's public directory (usually /usr/spool/uucppublic);
- (4) a partial pathname, which is prefixed by the current directory.

If the result is an erroneous pathname for the remote system, the copy will fail. If the destination-file is a directory, the last part of the source-file name is used.

Uucp preserves execute permissions across the transmission and gives 0666 read and write permissions (see chmod(2)).

The following options are interpreted by uucp.

- a Avoid doing a getwd to find the current directory. (This is sometimes used for efficiency.)
- c Use the source file when copying out rather than copying the file to the spool directory. (This is the default.)
- C Copy the source file to the spool directory and transmit the copy.

- d    Make all necessary directories for the file copy.  
      (This is the default.)
- f    Do not make intermediate directories for the file copy.
- ggrade  
      Grade is a single letter/number; lower ASCII sequence characters will cause a job to be transmitted earlier during a particular conversation. Default is `n'. By way of comparison, uux(1C) defaults to `A'; mail is usually sent at `C'.
- m    Send mail to the requester when the copy is complete.
- nuser  
      Notify user on remote system (i.e., send user mail) that a file was sent.
- r    Do not start the transfer, just queue the job.
- sspool  
      Use spool as the spool directory instead of the default.
- xdebug  
      Turn on the debugging at level debug.

#### FILES

/usr/spool/uucp - spool directory  
/etc/uucp/\* - data and configuration files

#### SEE ALSO

uux(1C), mail(1)

D. A. Nowitz and M. E. Lesk, A Dial-Up Network of UNIX Systems.

D. A. Nowitz, Uucp Implementation Description.

#### WARNING

The domain of remotely accessible files can (and for obvious security reasons, usually should) be severely restricted. You will very likely not be able to fetch files by pathname; ask a responsible person on the remote system to send them to you. For the same reasons you will probably not be able to send files to arbitrary pathnames.

#### BUGS

All files received by uucp will be owned by the uucp administrator (usually UID 5).

The `-m` option will only work sending files or receiving a single file. (Receiving multiple files specified by special shell characters `?*[]` will not activate the `-m` option.)

At present `uucp` cannot copy to a system several "hops" away, that is, a command of the form

```
uucp myfile system1!system2!system3!yourfile
```

is not permitted. Use `uusend(1C)` instead.

When invoking `uucp` from `cs(1)`, the ``!'` character must be prefixed by the ```\'` escape to inhibit `cs(1)`'s history mechanism. (Quotes are not sufficient.)

`Uucp` refuses to copy a file that does not give read access to ```other''`; that is, the file must have at least 0444 modes.

## NAME

uulog - display UUCP log files

## SYNOPSIS

uulog [ -s sys ] [ -u user ]

## DESCRIPTION

Uulog queries a log of uucp(1C) and uux(1C) transactions in the file /usr/spool/uucp/LOGFILE.

The options command uulog to print logging information:

-ssys    Print information about work involving system sys.

-uuser   Print information about work done for the specified user.

## FILES

/usr/spool/uucp/LOGFILE

## SEE ALSO

uucp(1C), uux(1C).

## NOTES

Very early releases of UUCP used separate log files for each of the UUCP utilities; uulog was used to merge the individual logs into a master file. This capability has not been necessary for some time and is no longer supported.

## BUGS

UUCP's recording of which user issued a request is unreliable.

Uulog is little more than an overspecialized version of grep(1).

## NAME

uuname - list names of UUCP hosts

## SYNOPSIS

uuname [ -l ]

## DESCRIPTION

Uuname lists the UUCP names of known systems. The -l option returns the local system name; this may differ from the hostname(1) for the system if the hostname is very long.

## SEE ALSO

uucp(1C), uux(1C).

## NAME

uuq - examine or manipulate the uucp queue

## SYNOPSIS

uuq [ -l ] [ -h ] [ -ssystem ] [ -uuser ] [ -djobno ] [ -rsdir ] [ -bbaud ]

## DESCRIPTION

Uuq is used to examine (and possibly delete) entries in the uucp queue.

When listing jobs, uuq uses a format reminiscent of ls. For the long format, information for each job listed includes job number, number of files to transfer, user who spooled the job, number of bytes to send, type of command requested (S for sending files, R for receiving files, X for remote uucp), and file or command desired.

Several options are available:

- h           Print only the summary lines for each system. Summary lines give system name, number of jobs for the system, and total number of bytes to send.
- l           Specifies a long format listing. The default is to list only the job numbers sorted across the page.
- ssystem     Limit output to jobs for systems whose system names begin with system.
- uuser       Limit output to jobs for users whose login names begin with user.
- djobno     Delete job number jobno (as obtained from a previous uuq command) from the uucp queue. Only the UUCP Administrator is permitted to delete jobs.
- rsdir       Look for files in the spooling directory sdir instead of the default directory.
- bbaud       Use baud to compute the transfer time instead of the default 1200 baud.

## FILES

/usr/spool/uucp/	Default spool directory
/usr/spool/uucp/C./C.*	Control files
/usr/spool/uucp/Dhostname./D.*	Outgoing data files
/usr/spool/uucp/X./X.*	Outgoing execution files

## SEE ALSO

uucp(1C), uux(1C), uulog(1C), uusnap(8C)



## BUGS

No information is available on work requested by the remote machine.

The user who requests a remote uucp command is unknown.

Uuq -l can be horrendously slow.

## AUTHOR

Lou Salkind, New York University

## NAME

uusend - send a file to a remote host

## SYNOPSIS

uusend [ -m mode ] sourcefile sys1!sys2!...!remotefile

## DESCRIPTION

Uusend sends a file to a given location on a remote system. The system need not be directly connected to the local system, but a chain of uucp(1) links needs to connect the two systems.

If the -m option is specified, the mode of the file on the remote end will be taken from the octal number given. Otherwise, the mode of the input file will be used.

The sourcefile can be ``-', meaning to use the standard input. Both of these options are primarily intended for internal use of uusend.

The remotefile can include the ~userid syntax.

## DIAGNOSTICS

If anything goes wrong any further away than the first system down the line, you will never hear about it.

## SEE ALSO

uux(1), uucp(1), uuencode(1)

## BUGS

This command should not exist, since uucp should handle it.

All systems along the line must have the uusend command available and allow remote execution of it.

Some uucp systems have a bug where binary files cannot be the input to a uux command. If this bug exists in any system along the line, the file will show up severely munged.

## NAME

uux - unix to unix command execution

## SYNOPSIS

```
uux [-] [-cClLnprz] [-aname] [-ggrade] [-xdebug]
command-string
```

## DESCRIPTION

Uux will gather zero or more files from various systems, execute a command on a specified system and then send standard output to a file on a specified system.

The command-string is made up of one or more arguments that look like a Shell command line, except that the command and file names may be prefixed by system-name!. A null system-name is interpreted as the local system.

File names may be one of

- (1) a full path name;
- (2) a path name preceded by ~user where user is a login name on the specified system and is replaced by that user's login directory;
- (3) a path name prefixed by ~; where ~ is expanded to the system's public directory (usually /usr/spool/uucppublic);
- (4) a partial pathname, which is prefixed by the current directory.

As an example, the command

```
uux "!diff usg!/usr/dan/file1 pwba!/a4/dan/file2 >
!~/dan/file.diff"
```

will get the file1 and file2 files from the ``usg'' and ``pwba'' machines, execute a diff(1) command and put the results in file.diff in the local /usr/spool/uucppublic/dan/ directory.

Any special shell characters, such as <>;|, should be quoted either by quoting the entire command-string, or quoting the special characters as individual arguments.

Uux will attempt to get all files to the execution system. For files that are output files, the file name must be escaped using parentheses. For example, the command

```
uux a!wc b!/usr/file1 \ (c!/usr/file2 \)
```

get /usr/file1 from system ``b'' and send it to system ``a'', perform a wc command on that file and send the result of the wc command to system ``c''.

Uux will notify you by mail if the requested command on the remote system was disallowed. This notification can be turned off by the -n option.

The following options are interpreted by uux:

- The standard input to uux is made the standard input to the command-string.
- aname  
    Use name as the user identification replacing the initiator user-id.
- c     Do not copy local file to the spool directory for transfer to the remote machine (this is the default).
- C     Force the copy of local files to the spool directory for transfer.
- ggrade  
    Grade is a single letter/number, from 0 to 9, A to Z, or a to z; 0 is the highest, and z is the lowest grade. The default is A; by comparison uucp(1) defaults to n and mail is usually sent at grade C. Lower grades should be specified for high-volume jobs, such as news.
- l     Try and make a link from the original file to the spool directory. If the link cannot be made, copy the file.
- n     Do not notify the user when the command completes.
- p     Same as -: The standard input to uux is made the standard input to the command-string.
- r     Do not start the file transfer, just queue the job.
- xdebug  
    Produce debugging output on stdout. The debug is a number between 0 and 9; higher numbers give more detailed information. Debugging is permitted only for privileged users (specifically, those with read access to L.sys(5)).
- z     Notify the user only if the command fails.
- L     Start up uucico with the -L flag. This will force calls to be made to local sites only (see uucico(8C)).

## FILES

/usr/spool/uucp        spool directories  
/etc/uucp/\*        UUCP configuration data.

## SEE ALSO

uucp(1), uucico(8), uuxqt(8).

## WARNING

For security reasons, many installations will limit the list of commands executable on behalf of an incoming request from uux. Many sites will permit little more than the receipt of mail (see mail(1)) via uux.

## BUGS

Only the first command of a shell pipeline may have a system-name!. All other commands are executed on the system of the first command.

The use of the shell metacharacter \* will probably not do what you want it to do.

The shell tokens << and >> are not implemented.

When invoking uux from csh(1), the `!' character must be prefixed by the `\' escape to inhibit csh's history mechanism. (Quotes are not sufficient.)

## NAME

vacation - return ``I am on vacation'' indication

## SYNOPSIS

```
vacation -I
vacation user
```

## DESCRIPTION

Vacation returns a message to the sender of a message telling that you are on vacation. The intended use is in a .forward file. For example, your .forward file might have:

```
\eric, "|vacation eric"
```

which would send messages to you (assuming your login name was eric) and send a message back to the sender.

Vacation expects a file .vacation.msg in your home directory containing a message to be sent back to each sender. It should be an entire message (including headers). For example, it might say:

```
From: eric@ucbmonet.Berkeley.EDU (Eric Allman)
Subject: I am on vacation
Delivered-By-The-Graces-Of: the Vacation program
```

```
I am on vacation until July 22. If you have something urgent,
please contact Joe Kalash <kalash@ucbingres.Berkeley.EDU>.
--eric
```

This message will only be sent once a week to each unique sender. The people who have sent you messages are kept in the files .vacation.pag and .vacation.dir in your home directory. The -I option initializes these files, and should be executed before you modify your .forward file.

If the -I flag is not specified, vacation reads the first line from the standard input for a UNIX-style ``From'' line to determine the sender. If this is not present, a nasty diagnostic is produced. Sendmail(8) includes the ``From'' line automatically.

No message is sent if the initial ``From'' line includes the string ``-REQUEST@'' or if a ``Precedence: bulk'' or ``Precedence: junk'' line is included in the header.

## SEE ALSO

sendmail(8)

## NAME

vgrind - grind nice listings of programs

## SYNOPSIS

vgrind [ -f ] [ - ] [ -t ] [ -n ] [ -x ] [ -W ] [ -sn ] [ -h  
header ] [ -d file ] [ -llanguage ] name ...

## DESCRIPTION

Vgrind formats the program sources which are arguments in a nice style using troff(1). Comments are placed in italics, keywords in bold face, and the name of the current function is listed down the margin of each page as it is encountered.

Vgrind runs in two basic modes, filter mode or regular mode. In filter mode vgrind acts as a filter in a manner similar to tbl(1). The standard input is passed directly to the standard output except for lines bracketed by the troff-like macros:

.vS - starts processing

.vE - ends processing

These lines are formatted as described above. The output from this filter can be passed to troff for output. There need be no particular ordering with eqn(1) or tbl(1).

In regular mode vgrind accepts input files, processes them, and passes them to troff(1) for output.

In both modes vgrind passes any lines beginning with a decimal point without conversion.

The options are:

-f forces filter mode

- forces input to be taken from standard input (default if -f is specified)

-t similar to the same option in troff causing formatted text to go to the standard output

-n forces no keyword bolding

-x outputs the index file in a ``pretty'' format. The index file itself is produced whenever vgrind is run with a file called index in the current directory. The index of function definitions can then be run off by giving vgrind the -x option and the file index as argument.

- W forces output to the (wide) Versatec printer rather than the (narrow) Varian
- s specifies a point size to use on output (exactly the same as the argument of a .ps)
- h specifies a particular header to put on every output page (default is the file name)
- d specifies an alternate language definitions file (default is /usr/share/misc/vgrindefs)
- l specifies the language to use. Currently known are PASCAL (-lp), MODEL (-lm), C (-lc or the default), CSH (-lcsh), SHELL (-lsh), RATFOR (-lrf), MODULA2 (-lmod2), YACC (-lyacc), ISP (-lisp), and ICON (-lI).

#### FILES

index file where source for index is created  
/usr/share/tmac/tmac.vgrindmacro package  
/usr/libexec/vfontedprpreprocessor  
/usr/share/misc/vgrindefslanguage descriptions

#### AUTHOR

Dave Presotto & William Joy

#### SEE ALSO

vlp(1), vtroff(1), vgrindefs(5)

#### BUGS

Vfontedpr assumes that a certain programming style is followed:

For C - function names can be preceded on a line only by spaces, tabs, or an asterisk. The parenthesized arguments must also be on the same line.

For PASCAL - function names need to appear on the same line as the keywords function or procedure.

For MODEL - function names need to appear on the same line as the keywords is beginproc.

If these conventions are not followed, the indexing and marginal function name comment mechanisms will fail.

More generally, arbitrary formatting styles for programs mostly look bad. The use of spaces to align source code fails miserably; if you plan to vgrind your program you should use tabs. This is somewhat inevitable since the font used by vgrind is variable width.



The mechanism of ctags in recognizing functions should be used here.

Filter mode does not work in documents using the -me or -ms macros. (So what use is it anyway?)

## NAME

vi - screen oriented (visual) display editor based on ex

## SYNOPSIS

vi [ -t tag ] [ -r ] [ +command ] [ -l ] [ -wn ] name ...

## DESCRIPTION

Vi (visual) is a display oriented text editor based on ex(1). Ex and vi run the same code; it is possible to get to the command mode of ex from within vi and vice-versa.

The Vi Quick Reference card and the Introduction to Display Editing with Vi provide full details on using vi.

## FILES

See ex(1).

## SEE ALSO

ex (1), edit (1), ``Vi Quick Reference'' card, ``An Introduction to Display Editing with Vi''.

## AUTHOR

William Joy  
Mark Horton added macros to visual mode and is maintaining version 3

## BUGS

Software tabs using ^T work only immediately after the autoindent.

Left and right shifts on intelligent terminals don't make use of insert and delete character operations in the terminal.

The wrapmargin option can be fooled since it looks at output columns when blanks are typed. If a long word passes through the margin and onto the next line without a break, then the line won't be broken.

Insert/delete within a line can be slow if tabs are present on intelligent terminals, since the terminals need help in doing this correctly.

Saving text on deletes in the named buffers is somewhat inefficient.

The source command does not work when executed as :source; there is no way to use the :append, :change, and :insert commands, since it is not possible to give more than one line of input to a : escape. To use these on a :global you must Q to ex command mode, execute them, and then reenter the screen editor with vi or open.

## NAME

vmstat - report virtual memory statistics (2.11BSD)

## SYNOPSIS

vmstat [ -fsip ] [ drives ] [ interval [ count ] ]

## DESCRIPTION

Vmstat delves into the system and normally reports certain statistics kept about process, virtual memory, disk, trap and cpu activity. If given a -f argument, it instead reports on the number of forks and vforks since system startup and the number of pages of virtual memory involved in each kind of fork. If given a -s argument, it instead prints the contents of the sum structure, giving the total number of several kinds of paging related events which have occurred since boot. If given a -i argument, it instead reports on the number of interrupts taken by each device since system startup.

If none of these options are given, vmstat will report in the first line a summary of the virtual memory activity since the system has been booted. If given a -p argument, a different set of information is shown with additional pdpl1 specific fields. If interval is specified, then successive lines are summaries over the last interval seconds.

`vmstat 5' will print what the system is doing every five seconds; this is a good choice of printing interval since this is how often some of the statistics are sampled in the system; others vary every second, running the output for a while will make it apparent which are recomputed every second. If a count is given, the statistics are repeated count times. The format fields are:

Procs: information about numbers of processes in various states.

r	in run queue
b	blocked for resources (i/o, paging, etc.)
w	runnable or short sleeper (< 20 secs) but swapped

Memory: information about the usage of virtual and real memory. Virtual pages are considered active if they belong to processes which are running or have run in the last 20 seconds. A ``page'' here is a disk block of 1024 bytes.

avm	active virtual memory
tx	fraction of active virtual memory used by text (-p only)
fre	size of the free list

Page: information about text table and swaping activity. These are averaged each five seconds, and given in units per second.

ti text table entries found in use/sticky  
tc text table entries found in cache  
pi pages swapped in  
po pages swapped out  
fr text table entries freed  
fc text table entries placed in cache  
ov user mode overlay changes

Swap: information about swaping activity (-p only)

i process swap in rate  
o process swap out rate

Disks: Disk I/O rates in kbytes/sec. The number under each of these is the unit number.

System: various rate averages per second over last 5 seconds.

pd pseudo-dma interrupts (-p only)  
in (non clock) device interrupts per second  
sy system calls per second  
tr traps/faults per second (-p only)  
ov user mode overlay changes (-p only)  
cs cpu context switch rate (switches/sec)

Cpu: breakdown of percentage usage of CPU time

us user time for normal processes (includes nice time if no -p)  
ni user time for low priority processes (-p only)  
sy system time  
id cpu idle

If more than 4 disk drives are configured in the system, vmstat displays only the first 4 drives, with priority given to Massbus disk drives (i.e. if both Unibus and Massbus drives are present and the total number of drives exceeds 4, then some number of Unibus drives will not be displayed in favor of the Massbus drives). To force vmstat to display specific drives, their names may be supplied on the command line.

#### FILES

/dev/kmem, /vmunix

#### SEE ALSO

systat(1), iostat(1)

The sections starting with ``Interpreting system activity'' in Installing and Operating 4.2bsd.

## NAME

vwidth - make troff width table for a font

## SYNOPSIS

```
vwidth fontfile pointsize > ftxx.c
cc -c ftxx.c mv ftxx.o /usr/share/font/ftxx
```

## DESCRIPTION

Vwidth translates from the width information stored in the vfont style format to the format expected by troff. Troff wants an object file in a.out(5) format. (This fact does not seem to be documented anywhere.) Troff should look directly in the font file but it doesn't.

Vwidth should be used after editing a font with fed(1). It is not necessary to use vwidth unless you have made a change that would affect the width tables. Such changes include numerically editing the width field, adding a new character, and moving or copying a character to a new position. It is not always necessary to use vwidth if the physical width of the glyph (e.g. the number of columns in the bit matrix) has changed, but if it has changed much the logical width should probably be changed and vwidth run.

Vwidth produces a C program on its standard output. This program should be run through the C compiler and the object (that is, the .o file) saved. The resulting file should be placed in /usr/share/font in the file ftxx where x is a one or two letter code that is the logical (internal to troff) font name. This name can be found by looking in the file /usr/share/fontinfo/fname\* where fname is the external name of the font.

## SEE ALSO

fed(1), vfont(5), troff(1), vtroff(1)

## BUGS

Produces the C file using obsolete syntax that the portable C compiler complains about.

## NAME

w, uptime - who is on and what they are doing; system time up

## SYNOPSIS

w [ -hswu ] [ user ]  
uptime

## DESCRIPTION

W prints a summary of the current activity on the system, including what each user is doing.

The uptime invocation prints only the header line.

The heading line shows the current time of day, how long the system has been up, the number of users logged into the system, and the load averages. The load average numbers give the number of jobs in the run queue averaged over 1, 5 and 15 minutes.

The fields output are: the users login name, the name of the tty the user is on, the time of day the user logged on, the number of minutes since the user last typed anything, the CPU time used by all processes and their children on that terminal, the CPU time used by the currently active processes, the name and arguments of the current process.

The -h flag suppresses the heading. The -s flag asks for a short form of output. In the short form, the tty is abbreviated, the login time and cpu times are left off, as are the arguments to commands.

The -w and -u flags force the w and uptime actions respectively, regardless of the name the program is invoked as.

If a user name is included, the output will be restricted to that user.

## FILES

/var/run/utmp	for login names
/dev/swap	secondary storage

## SEE ALSO

finger(1), ps(1), who(1)

## AUTHOR

Mark Horton

## BUGS

The notion of the ``current process'' is muddy. The current algorithm is ``the highest numbered process on the terminal that is not ignoring interrupts, or, if there is none, the

highest numbered process on the terminal''. This fails, for example, in critical sections of programs like the shell and editor, or when faulty programs running in the background fork and fail to ignore interrupts. (In cases where no process can be found, w prints ``-''.)

The CPU time is only an estimate, in particular, if someone leaves a background process running after logging out, the person currently on that terminal is ``charged'' with the time.

Background processes are not shown, even though they account for much of the load on the system.

Sometimes processes, typically those in the background, are printed with null or garbaged arguments. In these cases, the name of the command is printed in parentheses.

## NAME

wait - await completion of process

## SYNOPSIS

wait

## DESCRIPTION

Wait until all processes started with & have completed, and report on abnormal terminations.

Because the wait(2) system call must be executed in the parent process, the Shell itself executes wait, without creating a new process.

## SEE ALSO

sh(1)

## BUGS

Not all the processes of a 3- or more-stage pipeline are children of the Shell, and thus can't be waited for. (This bug does not apply to csh(1).)



## NAME

wall - write to all users

## SYNOPSIS

wall

## DESCRIPTION

Wall reads its standard input until an end-of-file. It then sends this message, preceded by 'Broadcast Message ...', to all logged in users.

The sender should be super-user to override any protections the users may have invoked.

## FILES

/dev/tty?  
/var/run/utmp

## SEE ALSO

mesg(1), write(1)

## DIAGNOSTICS

'Cannot send to ...' when the open on a user's tty file fails.

## NAME

wc - word count

## SYNOPSIS

wc [ -lwc ] [ name ... ]

## DESCRIPTION

Wc counts lines, words and characters in the named files, or in the standard input if no name appears. A word is a maximal string of characters delimited by spaces, tabs or newlines.

If an argument beginning with one of ``lwc'' is present, the specified counts (lines, words, or characters) are selected by the letters l, w, or c. The default is -lwc.

## BUGS

## NAME

what - show what versions of object modules were used to construct a file

## SYNOPSIS

what name ...

## DESCRIPTION

What reads each file and searches for sequences of the form ``@(#)' ' as inserted by the source code control system. It then prints the remainder of the string after this marker, up to a null character, newline, double quote, or ``>' ' character.

## BUGS

As SCCS is not licensed with UNIX/32V, this is a rewrite of the what command which is part of SCCS, and may not behave exactly the same as that command does.

## NAME

whatis - describe what a command is

## SYNOPSIS

whatis command ...

## DESCRIPTION

Whatis looks up a given command and gives the header line from the manual section. You can then run the man(1) command to get more information. If the line starts 'name(section) ...' you can do 'man section name' to get the documentation for it. Try 'whatis ed' and then you should do 'man 1 ed' to get the manual.

Whatis is actually just the -f option to the man(1) command.

## FILES

/usr/man/whatis      Data base

## SEE ALSO

man(1), catman(8)

## NAME

whereis - locate programs

## SYNOPSIS

whereis program ...

## DESCRIPTION

The whereis utility checks the standard binary directories for the specified programs, printing out the paths of any it finds.

The path searched is the string returned by the sysctl(8) utility for the ``user.cs\_path'' string.

## SEE ALSO

sysctl(8),

## COMPATIBILITY

The historic flags and arguments for the whereis utility are no longer available in this version.

## HISTORY

The whereis command appeared in 3.0BSD.

## NAME

which - locate a program file including aliases and paths  
(csh only)

## SYNOPSIS

which [ name ] ...

## DESCRIPTION

Which takes a list of names and looks for the files which would be executed had these names been given as commands. Each argument is expanded if it is aliased, and searched for along the user's path. Both aliases and path are taken from the user's .cshrc file.

## FILES

~/.cshrc source of aliases and path values

## DIAGNOSTICS

A diagnostic is given for names which are aliased to more than a single word, or if an executable file with the argument name was not found in the path.

## BUGS

Must be executed by a csh, since only csh's know about aliases.

## NAME

who - who is on the system

## SYNOPSIS

who [ who-file ] [ am I ]

## DESCRIPTION

Who, without an argument, lists the login name, terminal name, and login time for each current UNIX user.

Without an argument, who examines the /var/run/utmp file to obtain its information. If a file is given, that file is examined. Typically the given file will be /usr/adm/wtmp, which contains a record of all the logins since it was created. Then who lists logins, logouts, and crashes since the creation of the wtmp file. Each login is listed with user name, terminal name (with `/dev/' suppressed), and date and time. When an argument is given, logouts produce a similar line without a user name. Reboots produce a line with `x' in the place of the device name, and a fossil time indicative of when the system went down.

With two arguments, as in `who am I' (and also `who are you'), who tells who you are logged in as.

## FILES

/var/run/utmp

## SEE ALSO

getuid(2), utmp(5)

## NAME

whoami - display effective user id

## SYNOPSIS

whoami

## DESCRIPTION

The whoami utility has been obsoleted by the id(1) utility, and is equivalent to ``id -un''. The command ``id -p'' is suggested for normal interactive use.

The whoami utility displays your effective user ID as a name.

The whoami utility exits 0 on success, and >0 if an error occurs.

## SEE ALSO

id(1)



## NAME

whois - DARPA Internet user name directory service

## SYNOPSIS

whois name

## DESCRIPTION

whois help

Produces a helpful message similar to the following:

Please enter a name or a handle (``ident"), such as ``Smith" or ``SRI-NIC". Starting with a period forces a name-only search; starting with exclamation point forces handle-only. Examples:

Smith [looks for name or handle SMITH ]

!SRI-NIC [looks for handle SRI-NIC only ]

.Smith, John [looks for name JOHN SMITH only ]

Adding ``..." to the argument will match anything from that point, e.g. ``ZU..." will match ZUL, ZUM, etc.

To have the ENTIRE membership list of a group or organization, if you are asking about a group or org, shown with the record, use an asterisk character `\*' directly preceding the given argument. [CAUTION: If there are a lot of members this will take a long time!] You may of course use exclamation point and asterisk, or a period and asterisk together.

## SEE ALSO

RFC 812: Nicname/Whois

## NAME

window - window environment

## SYNOPSIS

```
window [-t] [-f] [-d] [-e escape-char] [-c command]
```

## DESCRIPTION

Window implements a window environment on ASCII terminals.

A window is a rectangular portion of the physical terminal screen associated with a set of processes. Its size and position can be changed by the user at any time. Processes communicate with their window in the same way they normally interact with a terminal--through their standard input, output, and diagnostic file descriptors. The window program handles the details of redirecting input and output to and from the windows. At any one time, only one window can receive input from the keyboard, but all windows can simultaneously send output to the display.

Windows can overlap and are framed as necessary. Each window is named by one of the digits ``1'' to ``9''. This one character identifier, as well as a user definable label string, are displayed with the window on the top edge of its frame. A window can be designated to be in the foreground, in which case it will always be on top of all normal, non-foreground windows, and can be covered only by other foreground windows. A window need not be completely within the edges of the terminal screen. Thus a large window (possibly larger than the screen) may be positioned to show only a portion of its full size.

Each window has a cursor and a set of control functions. Most intelligent terminal operations such as line and character deletion and insertion are supported. Display modes such as underlining and reverse video are available if they are supported by the terminal. In addition, similar to terminals with multiple pages of memory, each window has a text buffer which can have more lines than the window itself.

## OPTIONS

When window starts up, the commands (see long commands below) contained in the file .windowrc in the user's home directory are executed. If it does not exist, two equal sized windows spanning the terminal screen are created by default.

The command line options are

-t Turn on terse mode (see terse command below).

- f Fast. Don't perform any startup action.
- d Ignore .windowrc and create the two default windows instead.
- e escape-char  
Set the escape character to escape-char. Escape-char can be a single character, or in the form ^X where X is any character, meaning control-X.
- c command  
Execute the string command as a long command (see below) before doing anything else.

## PROCESS ENVIRONMENT

With each newly created window, a shell program is spawned with its process environment tailored to that window. Its standard input, output, and diagnostic file descriptors are bound to one end of either a pseudo-terminal (pty (4)) or a UNIX domain socket (socketpair (4)). If a pseudo-terminal is used, then its special characters and modes (see stty (1)) are copied from the physical terminal. A termcap (5) entry tailored to this window is created and passed as environment (environ (5)) variable TERMCAP. The termcap entry contains the window's size and characteristics as well as information from the physical terminal, such as the existence of underline, reverse video, and other display modes, and the codes produced by the terminal's function keys, if any. In addition, the window size attributes of the pseudo-terminal are set to reflect the size of this window, and updated whenever it is changed by the user. In particular, the editor vi (1) uses this information to redraw its display.

## OPERATION

During normal execution, window can be in one of two states: conversation mode and command mode. In conversation mode, the terminal's real cursor is placed at the cursor position of a particular window--called the current window--and input from the keyboard is sent to the process in that window. The current window is always on top of all other windows, except those in foreground. In addition, it is set apart by highlighting its identifier and label in reverse video.

Typing window's escape character (normally ^P) in conversation mode switches it into command mode. In command mode, the top line of the terminal screen becomes the command prompt window, and window interprets input from the keyboard as commands to manipulate windows.

There are two types of commands: short commands are usually one or two key strokes; long commands are strings either

typed by the user in the command window (see the ``:'' command below), or read from a file (see source below).

#### SHORT COMMANDS

Below, # represents one of the digits ``1'' to ``9'' corresponding to the windows 1 to 9. ^X means control-X, where X is any character. In particular, ^^ is control-^. Escape is the escape key, or ^[.

#        Select window # as the current window and return to conversation mode.

%#      Select window # but stay in command mode.

^^      Select the previous window and return to conversation mode. This is useful for toggling between two windows.

escape  
Return to conversation mode.

^P      Return to conversation mode and write ^P to the current window. Thus, typing two ^P's in conversation mode sends one to the current window. If the window escape is changed to some other character, that character takes the place of ^P here.

?        List a short summary of commands.

^L      Redraw the screen.

q        Exit window. Confirmation is requested.

^Z      Suspend window.

w        Create a new window. The user is prompted for the positions of the upper left and lower right corners of the window. The cursor is placed on the screen and the keys ``h'', ``j'', ``k'', and ``l'' move the cursor left, down, up, and right, respectively. The keys ``H'', ``J'', ``K'', and ``L'' move the cursor to the respective limits of the screen. Typing a number before the movement keys repeats the movement that number of times. Return enters the cursor position as the upper left corner of the window. The lower right corner is entered in the same manner. During this process, the placement of the new window is indicated by a rectangular box drawn on the screen, corresponding to where the new window will be framed. Typing escape at any point cancels this command.

This window becomes the current window, and is given the first available ID. The default buffer size is

used (see `nline` command below).

Only fully visible windows can be created this way.

`c#` Close window `#`. The process in the window is sent the hangup signal (see `kill (1)`). `Csh (1)` should handle this signal correctly and cause no problems.

`m#` Move window `#` to another location. A box in the shape of the window is drawn on the screen to indicate the new position of the window, and the same keys as those for the `w` command are used to position the box. The window can be moved partially off-screen.

`M#` Move window `#` to its previous position.

`s#` Change the size of window `#`. The user is prompted to enter the new lower right corner of the window. A box is drawn to indicate the new window size. The same keys used in `w` and `m` are used to enter the position.

`S#` Change window `#` to its previous size.

`^Y` Scroll the current window up by one line.

`^E` Scroll the current window down by one line.

`^U` Scroll the current window up by half the window size.

`^D` Scroll the current window down by half the window size.

`^B` Scroll the current window up by the full window size.

`^F` Scroll the current window down by the full window size.

`h` Move the cursor of the current window left by one column.

`j` Move the cursor of the current window down by one line.

`k` Move the cursor of the current window up by one line.

`l` Move the cursor of the current window right by one column.

`^S` Stop output in the current window.

`^Q` Start output in the current window.

`:` Enter a line to be executed as long commands. Normal line editing characters (erase character, erase word, erase line) are supported.

## LONG COMMANDS

Long commands are a sequence of statements parsed much like a programming language, with a syntax similar to that of C. Numeric and string expressions and variables are supported, as well as conditional statements.

There are two data types: string and number. A string is a sequence of letters or digits beginning with a letter. ``\_`` and ``.`` are considered letters. Alternately, non-alphanumeric characters can be included in strings by quoting them in ``"`` or escaping them with ``\``. In addition, the ``\`` sequences of C are supported, both inside and outside quotes (e.g., ``\n`` is a new line, ``\r`` a carriage return). For example, these are legal strings: abcde01234, "&#\$^\*&#", ab"\$#cd, ab\\$#cd, "/usr/ucb/window".

A number is an integer value in one of three forms: a decimal number, an octal number preceded by ``0``, or a hexadecimal number preceded by ``0x`` or ``0X``. The natural machine integer size is used (i.e., the signed integer type of the C compiler). As in C, a non-zero number represents a boolean true.

The character ``#`` begins a comment which terminates at the end of the line.

A statement is either a conditional or an expression. Expression statements are terminated with a new line or ``;``. To continue an expression on the next line, terminate the first line with ``\``.

## CONDITIONAL STATEMENT

Window has a single control structure: the fully bracketed if statement in the form

```
if <expr> then
 <statement>
 . . .
elsif <expr> then
 <statement>
 . . .
else
 <statement>
 . . .
endif
```

The else and elsif parts are optional, and the latter can be repeated any number of times. <Expr> must be numeric.

## EXPRESSIONS

Expressions in window are similar to those in the C language, with most C operators supported on numeric operands. In addition, some are overloaded to operate on strings.

When an expression is used as a statement, its value is discarded after evaluation. Therefore, only expressions with side effects (assignments and function calls) are useful as statements.

Single valued (no arrays) variables are supported, of both numeric and string values. Some variables are predefined. They are listed below.

The operators in order of increasing precedence:

`<expr1> = <expr2>`

Assignment. The variable of name `<expr1>`, which must be string valued, is assigned the result of `<expr2>`. Returns the value of `<expr2>`.

`<expr1> ? <expr2> : <expr3>`

Returns the value of `<expr2>` if `<expr1>` evaluates true (non-zero numeric value); returns the value of `<expr3>` otherwise. Only one of `<expr2>` and `<expr3>` is evaluated. `<Expr1>` must be numeric.

`<expr1> || <expr2>`

Logical or. Numeric values only. Short circuit evaluation is supported (i.e., if `<expr1>` evaluates true, then `<expr2>` is not evaluated).

`<expr1> && <expr2>`

Logical and with short circuit evaluation. Numeric values only.

`<expr1> | <expr2>`

Bitwise or. Numeric values only.

`<expr1> ^ <expr2>`

Bitwise exclusive or. Numeric values only.

`<expr1> & <expr2>`

Bitwise and. Numeric values only.

`<expr1> == <expr2>, <expr1> != <expr2>`

Comparison (equal and not equal, respectively). The boolean result (either 1 or 0) of the comparison is returned. The operands can be numeric or string valued. One string operand forces the other to be converted to a string in necessary.

`<expr1> < <expr2>, <expr1> > <expr2>, <expr1> <= <expr2>, <expr1> >= <expr2>`

Less than, greater than, less than or equal to, greater than or equal to. Both numeric and string values, with automatic conversion as above.

`<expr1> << <expr2>, <expr1> >> <expr2>`

If both operands are numbers, `<expr1>` is bit shifted left (or right) by `<expr2>` bits. If `<expr1>` is a string, then its first (or last) `<expr2>` characters are returned (if `<expr2>` is also a string, then its length is used in place of its value).

`<expr1> + <expr2>, <expr1> - <expr2>`

Addition and subtraction on numbers. For ```+''`, if one argument is a string, then the other is converted to a string, and the result is the concatenation of the two strings.

`<expr1> * <expr2>, <expr1> / <expr2>, <expr1> % <expr2>`

Multiplication, division, modulo. Numbers only.

`-<expr>, ~<expr>, !<expr>, $<expr>, $?<expr>`

The first three are unary minus, bitwise complement and logical complement on numbers only. The operator, ```$'`, takes `<expr>` and returns the value of the variable of that name. If `<expr>` is numeric with value `n` and it appears within an alias macro (see below), then it refers to the `n`th argument of the alias invocation. ```$?'` tests for the existence of the variable `<expr>`, and returns 1 if it exists or 0 otherwise.

`<expr>(<arglist>)`

Function call. `<Expr>` must be a string that is the unique prefix of the name of a builtin window function or the full name of a user defined alias macro. In the case of a builtin function, `<arglist>` can be in one of two forms:

`<expr1>, <expr2>, . . .`

`argname1 = <expr1>, argname2 = <expr2>, . . .`

The two forms can in fact be intermixed, but the result is unpredictable. Most arguments can be omitted; default values will be supplied for them. The argnames can be unique prefixes of the the argument names. The commas separating arguments are used only to disambiguate, and can usually be omitted.

Only the first argument form is valid for user defined aliases. Aliases are defined using the alias builtin function (see below). Arguments are accessed via a variant of the variable mechanism (see ```$'` operator above).

Most functions return value, but some are used for side effect only and so must be used as statements. When a function or an alias is used as a statement, the parenthesis surrounding the argument list may be omitted. Aliases return no value.



## BUILTIN FUNCTIONS

The arguments are listed by name in their natural order. Optional arguments are in square brackets (``[ ]'`). Arguments that have no names are in angle brackets (``< >'`).

`alias([<string>], [<string-list>])`

If no argument is given, all currently defined alias macros are listed. Otherwise, `<string>` is defined as an alias, with expansion `<string-list>`. The previous definition of `<string>`, if any, is returned. Default for `<string-list>` is no change.

`close(<window-list>)`

Close the windows specified in `<window-list>`. If `<window-list>` is the word `all`, then all windows are closed. No value is returned.

`cursormodes([modes])`

Set the window cursor to `modes`. `Modes` is the bitwise or of the mode bits defined as the variables `m_ul` (underline), `m_rev` (reverse video), `m_blk` (blinking), and `m_grp` (graphics, terminal dependent). Return value is the previous modes. Default is no change. For example, `cursor($m_rev|$m_blk)` sets the window cursors to blinking reverse video.

`echo([window], [<string-list>])`

Write the list of strings, `<string-list>`, to `window`, separated by spaces and terminated with a new line. The strings are only displayed in the window, the processes in the window are not involved (see `write` below). No value is returned. Default is the current window.

`escape([escapec])`

Set the escape character to `escape-char`. Returns the old escape character as a one character string. Default is no change. `Escapec` can be a string of a single character, or in the form `^X`, meaning control-`X`.

`foreground([window], [flag])`

Move window in or out of foreground. `Flag` can be one of `on`, `off`, `yes`, `no`, `true`, or `false`, with obvious meanings, or it can be a numeric expression, in which case a non-zero value is true. Returns the old foreground flag as a number. Default for `window` is the current window, default for `flag` is no change.

`label([window], [label])`

Set the label of window to `label`. Returns the old label as a string. Default for `window` is the current window, default for `label` is no change. To turn off a

label, set it to an empty string ("").

list()  
No arguments. List the identifiers and labels of all windows. No value is returned.

nline([nline])  
Set the default buffer size to nline. Initially, it is 48 lines. Returns the old default buffer size. Default is no change. Using a very large buffer can slow the program down considerably.

select([window])  
Make window the current window. The previous current window is returned. Default is no change.

shell(<string-list>))  
Set the default window shell program to <string-list>. Returns the first string in the old shell setting. Default is no change. Initially, the default shell is taken from the environment variable SHELL.

source(filename)  
Read and execute the long commands in filename. Returns -1 if the file cannot be read, 0 otherwise.

terse([flag])  
Set terse mode to flag. In terse mode, the command window stays hidden even in command mode, and errors are reported by sounding the terminal's bell. Flag can take on the same values as in foreground above. Returns the old terse flag. Default is no change.

unalias(alias)  
Undefine alias. Returns -1 if alias does not exist, 0 otherwise.

unset(variable)  
Undefine variable. Returns -1 if variable does not exist, 0 otherwise.

variables()  
No arguments. List all variables. No value is returned.

window([row], [column], [nrow], [ncol], [nline], [frame], [pty], [mapnl], [shell])  
Open a window with upper left corner at row, column and size nrow, ncol. If nline is specified, then that many lines are allocated for the text buffer. Otherwise, the default buffer size is used. Default values for row, column, nrow, and ncol are, respectively, the

upper, left-most, lower, or right-most extremes of the screen. Frame, pty, and mapnl are flag values interpreted in the same way as the argument to foreground (see above); they mean, respectively, put a frame around this window (default true), allocate pseudo-terminal for this window rather than socketpair (default true), and map new line characters in this window to carriage return and line feed (default true if socketpair is used, false otherwise). Shell is a list of strings that will be used as the shell program to place in the window (default is the program specified by shell, see below). The created window's identifier is returned as a number.

write([window], [<string-list>])

Send the list of strings, <string-list>, to window, separated by spaces but not terminated with a new line. The strings are actually given to the window as input. No value is returned. Default is the current window.

#### PREDEFINED VARIABLES

These variables are for information only. Redefining them does not affect the internal operation of window.

baud The baud rate as a number between 50 and 38400.

modes

The display modes (reverse video, underline, blinking, graphics) supported by the physical terminal. The value of modes is the bitwise or of some of the one bit values, m\_blk, m\_grp, m\_rev, and m\_ul (see below). These values are useful in setting the window cursors' modes (see cursormodes above).

m\_blk

The blinking mode bit.

m\_grp

The graphics mode bit (not very useful).

m\_rev

The reverse video mode bit.

m\_ul

The underline mode bit.

ncol The number of columns on the physical screen.

nrow The number of rows on the physical screen.

term The terminal type. The standard name, found in the second name field of the terminal's TERMCAP entry, is used.

## FILES

~/.windowrc      startup command file.  
/dev/[pt]ty[pq]?pseudo-terminal devices.

## DIAGNOSTICS

Should be self explanatory.

## BUGS

## NAME

write - write to another user

## SYNOPSIS

write user [ ttyname ]

## DESCRIPTION

Write copies lines from your terminal to that of another user. When first called, it sends the message

Message from yourname@yoursystem on yourttyname at time...

The recipient of the message should write back at this point. Communication continues until an end of file is read from the terminal or an interrupt is sent. At that point write writes `EOT' on the other terminal and exits.

If you want to write to a user who is logged in more than once, the ttyname argument may be used to indicate the appropriate terminal name.

Permission to write may be denied or granted by use of the mesg command. At the outset writing is allowed. Certain commands, in particular nroff and pr(1) disallow messages in order to prevent messy output.

If the character `!' is found at the beginning of a line, write calls the shell to execute the rest of the line as a command.

The following protocol is suggested for using write: when you first write to another user, wait for him to write back before starting to send. Each party should end each message with a distinctive signal-(o) for `over' is conventional-that the other may reply. (oo) for `over and out' is suggested when conversation is about to be terminated.

## FILES

/var/run/utmp to find user  
/bin/sh to execute `!'

## SEE ALSO

mesg(1), who(1), mail(1)

## NAME

xstr - extract strings from C programs to implement shared strings

## SYNOPSIS

xstr [ -c ] [ - ] [ file ]

## DESCRIPTION

Xstr maintains a file strings into which strings in component parts of a large program are hashed. These strings are replaced with references to this common area. This serves to implement shared constant strings, most useful if they are also read-only.

The command

xstr -c name

will extract the strings from the C source in name, replacing string references by expressions of the form (&xstr[number]) for some number. An appropriate declaration of xstr is prepended to the file. The resulting C text is placed in the file x.c, to then be compiled. The strings from this file are placed in the strings data base if they are not there already. Repeated strings and strings which are suffices of existing strings do not cause changes to the data base.

After all components of a large program have been compiled a file xs.c declaring the common xstr space can be created by a command of the form

xstr

This xs.c file should then be compiled and loaded with the rest of the program. If possible, the array can be made read-only (shared) saving space and swap overhead.

Xstr can also be used on a single file. A command

xstr name

creates files x.c and xs.c as before, without using or affecting any strings file in the same directory.

It may be useful to run xstr after the C preprocessor if any macro definitions yield strings or if there is conditional code which contains strings which may not, in fact, be needed. Xstr reads from its standard input when the argument '-' is given. An appropriate command sequence for running xstr after the C preprocessor is:

```
cc -E name.c | xstr -c -
cc -c x.c
mv x.o name.o
```

Xstr does not touch the file strings unless new items are added, thus make can avoid remaking xs.o unless truly necessary.

#### FILES

strings	Data base of strings
x.c	Massaged C source
xs.c	C source for definition of array `xstr'
/tmp/xs*	Temp file when `xstr name' doesn't touch strings

#### SEE ALSO

mkstr(1)

#### BUGS

If a string is a suffix of another string in the data base, but the shorter string is seen first by xstr both strings will be placed in the data base, when just placing the longer one there will do.

## NAME

yacc - yet another compiler-compiler

## SYNOPSIS

yacc [ -vd ] grammar

## DESCRIPTION

Yacc converts a context-free grammar into a set of tables for a simple automaton which executes an LR(1) parsing algorithm. The grammar may be ambiguous; specified precedence rules are used to break ambiguities.

The output file, `y.tab.c`, must be compiled by the C compiler to produce a program `yyparse`. This program must be loaded with the lexical analyzer program, `yylex`, as well as `main` and `yyerror`, an error handling routine. These routines must be supplied by the user; `Lex(1)` is useful for creating lexical analyzers usable by yacc.

If the `-v` flag is given, the file `y.output` is prepared, which contains a description of the parsing tables and a report on conflicts generated by ambiguities in the grammar.

If the `-d` flag is used, the file `y.tab.h` is generated with the define statements that associate the yacc-assigned 'token codes' with the user-declared 'token names'. This allows source files other than `y.tab.c` to access the token codes.

## FILES

`y.output`  
`y.tab.c`  
`y.tab.h` defines for token names  
`yacc.tmp`, `yacc.acts` temporary files  
`/usr/share/misc/yaccparparser` prototype for C programs

## SEE ALSO

`lex(1)`  
LR Parsing by A. V. Aho and S. C. Johnson, Computing Surveys, June, 1974.  
YACC - Yet Another Compiler Compiler by S. C. Johnson.

## DIAGNOSTICS

The number of reduce-reduce and shift-reduce conflicts is reported on the standard output; a more detailed report is found in the `y.output` file. Similarly, if some rules are not reachable from the start symbol, this is also reported.

## BUGS

Because file names are fixed, at most one yacc process can be active in a given directory at a time.



## NAME

yes - be repetitively affirmative

## SYNOPSIS

yes [ expletive ]

## DESCRIPTION

Yes repeatedly outputs "y", or if expletive is given, that is output repeatedly. Termination is by rubout.