

BSD Unix 2.11 man entries  
6/9/2019 - The ShadowTron Blog

Generated using the simh PDP-11/70 emulator with the PiDP11 Front Panel

PiDP11 - <https://obsolescence.wixsite.com/obsolescence/pidp-11>  
SimH - <http://simh.trailing-edge.com/>  
ShadowTronBlog - [https://www.youtube.com/channel/UCtUiwjYcRS\\_u6Egc8iTkHNg](https://www.youtube.com/channel/UCtUiwjYcRS_u6Egc8iTkHNg)  
<http://shadowtron.com>  
[shadowtronblog@gmail.com](mailto:shadowtronblog@gmail.com)

Manual Area covered

=====

1 Commands and Application Programs  
2 System Calls  
3 C Library Subroutines  
3F Fortran Library  
--> 4 Special Files  
5 File Formats  
6 Games  
7 Miscellaneous  
8 System Maintenance

\*\*\*\*\*

\*\*\*\* Manual 4 - Special Files \*\*\*\*

\*\*\*\*\*

intro	introduction to special files and hardware support
acc	ACC LH/DH IMP interface
ad	Data Translation A/D converter
arp	Address Resolution Protocol
autoconf	diagnostics from the autoconfiguration code
bk	line discipline for machine-machine communication (obsolete)
cons	VAX-11 console interface
crl	VAX 8600 console RL02 interface
css	DEC IMP-11A LH/DH IMP interface
ct	phototypesetter interface
ddn	DDN Standard Mode X.25 IMP interface
de	DEC DEUNA 10 Mb/s Ethernet interface
dh	DH-11/DM-11 communications multiplexer
dhu	DHU-11 communications multiplexer
dmc	DEC DMC-11/DMR-11 point-to-point communications device
dmf	DMF-32, terminal multiplexor
dmz	DMZ-32 terminal multiplexor
dn	DN-11 autocall unit interface
drum	paging device
dz	DZ-11 communications multiplexer
ec	3Com 10 Mb/s Ethernet interface
en	Xerox 3 Mb/s Ethernet interface
ex	Excelan 10 Mb/s Ethernet interface
fl	console floppy interface
hdh	ACC IF-11/HDH IMP interface
hk	RK6-11/RK06 and RK07 moving head disk
hp	MASSBUS disk interface
ht	TM-03/TE-16,TU-45,TU-77 MASSBUS magtape interface
hy	Network Systems Hyperchannel interface
icmp	Internet Control Message Protocol
idp	Xerox Internet Datagram Protocol
ik	Ikonas frame buffer, graphics device interface
il	Interlan NI1010 10 Mb/s Ethernet interface
imp	1822 network interface
imp	IMP raw socket interface
inet	Internet protocol family
ip	Internet Protocol
ix	Interlan Np100 10 Mb/s Ethernet interface
kg	KL-11/DL-11W line clock
lo	software loopback network interface
lp	line printer
mem	main memory
mt	TM78/TU-78 MASSBUS magtape interface
mtio	UNIX magtape interface
np	Interlan Np100 10 Mb/s Ethernet interface
ns	Xerox Network Systems(tm) protocol family
nsip	software network interface encapsulating ns packets in ip packets.
null	data sink
pcl	DEC CSS PCL-11 B Network Interface
ps	Evans and Sutherland Picture System 2 graphics device interface
pty	pseudo terminal driver
qe	DEC DEQNA Q-bus 10 Mb/s Ethernet interface
rx	DEC RX02 floppy disk interface
spp	Xerox Sequenced Packet Protocol

tb	line discipline for digitizing devices
tcp	Internet Transmission Control Protocol
tm	TM-11/TE-10 magtape interface
tmscp	DEC TMSCP magtape interface
ts	TS-11 magtape interface
tty	general terminal interface
tu	VAX-11/730 and VAX-11/750 TU58 console cassette interface
uda	UDA-50 disk controller interface
udp	Internet User Datagram Protocol
up	unibus storage module controller/drives
ut	UNIBUS TU45 tri-density tape drive interface
uu	TU58/DECTape II UNIBUS cassette interface
va	Benson-Varian interface
vp	Versatec interface
vv	Proteon proNET 10 Megabit ring

## NAME

intro - introduction to special files and hardware support

## DESCRIPTION

This section describes the special files, related driver functions, and networking support available in the system. In this part of the manual, the SYNOPSIS section of each configurable device gives a sample specification for use in constructing a system description for the /sys/conf/configscript, the autoconfig(8), program and describes the major and minor device numbers and their encoding. The DIAGNOSTICS section lists messages which may appear on the console and/or in the system error log /usr/adm/messages due to errors in device operation.

This section contains both devices which may be configured into the system, ``4'' entries, and network related information, ``4N'', ``4P'', and ``4F'' entries; The networking support is introduced in intro(4N).

## PDP DEVICE SUPPORT

This section describes the hardware supported on the DEC PDP-11. Software support for these devices comes in two forms. A hardware device may be supported with a character or block device driver, or it may be used within the networking subsystem and have a network interface driver. Block and character devices are accessed through files in the file system of a special type; c.f. mknod(8). Network interfaces are indirectly accessed through the interprocess communication facilities provided by the system; see socket(2).

A hardware device is identified to the system at configuration time and the appropriate device or network interface driver is then compiled into the system. When the resultant system is booted, the autoconfiguration facilities in the system probe for the device on either the UNIBUS (or Q-bus) or MASSBUS and, if found, enable the software support for it. If a device does not respond at autoconfiguration time it is not accessible at any time afterwards. To enable a device which did not autoconfigure, the system will have to be rebooted.

The autoconfiguration system is described in autoconfig(8). A list of the supported devices is given below.

## SEE ALSO

intro(4), networking(4), config(8).

## LIST OF DEVICES

The devices listed below are supported in this incarnation of the system. Pseudo-devices are not listed. Listed also

are devices which are in various stages of porting to 2.11BSD from 4.3BSD.

Devices are indicated by their functional interface. If second vendor products provide functionally identical interfaces they should be usable with the supplied software. (Beware, however, that we promise the software works ONLY with the hardware indicated on the appropriate manual page.) Occasionally, new devices of a similar type may be added simply by creating appropriate table entries in the driver.

The following are known to work:

acc	ACC LH/DH IMP communications interface
de	DEC DEUNA 10Mb/s Ethernet controller
dh	DH-11 emulators, terminal multiplexor
dhu	DHU-11 terminal multiplexor
dz	DZ-11 terminal multiplexor
ec	3Com 10Mb/s Ethernet controller
hk	RK6-11/RK06 and RK07 moving head disk
ht	TM03 MASSBUS tape drive interface (with TE-16, TU-45, TU-77)
il	Interlan 1010, 1010A, 2010A 10Mb/s Ethernet controller
lp	LP-11 parallel line printer interface
qe	DEC DEQNA Q-bus 10 Mb/s Ethernet interface
ra	DEC UDA-50, RQDX, KLESI disk controllers
rk	DEC RK05 disk controller
rl	DEC RL-11 disk controller
rx	DEC RX02 floppy interface
si	SI 9500 disk controller
tm	TM-11/TE-10 tape drive interface
tmscp	TMSCP-compatible tape controllers (e.g., TU81, TK50)
ts	TS-11 tape drive interface
vv	Proteon proNET 10Mb/s and 80Mb/s ring network interface
xp	General purpose SMD disk controller

The following should work:

dr	DR-11W general purpose DMA UNIBUS interface
----	---

The following worked in the past but will probably require work:

css	DEC IMP-11A communications interface
dmc	DEC DMC-11/DMR-11 point-to-point communications device
en	Xerox 3Mb/s Ethernet controller (obsolete)
sri	DR-11C IMP interface

It should be possible to port these from 4.3BSD:

ex	Excelan 10Mb/s Ethernet controller
ix	Interlan NP-100 10Mb/s Ethernet controller
np	Interlan NP-100 10Mb/s Ethernet controller (intelligent mode)
pcl	DEC PCL-11 communications interface

No idea whether any of the following could be ported:

ad	Data translation A/D interface
ct	C/A/T or APS phototypesetter
ddn	ACC ACP625 DDN Standard Mode X.25 IMP interface
dmf	DEC DMF-32 terminal multiplexor and parallel printer interface
dmz	DEC DMZ-32 terminal multiplexor
dn	DEC DN-11 autodialer interface
hdh	ACC IF-11/HDH IMP interface
hp	MASSBUS disk interface (with RP06, RM03, RM05, etc.)
hy	DR-11B or GI-13 interface to an NSC Hyperchannel
kg	KL-11/DL-11W line clock
mt	TM78 MASSBUS tape drive interface
tu	VAX-11/730 TU58 console cassette interface
un	DR-11W interface to Ungermann-Bass
up	Emulex SC-21V, SC-31 UNIBUS disk controller
ut	UNIBUS TU-45 tape drive interface
uu	TU58 dual cassette drive interface (DL11)
va	Benson-Varian printer/plotter interface
vp	Versatec printer/plotter interface

## NAME

acc - ACC LH/DH IMP interface

## SYNOPSIS

```
/sys/conf/SYSTEM:
  NACC 0      # ACC LH/DH ARPAnet IMP interface
  PLI  YES    # LH/DH is connected to a PLI
```

## DESCRIPTION

The acc device provides a Local Host/Distant Host interface to an IMP. It is normally used when participating in the DARPA Internet. The controller itself is not accessible to users, but instead provides the hardware support to the IMP interface described in imp(4). When configuring, the imp(NIMP) pseudo-device must also be included.

## DIAGNOSTICS

acc%d: not alive. The initialization routine was entered even though the device did not autoconfigure. This indicates a system problem.

acc%d: can't initialize. Insufficient UNIBUS resources existed to initialize the device. This is likely to occur when the device is run on a buffered data path on an 11/750 and other network interfaces are also configured to use buffered data paths, or when it is configured to use buffered data paths on an 11/730 (which has none).

acc%d: imp doesn't respond, icsr=%b. The driver attempted to initialize the device, but the IMP failed to respond after 500 tries. Check the cabling.

acc%d: stray xmit interrupt, csr=%b. An interrupt occurred when no output had previously been started.

acc%d: output error, ocsr=%b, icsr=%b. The device indicated a problem sending data on output.

acc%d: input error, csr=%b. The device indicated a problem receiving data on input.

acc%d: bad length=%d. An input operation resulted in a data transfer of less than 0 or more than 1008 bytes of data into memory (according to the word count register). This should never happen as the maximum size of a host-IMP message is 1008 bytes.

## NAME

arp - Address Resolution Protocol

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NETHER    1      # ether pseudo-device
```

## DESCRIPTION

ARP is a protocol used to dynamically map between DARPA Internet and 10Mb/s Ethernet addresses. It is used by all the 10Mb/s Ethernet interface drivers. It is not specific to Internet protocols or to 10Mb/s Ethernet, but this implementation currently supports only that combination.

ARP caches Internet-Ethernet address mappings. When an interface requests a mapping for an address not in the cache, ARP queues the message which requires the mapping and broadcasts a message on the associated network requesting the address mapping. If a response is provided, the new mapping is cached and any pending message is transmitted. ARP will queue at most one packet while waiting for a mapping request to be responded to; only the most recently ``transmitted'' packet is kept.

To facilitate communications with systems which do not use ARP, ioctls are provided to enter and delete entries in the Internet-to-Ethernet tables. Usage:

```
#include <sys/ioctl.h>
#include <sys/socket.h>
#include <net/if.h>
struct arpreq arpreq;
```

```
ioctl(s, SIOCSARP, (caddr_t)&arpreq);
ioctl(s, SIOCGARP, (caddr_t)&arpreq);
ioctl(s, SIOCDAEP, (caddr_t)&arpreq);
```

Each ioctl takes the same structure as an argument. SIOCSARP sets an ARP entry, SIOCGARP gets an ARP entry, and SIOCDAEP deletes an ARP entry. These ioctls may be applied to any socket descriptor s, but only by the super-user. The arpreq structure contains:

```
/*
 * ARP ioctl request
 */
struct arpreq {
    struct sockaddr  arp_pa; /* protocol address */
    struct sockaddr  arp_ha; /* hardware address */
    int              arp_flags; /* flags */
};
/* arp_flags field values */
#define ATF_COM      0x02 /* completed entry (arp_ha valid) */
```



```
#define ATF_PERM      0x04      /* permanent entry */
#define ATF_PUBL      0x08      /* publish (respond for other host) */
#define ATF_USETRAILERS 0x10      /* send trailer packets to host */
```

The address family for the `arp_pa` `sockaddr` must be `AF_INET`; for the `arp_ha` `sockaddr` it must be `AF_UNSPEC`. The only flag bits which may be written are `ATF_PERM`, `ATF_PUBL` and `ATF_USETRAILERS`. `ATF_PERM` causes the entry to be permanent if the `ioctl` call succeeds. The peculiar nature of the ARP tables may cause the `ioctl` to fail if more than 8 (permanent) Internet host addresses hash to the same slot. `ATF_PUBL` specifies that the ARP code should respond to ARP requests for the indicated host coming from other machines. This allows a host to act as an ``ARP server,' ' which may be useful in convincing an ARP-only machine to talk to a non-ARP machine.

ARP is also used to negotiate the use of trailer IP encapsulations; trailers are an alternate encapsulation used to allow efficient packet alignment for large packets despite variable-sized headers. Hosts which wish to receive trailer encapsulations so indicate by sending gratuitous ARP translation replies along with replies to IP requests; they are also sent in reply to IP translation replies. The negotiation is thus fully symmetrical, in that either or both hosts may request trailers. The `ATF_USETRAILERS` flag is used to record the receipt of such a reply, and enables the transmission of trailer packets to that host.

ARP watches passively for hosts impersonating the local host (i.e. a host which responds to an ARP mapping request for the local host's address).

#### DIAGNOSTICS

duplicate IP address!! sent from ethernet address:  
%x:%x:%x:%x:%x:%x. ARP has discovered another host on the local network which responds to mapping requests for its own Internet address.

#### SEE ALSO

`ec(4)`, `de(4)`, `il(4)`, `inet(4F)`, `arp(8C)`, `ifconfig(8C)`  
``An Ethernet Address Resolution Protocol,' ' RFC826, Dave Plummer, Network Information Center, SRI.  
``Trailer Encapsulations,' ' RFC893, S.J. Leffler and M.J. Karels, Network Information Center, SRI.

#### BUGS

ARP packets on the Ethernet use only 42 bytes of data; however, the smallest legal Ethernet packet is 60 bytes (not including CRC). Some systems may not enforce the minimum packet size, others will.

## NAME

bk - line discipline for machine-machine communication  
(obsolete)

## SYNOPSIS

Not currently supported under 2.11BSD

## DESCRIPTION

This line discipline provides a replacement for the old and new tty drivers described in tty(4) when high speed output to and especially input from another machine is to be transmitted over a asynchronous communications line. The discipline was designed for use by the Berkeley network. It may be suitable for uploading of data from microprocessors into the system. If you are going to send data over asynchronous communications lines at high speed into the system, you must use this discipline, as the system otherwise may detect high input data rates on terminal lines and disables the lines; in any case the processing of such data when normal terminal mechanisms are involved saturates the system.

The line discipline is enabled by a sequence:

```
#include <sgtty.h>
int ldisc = NETLDISC, fildes; ...
ioctl(fildes, TIOCSETD, &ldisc);
```

A typical application program then reads a sequence of lines from the terminal port, checking header and sequencing information on each line and acknowledging receipt of each line to the sender, who then transmits another line of data. Typically several hundred bytes of data and a smaller amount of control information will be received on each handshake.

The old standard teletype discipline can be restored by doing:

```
ldisc = OTTYDISC;
ioctl(fildes, TIOCSETD, &ldisc);
```

While in networked mode, normal teletype output functions take place. Thus, if an 8 bit output data path is desired, it is necessary to prepare the output line by putting it into RAW mode using ioctl(2). This must be done before changing the discipline with TIOCSETD, as most ioctl(2) calls are disabled while in network line-discipline mode.

When in network mode, input processing is very limited to reduce overhead. Currently the input path is only 7 bits wide, with newline the only recognized character, terminating an input record. Each input record must be read and acknowledged before the next input is read as the system

refuses to accept any new data when there is a record in the buffer. The buffer is limited in length, but the system guarantees to always be willing to accept input resulting in 512 data characters and then the terminating newline.

User level programs should provide sequencing and checksums on the information to guarantee accurate data transfer.

#### SEE ALSO

tty(4)

#### DIAGNOSTICS

None.

#### BUGS

The Purdue uploading line discipline, which provides 8 bits and uses timeout's to terminate uploading should be incorporated into the standard system, as it is much more suitable for microprocessor connections.

## NAME

br - EATON BR1537/BR1711 1538[A,B,C,D] moving head disk

## SYNOPSIS

/sys/conf/SYSTEM:

NBR br\_drives # EATON 1537/1711, EATON 1538A, B, C, D

/etc/dtab:

#Name	Unit#	Addr	Vector	Br	Handler(s)	# Comments
br	?176710		254	5	brintr	# Eaton 1537/8

major device number(s):

raw: 20

block: 11

minor device encoding:

bits 0007 specify partition of BR drive

bits 0070 specify BR drive

## DESCRIPTION

Files with minor device numbers 0 through 7 refer to various portions of drive 0; minor devices 8 through 15 refer to drive 1, etc. There are four drive types supported by the Eaton BR1537 and BR1711 controllers, these are the 1538A (50 Mb), 1538B (80 Mb), 1538C (200 Mb) and 1538D (300 Mb). Capacities are unformatted megabytes. The standard device names begin with ``br'' followed by the drive number and then a letter a-h for partitions 0-7 respectively. The character ? stands here for a drive number in the range 0-7.

The block files access the disk via the system's normal buffering mechanism and may be read and written without regard to physical disk records. There is also a `raw' interface which provides for direct transmission between the disk and the user's read or write buffer. A single read or write call results in exactly one I/O operation and therefore raw I/O is considerably more efficient when many words are transmitted. The names of the raw files conventionally begin with an extra `r.'

In raw I/O the buffer must begin on a word (even) boundary, and counts should be a multiple of 512 bytes (a disk sector). Likewise seek calls should specify a multiple of 512 bytes.

## DISK SUPPORT

The size (in sectors) of the pseudo-disks on each drive are as follows:

## 1538A partitions:

disk	length	cyls	comments
br?a	18260	0 - 165	1538A has 22 sec/trk, 5 trk/cyl
br?b	12210	166 - 276	
br?c	59180	277 - 814	
br?d	unused		
br?e	unused		
br?f	unused		
br?g	unused		
br?h	89650	0 - 814	

## 1538B partitions:

disk	length	cyls	comments
br?a	18400	0 - 114	1538B has 32 sec/trk, 5 trk/cyl
br?b	12320	115 - 190	
br?c	99840	191 - 814	
br?d	unused		
br?e	unused		
br?f	unused		
br?g	unused		
br?h	99840	0 - 814	

## 1538C partitions:

disk	length	cyls	comments
br?a	18392	0 - 43	1538C has 22 sec/trk, 19 trk/cyl
br?b	12122	44 - 72	
br?c	231990	73 - 627	
br?d	78166	628 - 814	
br?e	unused		
br?f	unused		
br?g	unused		
br?h	340670	0 - 814	

## 1538D partitions:

disk	length	cyls	comments
br?a	18240	0 - 29	1538D has 32 sec/trk, 19 trk/cyl
br?b	12160	30 - 49	
br?c	232256	50 - 431	
br?d	232256	432 - 813	
br?e	unused		
br?f	unused		
br?g	unused		
br?h	495520	0 - 814	

## FILES

/dev/br[0-7][a-h] block files  
 /dev/rbr[0-7][a-h] raw files  
 /dev/MAKEDEV script to create special files  
 /dev/MAKEDEV.local script to localize special files

## SEE ALSO

ra(4), ram(4), rk(4), rl(4), rx(4), si(4), xp(4), dtab(5),

autoconfig(8)

#### DIAGNOSTICS

br%d%c: hard error sn%d cs2=%b ds=%b er=%b. An unrecoverable error occurred during transfer of the specified sector of the specified disk partition. The contents of the cs2, ds and er registers are printed in octal and symbolically with bits decoded. The error was either unrecoverable, or a large number of retry attempts (including offset positioning and drive recalibration) could not recover the error.

#### BUGS

In raw I/O read and write(2) truncate file offsets to 512-byte block boundaries, and write scribbles on the tail of incomplete blocks. Thus, in programs that are likely to access raw devices, read, write and lseek(2) should always deal in 512-byte multiples.

DEC-standard error logging should be supported.

A program to analyze the logged error information (even in its present reduced form) is needed.

The partition tables for the file systems should be read off of each pack, as they are never quite what any single installation would prefer, and this would make packs more portable.

Only the 1538D (300Mb) disk has been used with this driver, the disktab(5) file and newfs(8) program only know about the 1538D.

## NAME

cons - PDP-11 console interface

## DESCRIPTION

The console is available to the processor through the console registers. It acts like a normal terminal, except that typing control-P on the consoles of some PDP-11's puts the console in local console mode with an ODT (console debugger mode) prompt. The operation of the console in this mode varies per-processor. See the PDP-11 Processor Handbook for your particular processor.

On an 11/44 the console can be put into local mode by typing a control-P. The ODT prompt is ``>>>'. The processor is not stopped by entering local console mode. The CPU may be halted with the ``halt' command, which may be abbreviated to ``h.' Conversational mode is re-entered by using the command ``c' (continue).

On 11/73 the processor's front panel halt switch must be pressed to put the console into ODT mode. The ODT prompt is an ``@'. No characters typed on the console are special while UNIX is in control of the console.

## FILES

/dev/console

## SEE ALSO

tty(4), reboot(8)  
PDP-11 Hardware Handbook specific to hardware

## NAME

css - DEC IMP-11A LH/DH IMP interface

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NCSS css_controllers      # DEC/CSS IMP11-A ARPAnet IMP interface
```

## DESCRIPTION

The css device provides a Local Host/Distant Host interface to an IMP. It is normally used when participating in the DARPA Internet. The controller itself is not accessible to users, but instead provides the hardware support to the IMP interface described in imp(4). When configuring, the imp(NIMP) pseudo-device is also included.

## DIAGNOSTICS

css%d: not alive. The initialization routine was entered even though the device did not autoconfigure. This indicates a system problem.

css%d: can't initialize. Insufficient UNIBUS resources existed to initialize the device. This is likely to occur when the device is run on a buffered data path on an 11/750 and other network interfaces are also configured to use buffered data paths, or when it is configured to use buffered data paths on an 11/730 (which has none).

css%d: imp doesn't respond, icsr=%b. The driver attempted to initialize the device, but the IMP failed to respond after 500 tries. Check the cabling.

css%d: stray output interrupt csr=%b. An interrupt occurred when no output had previously been started.

css%d: output error, ocsr=%b icsr=%b. The device indicated a problem sending data on output.

css%d: recv error, csr=%b. The device indicated a problem receiving data on input.

css%d: bad length=%d. An input operation resulted in a data transfer of less than 0 or more than 1008 bytes of data into memory (according to the word count register). This should never happen as the maximum size of a host-IMP message is 1008 bytes.



## NAME

de - DEC DEUNA 10 Mb/s Ethernet interface

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NDE  de_controllers # DEUNA
```

## DESCRIPTION

The de interface provides access to a 10 Mb/s Ethernet network through a Digital Equipment UNIBUS Network Adapter (DEUNA).

Each of the host's network addresses is specified at boot time with an SIOCSIFADDR ioctl. The de interface employs the address resolution protocol described in arp(4P) to dynamically map between Internet and Ethernet addresses on the local network.

The interface normally tries to use a ``trailer'' encapsulation to minimize copying data on input and output. The use of trailers is negotiated with ARP. This negotiation may be disabled, on a per-interface basis, by setting the IFF\_NOTRAILERS flag with an SIOCSIFFLAGS ioctl.

## DIAGNOSTICS

de%d: hardware address %s. This is a normal autoconfiguration message noting the 6 byte physical ethernet address of the adapter.

de%d: oerror, flags=%b tdrerr=%b (len=%d). The hardware indicated an error in transmitting a packet to the cable. The status and error flags are reported.

de%d: ierror, flags=%b lenerr=%b (len=%d). The hardware indicated an error in reading a packet from the cable. The status and error flags are reported.

de%d: can't handle af%d. The interface was handed a message with addresses formatted in an unsuitable address family; the packet was dropped.

de%d: buffer unavailable. The interface received more packets than it had buffers allocated to receive them.

de%d: address change failed, csr0=%b csr1=%b. The interface was unable to reprogram its physical ethernet address. This may happen with very early models of the interface. This facility is used only when the controller is not the first network interface configured for XNS.

The following messages indicate a probable hardware error performing the indicated operation during autoconfiguration

or initialization. The two control and status registers should indicate the nature of the failure. See the hardware manual for details.

de%d: reset failed, csr0=%b csr1=%b.

de%d: ppcb failed, csr0=%b csr1=%b.

de%d: read addr failed, csr0=%b csr1=%b.

de%d: wtring failed, csr0=%b csr1=%b.

de%d: wtmode failed, csr0=%b csr1=%b.

SEE ALSO

intro(4N), inet(4F), arp(4P)

## NAME

dh - DH-11/DM-11 communications multiplexer

## SYNOPSIS

```
/sys/conf/SYSTEM:
```

```
NDH  dh_units  # DH11; NDH is in units of boards (16 each)
```

```
NDM  dm_units  # DM11; NDM is in units of boards (16 each)
```

```
/etc/dtab:
```

#Name	Unit#	Addr	Vector	Br	Handler(s)	# Comments
dm	?	170500	310	4	dmintr	# dm11 dh modem control
dh	?	160020	320	5	dhrintr dhxint	# dh11 terminal mux

```
major device number(s):
```

```
raw: 3
```

```
minor device encoding:
```

```
bits 0017 specify line on DH unit
```

```
bits 0060 specify DH unit
```

```
bit 0100 specify RTS/CTS ('`hardware'`) flowcontrol
```

```
bit 0200 specifies non-blocking open ('`CD always on'`')
```

## DESCRIPTION

A dh-11 provides 16 communication lines; dm-11's may be optionally paired with dh-11's to provide modem control for the lines. If there are fewer dm-11's than dh-11's, the dm-11's are assumed to be paired with the first dm\_units dh-11's.

Each line attached to the DH-11 communications multiplexer behaves as described in tty(4). Input and output for each line may independently be set to run at any of 16 speeds; see tty(4) for the encoding.

Bit 0200 of the minor device number for DH lines may be set to say that a line is not properly connected, and that the line should be treated as hard-wired with carrier always present. Thus creating the special character device node "3, 130" via mknod /dev/ttyh2 c 3 130 would cause line ttyh2 to be treated in this way. Bit 0100 of the minor device number enables RTS/CTS (also called "`hardware'`) flow control. It is enabled by adding 64 to the minor device number: mknod /dev/ttyh2 c 3 194

The dh driver monitors the rate of input on each board, and switches between the use of character-at-a-time interrupts and input silos. While the silo is enabled during periods of high-speed input, the driver polls for input 30 times per second.

## FILES

```
/dev/tty[h-k][0-9a-f]
```

```
/dev/ttyd[0-9a-f]
```

/dev/MAKEDEV script to create special files  
/dev/MAKEDEV.local script to localize special files

SEE ALSO

tty(4), dtab(5), autoconfig(8)

DIAGNOSTICS

dh%d NXM. No response from UNIBUS on a dma transfer within a timeout period. This has never been observed on a PDP-11 and is a carryover from the VAX driver when it was ported. It is not serious.

dh%d %d overruns. The character input silo overflowed before it could be serviced. This message is only printed at line close time. It is not serious but does indicate that the system was not able to keep up with the data flow.

## NAME

dhu - DHU-11 communications multiplexer

## SYNOPSIS

```
/sys/conf/SYSTEM:
```

```
NDHU dhu_units # DHU11
```

```
/etc/dtab:
```

```
#Name Unit# Addr Vector Br Handler(s) # Comments
du ? 160020 310 5 dhurint dhuxint # dhull terminal mux
```

```
major device number(s):
```

```
raw: 4
```

```
minor device encoding:
```

```
bits 0017 specify line on DHU unit
```

```
bits 0060 specify DHU unit
```

```
bit 0100 specify RTS/CTS flow control
```

```
bit 0200 specifies non-blocking open ('`CD always on''')
```

## DESCRIPTION

A DHU-11 provides 16 communication lines.

Each line attached to the DHU-11 communications multiplexer behaves as described in tty(4). Input and output for each line may independently be set to run at any of 13 speeds (50, 200 and 38400 baud are not available); see tty(4) for the encoding.

Bit 0200 of the minor device number for DHU lines may be set to say that a line is not properly connected, and that the line should be treated as hard-wired with carrier always present. Thus creating the special character device node "4, 130" via mknod /dev/ttyS2 c 4 130 would cause line ttyS2 to be treated in this way. Turning on bit 6 (adding 64) to the minor device number via mknod /dev/ttyS2 c 4 194 enables RTS/CTS flow control.

The DHU-11 driver normally uses input silos and delays receiver interrupts by 20 milliseconds rather than taking an interrupt on each input character.

## FILES

```
/dev/tty[S-Z][0-9a-f]
```

## SEE ALSO

```
tty(4)
```

## DIAGNOSTICS

dhu(%d,%d) NXM. No response from UNIBUS on a DMA transfer within a timeout period. This has never been observed on a PDP-11 and is believed to be a carryover from the VAX driver when it was ported. It is not serious.

dhu%d %d overruns. The character input silo overflowed before it could be serviced. This message is printed only at line close time rather than on each overrun error. Kernel printf's are not interrupt driven and caused more overruns by blocking interrupts for lengthy periods of time.

#### NOTES

The driver currently does not make full use of the hardware capabilities of the DHU-11, for dealing with XON/XOFF flow-control or hard-wired lines for example.

The 4 40-way cables are a pain.

## NAME

dhv - DHV-11 communications multiplexer

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NDHV dhv_units # DHV11

/etc/dtab:
    #Name Unit# Addr   Vector Br Handler(s)      # Comments
    dhv ?      160440 310    5   dhvrint dhvxint # dhv terminal mux

major device number(s):
    raw: 24

minor device encoding:
    bits 0007 specify line on DHV unit
    bits 0070 specify DHV unit
    bit 0100 specifies RTS/CTS ('`hardware'`) flowcontrol
    bit 0200 specifies non-blocking open ('`CD always on'`')
```

## DESCRIPTION

A DHV-11 provides 8 communication lines.

Each line attached to the DHV-11 communications multiplexer behaves as described in tty(4). Input and output for each line may independently be set to run at any of 13 speeds (50 and 200 baud are not available). While 38400 is available the underlying hardware is not fast enough to handle it and there will be pauses/gaps between characters.

Bit 0200 of the minor device number for DHV lines may be set to say that a line is not properly connected, and that the line should be treated as hard-wired with carrier always present. Thus creating the special character device node "4, 130" via `mknod /dev/ttyS2 c 4 130` would cause line ttyS2 to be treated in this way. Turning on bit 6 in the minor device number via `mknod /dev/ttyS2 c 4 194` would enable RTS/CTS flow control.

The DHV-11 has an input silo but does not have the programmable receiver delay that the DHU (and DHQ) have. Thus system services more interrupts (i.e. gets fewer characters per interrupt on average) with a DHV-11 than with a DHQ (in DHU mode).

## FILES

/dev/tty[S-Z][0-9a-f]

## SEE ALSO

tty(4)

## DIAGNOSTICS

dhv%d,%d NXM. No response from QBUS on a DMA transfer

within a timeout period. This error has never been observed on a PDP-11 and is a carryover from the VAX driver when that was ported to 2BSD.

dhv%d diag %o. Diagnostic information from the DHV11. This has never been observed. The DHV-11 hardware manual will be required to decode the value printed out.

dhv%d: %d overruns. The character input silo overflowed before it could be serviced. This message is printed only when the line is closed. By only printing this when the line is closed further silo overruns are avoided (kernel printf statements are not interrupt driven).

#### NOTES

The DHV lacks the receiver delay that the DHU (and DHQ) have. Thus it is extremely easy (indeed it's almost certain at higher data rates) for a DHV-11 to enter an interrupt per character mode and use 70-80% of the cpu.



## NAME

dmc - DEC DMC-11/DMR-11 point-to-point communications device

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NDMC dmc_controllers      # DMC11
```

## DESCRIPTION

The dmc interface provides access to a point-to-point communications device which runs at either 1 Mb/s or 56 Kb/s. DMC-11's communicate using the DEC DDCMP link layer protocol.

The dmc interface driver also supports a DEC DMR-11 providing point-to-point communication running at data rates from 2.4 Kb/s to 1 Mb/s. DMR-11's are a more recent design and thus are preferred over DMC-11's. The NXMT and NRCV constants in the driver should be increased in this case, as the DMR can accept up to 64 transmit and receive buffers, as opposed to 7 for the DMC.

The configuration flags specify how to set up the device,

- 0 -- full duplex DDCMP (normal mode)
- 1 -- DDCMP Maintenance mode (generally useless)
- 2 -- DDCMP Half Duplex, primary station
- 3 -- DDCMP Half Duplex, secondary station

Several device error counters are available via "adb", for more information see the adb script /usr/share/adb/dmcstats, or the DMC11 technical manual.

The host's address must be specified with an SIOCSIFADDR ioctl, and the destination address specified with a SIOC-SIFDSTADDR ioctl, before the interface will transmit or receive any packets.

## ROUTING

The driver places a HOST entry in the kernel routing tables for the address given in the SIOCSIFDSTADDR ioctl. To use the DMC as a link between local nets, the route to the remote net must be added manually with the route(8) command, or by the use of the routing process routed(8) on each end of the link.

## DIAGNOSTICS

dmc%d: bad control %o. A bad parameter was passed to the dmcload routine.

dmc%d: unknown address type %d. An input packet was received which contained a type of address unknown to the driver.

DMC fatal error 0%. A fatal error in DDMCP occurred, causing the device to be restarted.

DMC soft error 0%. A non-fatal error in DDMCP has occurred.

dmc%d: af%d not supported. The interface was handed a message which has addresses formatted in an unsuitable address family.

#### SEE ALSO

intro(4N), inet(4F)

#### BUGS

The current version of the driver uses a link-level encapsulation so that multiple protocol types may be used. It is thus incompatible with earlier drivers, including the 4.2BSD version.

## NAME

dr - DR11-W general purpose interface driver

## SYNOPSIS

```
/sys/conf/SYSTEM:
```

```
    NDR    dr11-w_units    # DR11-W
```

```
/etc/dtab:
```

#Name	Unit#	Addr	Vector	Br	Handler(s)	# Comments
dr	?	172410	124	5	drintr	# dr11-w (b-bus emulator)
dr	?	172430	130	5	drintr	# dr11-w (gc-bus emulator)
dr	?	172450	134	5	drintr	# dr11-w (em-2 interface)

```
major device number(s):
```

```
    raw: 12
```

```
minor device encoding:
```

```
    bits 0007 specify DR unit
```

## DESCRIPTION

The Digital Equipment 1981-82 Peripherals Handbook states the following about the DR11-W:

The The DR11-W is a general purpose, direct memory access (DMA) interface to the PDP-11 UNIBUS or VAX UNIBUS. The DR11-W moves data directly between memory and the UNIBUS to and from the user's peripheral.

It features:

- + Word or byte transfers.
- + Programmed or direct memory access (DMA) block transfers.
- + Burst data transfers.
- + User-controlled transfer rates up to memory speed.

The DR11-W is a 53-line direct memory access (DMA) interface to the PDP-11 UNIBUS or VAX UBA, which allows the user to control data transfers between the host processor and a peripheral. The DR11-W has 32 data lines (for transferring 16-bit parallel data) and 21 control lines (for transferring control and status information).

When used as an interprocessor buffer (IPB), the DR11-W allows data transfers between two processors. Inter-processor communication is accomplished by attaching one DR11-W to each processor UNIBUS or UBA and then cabling the two DR11-Ws together.

There. Now you know as much about it as I do.

A number of `ioctl(2)` calls apply to the dr devices, and have the form

```
#include <pdpuba/drreg.h>
ioctl(fildes, code, arg)
int *arg;
```

The applicable codes are:

DRGTTY	Get DR11-W status.
DRSTTY	Set flags and function.
DRSFUN	Set function.
DRSFLAG	Set flags.
DRGCSR	Get CSR and WCR.
DRSSIG	Set signal for ATTN interrupt.
DRESET	Reset DR11-W interface.
DRSTIME	Set timeout.
DRCTIME	Set timeout inactive.
DROUTPUT	Put word in output data register.
DRINPUT	Get word from input data register.
DRITIME	Don't ignore errors on timeout.

#### FILES

```
/dev/dr[0-7] device special files
/dev/MAKEDEV script to create special files
/dev/MAKEDEV.local script to localize special files
```

#### SEE ALSO

`dtab(5)`, `autoconfig(8)`

#### DIAGNOSTICS

```
dr%d: error csr=%b, eir=%b

dr%d: timeout error
```

#### BUGS

This interface is only available under 2.9BSD and 2.11BSD.  
No documentation exists on how to use it.

## NAME

dz - DZ-11 communications multiplexer

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NDZ  dz_units  # DZ11; NDZ is in units of boards (8 each)

/etc/dtab:
    #Name Unit# Addr  Vector Br Handler(s)      # Comments
    dz  ?      160100 310    5  dzrint  dzdma    # dz11 terminal mux

major device number:
    raw: 2

minor device encoding:
    bits 0007 specify line on DZ unit
    bits 0170 specify DZ unit
    bit 0200 specifies non-blocking open ('`CD always on')
```

## DESCRIPTION

A DZ11 provides 8 communication lines with partial modem control, adequate for UNIX dialup use. Each line attached to the DZ11 communications multiplexer behaves as described in tty(4) and may be set to run at any of 16 speeds; see tty(4) for the encoding.

Bit 0200 of the minor device number for DZ lines may be set to say that a line is not properly connected, and that the line should be treated as hard-wired with carrier always present. Thus creating the special character device node "2, 130" via `mknod /dev/tty02 c 2 130` would cause line tty02 to be treated in this way.

The dz driver monitors the rate of input on each board, and switches between the use of character-at-a-time interrupts and input silos. While the silo is enabled during periods of high-speed input, the driver polls for input 30 times per second.

## FILES

```
/dev/tty[0-9][0-9]
/dev/ttyd[0-9a-f]  dialups
/dev/MAKEDEV  script to create special files
/dev/MAKEDEV.local  script to localize special files
```

## SEE ALSO

tty(4), dtab(5), autoconfig(8)

## DIAGNOSTICS

dz%d: silo overflow. The 64 character input silo overflowed before it could be serviced. This can happen if a hard error occurs when the CPU is running with elevated priority, as the system will then print a message on the console with

interrupts disabled. It is not serious.

## NAME

ec - 3Com 10 Mb/s Ethernet interface

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NEC  ec_controllers # 3Com Ethernet
```

## DESCRIPTION

The ec interface provides access to a 10 Mb/s Ethernet network through a 3com controller.

The hardware has 32 kilobytes of dual-ported memory on the UNIBUS. This memory is used for internal buffering by the board, and the interface code reads the buffer contents directly through the UNIBUS. The address of this memory is given in the flags field in the configuration file. The first interface normally has its memory at Unibus address 0.

Each of the host's network addresses is specified at boot time with an SIOCSIFADDR ioctl. The ec interface employs the address resolution protocol described in arp(4P) to dynamically map between Internet and Ethernet addresses on the local network.

The interface normally tries to use a ``trailer'' encapsulation to minimize copying data on input and output. The use of trailers is negotiated with ARP. This negotiation may be disabled, on a per-interface basis, by setting the IFF\_NOTRAILERS flag with an SIOCSIFFLAGS ioctl.

The interface software implements an exponential backoff algorithm when notified of a collision on the cable. This algorithm utilizes a 16-bit mask and the VAX-11's interval timer in calculating a series of random backoff values. The algorithm is as follows:

1. Initialize the mask to be all 1's.
2. If the mask is zero, 16 retries have been made and we give up.
3. Shift the mask left one bit and formulate a backoff by masking the interval timer with the smaller of the complement of this mask and a 5-bit mask, resulting in a pseudo-random number between 0 and 31. This produces the number of slot times to delay, where a slot is 51 microseconds.
4. Use the value calculated in step 3 to delay before retransmitting the packet. The delay is done in a software busy loop.

## DIAGNOSTICS

ec%d: send error. After 16 retransmissions using the exponential backoff algorithm described above, the packet was dropped.

ec%d: input error (offset=%d). The hardware indicated an error in reading a packet off the cable or an illegally sized packet. The buffer offset value is printed for debugging purposes.

ec%d: can't handle af%d. The interface was handed a message with addresses formatted in an unsuitable address family; the packet was dropped.

## SEE ALSO

intro(4N), inet(4F), arp(4P)

## BUGS

The hardware is not capable of talking to itself. The software implements local sending and broadcast by sending such packets to the loop interface. This is a kludge.

Backoff delays are done in a software busy loop. This can degrade the system if the network experiences frequent collisions.



## NAME

en - Xerox 3 Mb/s Ethernet interface

## SYNOPSIS

```
/sys/conf/SYSTEM:  
    NEN  en_controllers # Xerox prototype (3 Mb) Ethernet
```

## DESCRIPTION

The en interface provides access to a 3 Mb/s Ethernet network. Due to limitations in the hardware, DMA transfers to and from the network must take place in the lower 64K bytes of the UNIBUS address space, and thus this must be among the first UNIBUS devices enabled after boot.

Each of the host's network addresses is specified at boot time with an SIOCSIFADDR ioctl. The station address is discovered by probing the on-board Ethernet address registers, and is used to verify the protocol addresses. No packets will be sent or accepted until a network address is supplied.

The interface software implements an exponential backoff algorithm when notified of a collision on the cable. This algorithm utilizes a 16-bit mask and the VAX-11's interval timer in calculating a series of random backoff values. The algorithm is as follows:

1. Initialize the mask to be all 1's.
2. If the mask is zero, 16 retries have been made and we give up.
3. Shift the mask left one bit and formulate a backoff by masking the interval timer with the mask (this is actually the two's complement of the value).
4. Use the value calculated in step 3 to delay before retransmitting the packet.

The interface handles both Internet and NS protocol families. It normally tries to use a ``trailer'' encapsulation to minimize copying data on input and output. The use of trailers is negotiated with ARP. This negotiation may be disabled, on a per-interface basis, by setting the IFF\_NOTRAILERS flag with an SIOCSIFFLAGS ioctl.

## DIAGNOSTICS

en%d: output error. The hardware indicated an error on the previous transmission.

en%d: send error. After 16 retransmissions using the exponential backoff algorithm described above, the packet

was dropped.

en%d: input error. The hardware indicated an error in reading a packet off the cable.

en%d: can't handle af%d. The interface was handed a message with addresses formatted in an unsuitable address family; the packet was dropped.

#### SEE ALSO

intro(4N), inet(4F)

#### BUGS

The device has insufficient buffering to handle back to back packets. This makes use in a production environment painful.

The hardware does word at a time DMA without byte swapping. To compensate, byte swapping of user data must either be done by the user or by the system. A kludge to byte swap only IP packets is provided if the ENF\_SWABIPS flag is defined in the driver and set at boot time with an SIOCSIF-FLAGS ioctl.

## NAME

fd, stdin, stdout, stderr file descriptor files

## DESCRIPTION

The files /dev/fd/0 through /dev/fd/# refer to file descriptors which can be accessed through the file system. If the file descriptor is open and the mode the file is being opened with is a subset of the mode of the existing descriptor, the call:

```
fd = open("/dev/fd/0", mode);
```

and the call:

```
fd = fcntl(0, F_DUPFD, 0);
```

are equivalent.

Opening the files /dev/stdin, /dev/stdout and /dev/stderr is equivalent to the following calls:

```
fd = fcntl(STDIN_FILENO, F_DUPFD, 0);
fd = fcntl(STDOUT_FILENO, F_DUPFD, 0);
fd = fcntl(STDERR_FILENO, F_DUPFD, 0);
```

Flags to the open(2) call other than O\_RDONLY, O\_WRONLY and O\_RDWR are ignored.

## FILES

```
/dev/fd/#
/dev/stdin
/dev/stdout
/dev/stderr
```

## SEE ALSO

tty(4)

## NAME

hk - RK6-11/RK06 and RK07 moving head disk

## SYNOPSIS

```
/sys/conf/SYSTEM:
```

```
NHK  hk_drives # RK611, RK06/07
```

```
/etc/dtab:
```

#Name	Unit#	Addr	Vector	Br	Handler(s)	# Comments
hk	?	177440	210	5	hkintr	# rk611/711 rk06/07

```
major device number(s):
```

```
raw: 13
```

```
block: 4
```

```
minor device encoding:
```

```
bits 0007 specify partition of HK drive
```

```
bits 0070 specify HK drive
```

## DESCRIPTION

Files with minor device numbers 0 through 7 refer to various portions of drive 0; minor devices 8 through 15 refer to drive 1, etc. The standard device names begin with ``hk'' followed by the drive number and then a letter a-h for partitions 0-7 respectively. The character ? stands here for a drive number in the range 0-7.

The block files access the disk via the system's normal buffering mechanism and may be read and written without regard to physical disk records. There is also a `raw' interface which provides for direct transmission between the disk and the user's read or write buffer. A single read or write call results in exactly one I/O operation and therefore raw I/O is considerably more efficient when many words are transmitted. The names of the raw files conventionally begin with an extra `r.'

In raw I/O the buffer must begin on a word (even) boundary, and counts should be a multiple of 512 bytes (a disk sector). Likewise seek calls should specify a multiple of 512 bytes.

## DISK SUPPORT

Disks must be labeled using either the standalone disklabel program on the boot tape or with the disklabel(8) program. There are no partition tables coded into the hk drivers, these must be placed on the drive with disklabel.

Traditionally the hk?a partition is used for the root filesystem, the b partition as a swap area and the c partition for disk to disk copying (it spans the entire disk).

The kernel uses the c partition to access the bad block information stored at the end of some packs. Extreme care must be taken when creating file systems on this partition to avoid overwriting any bad block information present. User data should use the h partition which should be at least one track (one cylinder might be best) shorter than the c partition.

#### FILES

/dev/hk[0-7][a-h] block files  
/dev/rhk[0-7][a-h] raw files  
/dev/MAKEDEV script to create special files  
/dev/MAKEDEV.local script to localize special files

#### SEE ALSO

ra(4), ram(4), rk(4), rl(4), rx(4), si(4), xp(4), dtab(5),  
autoconfig(8), disklabel(8)

#### DIAGNOSTICS

hk%d%c: hard error sn%d cs2=%b ds=%b er=%b. An unrecoverable error occurred during transfer of the specified sector of the specified disk partition. The contents of the cs2, ds and er registers are printed in octal and symbolically with bits decoded. The error was either unrecoverable, or a large number of retry attempts (including offset positioning and drive recalibration) could not recover the error.

hk%d: write locked. The write protect switch was set on the drive when a write was attempted. The write operation is not recoverable.

hk%d: not ready. The drive was spun down or off line when it was accessed. The i/o operation is not recoverable.

hk%d: not ready (came back!). The drive was not ready, but after printing the message about being not ready (which takes a fraction of a second) was ready. The operation is recovered if no further errors occur.

hk%d%c: soft ecc sn%d. A recoverable ECC error occurred on the specified sector of the specified disk partition. This happens normally a few times a week. If it happens more frequently than this the sectors where the errors are occurring should be checked to see if certain cylinders on the pack, spots on the carriage of the drive or heads are indicated.

#### BUGS

In raw I/O read and write(2) truncate file offsets to 512-byte block boundaries, and write scribbles on the tail of incomplete blocks. Thus, in programs that are likely to access raw devices, read, write and lseek(2) should always

deal in 512-byte multiples.

DEC-standard error logging should be supported.

A program to analyze the logged error information (even in its present reduced form) is needed.

## NAME

ht - TM-03/TE-16,TU-45,TU-77 MASSBUS magtape interface

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NHT ht_drives # TE16, TU45, TU77

/etc/dtab:
    #Name Unit# Addr    Vector Br Handler(s)      # Comments
    ht  ?      172440 224      5  htintr      # tu 16 massbus tape

major device number(s):
    raw: 6
    block: 0
minor device encoding:
    bits 0003 specify HT drive
    bit 0004 specifies no-rewind operation
    bit 0010 specifies 1600BPI recording density instead of 800BPI
```

## DESCRIPTION

The tm-03/transport combination provides a standard tape drive interface as described in mtio(4). All drives provide both 800 and 1600 bpi; the TE-16 runs at 45 ips, the TU-45 at 75 ips, while the TU-77 runs at 125 ips and autoloads tapes.

## FILES

```
/dev/MAKEDEV script to create special files
/dev/MAKEDEV.local script to localize special files
```

## SEE ALSO

mt(1), tar(1), tp(1), mtio(4), tm(4), ts(4), dtab(5), auto-config(8)

## DIAGNOSTICS

tu%d: no write ring. An attempt was made to write on the tape drive when no write ring was present; this message is written on the terminal of the user who tried to access the tape.

tu%d: not online. An attempt was made to access the tape while it was offline; this message is written on the terminal of the user who tried to access the tape.

tu%d: can't change density in mid-tape. An attempt was made to write on a tape at a different density than is already recorded on the tape. This message is written on the terminal of the user who tried to switch the density.

tu%d: hard error bn%d er=%b ds=%b. A tape error occurred at block bn; the ht error register and drive status register are printed in octal with the bits symbolically decoded.

Any error is fatal on non-raw tape; when possible the driver will have retried the operation which failed several times before reporting the error.

#### BUGS

If any non-data error is encountered on non-raw tape, it refuses to do anything more until closed.

The system should remember which controlling terminal has the tape drive open and write error messages to that terminal rather than on the console.



## NAME

hy - Network Systems Hyperchannel interface

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NHY  hy_controllers # Hyperchannel
```

## DESCRIPTION

The hy interface provides access to a Network Systems Corporation Hyperchannel Adapter.

The network to which the interface is attached is specified at boot time with an SIOCSIFADDR ioctl. The host's address is discovered by reading the adapter status register. The interface will not transmit or receive packets until the network number is known.

## DIAGNOSTICS

hy%d: unit number 0x%x port %d type %x microcode level 0x%x.  
Identifies the device during autoconfiguration.

hy%d: can't handle af%d. The interface was handed a message with addresses formatted in an unsuitable address family; the packet was dropped.

hy%d: can't initialize. The interface was unable to allocate UNIBUS resources. This is usually due to having too many network devices on an 11/750 where there are only 3 buffered data paths.

hy%d: NEX - Non Existent Memory. Non existent memory error returned from hardware.

hy%d: BAR overflow. Bus address register overflow error returned from hardware.

hy%d: Power Off bit set, trying to reset. Adapter has lost power, driver will reset the bit and see if power is still out in the adapter.

hy%d: Power Off Error, network shutdown. Power was really off in the adapter, network connections are dropped. Software does not shut down the network unless power has been off for a while.

hy%d: RECVD MP > MPSIZE (%d). A message proper was received that is too big. Probable a driver bug. Shouldn't happen.

hy%d: xmit error - len > hy\_olen [%d > %d]. Probable driver error. Shouldn't happen.

hy%d: DRIVER BUG - INVALID STATE %d. The driver state machine reached a non-existent state. Definite driver bug.

hy%d: watchdog timer expired. A command in the adapter has taken too long to complete. Driver will abort and retry the command.

hy%d: adapter power restored. Software was able to reset the power off bit, indicating that the power has been restored.

#### SEE ALSO

intro(4N), inet(4F)

#### BUGS

If the adapter does not respond to the status command issued during autoconfigure, the adapter is assumed down. A reboot is required to recognize it.

The adapter power fail interrupt seems to occur sporadically when power has, in fact, not failed. The driver will believe that power has failed only if it can not reset the power fail latch after a ``reasonable'' time interval. These seem to appear about 2-4 times a day on some machines. There seems to be no correlation with adapter rev level, number of ports used etc. and whether a machine will get these ``bogus powerfails''. They don't seem to cause any real problems so they have been ignored.

## NAME

icmp - Internet Control Message Protocol

## SYNOPSIS

```
#include <sys/socket.h>
#include <netinet/in.h>
```

```
s = socket(AF_INET, SOCK_RAW, proto);
```

## DESCRIPTION

ICMP is the error and control message protocol used by IP and the Internet protocol family. It may be accessed through a "raw socket" for network monitoring and diagnostic functions. The proto parameter to the socket call to create an ICMP socket is obtained from getprotobyname(3N). ICMP sockets are connectionless, and are normally used with the sendto and recvfrom calls, though the connect(2) call may also be used to fix the destination for future packets (in which case the read(2) or recv(2) and write(2) or send(2) system calls may be used).

Outgoing packets automatically have an IP header prepended to them (based on the destination address). Incoming packets are received with the IP header and options intact.

## DIAGNOSTICS

A socket operation may fail with one of the following errors returned:

[EISCONN]           when trying to establish a connection on a socket which already has one, or when trying to send a datagram with the destination address specified and the socket is already connected;

[ENOTCONN]          when trying to send a datagram, but no destination address is specified, and the socket hasn't been connected;

[ENOBUFS]           when the system runs out of memory for an internal data structure;

[EADDRNOTAVAIL]     when an attempt is made to create a socket with a network address for which no network interface exists.

## SEE ALSO

send(2), recv(2), intro(4N), inet(4F), ip(4P)

## NAME

idp - Xerox Internet Datagram Protocol

## SYNOPSIS

```
#include <sys/socket.h>
#include <netns/ns.h>
#include <netns/idp.h>

s = socket(AF_NS, SOCK_DGRAM, 0);
```

## DESCRIPTION

IDP is a simple, unreliable datagram protocol which is used to support the SOCK\_DGRAM abstraction for the Internet protocol family. IDP sockets are connectionless, and are normally used with the sendto and recvfrom calls, though the connect(2) call may also be used to fix the destination for future packets (in which case the recv(2) or read(2) and send(2) or write(2) system calls may be used).

Xerox protocols are built vertically on top of IDP. Thus, IDP address formats are identical to those used by SPP. Note that the IDP port space is the same as the SPP port space (i.e. a IDP port may be "connected" to a SPP port, with certain options enabled below). In addition broadcast packets may be sent (assuming the underlying network supports this) by using a reserved "broadcast address"; this address is network interface dependent.

## DIAGNOSTICS

A socket operation may fail with one of the following errors returned:

- [EISCONN]           when trying to establish a connection on a socket which already has one, or when trying to send a datagram with the destination address specified and the socket is already connected;
- [ENOTCONN]          when trying to send a datagram, but no destination address is specified, and the socket hasn't been connected;
- [ENOBUFS]           when the system runs out of memory for an internal data structure;
- [EADDRINUSE]        when an attempt is made to create a socket with a port which has already been allocated;
- [EADDRNOTAVAIL]     when an attempt is made to create a socket with a network address for which no network interface exists.

## SOCKET OPTIONS

## [SO\_HEADERS\_ON\_INPUT]

When set, the first 30 bytes of any data returned from a read or recv from will be the initial 30 bytes of the IDP packet, as described by

```
struct idp {
    u_short    idp_sum;
    u_short    idp_len;
    u_char      idp_tc;
    u_char      idp_pt;
    struct ns_addr idp_dna;
    struct ns_addr idp_sna;
};
```

This allows the user to determine the packet type, and whether the packet was a multi-cast packet or directed specifically at the local host. When requested, gives the current state of the option, (NSP\_RAWIN or 0).

## [SO\_HEADERS\_ON\_OUTPUT]

When set, the first 30 bytes of any data sent will be the initial 30 bytes of the IDP packet. This allows the user to determine the packet type, and whether the packet should be multi-cast packet or directed specifically at the local host. You can also misrepresent the sender of the packet. When requested, gives the current state of the option. (NSP\_RAWOUT or 0).

## [SO\_DEFAULT\_HEADERS]

The user provides the kernel an IDP header, from which it gleans the Packet Type. When requested, the kernel will provide an IDP header, showing the default packet type, and local and foreign addresses, if connected.

## [SO\_ALL\_PACKETS]

When set, this option defeats automatic processing of Error packets, and Sequence Protocol packets.

## [SO\_SEQNO]

When requested, this returns a sequence number which is not likely to be repeated until the machine crashes or a very long time has passed. It is useful in constructing Packet Exchange Protocol packets.

## SEE ALSO

send(2), recv(2), intro(4N), ns(4F)

## NAME

il - Interlan NI1010 10 Mb/s Ethernet interface

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NIL  il_controllers # Interlan Ethernet
```

## DESCRIPTION

The il interface provides access to a 10 Mb/s Ethernet network through an Interlan 1010 or 1010A controller.

Each of the host's network addresses is specified at boot time with an SIOCSIFADDR ioctl. The il interface employs the address resolution protocol described in arp(4P) to dynamically map between Internet and Ethernet addresses on the local network.

The interface normally tries to use a ``trailer'' encapsulation to minimize copying data on input and output. The use of trailers is negotiated with ARP. This negotiation may be disabled, on a per-interface basis, by setting the IFF\_NOTRAILERS flag with an SIOCSIFFLAGS ioctl.

## DIAGNOSTICS

il%d: input error. The hardware indicated an error in reading a packet off the cable or an illegally sized packet.

il%d: can't handle af%d. The interface was handed a message with addresses formatted in an unsuitable address family; the packet was dropped.

il%d: setaddr didn't work. The interface was unable to reprogram its physical ethernet address. This may happen with very early models of the interface. This facility is used only when the controller is not the first network interface configured for XNS. The oldest interface tested (2.7.1.0.1.45) has never failed in this way.

The following messages indicate a probable hardware error performing the indicated operation during autoconfiguration or initialization. The status field in the control and status register (the low-order four bits) should indicate the nature of the failure. See the hardware manual for details.

il%d: reset failed, csr=%b.

il%d: status failed, csr=%b.

il%d: hardware diag failed, csr=%b.

il%d: verifying setaddr, csr=%b.

il%d: stray xmit interrupt, csr=%b.

il%d: can't initialize.

SEE ALSO

intro(4N), inet(4F), arp(4P)

## NAME

imp - IMP raw socket interface

## SYNOPSIS

```
#include <sys/socket.h>
#include <netinet/in.h>
#include <netimp/if_imp.h>

s = socket(AF_IMPLINK, SOCK_RAW, proto);
```

## DESCRIPTION

The raw imp socket provides direct access to the imp(4) network interface. Users send packets through the interface using the send(2) calls, and receive packets with the recv(2), calls. All outgoing packets must have an 1822 96-bit leader on the front. Likewise, packets received by the user will have this leader on the front. The 1822 leader and the legal values for the various fields are defined in the include file <netimp/if\_imp.h>. The raw imp interface automatically installs the length and destination address in the 1822 leader of all outgoing packets; these need not be filled in by the user.

If the protocol selected, proto, is zero, the socket will receive all IMP messages except RFNM and incompletes which are not input data for a kernel protocol. If proto is non-zero, only messages for the specified link type will be received.

## DIAGNOSTICS

An operation on a socket may fail with one of the following errors:

- [EISCONN]           when trying to establish a connection on a socket which already has one, or when trying to send a datagram with the destination address specified and the socket is already connected;
- [ENOTCONN]          when trying to send a datagram, but no destination address is specified, and the socket hasn't been connected;
- [ENOBUFS]           when the system runs out of memory for an internal data structure;
- [ENOBUFS]           eight messages to the destination host are outstanding, and another eight are already queued for output;
- [EADDRNOTAVAIL]     when an attempt is made to create a socket



with a network address for which no network interface exists.

SEE ALSO

intro(4N), inet(4F), imp(4)

## NAME

imp - 1822 network interface

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NIMP count      # ARPAnet IMP 1822 interface
```

## DESCRIPTION

The imp interface, as described in BBN Report 1822, provides access to an intelligent message processor normally used when participating in the Department of Defense ARPA network. The network interface communicates through a device controller, usually an ACC LH/DH or HDH or a DEC IMP-11A, with the IMP. The interface is "reliable" and "flow-controlled" by the host-IMP protocol.

To configure IMP support, at least one of acc(4), css(4) or hdh(4) must be included. The count specifies the total number of IMP connections. The network number on which the interface resides is specified at boot time using the SIOC-SIFADDR ioctl. The host number is discovered through receipt of NOOP messages from the IMP.

The network interface is always in one of four states: up, down, initializing, or going down. When the system is booted, the interface is marked down. If the hardware controller is successfully probed, the interface enters the initializing state and transmits three NOOP messages to the IMP. It then waits for the IMP to respond with two or more NOOP messages in reply. When it receives these messages it enters the up state. The "going down" state is entered only when notified by the IMP of an impending shutdown. Packets may be sent through the interface only while it is in the up state. Outgoing packets are dropped with the error ENETDOWN returned to the caller if the interface is in any other state.

## DIAGNOSTICS

imp%d: not configured. A hardware interface could not be attached during autoconfiguration because too few IMP pseudo-devices were configured.

imp%d: leader error. The IMP reported an error in a leader (1822 message header). This causes the interface to be reset and any packets queued up for transmission to be purged.

imp%d: going down in 30 seconds.

imp%d: going down for hardware PM.

imp%d: going down for reload software.

imp%d: going down for emergency reset. The Network Control Center (NCC) is manipulating the IMP. By convention these

messages are reported to all hosts on an IMP.

imp?: host %x, lost %d rfnms. The IMP had messages outstanding to the host listed, but no RFNM (Request for Next Message) messages were received from the IMP in 127 seconds. The software state for that host is reinitialized.

imp%d: interface reset. The host has received an interface reset message from the IMP.

imp%d: address reset to x%x (%d/%d). The host has received a NOOP message which caused it to reset its notion of its current address. The Internet address is printed in hexadecimal, with the host and IMP numbers following. This indicates that the address originally set by ifconfig(8) was incorrect, that the IMP has undergone an identity crisis, or that communication between the IMP and the host is being garbled.

imp%d: data error. The IMP noted an error in data transmitted. The host-IMP interface is reset and the host enters the init state (awaiting NOOP messages).

imp%d: interface reset. The reset process has been completed.

imp%d: marked down. After receiving a "going down in 30 seconds" message, and waiting 30 seconds, the host has marked the IMP unavailable. Before packets may be sent to the IMP again, the IMP must notify the host, through a series of NOOP messages, that it is back up.

imp%d: can't handle af%d. The interface was handed a message with addresses formatting in an unsuitable address family; the packet was dropped.

SEE ALSO

intro(4N), inet(4F), acc(4), css(4), hdh(4), implog(8),  
implogd(8)

## NAME

inet - Internet protocol family

## SYNOPSIS

```
#include <sys/types.h>
#include <netinet/in.h>
```

## DESCRIPTION

The Internet protocol family is a collection of protocols layered atop the Internet Protocol (IP) transport layer, and utilizing the Internet address format. The Internet family provides protocol support for the SOCK\_STREAM, SOCK\_DGRAM, and SOCK\_RAW socket types; the SOCK\_RAW interface provides access to the IP protocol.

## ADDRESSING

Internet addresses are four byte quantities, stored in network standard format: layed out as highest to lowest order bytes in memory or ``Big Endian'' (the VAX is word and byte reversed, or ``Little Endian''; the PDP-11 is byte reversed within each word, or ``Middle Endian'). The include file <netinet/in.h> defines this address as a discriminated union.

Sockets bound to the Internet protocol family utilize the following addressing structure,

```
struct sockaddr_in {
    short      sin_family;
    u_short    sin_port;
    struct      in_addr sin_addr;
    char       sin_zero[8];
};
```

Sockets may be created with the local address INADDR\_ANY to effect "wildcard" matching on incoming messages. The address in a connect(2) or sendto(2) call may be given as INADDR\_ANY to mean ``this host.''. The distinguished address INADDR\_BROADCAST is allowed as a shorthand for the broadcast address on the primary network if the first network configured supports broadcast.

## PROTOCOLS

The Internet protocol family is comprised of the IP transport protocol, Internet Control Message Protocol (ICMP), Transmission Control Protocol (TCP), and User Datagram Protocol (UDP). TCP is used to support the SOCK\_STREAM abstraction while UDP is used to support the SOCK\_DGRAM abstraction. A raw interface to IP is available by creating an Internet socket of type SOCK\_RAW. The ICMP message protocol is accessible from a raw socket.

The 32-bit Internet address contains both network and host parts. It is frequency-encoded; the most-significant bit is clear in Class A addresses, in which the high-order 8 bits are the network number. Class B addresses use the high-order 16 bits as the network field, and Class C addresses have a 24-bit network part. Sites with a cluster of local networks and a connection to the DARPA Internet may chose to use a single network number for the cluster; this is done by using subnet addressing. The local (host) portion of the address is further subdivided into subnet and host parts. Within a subnet, each subnet appears to be an individual network; externally, the entire cluster appears to be a single, uniform network requiring only a single routing entry. Subnet addressing is enabled and examined by the following `ioctl(2)` commands on a datagram socket in the Internet domain; they have the same form as the `SIOCIFADDR` command (see `intro(4N)`).

`SIOCSIFNETMASK`      Set interface network mask. The network mask defines the network part of the address; if it contains more of the address than the address type would indicate, then subnets are in use.

`SIOCGIFNETMASK`      Get interface network mask.

SEE ALSO

`ioctl(2)`, `socket(2)`, `intro(4N)`, `tcp(4P)`, `udp(4P)`, `ip(4P)`, `icmp(4P)`

An Introductory 4.3BSD Interprocess Communication Tutorial (PS1:7).

An Advanced 4.3BSD Interprocess Communication Tutorial (PS1:8).

CAVEAT

The Internet protocol support is subject to change as the Internet protocols develop. Users should not depend on details of the current implementation, but rather the services exported.

## NAME

ip - Internet Protocol

## SYNOPSIS

```
#include <sys/socket.h>
#include <netinet/in.h>
```

```
s = socket(AF_INET, SOCK_RAW, proto);
```

## DESCRIPTION

IP is the transport layer protocol used by the Internet protocol family. Options may be set at the IP level when using higher-level protocols that are based on IP (such as TCP and UDP). It may also be accessed through a "raw socket" when developing new protocols, or special purpose applications.

A single generic option is supported at the IP level, `IP_OPTIONS`, that may be used to provide IP options to be transmitted in the IP header of each outgoing packet. Options are set with `setsockopt(2)` and examined with `getsockopt(2)`. The format of IP options to be sent is that specified by the IP protocol specification, with one exception: the list of addresses for Source Route options must include the first-hop gateway at the beginning of the list of gateways. The first-hop gateway address will be extracted from the option list and the size adjusted accordingly before use. IP options may be used with any socket type in the Internet family.

Raw IP sockets are connectionless, and are normally used with the `sendto` and `recvfrom` calls, though the `connect(2)` call may also be used to fix the destination for future packets (in which case the `read(2)` or `recv(2)` and `write(2)` or `send(2)` system calls may be used).

If `proto` is 0, the default protocol `IPPROTO_RAW` is used for outgoing packets, and only incoming packets destined for that protocol are received. If `proto` is non-zero, that protocol number will be used on outgoing packets and to filter incoming packets.

Outgoing packets automatically have an IP header prepended to them (based on the destination address and the protocol number the socket is created with). Incoming packets are received with IP header and options intact.

## DIAGNOSTICS

A socket operation may fail with one of the following errors returned:

[EISCONN]           when trying to establish a connection on a socket which already has one, or when trying

to send a datagram with the destination address specified and the socket is already connected;

[ENOTCONN]        when trying to send a datagram, but no destination address is specified, and the socket hasn't been connected;

[ENOBUFFS]        when the system runs out of memory for an internal data structure;

[EADDRNOTAVAIL]    when an attempt is made to create a socket with a network address for which no network interface exists.

The following errors specific to IP may occur when setting or getting IP options:

[EINVAL]            An unknown socket option name was given.

[EINVAL]            The IP option field was improperly formed; an option field was shorter than the minimum value or longer than the option buffer provided.

SEE ALSO

getsockopt(2), send(2), recv(2), intro(4N), icmp(4P),  
inet(4F)

## NAME

lp - line printer

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NLP      lp_printers  # Line Printer
    LP_MAXCOL 132      # Maximum number of columns on line printers

/etc/dtab:
    #Name Unit# Addr   Vector Br Handler(s)      # Comments
    lp ?      177514 200    4  lpintr      # lp-11 line printer

major device number(s):
    raw: 5
minor device encoding:
    bit 01 specifies 64-character set (instead of 96-character set)
    bits 06 are reserved
    bits 0370 specify line printer unit: <lp_unit> * 8
```

## DESCRIPTION

lp provides the interface to any of the standard DEC line printers on an LP-11 parallel interface. When it is opened or closed, a suitable number of page ejects is generated. Bytes written are printed.

The unit number of the printer is specified by the minor device after removing the low 3 bits, which act as per-device parameters. Currently only the lowest of the low three bits is interpreted: if it is set, the device is treated as having a 64-character set, rather than a full 96-character set. In the resulting half-ASCII mode, lower case letters are turned into upper case and certain characters are escaped according to the following table:

{	(
}	)
`	'
	!
~	^

The driver correctly interprets carriage returns, backspaces, tabs, and form feeds. Lines longer than the maximum page width are truncated. The page width is specified via the LP\_MAXCOL definition, set to 132 in the GENERIC distribution.

## FILES

```
/dev/lp
/dev/MAKEDEV script to create special files
/dev/MAKEDEV.local script to localize special files
```



## SEE ALSO

lpr(1), dtab(5), autoconfig(8)

## DIAGNOSTICS

None.

## BUGS

Although the driver supports multiple printers this has never been tried.

## NAME

mem, kmem - main memory

## SYNOPSIS

major device number(s):  
raw: 1  
minor device encoding:  
mem: 0; kmem: 1; null: 2

## DESCRIPTION

Mem is a special file that is an image of the main memory of the computer. It may be used, for example, to examine (and even to patch) the system.

Byte addresses in mem are interpreted as physical memory addresses. References to non-existent locations cause errors to be returned.

The file kmem is the same as mem except that kernel virtual memory rather than physical memory is accessed. Only kernel virtual addresses that are mapped to memory are allowed. Examining and patching device registers is likely to lead to unexpected results when read-only or write-only bits are present.

On PDP-11s, the I/O page begins at location 0160000 of kmem and the per-process data segment for the current process begins at 0140000 and is USIZE clicks (64 bytes each) long.

## FILES

/dev/mem  
/dev/kmem  
/dev/MAKEDEV script to create special files  
/dev/MAKEDEV.local script to localize special files

## BUGS

On PDP-11's, specifying an odd kernel or user address, or an odd transfer count is [generally] slower than using all even parameters.

On machines with ENABLE/34(tm) memory mapping boards the I/O page can be accessed only through kmem.

## NAME

mtio - UNIX magtape interface

## DESCRIPTION

The files mt0, ..., mt15 refer to the UNIX magtape drives, which may be on the MASSBUS using the TM03 formatter ht(4), or on the UNIBUS using either the TM11 or TS11 formatters tm(4) or ts(4). The following description applies to any of the transport/controller pairs.

The files mt0, ..., mt7 are 800bpi (or the transport's lowest density), mt8, ..., mt15 are 1600bpi (or the transport's second density), and mt16, ..., mt23 are 6250bpi (or the transport's third density). (But note that only 1600 bpi is available with the TS11.) The files mt0, ..., mt3, mt8, ..., mt11, and mt16, ..., mt19 are rewound when closed; the others are not. When a file open for writing is closed, two end-of-files are written. If the tape is not to be rewound it is positioned with the head between the two tapemarks.

A standard tape consists of a series of 1024 byte records terminated by an end-of-file. To the extent possible, the system makes it possible, if inefficient, to treat the tape like any other file. Seeks have their usual meaning and it is possible to read or write a byte at a time. Writing in very small units is inadvisable, however, because it uses most of the tape in record gaps.

The mt files discussed above are useful when it is desired to access the tape in a way compatible with ordinary files. When foreign tapes are to be dealt with, and especially when long records are to be read or written, the 'raw' interface is appropriate. The associated files are named rmt0, ..., rmt23, but the same minor-device considerations as for the regular files still apply. A number of other ioctl operations are available on raw magnetic tape. The following definitions are from <sys/mtio.h>:

```
/*
 * Structures and definitions for mag tape io control commands
 */

/* structure for MTIOCTOP - mag tape op command */
struct mtop {
    short    mt_op;          /* operations defined below */
    daddr_t  mt_count; /* how many of them */
};

/* operations */
#define MTWEOF 0 /* write an end-of-file record */
#define MTFSF 1 /* forward space file */
```

```

#define MTBSF 2 /* backward space file */
#define MTFSR 3 /* forward space record */
#define MTBSR 4 /* backward space record */
#define MTREW 5 /* rewind */
#define MTOFFL 6 /* rewind and put the drive offline */
#define MTNOP 7 /* no operation, sets status only */
#define MTCACHE 8 /* enable controller cache */
#define MTNOCACHE 9 /* disable controller cache */

/* structure for MTIOCGGET - mag tape get status command */

struct mtget {
    short mt_type; /* type of magtape device */
    /* the following two registers are grossly device dependent */
    short mt_dsreg; /* ``drive status'' register */
    short mt_erreg; /* ``error'' register */
    /* end device-dependent registers */
    short mt_resid; /* residual count */
    /* the following two are not yet implemented */
    daddr_t mt_fileno; /* file number of current position */
    daddr_t mt_blkno; /* block number of current position */
    /* end not yet implemented */
};

/*
 * Constants for mt_type byte. These are the same
 * for other controllers compatible with the types listed.
 */
#define MT_ISTS 0x01 /* TS-11 */
#define MT_ISHT 0x02 /* TM03 Massbus: TE16, TU45, TU77 */
#define MT_ISTM 0x03 /* TM11/TE10 Unibus */
#define MT_ISMT 0x04 /* TM78/TU78 Massbus */
#define MT_ISUT 0x05 /* SI TU-45 emulation on Unibus */
#define MT_ISCPC 0x06 /* SUN */
#define MT_ISAR 0x07 /* SUN */
#define MT_ISTMSCP 0x08 /* DEC TMSCP protocol (TU81, TK50) */

/* mag tape io control commands */
#define MTIOCTOP _IOW(m, 1, struct mtop) /* do a mag tape op */
#define MTIOCGGET _IOR(m, 2, struct mtget) /* get tape status */
#define MTIOCIEOT _IO(m, 3) /* ignore EOT error */
#define MTIOCEEOT _IO(m, 4) /* enable EOT error */

#ifndef KERNEL
#define DEFTAPE "/dev/rmt12"
#endif

```

Each read or write call reads or writes the next record on the tape. In the write case the record has the same length as the buffer given. During a read, the record size is passed back as the number of bytes read, provided it is no greater than the buffer size; if the record is long, an

error is indicated. In raw tape I/O seeks are ignored. A zero byte count is returned when a tape mark is read, but another read will fetch the first record of the new tape file.

The MTCACHE and MTNOCACHE commands do not apply to all tape drivers. At present only the TMSCP driver supports those commands and then only for drives such as the TU81+.

#### FILES

/dev/mt?  
/dev/rmt?

#### SEE ALSO

mt(1), tar(1), tp(1), ht(4), tm(4), ts(4)

#### BUGS

The status should be returned in a device independent format.

The special file naming should be redone in a more consistent and understandable manner.

## NAME

networking - introduction to networking facilities

## SYNOPSIS

```
#include <sys/socket.h>
#include <net/route.h>
#include <net/if.h>
```

## DESCRIPTION

This section briefly describes the networking facilities available in the system. Documentation in this part of section 4 is broken up into three areas: protocol families (domains), protocols, and network interfaces. Entries describing a protocol family are marked ``4F,' ' while entries describing protocol use are marked ``4P.' ' Hardware support for network interfaces are found among the standard ``4' ' entries.

All network protocols are associated with a specific protocol family. A protocol family provides basic services to the protocol implementation to allow it to function within a specific network environment. These services may include packet fragmentation and reassembly, routing, addressing, and basic transport. A protocol family may support multiple methods of addressing, though the current protocol implementations do not. A protocol family is normally comprised of a number of protocols, one per socket(2) type. It is not required that a protocol family support all socket types. A protocol family may contain multiple protocols supporting the same socket abstraction.

A protocol supports one of the socket abstractions detailed in socket(2). A specific protocol may be accessed either by creating a socket of the appropriate type and protocol family, or by requesting the protocol explicitly when creating a socket. Protocols normally accept only one type of address format, usually determined by the addressing structure inherent in the design of the protocol family/network architecture. Certain semantics of the basic socket abstractions are protocol specific. All protocols are expected to support the basic model for their particular socket type, but may, in addition, provide non-standard facilities or extensions to a mechanism. For example, a protocol supporting the SOCK\_STREAM abstraction may allow more than one byte of out-of-band data to be transmitted per out-of-band message.

A network interface is similar to a device interface. Network interfaces comprise the lowest layer of the networking subsystem, interacting with the actual transport hardware. An interface may support one or more protocol families and/or address formats. The SYNOPSIS section of each

network interface entry gives a sample specification of the related drivers for use in providing a system description to the /sys/conf/config script. The DIAGNOSTICS section lists messages which may appear on the console and/or in the system error log, /usr/adm/messages (see syslogd(8)), due to errors in device operation.

## PROTOCOLS

The system currently supports the DARPA Internet protocols and the Xerox Network Systems(tm) protocols. Raw socket interfaces are provided to the IP protocol layer of the DARPA Internet, to the IMP link layer (1822), and to the IDP protocol of Xerox NS. Consult the appropriate manual pages in this section for more information regarding the support for each protocol family.

## ADDRESSING

Associated with each protocol family is an address format. The following address formats are used by the system (and additional formats are defined for possible future implementation):

```
#define AF_UNIX          1      /* local to host (pipes, portals) */
#define AF_INET          2      /* internetwork: UDP, TCP, etc. */
#define AF_IMPLINK       3      /* arpanet imp addresses */
#define AF_PUP           4      /* pup protocols: e.g. BSP */
#define AF_NS            6      /* Xerox NS protocols */
#define AF_HYLINK        15     /* NSC Hyperchannel */
```

## ROUTING

The network facilities provided limited packet routing. A simple set of data structures comprise a ``routing table'' used in selecting the appropriate network interface when transmitting packets. This table contains a single entry for each route to a specific network or host. A user process, the routing daemon, maintains this data base with the aid of two socket-specific ioctl(2) commands, SIOCADDRT and SIOCDELRT. The commands allow the addition and deletion of a single routing table entry, respectively. Routing table manipulations may only be carried out by super-user.

A routing table entry has the following form, as defined in <net/route.h>;

```
struct rtentry {
    u_long    rt_hash;
    struct    sockaddr rt_dst;
    struct    sockaddr rt_gateway;
    short     rt_flags;
    short     rt_refcnt;
    u_long    rt_use;
    struct    ifnet *rt_ifp;
```

```
};
```

with `rt_flags` defined from,

```
#define RTF_UP          0x1    /* route usable */
#define RTF_GATEWAY     0x2    /* destination is a gateway */
#define RTF_HOST        0x4    /* host entry (net otherwise) */
#define RTF_DYNAMIC     0x10   /* created dynamically (by redirect) */
```

Routing table entries come in three flavors: for a specific host, for all hosts on a specific network, for any destination not matched by entries of the first two types (a wildcard route). When the system is booted and addresses are assigned to the network interfaces, each protocol family installs a routing table entry for each interface when it is ready for traffic. Normally the protocol specifies the route through each interface as a ``direct'' connection to the destination host or network. If the route is direct, the transport layer of a protocol family usually requests the packet be sent to the same host specified in the packet. Otherwise, the interface is requested to address the packet to the gateway listed in the routing entry (i.e. the packet is forwarded).

Routing table entries installed by a user process may not specify the hash, reference count, use, or interface fields; these are filled in by the routing routines. If a route is in use when it is deleted (`rt_refcnt` is non-zero), the routing entry will be marked down and removed from the routing table, but the resources associated with it will not be reclaimed until all references to it are released. The routing code returns `EEXIST` if requested to duplicate an existing entry, `ESRCH` if requested to delete a non-existent entry, or `ENOBUFS` if insufficient resources were available to install a new route. User processes read the routing tables through the `/dev/kmem` device. The `rt_use` field contains the number of packets sent along the route.

When routing a packet, the kernel will first attempt to find a route to the destination host. Failing that, a search is made for a route to the network of the destination. Finally, any route to a default (``wildcard'') gateway is chosen. If multiple routes are present in the table, the first route found will be used. If no entry is found, the destination is declared to be unreachable.

A wildcard routing entry is specified with a zero destination address value. Wildcard routes are used only when the system fails to find a route to the destination host and network. The combination of wildcard routes and routing redirects can provide an economical mechanism for routing traffic.



## INTERFACES

Each network interface in a system corresponds to a path through which messages may be sent and received. A network interface usually has a hardware device associated with it, though certain interfaces such as the loopback interface, `lo(4)`, do not.

The following `ioctl` calls may be used to manipulate network interfaces. The `ioctl` is made on a socket (typically of type `SOCK_DGRAM`) in the desired domain. Unless specified otherwise, the request takes an `ifrequest` structure as its parameter. This structure has the form

```
struct    ifreq {
#define    IFNAMSIZ 16
char ifr_name[IFNAMSIZ];    /* if name, e.g. "en0" */
union {
    struct    sockaddr ifru_addr;
    struct    sockaddr ifru_dstaddr;
    struct    sockaddr ifru_broadaddr;
    short     ifru_flags;
    int        ifru_metric;
    caddr_t    ifru_data;
} ifr_ifru;
#define    ifr_addr    ifr_ifru.ifru_addr /* address */
#define    ifr_dstaddr    ifr_ifru.ifru_dstaddr /* other end of p-to-p link */
#define    ifr_broadaddr    ifr_ifru.ifru_broadaddr /* broadcast address */
#define    ifr_flags    ifr_ifru.ifru_flags /* flags */
#define    ifr_metric    ifr_ifru.ifru_metric /* metric */
#define    ifr_data    ifr_ifru.ifru_data /* for use by interface */
};
```

## SIOCSIFADDR

Set interface address for protocol family. Following the address assignment, the ``initialization'' routine for the interface is called.

## SIOCGIFADDR

Get interface address for protocol family.

## SIOCSIFDSTADDR

Set point to point address for protocol family and interface.

## SIOCGIFDSTADDR

Get point to point address for protocol family and interface.

## SIOCSIFBRDADDR

Set broadcast address for protocol family and interface.

**SIOCGIFBRDADDR**

Get broadcast address for protocol family and interface.

**SIOCSIFFLAGS**

Set interface flags field. If the interface is marked down, any processes currently routing packets through the interface are notified; some interfaces may be reset so that incoming packets are no longer received. When marked up again, the interface is reinitialized.

**SIOCGIFFLAGS**

Get interface flags.

**SIOCSIFMETRIC**

Set interface routing metric. The metric is used only by user-level routers.

**SIOCGIFMETRIC**

Get interface metric.

**SIOCGIFCONF**

Get interface configuration list. This request takes an ifconf structure (see below) as a value-result parameter. The ifc\_len field should be initially set to the size of the buffer pointed to by ifc\_buf. On return it will contain the length, in bytes, of the configuration list.

```
/*
 * Structure used in SIOCGIFCONF request.
 * Used to retrieve interface configuration
 * for machine (useful for programs which
 * must know all networks accessible).
 */
struct    ifconf {
    int    ifc_len;          /* size of associated buffer */
    union {
        caddr_t    ifcu_buf;
        struct      ifreq *ifcu_req;
    } ifc_ifcu;
#define    ifc_buf    ifc_ifcu.ifcu_buf    /* buffer address */
#define    ifc_req    ifc_ifcu.ifcu_req    /* array of structures returned */
};
```

**SEE ALSO**

socket(2), ioctl(2), intro(4), config(8), routed(8C)

## NAME

ns - Xerox Network Systems(tm) protocol family

## SYNOPSIS

Not currently supported under 2.11BSD

## DESCRIPTION

The NS protocol family is a collection of protocols layered atop the Internet Datagram Protocol (IDP) transport layer, and using the Xerox NS address formats. The NS family provides protocol support for the SOCK\_STREAM, SOCK\_DGRAM, SOCK\_SEQPACKET, and SOCK\_RAW socket types; the SOCK\_RAW interface is a debugging tool, allowing you to trace all packets entering, (or with toggling kernel variable, additionally leaving) the local host.

## ADDRESSING

NS addresses are 12 byte quantities, consisting of a 4 byte Network number, a 6 byte Host number and a 2 byte port number, all stored in network standard format. (on the VAX these are word and byte reversed; on the Sun they are not reversed). The include file <netns/ns.h> defines the NS address as a structure containing unions (for quicker comparisons).

Sockets in the Internet protocol family use the following addressing structure:

```
struct sockaddr_ns {
    short      sns_family;
    struct ns_addr sns_addr;
    char       sns_zero[2];
};
```

where an ns\_addr is composed as follows:

```
union ns_host {
    u_char    c_host[6];
    u_short   s_host[3];
};
```

```
union ns_net {
    u_char    c_net[4];
    u_short   s_net[2];
};
```

```
struct ns_addr {
    union ns_net    x_net;
    union ns_host   x_host;
    u_short         x_port;
};
```

Sockets may be created with an address of all zeroes to effect ``wildcard'' matching on incoming messages. The local port address specified in a `bind(2)` call is restricted to be greater than `NSPORT_RESERVED` (=3000, in `<netns/ns.h>`) unless the creating process is running as the super-user, providing a space of protected port numbers.

#### PROTOCOLS

The NS protocol family supported by the operating system is comprised of the Internet Datagram Protocol (IDP) `idp(4P)`, Error Protocol (available through IDP), and Sequenced Packet Protocol (SPP) `spp(4P)`.

SPP is used to support the `SOCK_STREAM` and `SOCK_SEQPACKET` abstraction, while IDP is used to support the `SOCK_DGRAM` abstraction. The Error protocol is responded to by the kernel to handle and report errors in protocol processing; it is, however, only accessible to user programs through heroic actions.

#### SEE ALSO

`intro(3)`, `byteorder(3N)`, `gethostbyname(3N)`, `getnetent(3N)`,  
`getprotoent(3N)`, `getservent(3N)`, `ns(3N)`, `intro(4N)`, `spp(4P)`,  
`idp(4P)`, `nsip(4)`  
Internet Transport Protocols, Xerox Corporation document  
XSIS-028112  
An Advanced 4.3BSD Interprocess Communication Tutorial

## NAME

nsip - software network interface encapsulating ns packets in ip packets.

## SYNOPSIS

```
#include <netns/ns_if.h>
```

Not currently supported under 2.11BSD

## DESCRIPTION

The nsip interface is a software mechanism which may be used to transmit Xerox NS(tm) packets through otherwise uncooperative networks. It functions by prepending an IP header, and resubmitting the packet through the unix IP machinery.

The super-user can advise the operating system of a willing partner by naming an IP address to be associated with an NS address. Presently, only specific hosts pairs are allowed, and for each host pair, an artificial point-to-point interface is constructed. At some future date, IP broadcast addresses or hosts may be paired with NS networks or hosts.

Specifically, a socket option of SO\_NSIP\_ROUTE is set on a socket of family AF\_NS, type SOCK\_DGRAM, passing the following structure:

```
struct nsip_req {
    struct sockaddr rq_ns;    /* must be ns format destination */
    struct sockaddr rq_ip;    /* must be ip format gateway */
    short rq_flags;
};
```

## DIAGNOSTICS

nsip%d: can't handle af%d. The interface was handed a message with addresses formatted in an unsuitable address family; the packet was dropped.

## SEE ALSO

intro(4N), ns(4F)

## BUGS

It is absurd to have a separate pseudo-device for each pt-to-pt link. There is no way to change the IP address for an NS host once the the encapsulation interface is set up. The request should honor flags of RTF\_GATEWAY to indicate remote networks, and the absence of RTF\_UP should be a clue to remove that partner. This was intended to postpone the necessity of rewriting reverse ARP for the en device, and to allow passing XNS packets through an Arpanet-Milnet gateway, to facilitate testing between some co-operating

universities.

## NAME

    null - data sink

## SYNOPSIS

    major device number(s):  
        raw: 1  
    minor device encoding:  
        mem: 0; kmem: 1; null: 2

## DESCRIPTION

    Data written on a null special file is discarded.

    Reads from a null special file always return 0 bytes.

## FILES

    /dev/null  
    /dev/MAKEDEV script to create special files  
    /dev/MAKEDEV.local script to localize special files

## SEE ALSO

    mem(4)

## NAME

pty - pseudo terminal driver

## SYNOPSIS

/sys/conf/SYSTEM:

NPTY ptys # pseudo-terminals, in groups of 8

## DESCRIPTION

The pty driver provides support for a device-pair termed a pseudo terminal. A pseudo terminal is a pair of character devices, a master device and a slave device. The slave device provides processes an interface identical to that described in tty(4). However, whereas all other devices which provide the interface described in tty(4) have a hardware device of some sort behind them, the slave device has, instead, another process manipulating it through the master half of the pseudo terminal. That is, anything written on the master device is given to the slave device as input and anything written on the slave device is presented as input on the master device.

In configuring, ptys specifies the number of pseudo terminal pairs are configured.

The following ioctl calls apply only to pseudo terminals:

## TIOCSTOP

Stops output to a terminal (e.g. like typing ^S).

Takes no parameter.

## TIOCSTART

Restarts output (stopped by TIOCSTOP or by typing ^S).

Takes no parameter.

## TIOCPKT

Enable/disable packet mode. Packet mode is enabled by specifying (by reference) a nonzero parameter and disabled by specifying (by reference) a zero parameter. When applied to the master side of a pseudo terminal, each subsequent read from the terminal will return data written on the slave part of the pseudo terminal preceded by a zero byte (symbolically defined as TIOCPKT\_DATA), or a single byte reflecting control status information. In the latter case, the byte is an inclusive-or of zero or more of the bits:

## TIOCPKT\_FLUSHREAD

whenever the read queue for the terminal is flushed.

## TIOCPKT\_FLUSHWRITE

whenever the write queue for the terminal is



flushed.

#### TIOCPKT\_STOP

whenever output to the terminal is stopped a la  
^S.

#### TIOCPKT\_START

whenever output to the terminal is restarted.

#### TIOCPKT\_DOSTOP

whenever t\_stopc is ^S and t\_startc is ^Q.

#### TIOCPKT\_NOSTOP

whenever the start and stop characters are not  
^S/^Q.

While this mode is in use, the presence of control status information to be read from the master side may be detected by a select for exceptional conditions.

This mode is used by rlogin(1C) and rlogind(8C) to implement a remote-echoed, locally ^S/^Q flow-controlled remote login with proper back-flushing of output; it can be used by other similar programs.

#### TIOCUCNTL

Enable/disable a mode that allows a small number of simple user ioctl commands to be passed through the pseudo-terminal, using a protocol similar to that of TIOCPKT. The TIOCUCNTL and TIOCPKT modes are mutually exclusive. This mode is enabled from the master side of a pseudo terminal by specifying (by reference) a nonzero parameter and disabled by specifying (by reference) a zero parameter. Each subsequent read from the master side will return data written on the slave part of the pseudo terminal preceded by a zero byte, or a single byte reflecting a user control operation on the slave side. A user control command consists of a special ioctl operation with no data; the command is given as UIOCCMD(n), where n is a number in the range 1-255. The operation value n will be received as a single byte on the next read from the master side. The ioctl UIOCCMD(0) is a no-op that may be used to probe for the existence of this facility. As with TIOCPKT mode, command operations may be detected with a select for exceptional conditions.

#### TIOCREMOTE

A mode for the master half of a pseudo terminal, independent of TIOCPKT. This mode causes input to the pseudo terminal to be flow controlled and not input edited (regardless of the terminal mode). Each write

to the control terminal produces a record boundary for the process reading the terminal. In normal usage, a write of data is like the data typed as a line on the terminal; a write of 0 bytes is like typing an end-of-file character. TIOCREMOTE can be used when doing remote line editing in a window manager, or whenever flow controlled input is required.

#### FILES

/dev/pty[p-r][0-9a-f]	master pseudo terminals
/dev/tty[p-r][0-9a-f]	slave pseudo terminals

#### DIAGNOSTICS

None.

## NAME

qe - DEC DEQNA Q-bus 10 Mb/s Ethernet interface

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NQE  qe_controllers # DEQNA
```

## DESCRIPTION

The qe interface provides access to a 10 Mb/s Ethernet network through the DEC DEQNA Q-bus controller.

Each of the host's network addresses is specified at boot time with an SIOCSIFADDR ioctl. The qe interface employs the address resolution protocol described in arp(4P) to map dynamically between Internet and Ethernet addresses on the local network.

The interface normally tries to use a ``trailer'' encapsulation to minimize copying data on input and output. The use of trailers is negotiated with ARP. This negotiation may be disabled, on a per-interface basis, by setting the IFF\_NOTRAILERS flag with an SIOCSIFFLAGS ioctl.

## DIAGNOSTICS

## SEE ALSO

inet(4F), intro(4N), arp(4P)

## NAME

ra - MSCP disk controller interface

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NRAC ra_controllers # NRAD controllers
    NRAD ra_drives      # RX33/50, RC25, RD51/52/53/54, RA60/80/81/82

/etc/dtab:
    #Name Unit# Addr   Vector Br Handler(s)      # Comments
    ra  ?      172150 154    5  raintr      # uda50, rqdx1/2/3

major device number(s):
    raw: 14
    block: 5
minor device encoding:
    bits 0007 specify partition of RA drive
    bits 0070 specify RA drive
    bits 0300 specify RA controller
```

## DESCRIPTION

This is a driver for the DEC UDA-50 disk controller and for other compatible controllers. The UDA-50 communicates with the host through a packet oriented protocol termed the Mass Storage Control Protocol (MSCP). Consult the file `<pdp/mscp.h>` for a detailed description of this protocol.

Files with minor device numbers 0 through 7 refer to various portions of drive 0; minor devices 8 through 15 refer to drive 1, etc. The standard device names begin with ``ra'` followed by the drive number and then a letter a-h for partitions 0-7 respectively. The character `?` stands here for a drive number in the range 0-7.

The block files access the disk via the system's normal buffering mechanism and may be read and written without regard to physical disk records. There is also a ``raw'` interface which provides for direct transmission between the disk and the user's read or write buffer. A single read or write call results in exactly one I/O operation and therefore raw I/O is considerably more efficient when many words are transmitted. The names of the raw files conventionally begin with an extra ``r.'`

In raw I/O the buffer must begin on a word (even) boundary, and counts should be a multiple of 512 bytes (a disk sector). Likewise seek calls should specify a multiple of 512 bytes.

## DISK SUPPORT

This driver configures the drive type of each drive when it is first opened. Partition information is read from the

disklabel. If there is no label or the label is corrupt then the 'a' partition is used to span the entire drive.

The ra?a partition is normally used for the root file system, the ra?b partition as a swap area, and the ra?c partition for pack-pack copying (it maps the entire disk).

#### FILES

/dev/ra[0-7][a-h]  
 /dev/rra[0-7][a-h]  
 /dev/MAKEDEV script to create special files

#### SEE ALSO

hk(4), ram(4), rk(4), rl(4), rp(4), rx(4), si(4), xp(4),  
 dtab(5), autoconfig(8), disklabel(8)

#### DIAGNOSTICS

rasa %o, state %d. (Additional status information given after a hard i/o error.) The values of the UDA-50 status register and the internal driver state are printed.

ra%d: interrupt in unknown state %d ignored. An interrupt was received when the driver was in an unknown internal state. Indicates a hardware problem or a driver bug.

ra%d: fatal error (%o). The UDA-50 indicated a ``fatal error'' in the status returned to the host. The contents of the status register are displayed.

ra%d,%d: OFFLINE. (Additional status information given after a hard i/o error.) A hard i/o error occurred because the drive was not on-line. The attached unit number and the MSCP unit numbers are printed.

status %o. (Additional status information given after a hard i/o error.) The status information returned from the UDA-50 is tacked onto the end of the hard error message printed on the console.

ra: unknown packet opcode=0%o. An MSCP packet of unknown type was received from the UDA-50. Check the cabling to the controller.

The following errors are interpretations of MSCP error messages returned by the UDA-50 to the host.

ra: %s error, controller error, event 0%o.

ra: %s error, host memory access error, event 0%o, addr 0%o.

ra: %s error, disk transfer error, unit %d, grp 0x%x, hdr 0x%x.

ra: %s error, SDI error, unit %d, event 0%o, hdr 0x%x.

ra: %s error, small disk error, unit %d, event 0%o, cyl %d.

ra: %s error, unknown error, unit %d, format 0%o, event 0%o.

#### BUGS

In raw I/O read and write(2) truncate file offsets to 512-byte block boundaries, and write scribbles on the tail of incomplete blocks. Thus, in programs that are likely to access raw devices, read, write and lseek(2) should always deal in 512-byte multiples.

## NAME

ram - ram disk driver

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NRAM ram_size # RAM disk size (512-byte blocks)

major device number(s):
    block: 3
minor device encoding:
    must be zero (0)
```

## DESCRIPTION

The ram pseudo-device provides a very fast extended memory store. Its use is intended for file systems like /tmp and applications which need to access a reasonably large amount of data quickly.

The amount of memory dedicated to the ram device is controlled by the NRAM definition in units of 512-byte blocks. This is also patchable in the system binary through the variable ram\_size (though a patched system would have to be rebooted before any change took effect; see adb(1)). This makes it easy to test the effects of different ram disk sizes on system performance. It's important to note that any space given to the ram device is permanently allocated at system boot time. Dedicating too much memory can adversely affect system performance by forcing the system to swap heavily as in a memory poor environment.

The block file accesses the ram disk via the system's buffering mechanism through a buffer sharing arrangement with the buffer cache. It may be read and written without regard to physical disk records. There is no 'raw' interface since no speed advantage is gained by such an interface with the ram disk.

## DISK SUPPORT

The ram driver does not support pseudo-disks (partitions). The special files refer to the entire 'drive' as a single sequentially addressed file.

A typical use for the ram disk would be to mount /tmp on it. Note that if this arrangement is recorded in /etc/fstab then /etc/rc will have to be modified slightly to do a mkfs(8) on the ram disk before the standard file system checks are done.

## FILES

```
/dev/ram          block file
/dev/MAKEDEV      script to create special files
/dev/MAKEDEV.local script to localize special files
```

## SEE ALSO

hk(4), ra(4), rl(4), rk(4), rp(4), rx(4), si(4), xp(4)  
dtab(5), autoconfig(8)

## DIAGNOSTICS

ram: no space. There is not enough memory to allocate the space needed by the ram disk. The ram disk is disabled. Any attempts to access it will return an error.

ram: not allocated. No memory was allocated to the ram disk and an attempt was made to open it. Either not enough memory was available at boot time or the kernel variable ram\_size was set to zero.

## BUGS

The ram driver is only available under 2.11BSD.



## NAME

rk - RK-11/RK05 disk

## SYNOPSIS

/sys/conf/SYSTEM:

NRK rk\_drives # RK05

/etc/dtab:

#Name	Unit#	Addr	Vector	Br	Handler(s)	# Comments
rk ?		177400	220	5	rkintr	# rk05

major device number(s):

raw: 15

block: 6

minor device encoding:

specifies drive: <rk\_drive>

## DESCRIPTION

Minor device numbers are drive numbers on one controller. The standard device names begin with ``rk'' followed by the drive number and then the letter "h". The character ? stands here for a drive number in the range 0-7.

The block files access the disk via the system's normal buffering mechanism and may be read and written without regard to physical disk records. There is also a `raw' interface which provides for direct transmission between the disk and the user's read or write buffer. A single read or write call results in exactly one I/O operation and therefore raw I/O is considerably more efficient when many words are transmitted. The names of the raw files conventionally begin with an extra `r.'

In raw I/O the buffer must begin on a word (even) boundary, and counts should be a multiple of 512 bytes (a disk sector). Likewise seek calls should specify a multiple of 512 bytes.

## DISK SUPPORT

The rk driver does not support pseudo-disks (partitions). Each file rk?h refers to the entire drive as a single sequentially addressed file. Each drive has 4872 512-byte blocks.

It's not clear what one would do with one of these drives if one had one ...

## FILES

/dev/rk[0-7]h        block files  
 /dev/rrk[0-7]h     raw files  
 /dev/MAKEDEV script to create special files  
 /dev/MAKEDEV.local script to localize special files

## SEE ALSO

hk(4), ra(4), ram(4), rl(4), rp(4), rx(4), si(4), xp(4),  
dtab(5), autoconfig(8)

## DIAGNOSTICS

rk%d: hard error sn%d er=%b ds=%b. An unrecoverable error occurred during transfer of the specified sector of the specified disk. The contents of the two error registers are also printed in octal and symbolically with bits decoded. The error was either unrecoverable, or a large number of retry attempts could not recover the error.

rk%d: write locked. The write protect switch was set on the drive when a write was attempted. The write operation is not recoverable.

## BUGS

In raw I/O read and write(2) truncate file offsets to 512-byte block boundaries, and write scribbles on the tail of incomplete blocks. Thus, in programs that are likely to access raw devices, read, write and lseek(2) should always deal in 512-byte multiples.

DEC-standard error logging should be supported.

A program to analyze the logged error information (even in its present reduced form) is needed.

## NAME

rl - RL-11/RL01, RL02 moving-head disk

## SYNOPSIS

```
/sys/conf/SYSTEM:
```

```
NRL  rl_drives # RL01/02
```

```
/etc/dtab:
```

#Name	Unit#	Addr	Vector	Br	Handler(s)	# Comments
rl	?	174400	160	5	rlintr	# rl-01/02

```
major device number(s):
```

```
raw: 16
```

```
block: 7
```

```
minor device encoding:
```

```
bits 0007 specify partition of RL drive
```

```
bits 0070 specify RL drive
```

## DESCRIPTION

The block files access the disk via the system's normal buffering mechanism and may be read and written without regard to physical disk records. There is also a 'raw' interface which provides for direct transmission between the disk and the user's read or write buffer. A single read or write call results in exactly one I/O operation and therefore raw I/O is considerably more efficient when many words are transmitted. The names of the raw files conventionally begin with an extra 'r.'

In raw I/O the buffer must begin on a word (even) boundary, and counts should be a multiple of 512 bytes (a disk sector). Likewise seek calls should specify a multiple of 512 bytes.

## DISK SUPPORT

The RL01 drives are each 10240 blocks long and the RL02 drives are 20480 blocks long.

On a RL02 there is room for a full sized root ('a') partition and a reasonable sized swap ('b') partition. The RL01 can only (realistically) have a single 5mb partition.

## FILES

```
/dev/rl[0-3][a-h]  block files
```

```
/dev/rrl[0-3][a-h] raw files
```

```
/dev/MAKEDEV  script to create special files
```

```
/dev/MAKEDEV.local  script to localize special files
```

## SEE ALSO

```
hk(4), ra(4), ram(4), rk(4), rp(4), rx(4), si(4), xp(4),  
dtab(5), autoconfig(8)
```

## DIAGNOSTICS

rl%d: hard error sn%d cs=%b da=%b. An unrecoverable error occurred during transfer of the specified sector of the specified disk. The contents of the two error registers are also printed in octal and symbolically with bits decoded. The error was either unrecoverable, or a large number of retry attempts could not recover the error.

rl%d: hard error sn%d mp=%b da=%b. An unrecoverable drive error occurred during transfer of the specified sector of the specified disk. The contents of the two error registers are also printed in octal and symbolically with bits decoded. The error was either unrecoverable, or a large number of retry attempts could not recover the error.

rl%d: write locked. The write protect switch was set on the drive when a write was attempted. The write operation is not recoverable.

rl%d: can't get status. A ``get status'' command on the specified drive failed. The error is unrecoverable.

## BUGS

In raw I/O read and write(2) truncate file offsets to 512-byte block boundaries, and write scribbles on the tail of incomplete blocks. Thus, in programs that are likely to access raw devices, read, write and lseek(2) should always deal in 512-byte multiples.

DEC-standard error logging should be supported.

A program to analyze the logged error information (even in its present reduced form) is needed.

## NAME

rx, rb, rz - XMODEM, YMODEM, ZMODEM (Batch) file receive

## SYNOPSIS

```
rz [- +abepqtuvy]
rb [- +abqtuvy]
rx [- abceqtuv] file
gz file ...
[-][v]rzCOMMAND
```

## DESCRIPTION

This program uses error correcting protocols to receive files over a dial-in serial port from a variety of programs running under PC-DOS, CP/M, Unix, and other operating systems. It is invoked from a shell prompt manually, or automatically as a result of an "sz file ..." command given to the calling program.

Rz is not intended be called from cu(1), or other communications programs. Unix flavors of Omen Technology's Professional-YAM communications software are available for dial-out applications. (Reg.)

Rz (Receive ZMODEM) receives files with the ZMODEM batch protocol. Pathnames are supplied by the sending program, and directories are made if necessary (and possible). Normally, the "rz" command is automatically issued by the calling ZMODEM program, but some defective ZMODEM implementations may require starting rz the old fashioned way. Rz does not support ZMODEM Crash Recovery and certain other ZMODEM features. Unix flavors of Professional-YAM may be linked to "rz" and used in place of this program to support these ZMODEM features.

Rb receives file(s) with YMODEM, accepting either standard 128 byte sectors or 1024 byte sectors (YAM sb -k option). The user should determine when the 1024 byte block length actually improves throughput without causing lost data or even system crashes.

If True YMODEM (Omen Technology trademark) file information (file length, etc.) is received, the file length controls the number of bytes written to the output dataset, and the modify time and file mode (iff non zero) are set accordingly.

If no True YMODEM file information is received, slashes in the pathname are changed to underscore, and any trailing period in the pathname is eliminated. This conversion is useful for files received from CP/M systems. With YMODEM,

each file name is converted to lower case unless it contains one or more lower case letters.

Rx receives a single file with XMODEM or XMODEM-1k protocol. The user should determine when the 1024 byte block length actually improves throughput without causing problems. The user must supply the file name to both sending and receiving programs. Up to 1023 garbage characters may be added to the received file.

Gz is a shell script which calls sz to command Pro-YAM or ZCOMM to transmit the specified files. Pathnames used with gz must be escaped if they have special significance to the Unix shell.

EXAMPLE: gz "-a C:\*.c D:\*.h"

Rz may be invoked as rzCOMMAND (with an optional leading - as generated by login(1)). For each received file, rz will pipe the file to ``COMMAND filename'' where filename is the name of the transmitted file with the file contents as standard input.

Each file transfer is acknowledged when COMMAND exits with 0 status. A non zero exit status terminates transfers.

A typical use for this form is rzmmail which calls rmail(1) to post mail to the user specified by the transmitted file name. For example, sending the file "caf" from a PC-DOS system to rzmmail on a Unix system would result in the contents of the DOS file "caf" being mailed to user "caf".

On some Unix systems, the login directory must contain a link to COMMAND as login sets SHELL=rsh which disallows absolute pathnames. If invoked with a leading ``v'', rz will report progress to /tmp/rzlog. The following entry works for Unix SYS III/V:

```
    rzmmail::5:1::/bin:/usr/local/rzmmail
```

If the SHELL environment variable includes rsh or rksh (restricted shell), rz will not accept absolute pathnames or references to a parent directory, will not modify an existing file, and removes any files received in error.

If rz is invoked with stdout and stderr to different datasets, Verbose is set to 2, causing frame by frame progress reports to stderr. This may be disabled with the q option.

The meanings of the available options are:

- a Convert files to Unix conventions by stripping carriage returns and all characters beginning with the first Control Z (CP/M end of file).
- b Binary (tell it like it is) file transfer override.
- c Request 16 bit CRC. XMODEM file transfers default to 8 bit checksum. YMODEM and ZMODEM normally use 16 bit CRC.
- D Output file data to /dev/null; for testing. (Unix only)
- e Force sender to escape all control characters; normally XON, XOFF, DLE, CR-@-CR, and Ctrl-X are escaped.
- p (ZMODEM) Protect: skip file if destination file exists.
- q Quiet suppresses verbosity.
- t tim  
Change timeout to tim tenths of seconds.
- v Verbose causes a list of file names to be appended to /tmp/rzlog . More v's generate more output.
- y Yes, clobber any existing files with the same name.

#### EXAMPLES

```
(Pro-YAM command)
<ALT-2>
Pro-YAM Command: sz *.h *.c
(This automatically invokes rz on the connected system.)
```

#### SEE ALSO

ZMODEM.DOC, YMODEM.DOC, Professional-YAM, crc(omen),  
sz(omen), usq(omen), undos(omen)

Compile time options required for various operating systems are described in the source file.

#### NOTES

ZMODEM's support of XOFF/XON flow control allows proper operation in many environments that do not support XMODEM uploads. Unfortunately, not all timesharing systems support input flow control. The TTY input buffering on some systems may not adequately buffer long blocks or streaming input at high speed. You should suspect this problem when you can't send data to the Unix system at high speeds using ZMODEM, but YMODEM-1k or XMODEM-1k, when YMODEM with 128 byte blocks works properly.

The DSZ or Pro-YAM zmodem l numeric parameter may be set to a value between 64 and 1024 to limit the burst length ("zmodem pll28"). Although this compromises ZMODEM's throughput, ZMODEM's superior reliability remains intact.

If a program that does not properly implement the specified file transfer protocol causes rz to "hang" the port after a failed transfer, either wait for rz to time out or keyboard a dozen Ctrl-X characters. Every reported instance of this

problem has been corrected by using ZCOMM, Pro-YAM, DSZ, or other program with a correct implementation of the specified protocol.

Many programs claiming to support YMODEM only support XMODEM with 1k blocks, and they often don't get that quite right.

In the case of a few poorly designed microcomputers, sending serial data to a tty port at sustained high speeds has been known to cause lockups, system halts, kernel panics, and occasional antisocial behaviour. This problem is not unique to rz; CRT terminals with block mode transmission and line noise have the same effect. When experimenting with high speed input to a system, consider rebooting the system if the file transfers are not successful, especially if the personality of the system appears altered.

The Unix "ulimit" parameter must be set high enough to permit large file transfers to Unix.

32 bit CRC code courtesy Gary S. Brown. Directory creation code from John Gilmore's PD TAR program.

#### BUGS

Rz is not intended be called from cu(1), or other communications programs. Unix flavors of Omen Technology's Professional-YAM communications software are available for dial-out applications.

Rz does not support ZMODEM Crash Recovery and many other ZMODEM features. Unix flavors of Professional-YAM may be linked to "rz" to support these features.

Pathnames are restricted to 127 characters. In XMODEM single file mode, the pathname given on the command line is still processed as described above. The ASCII option's CR/LF to NL translation merely deletes CR's; undos(omen) performs a more intelligent translation.

#### VMS VERSION

The VMS version does not set the file time.

VMS C Standard I/O and RMS may interact to modify file contents unexpectedly.

The VMS version does not support invocation as rzCOMMAND. The current VMS version does not support XMODEM, XMODEM-1k, or YMODEM.

According to the VMS documentation, the buffered input routine used on the VMS version of rz introduces a delay of up to one second for each protocol transaction. This delay may



be significant for very short files. Removing the "#define BUFSIZE" line from rz.c will eliminate this delay at the expense of increased CPU utilization.

For high speed operation, try increasing the SYSGEN parameter TTY\_TYPAHDSZ to 256.

The VMS version causes DCL to generate a random off the wall error message under some error conditions; this is a result of the incompatibility of the VMS "exit" function with the Unix/MSDOS standard.

#### ZMODEM CAPABILITIES

Rz supports incoming ZMODEM binary (-b), ASCII (-a), protect (-p), clobber (-y), and append (-+) requests. Other options sent by the sender are ignored. The default is protect (-p) and binary (-b).

The Unix versions support ZMODEM command execution.

#### FILES

rz.c, crctab.c, rbsb.c, zm.c, zmodem.h Unix source files.

rz.c, crctab.c, vrzsz.c, zm.c, zmodem.h, vmodem.h, vvmodem.c, VMS source files.

/tmp/rzlog stores debugging output generated with -vv option (rzlog on VMS).

## NAME

rx - DEC RX02 floppy disk

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NRX rx_units # RX02

/etc/dtab:
    #Name Unit# Addr    Vector Br Handler(s)      # Comments
    rx ?      177170 264    5  rxintr      # rx02

major device number(s):
    raw: 17
    block: 8
minor device encoding:
    bit 01 specifies drive: <rx_drive>
    bit 02 specifies density: single: 0, double: 1
```

## DESCRIPTION

The rx device provides access to a DEC RX02 floppy disk unit with M8256 interface module (RX211 configuration). The RX02 uses 8-inch, single-sided, soft-sectored floppy disks (with pre-formatted industry-standard headers) in either single or double density.

Floppy disks handled by the RX02 contain 77 tracks, each with 26 sectors (for a total of 2,002 sectors). The sector size is 128 bytes for single density, 256 bytes for double density. Single density disks are compatible with the RX01 floppy disk unit and with IBM 3740 Series Diskette 1 systems.

Files with minor device numbers 0 and 1 refer to drives 0 and 1 in single density mode; minor devices 2 and 3 refer to drives 0 and 1 in double density mode. The standard device names begin with `rx' followed by the drive number and then a letter a-b for single and double density access to the drive respectively. The character ? stands here for a drive number in the range 0-1.

The block files access the disk via the system's normal buffering mechanism and may be read and written without regard to physical disk records. There is also a `raw' interface which provides for direct transmission between the disk and the user's read or write buffer. A single read or write call results in exactly one I/O operation and therefore raw I/O is considerably more efficient when many words are transmitted. The names of the raw files conventionally begin with an extra `r.'

In raw I/O the buffer must begin on a word (even) boundary, and counts should be a multiple of the disk's sector size

(either 128 or 256 bytes depending on selected density.) Likewise seek calls should specify a multiple of the disk's sector size.

In addition to normal ('block' and 'raw') i/o, the driver supports formatting of disks for either density.

#### DISK SUPPORT

The rx driver does not support pseudo-disks (partitions). Each file rx?[ab] refers to the entire drive as a single sequentially addressed file. The physical disk sector size is 128 bytes for single density and 256 bytes for double density. The logical block size is 1024 bytes. Each diskette has 250 logical blocks in single density and 500 logical blocks in double density.

The size and density of the disks are specified as follows:

disk	minor	device	unit	density
rx0a	0	0		single
rx1a	1	1		single
rx0b	2	0		double
rx1b	3	1		double

Even though the storage capacity on a floppy disk is quite small, it is possible to make file systems on them. For example, the commands

```
% rxformat /dev/rrx0a
```

```
% newfs /dev/rrx0a
```

and

```
% rxformat /dev/rrx1b
```

```
% newfs /dev/rrx1b
```

format and make file systems on the single density disk in the left drive and the double density disk in the right drive with 241 and 486 kbytes available, respectively, for file storage. Using tar(1) gives somewhat more efficient utilization of the available space for file storage. The RX02 disks are primarily useful for small archives and transfer of small systems or distributions.

An ioctl(2) call is available to format RX02 disks:

```
#include <pdpuba/rxreg.h>
```

```
ioctl(fildes, RXIOC_FORMAT)
```

The density is determined by the device opened.

#### FILES

```
/dev/rx[01][ab]    block files
```

```
/dev/rrx[01][ab]   raw files
```

```
/dev/MAKEDEV       script to create special files
```

```
/dev/MAKEDEV.local  script to localize special files
```

## SEE ALSO

hk(4), ra(4), ram(4), rk(4), rl(4), rp(4), si(4), xp(4),  
dtab(5), autoconfig(8), rxformat(8V)

## DIAGNOSTICS

rx2%d: hard error sn%d cs=%b er=%b. An unrecoverable error occurred during transfer of the specified sector of the specified disk. The contents of the two error registers are also printed in octal and symbolically with bits decoded. The error was either unrecoverable, or a large number of retry attempts could not recover the error.

## BUGS

In raw I/O read and write(2) truncate file offsets to disk sector size block boundaries (either 128 or 256 bytes depending on the selected density), and write scribbles on the tail of incomplete blocks. Thus, in programs that are likely to access raw devices, read, write and lseek(2) should always deal in disk sector size multiples.

DEC-standard error logging should be supported.

A program to analyze the logged error information (even in its present reduced form) is needed.

The 4.3BSD rx driver which supports more capabilities should be ported to 2.11BSD.

## NAME

si - SI 9500/CDC 9766 moving head disk

## SYNOPSIS

/sys/conf/SYSTEM:

NSI si\_drives # SI 9500 driver for CDC 9766 disks

/etc/dtab:

#Name	Unit#	Addr	Vector	Br	Handler(s)	# Comments
si	?	176700	170	5	siintr	# si9500

major device number(s):

raw: 18

block: 9

minor device encoding:

bits 0007 specify partition of SI drive

bits 0070 specify SI drive

## DESCRIPTION

Files with minor device numbers 0 through 7 refer to various portions of drive 0; minor devices 8 through 15 refer to drive 1, etc. The standard device names begin with ``si'' followed by the drive number and then a letter a-h for partitions 0-7 respectively. The character ? stands here for a drive number in the range 0-7.

The block files access the disk via the system's normal buffering mechanism and may be read and written without regard to physical disk records. There is also a `raw' interface which provides for direct transmission between the disk and the user's read or write buffer. A single read or write call results in exactly one I/O operation and therefore raw I/O is considerably more efficient when many words are transmitted. The names of the raw files conventionally begin with an extra `r.'

In raw I/O the buffer must begin on a word (even) boundary, and counts should be a multiple of 512 bytes (a disk sector). Likewise seek calls should specify a multiple of 512 bytes.

## DISK SUPPORT

The origin and size (in sectors) of the pseudo-disks on each drive are as follows:

SI 9500/CDC9766 partitions:

disk	start	length	cyls	comments
xp?a	0	9120	0 - 14	/
xp?b	9120	9120	15 - 29	swap
xp?c	18240	234080	30 - 414	
xp?d	252320	247906	415 - 822*	
xp?e	18240	164160	30 - 299	/usr
xp?f	182400	152000	300 - 549	
xp?g	334400	165826	550 - 822*	
xp?h	0	500384	0 - 822	whole pack

Those partitions marked with an asterisk (``\*') actually stop short of the indicated ending cylinder to protect any bad block forwarding information on the packs. The indicated lengths are correct. Partition ``h' must be used to access the bad block forwarding area. N.B.: the si driver does not support bad block forwarding; the space is reserved in the event bad block forwarding is ever added to the driver.

#### FILES

```
/dev/si[0-7][a-h]  block files
/dev/rsi[0-7][a-h] raw files
/dev/MAKEDEV      script to create special files
/dev/MAKEDEV.local script to localize special files
```

#### SEE ALSO

```
hk(4), ra(4), ram(4), rk(4), rl(4), rp(4), rx(4), xp(4),
dtab(5), autoconfig(8)
```

#### DIAGNOSTICS

si%d%c: hard error sn%d cnr=%b err=%b. An unrecoverable error occurred during transfer of the specified sector of the specified disk partition. The contents of the two error registers are also printed in octal and symbolically with bits decoded. The error was either unrecoverable, or a large number of retry attempts (including offset positioning and drive recalibration) could not recover the error.

si%d%c: hard error sn%d ssr=%b err=%b. An unrecoverable error occurred during transfer of the specified sector of the specified disk partition. The contents of the two error registers are also printed in octal and symbolically with bits decoded. The error was either unrecoverable, or a large number of retry attempts (including offset positioning and drive recalibration) could not recover the error.

#### BUGS

In raw I/O read and write(2) truncate file offsets to 512-byte block boundaries, and write scribbles on the tail of incomplete blocks. Thus, in programs that are likely to access raw devices, read, write and lseek(2) should always

deal in 512-byte multiples.

The partition tables for the file systems should be read off of each pack, as they are never quite what any single installation would prefer, and this would make packs more portable.

## NAME

spp - Xerox Sequenced Packet Protocol

## SYNOPSIS

```
#include <sys/socket.h>
#include <netns/ns.h>
s = socket(AF_NS, SOCK_STREAM, 0);

#include <netns/sp.h>
s = socket(AF_NS, SOCK_SEQPACKET, 0);
```

## DESCRIPTION

The SPP protocol provides reliable, flow-controlled, two-way transmission of data. It is a byte-stream protocol used to support the SOCK\_STREAM abstraction. SPP uses the standard NS(tm) address formats.

Sockets utilizing the SPP protocol are either "active" or "passive". Active sockets initiate connections to passive sockets. By default SPP sockets are created active; to create a passive socket the listen(2) system call must be used after binding the socket with the bind(2) system call. Only passive sockets may use the accept(2) call to accept incoming connections. Only active sockets may use the connect(2) call to initiate connections.

Passive sockets may "underspecify" their location to match incoming connection requests from multiple networks. This technique, termed "wildcard addressing", allows a single server to provide service to clients on multiple networks. To create a socket which listens on all networks, the NS address of all zeroes must be bound. The SPP port may still be specified at this time; if the port is not specified the system will assign one. Once a connection has been established the socket's address is fixed by the peer entity's location. The address assigned the socket is the address associated with the network interface through which packets are being transmitted and received. Normally this address corresponds to the peer entity's network.

If the SOCK\_SEQPACKET socket type is specified, each packet received has the actual 12 byte sequenced packet header left for the user to inspect:

```
struct sphdr {
    u_char    sp_cc;           /* connection control */
#define SP_EM 0x10           /* end of message */
    u_char    sp_dt;           /* datastream type */
    u_short   sp_sid;
    u_short   sp_did;
    u_short   sp_seq;
    u_short   sp_ack;
    u_short   sp_alo;
```



```
};
```

This facilitates the implementation of higher level Xerox protocols which make use of the data stream type field and the end of message bit. Conversely, the user is required to supply a 12 byte header, the only part of which inspected is the data stream type and end of message fields.

For either socket type, packets received with the Attention bit sent are interpreted as out of band data. Data sent with `send(..., ..., ..., MSG_OOB)` cause the attention bit to be set.

#### DIAGNOSTICS

A socket operation may fail with one of the following errors returned:

- [EISCONN]               when trying to establish a connection on  
                          a socket which already has one;
- [ENOBUFFS]             when the system runs out of memory for  
                          an internal data structure;
- [ETIMEDOUT]            when a connection was dropped due to  
                          excessive retransmissions;
- [ECONNRESET]           when the remote peer forces the connec-  
                          tion to be closed;
- [ECONNREFUSED]         when the remote peer actively refuses  
                          connection establishment (usually  
                          because no process is listening to the  
                          port);
- [EADDRINUSE]           when an attempt is made to create a  
                          socket with a port which has already  
                          been allocated;
- [EADDRNOTAVAIL]        when an attempt is made to create a  
                          socket with a network address for which  
                          no network interface exists.

#### SOCKET OPTIONS

- SO\_DEFAULT\_HEADERS    when set, this determines the data  
                          stream type and whether the end of mes-  
                          sage bit is to be set on every ensuing  
                          packet.
- SO\_MTU                 This specifies the maximum ammount of  
                          user data in a single packet.        The  
                          default is 576 bytes - sizeof(struct  
                          spidp). This quantity affects windowing  
                          -- increasing it without increasing the

amount of buffering in the socket will lower the number of unread packets accepted. Anything larger than the default will not be forwarded by a bona fide XEROX product internetwork router. The data argument for the setsockopt call must be an unsigned short.

SEE ALSO

intro(4N), ns(4F)

BUGS

There should be some way to reflect record boundaries in a stream. For stream mode, there should be an option to get the data stream type of the record the user process is about to receive.

## NAME

sri - DR11-C IMP interface

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NSRI sri_controllers      # SRI DR11c ARPAnet IMP
```

## DESCRIPTION

## DIAGNOSTICS

sri%d: not alive. The initialization routine was entered even though the device did not autoconfigure. This indicates a system problem.

sri%d: can't initialize. Insufficient UNIBUS resources existed to initialize the device. This is likely to occur when the device is run on a buffered data path on an 11/750 and other network interfaces are also configured to use buffered data paths, or when it is configured to use buffered data paths on an 11/730 (which has none).

sri%d: imp doesn't respond, ibf=%b. The driver attempted to initialize the device, but the IMP failed to respond after 5 tries. Check the cabling.

sri%d: stray xmit interrupt, csr=%b. An interrupt occurred when no output had previously been started.

sri%d: output error, csr=%b. The device indicated a problem sending data on output.

sri%d: input error, ibf=%b. The device indicated a problem receiving data on input.

sri%d: bad length=%d. An input operation resulted in a data transfer of less than 10 or more than 1008 bytes of data into memory (according to the word count register). This should never happen as the maximum size of a host-IMP message is 1018 bytes.

## NAME

swap - swap device

## DESCRIPTION

This file refers to the swap device in use by the system. It should be linked to the special file for the disk or disk partition on which the system swaps. It is used by programs that wish to read from the swap device without knowing its real name. The script /dev/MAKEDEV.local normally is edited by the system builder to set /dev/swap up correctly.

## FILES

/dev/swap  
/dev/MAKEDEV script to create special files  
/dev/MAKEDEV.local script to localize special files

## SEE ALSO

ps(1), w(1)

## NAME

tb - line discipline for digitizing devices

## SYNOPSIS

```
/sys/conf/SYSTEM:
NTB 1      # RS232 interface for Genisco/Hitachi tablets
```

## DESCRIPTION

This line discipline provides a polled interface to many common digitizing devices which are connected to a host through a serial line. When these devices stream data at high speed, the use of the line discipline is critical in minimizing the number of samples that would otherwise be lost due to buffer exhaustion in the tty(4) handler.

The line discipline is enabled by a sequence:

```
#include <sys/tablet.h>
int ldisc = TBLDISC, fildes; ...
ioctl(fildes, TIOCSETD, &ldisc);
```

A typical application program then polls the digitizing device by reading a binary data structure which contains: the current X and Y positions (in the device coordinate space), up-down status of the buttons or pen stylus, proximity information (when available), and a count of the number of samples received from the input device since it was opened. In addition, devices such as the GTCO append tilt and pressure information to the end of the aforementioned structure. For the Polhemus 3-D digitizer the structure read is completely different. Refer to the include file for a complete description.

While in tablet mode, normal teletype input and output functions take place. Thus, if an 8 bit output data path is desired, it is necessary to prepare the output line by putting it into RAW mode using ioctl(2). This must be done before changing the discipline with TIOCSETD, as most ioctl(2) calls are disabled while in tablet line-discipline mode.

The line discipline supports ioctl(2) requests to get/set the operating mode, and to get/set the tablet type and operating mode by or-ing the two values together.

The line discipline supports digitizing devices which are compatible with Hitachi, GTCO, or Polhemus protocol formats. For Hitachi there are several formats with that used in the newer model HDG-1111B the most common.

## SEE ALSO

tty(4)

DIAGNOSTICS  
None.

## NAME

tcp - Internet Transmission Control Protocol

## SYNOPSIS

```
#include <sys/socket.h>
#include <netinet/in.h>
```

```
s = socket(AF_INET, SOCK_STREAM, 0);
```

## DESCRIPTION

The TCP protocol provides reliable, flow-controlled, two-way transmission of data. It is a byte-stream protocol used to support the SOCK\_STREAM abstraction. TCP uses the standard Internet address format and, in addition, provides a per-host collection of "port addresses". Thus, each address is composed of an Internet address specifying the host and network, with a specific TCP port on the host identifying the peer entity.

Sockets utilizing the tcp protocol are either "active" or "passive". Active sockets initiate connections to passive sockets. By default TCP sockets are created active; to create a passive socket the listen(2) system call must be used after binding the socket with the bind(2) system call. Only passive sockets may use the accept(2) call to accept incoming connections. Only active sockets may use the connect(2) call to initiate connections.

Passive sockets may "underspecify" their location to match incoming connection requests from multiple networks. This technique, termed "wildcard addressing", allows a single server to provide service to clients on multiple networks. To create a socket which listens on all networks, the Internet address INADDR\_ANY must be bound. The TCP port may still be specified at this time; if the port is not specified the system will assign one. Once a connection has been established the socket's address is fixed by the peer entity's location. The address assigned the socket is the address associated with the network interface through which packets are being transmitted and received. Normally this address corresponds to the peer entity's network.

TCP supports one socket option which is set with setsockopt(2) and tested with getsockopt(2). Under most circumstances, TCP sends data when it is presented; when outstanding data has not yet been acknowledged, it gathers small amounts of output to be sent in a single packet once an acknowledgement is received. For a small number of clients, such as window systems that send a stream of mouse events which receive no replies, this packetization may cause significant delays. Therefore, TCP provides a boolean option, TCP\_NODELAY (from <netinet/tcp.h>, to defeat this

algorithm. The option level for the `setsockopt` call is the protocol number for TCP, available from `getprotobyname(3N)`.

Options at the IP transport level may be used with TCP; see `ip(4P)`. Incoming connection requests that are source-routed are noted, and the reverse source route is used in responding.

#### DIAGNOSTICS

A socket operation may fail with one of the following errors returned:

- [EISCONN]               when trying to establish a connection on  
                          a socket which already has one;
- [ENOBUFFS]             when the system runs out of memory for  
                          an internal data structure;
- [ETIMEDOUT]            when a connection was dropped due to  
                          excessive retransmissions;
- [ECONNRESET]           when the remote peer forces the connec-  
                          tion to be closed;
- [ECONNREFUSED]         when the remote peer actively refuses  
                          connection establishment (usually  
                          because no process is listening to the  
                          port);
- [EADDRINUSE]           when an attempt is made to create a  
                          socket with a port which has already  
                          been allocated;
- [EADDRNOTAVAIL]        when an attempt is made to create a  
                          socket with a network address for which  
                          no network interface exists.

#### SEE ALSO

`getsockopt(2)`, `socket(2)`, `intro(4N)`, `inet(4F)`, `ip(4P)`



## NAME

tm - TM-11/TE-10 magtape interface

## SYNOPSIS

/sys/conf/SYSTEM:

```
# Setting AVIVTM configures the TM driver for the AVIV 800/1600/6250
# controller.  For more details, see /sys/pdpuba/tm.c.
NTM tm_drives # TM11
AVIVTM YES      # AVIV 800/1600/6250 controller
```

/etc/dtab:

#Name	Unit#	Addr	Vector	Br	Handler(s)	# Comments
tm	?	172520	224	5	tmintr	# tm11 driver

major device number(s):

raw: 7

block: 1

minor device encoding:

bits 0003 specify TS drive

bit 0004 specifies no-rewind operation

bits 0030 specify recording density:

0000: 800BPI

0010: 1600BPI (AVIVTM and some other controllers)

0020: 6250BPI (AVIVTM only)

## DESCRIPTION

The tm-11/te-10 combination provides a standard tape drive interface as described in mtio(4). The standard DEC tm-11 operates only at 800 bpi. Other controllers of this type may also allow operation at 1600 bpi, under software control or by switching manually.

## FILES

/dev/MAKEDEV script to create special files

/dev/MAKEDEV.local script to localize special files

## SEE ALSO

mt(1), tar(1), tp(1), mtio(4), ht(4), ts(4), mt(4), dtab(5),  
autoconfig(8)

## DIAGNOSTICS

te%d: no write ring. An attempt was made to write on the tape drive when no write ring was present; this message is written on the terminal of the user who tried to access the tape.

te%d: not online. An attempt was made to access the tape while it was offline; this message is written on the terminal of the user who tried to access the tape.

te%d: can't change density in mid-tape. An attempt was made to write on a tape at a different density than is already

recorded on the tape. This message is written on the terminal of the user who tried to switch the density.

te%d: hard error bn%d er=%b. A tape error occurred at block bn; the tm error register is printed in octal with the bits symbolically decoded. Any error is fatal on non-raw tape; when possible the driver will have retried the operation which failed several times before reporting the error.

te%d: lost interrupt. A tape operation did not complete within a reasonable time, most likely because the tape was taken off-line during rewind or lost vacuum. The controller should, but does not, give an interrupt in these cases. The device will be made available again after this message, but any current open reference to the device will return an error as the operation in progress aborts.

#### BUGS

If any non-data error is encountered on non-raw tape, it refuses to do anything more until closed.

## NAME

tmscp - DEC TMSCP magtape interface

## SYNOPSIS

/sys/conf/SYSTEM:

```

NTMSCP  1  # TMSCP controllers
NTMS     1  # TMSCP drives
TMSCP_DEBUG NO  # debugging code in in TMSCP drive (EXPENSIVE)

```

/etc/dtab:

#Name	Unit#	Addr	Vector	Br	Handler(s)	# Comments
tms ?		174500	260	5	tmsintr	# tmscp driver
tms ?		164334	0	5	tmsintr	# alternate

major device number(s):

```

raw: 23
block: 12

```

minor device encoding:

```

bit:  | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
      -----
      C  C  X  D  D  N  U  U

```

C = Controller # (max of 4 controllers)

D = Density

N = Norewind on close

U = Unit (drive) number (max of 4 drives per controller)

## DESCRIPTION

Tape controllers compatible with the DEC Tape Mass Storage Control Protocol (TMSCP) architecture such as the TU81 and the TK50 provide a standard tape drive interface as described in mtio(4). The controller communicates with the host through a packet oriented protocol. Consult the file <pdp/tmscp.h> for a detailed description of this protocol.

## FILES

```

/dev/MAKEDEV script to create special files
/dev/MAKEDEV.local script to localize special files

```

The vector specified in /etc/dtab may be given as an explicit value in which case autoconfig(8) will attempt to allocate the specified vector. The vector may also be (if the system is at revision level 100 or higher) given as 0 - in which case autoconfig(8) will request the kernel to allocate the next available vector (beginning at 01000 and proceeding downward) which autoconfig(8) will initialize.

Multiple drives on a single controller are supported. Multiple controllers are also supported.

Density selection is specified by bits 3 and 4 in the minor device number. A value of 0 requests the lowest density of a drive and a value of 3 the highest density. Values of 1 and 2 are essentially equivalent (because because TMSCP only defines values for three densities) and request the middle density (for tri-density tape drives).

#### DIAGNOSTICS

tms: !drives. Not enough drives were declared when the kernel was built. The NTMS parameter in the kernel config file needs to be increased by at least one.

tms%d stepN init failed: sa %x. Step N of the 4 step initialization sequence has failed.

tms%d: random intr. An unexpected interrupt was received. This is indicative of some other device using the same interrupt vector as the TMSCP controller. The interrupt is ignored.

tms%d Ver %d Mod %d. The version and model number of the controller are displayed when the controller is initialized. This is an information message and not an error.

tms%d: state %d. The controller state is not one of the 4 initialization states or the RUN state. This indicates a serious problem in the driver and possibly the hardware.

tms%d: fatal error %x. The controller detected a ``fatal error'' in the status returned to the host. The contents of the status register are displayed.

tms%d init fail. The controller failed to initialize. Indicative of a hardware problem.

tms%d,%d flush fail. The cache failed to flush during a close operation. Data loss is possible in this case.

tms%d,%d: sa %x state %d. A fatal error. The controller will be reset in an attempt to resume operation.

tms ioctl. An invalid internal ioctl function call has been made. This is a driver bug.

tms%d,%d inv end. An invalid end (completion) code has been detected. A drive has returned 0 as the opcode originally issued. This is a hardware problem.

tms%d,%d bad rsp: %x. An unrecognized response has been received. This is a driver bug.

tms%d,%d cache lost The cache on the drive has been lost. Data loss is likely. Usually due to a hardware problem.

The following error is produced when a TMSCP error log datagram is received:

tms%d,%d dgram fmt=%x evt=%x grp=%x flg=%x pos=%D.

Earlier versions of the driver consumed a noticeable amount of kernel D-space decoding and pretty printing more detailed information. This has been removed in favor of a shorter message. In the future an error log daemon will be written and the datagrams from the MSCP and TMSCP drivers passed to it for analysis.

#### SEE ALSO

mt(1), tar(1), tp(1), mtio(4), tm(4), ts(4), ut(4),  
dmesg(8), dtab(5), autoconfig(8)

#### BUGS

If any non-data error is encountered on non-raw tape, it refuses to do anything more until closed.

On quad-density tape drives (the Kennedy 9662 for example) the middle density of 3200bpi is not host selectable (it can be manually selected from the drive's front panel) because TMSCP only defines 800, 1600 and 6250bpi.

## NAME

ts - TS-11 magtape interface

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NTS  ts_drives      # TS11

/etc/dtab:
    #Name Unit# Addr   Vector Br Handler(s)      # Comments
    ts  ?      172520 224    5  tsintr        # ts11 driver

major device number(s):
    raw: 8
    block: 2
minor device encoding:
    bits 0003 specify TS drive
    bit 0004 specifies no-rewind operation
    bit 0010 ignored
```

## DESCRIPTION

The ts-11 combination provides a standard tape drive interface as described in mtio(4). The ts-11 operates only at 1600 bpi, and only one transport is possible per controller.

## FILES

```
/dev/MAKEDEV script to create special files
/dev/MAKEDEV.local script to localize special files
```

## SEE ALSO

mt(1), tar(1), tp(1), mtio(4), ht(4), tm(4), dtab(5), auto-config(8)

## DIAGNOSTICS

ts%d: no write ring. An attempt was made to write on the tape drive when no write ring was present; this message is written on the terminal of the user who tried to access the tape.

ts%d: not online. An attempt was made to access the tape while it was offline; this message is written on the terminal of the user who tried to access the tape.

ts%d: hard error bn%d xs0=%b xs1=%b xs2=%b xs3=%b. A hard error occurred on the tape at block bn; status registers 0-3 are printed in octal and symbolically decoded as bits.

ts%d: addr mod 4 != 0. The address of a TS-11 command packet was not on an even longword boundary.

## BUGS

If any non-data error is encountered on non-raw tape, it refuses to do anything more until closed.

The device lives at the same address as a tm-11 tm(4).

## NAME

tty - general terminal interface

## SYNOPSIS

```
#include <sgtty.h>
```

## DESCRIPTION

This section describes both a particular special file /dev/tty and the terminal drivers used for conversational computing.

Line disciplines.

The system provides different line disciplines for controlling communications lines. In this version of the system there are two disciplines available for use with terminals:

old        The old (Version 7) terminal driver. This is sometimes used when using the standard shell sh(1).

new        The standard Berkeley terminal driver, with features for job control; this must be used when using csh(1).

Line discipline switching is accomplished with the TIOCSSTD ioctl:

```
int ldisc = LDISC;
ioctl(f, TIOCSSTD, &ldisc);
```

where LDISC is OTTYDISC for the standard tty driver and NTTYDISC for the ``new'' driver. The standard (currently old) tty driver is discipline 0 by convention. Other disciplines may exist for special purposes, such as use of communications lines for network connections. The current line discipline can be obtained with the TIOCGETD ioctl. Pending input is discarded when the line discipline is changed.

All of the low-speed asynchronous communications ports can use any of the available line disciplines, no matter what hardware is involved. The remainder of this section discusses the "old" and "new" disciplines.

The control terminal.

When a terminal file is opened, it causes the process to wait until a connection is established. In practice, user programs seldom open these files; they are opened by getty(8) or rlogind(8C) and become a user's standard input and output file.



If a process which has no control terminal opens a terminal file, then that terminal file becomes the control terminal for that process. The control terminal is thereafter inherited by a child process during a `fork(2)`, even if the control terminal is closed.

The file `/dev/tty` is, in each process, a synonym for a control terminal associated with that process. It is useful for programs that wish to be sure of writing messages on the terminal no matter how output has been redirected. It can also be used for programs that demand a file name for output, when typed output is desired and it is tiresome to find out which terminal is currently in use.

A process can remove the association it has with its controlling terminal by opening the file `/dev/tty` and issuing an

```
ioctl(f, TIOCNOTTY, 0);
```

This is often desirable in server processes.

Process groups.

Command processors such as `csh(1)` can arbitrate the terminal between different jobs by placing related jobs in a single process group and associating this process group with the terminal. A terminal's associated process group may be set using the `TIOCSGRP` `ioctl(2)`:

```
ioctl(fildes, TIOCSGRP, &pgrp);
```

or examined using `TIOCGGRP`, which returns the current process group in `pgrp`. The new terminal driver aids in this arbitration by restricting access to the terminal by processes which are not in the current process group; see Job access control below.

Modes.

The terminal drivers have three major modes, characterized by the amount of processing on the input and output characters:

**cooked**     The normal mode. In this mode lines of input are collected and input editing is done. The edited line is made available when it is completed by a newline, or when the `t_brkc` character (normally undefined) or `t_eofc` character (normally an EOT, control-D, hereafter `^D`) is entered. A carriage return is usually made synonymous with newline in this mode, and replaced with a newline whenever it

is typed. All driver functions (input editing, interrupt generation, output processing such as tab expansion, etc.) are available in this mode.

**CBREAK** This mode eliminates the character, word, and line editing input facilities, making the input character available to the user program as it is typed. Flow control, literal-next and interrupt processing are still done in this mode. Output processing is done.

**RAW** This mode eliminates all input processing and makes all input characters available as they are typed; no output processing is done either.

The style of input processing can also be very different when the terminal is put in non-blocking I/O mode; see the **FNDELAY** flag described in **fcntl(2)**. In this case a **read(2)** from the control terminal will never block, but rather return an error indication (**EWOULDBLOCK**) if there is no input available.

A process may also request that a **SIGIO** signal be sent it whenever input is present and also whenever output queues fall below the low-water mark. To enable this mode the **FASYNC** flag should be set using **fcntl(2)**.

Input editing.

A UNIX terminal ordinarily operates in full-duplex mode. Characters may be typed at any time, even while output is occurring, and are only lost when the system's character input buffers become completely choked, which is rare, or when the user has accumulated the maximum allowed number of input characters that have not yet been read by some program. Currently this limit is 256 characters. In **RAW** mode, the terminal driver throws away all input and output without notice when the limit is reached. In **CBREAK** or cooked mode it refuses to accept any further input and, if in the new line discipline, rings the terminal bell.

Input characters are normally accepted in either even or odd parity with the parity bit being stripped off before the character is given to the program. By clearing either the **EVEN** or **ODD** bit in the flags word it is possible to have input characters with that parity discarded (see the Summary below.)

In all of the line disciplines, it is possible to simulate terminal input using the **TIOCSTI** ioctl, which takes, as its third argument, the address of a character. The system pretends that this character was typed on the argument

terminal, which must be the control terminal except for the super-user (this call is not in standard version 7 UNIX).

Input characters are normally echoed by putting them in an output queue as they arrive. This may be disabled by clearing the ECHO bit in the flags word using the `stty(3C)` call or the `TIOCSETN` or `TIOCSETP` ioctls (see the Summary below).

In cooked mode, terminal input is processed in units of lines. A program attempting to read will normally be suspended until an entire line has been received (but see the description of `SIGTTIN` in Job access control and of `FIONREAD` in Summary, both below.) No matter how many characters are requested in the read call, at most one line will be returned. It is not, however, necessary to read a whole line at once; any number of characters may be requested in a read, even one, without losing information.

During input, line editing is normally done, with the erase character `sg_erase` (by default, `DELETE`) logically erasing the last character typed and the `sg_kill` character (default, `^U`: control-U) logically erasing the entire current input line. These characters never erase beyond the beginning of the current input line or an eof.

The drivers normally treat either a carriage return or a newline character as terminating an input line, replacing the return with a newline and echoing a return and a line feed. If the `CRMOD` bit is cleared in the local mode word then the processing for carriage return is disabled, and it is simply echoed as a return, and does not terminate cooked mode input.

In the new driver there is a literal-next character (normally `^V`) which can be typed in both cooked and `CBREAK` mode preceding any character to prevent its special meaning to the terminal handler.

The new terminal driver also provides two other editing characters in normal mode. The word-erase character, normally `^W`, erases the preceding word, but not any spaces before it. For the purposes of `^W`, a word is defined as a sequence of non-blank characters, with tabs counted as blanks. Finally, the reprint character, normally `^R`, retypes the pending input beginning on a new line. Retyping occurs automatically in cooked mode if characters which would normally be erased from the screen are fouled by program output.

Input echoing and redisplay

The terminal driver has several modes for handling the echoing of terminal input, controlled by bits in a local mode word.

Hardcopy terminals. When a hardcopy terminal is in use, the LPRTERA bit is normally set in the local mode word. Characters which are logically erased are then printed out backwards preceded by `\' and followed by `/' in this mode.

CRT terminals. When a CRT terminal is in use, the LCRTBS bit is normally set in the local mode word. The terminal driver then echoes the proper number of erase characters when input is erased; in the normal case where the erase character is a ^H this causes the cursor of the terminal to back up to where it was before the logically erased character was typed. If the input has become fouled due to interspersed asynchronous output, the input is automatically retyped.

Erasing characters from a CRT. When a CRT terminal is in use, the LCRTERA bit may be set to cause input to be erased from the screen with a "backspace-space-backspace" sequence when character or word deleting sequences are used. A LCRTKIL bit may be set as well, causing the input to be erased in this manner on line kill sequences as well.

Echoing of control characters. If the LCTLECH bit is set in the local state word, then non-printing (control) characters are normally echoed as ^X (for some X) rather than being echoed unmodified; delete is echoed as ^?.

The normal modes for use on CRT terminals are speed dependent. At speeds less than 1200 baud, the LCRTERA and LCRTKILL processing is painfully slow, and stty(1) normally just sets LCRTBS and LCTLECH; at speeds of 1200 baud or greater all of these bits are normally set. Stty(1) summarizes these option settings and the use of the new terminal driver as "newcrt."

Output processing.

When one or more characters are written, they are actually transmitted to the terminal as soon as previously-written characters have finished typing. (As noted above, input characters are normally echoed by putting them in the output queue as they arrive.) When a process produces characters more rapidly than they can be typed, it will be suspended when its output queue exceeds some limit. When the queue has drained down to some threshold the program is resumed. Even parity is normally generated on output. The EOT character is not transmitted in cooked mode to prevent terminals that respond to it from hanging up; programs using RAW or CBREAK mode should be careful.

The terminal drivers provide necessary processing for cooked and CBREAK mode output including parity generation. The driver will also optionally expand tabs into spaces, where the tab stops are assumed to be set every eight columns, and optionally convert newlines to carriage returns followed by newline. These functions are controlled by bits in the `tty flags` word; see Summary below.

Finally, in the new terminal driver, there is a output flush character, normally `^O`, which sets the `LFLUSHO` bit in the local mode word, causing subsequent output to be flushed until it is cleared by a program or more input is typed. This character has effect in both cooked and CBREAK modes and causes pending input to be retyped if there is any pending input. An `ioctl` to flush the characters in the input or output queues, `TIOCFLUSH`, is also available.

#### Flow control.

There are two characters (the stop character, normally `^S`, and the start character, normally `^Q`) which cause output to be suspended and resumed respectively. Extra stop characters typed when output is already stopped have no effect, unless the start and stop characters are made the same, in which case output resumes.

A bit in the `flags` word may be set to put the terminal into TANDEM mode. In this mode the system produces a stop character (default `^S`) when the input queue is in danger of overflowing, and a start character (default `^Q`) when the input has drained sufficiently. This mode is useful when the terminal is actually another machine that obeys those conventions.

If software flow control is not sufficient (it begins failing above 4800 bits per second on slow systems) RTS/CTS flow control is now available. Hardware flow control causes RTS to be dropped when the remote system (modem) should cease sending and to be raised when additional data can be accepted. If the remote system drops CTS then the local host ceases sending data until the CTS is raised. Hardware flow control is enabled via the `TIOCLBIS` (or `TIOCLBIC` or `TIOCLSET`) function and the bit `LRTSCTS`.

#### Line control and breaks.

There are several `ioctl` calls available to control the state of the terminal line. The `TIOCSBRK` `ioctl` will set the break bit in the hardware interface causing a break condition to exist; this can be cleared (usually after a delay with `sleep(3)`) by `TIOCCBRK`. Break conditions in the input are reflected as a null character in RAW mode or as the

interrupt character in cooked or CBREAK mode. The TIOCCDTR ioctl will clear the data terminal ready condition; it can be set again by TIOCSDTR.

When the carrier signal from the dataset drops (usually because the user has hung up his terminal) a SIGHUP hangup signal is sent to the processes in the distinguished process group of the terminal; this usually causes them to terminate. The SIGHUP can be suppressed by setting the LNOHANG bit in the local state word of the driver. Access to the terminal by other processes is then normally revoked, so any further reads will fail, and programs that read a terminal and test for end-of-file on their input will terminate appropriately.

It is possible to ask that the phone line be hung up on the last close with the TIOCHPCL ioctl; this is normally done on the outgoing lines and dialups.

Interrupt characters.

There are several characters that generate interrupts in cooked and CBREAK mode; all are sent to the processes in the control group of the terminal, as if a TIOCGPGRP ioctl were done to get the process group and then a killpg(2) system call were done, except that these characters also flush pending input and output when typed at a terminal ('a`'la TIOCFLUSH). The characters shown here are the defaults; the field names in the structures (given below) are also shown. The characters may be changed.

^C    t\_intrc (ETX) generates a SIGINT signal. This is the normal way to stop a process which is no longer interesting, or to regain control in an interactive program.

^\    t\_quitc (FS) generates a SIGQUIT signal. This is used to cause a program to terminate and produce a core image, if possible, in the file core in the current directory.

^Z    t\_suspc (EM) generates a SIGTSTP signal, which is used to suspend the current process group.

^Y    t\_dsuspc (SUB) generates a SIGTSTP signal as ^Z does, but the signal is sent when a program attempts to read the ^Y, rather than when it is typed.

Job access control.

When using the new terminal driver, if a process which is not in the distinguished process group of its control

terminal attempts to read from that terminal its process group is sent a SIGTTIN signal. This signal normally causes the members of that process group to stop. If, however, the process is ignoring SIGTTIN, has SIGTTIN blocked, or is in the middle of process creation using `vfork(2)`, the read will return -1 and set `errno` to `EIO`.

When using the new terminal driver with the `LTOSTOP` bit set in the local modes, a process is prohibited from writing on its control terminal if it is not in the distinguished process group for that terminal. Processes which are holding or ignoring SIGTTOU signals or which are in the middle of a `vfork(2)` are excepted and allowed to produce output.

Terminal/window sizes. In order to accommodate terminals and workstations with variable-sized windows, the terminal driver provides a mechanism for obtaining and setting the current terminal size. The driver does not use this information internally, but only stores it and provides a uniform access mechanism. When the size is changed, a SIGWINCH signal is sent to the terminal's process group so that knowledgeable programs may detect size changes. This facility was added in 4.3BSD and is not available in earlier versions of the system.

Summary of modes.

Unfortunately, due to the evolution of the terminal driver, there are 4 different structures which contain various portions of the driver data. The first of these (`sgttyb`) contains that part of the information largely common between version 6 and version 7 UNIX systems. The second contains additional control characters added in version 7. The third is a word of local state added in 4BSD, and the fourth is another structure of special characters added for the new driver. In the future a single structure may be made available to programs which need to access all this information; most programs need not concern themselves with all this state.

Basic modes: `sgtty`.

The basic `ioctl`s use the structure defined in `<sgtty.h>`:

```
struct sgttyb {
    char sg_ispeed;
    char sg_ospeed;
    char sg_erase;
    char sg_kill;
    short sg_flags;
};
```

The `sg_ispeed` and `sg_ospeed` fields describe the input and output speeds of the device according to the following table, which corresponds to the DEC DH-11 interface. If other hardware is used, impossible speed changes are ignored. Symbolic values in the table are as defined in `<sgtty.h>`.

B0	0	(hang up dataphone)
B50	1	50 baud
B75	2	75 baud
B110	3	110 baud
B134	4	134.5 baud
B150	5	150 baud
B200	6	200 baud
B300	7	300 baud
B600	8	600 baud
B1200	9	1200 baud
B1800	10	1800 baud
B2400	11	2400 baud
B4800	12	4800 baud
B9600	13	9600 baud
EXTA	14	External A
EXTB	15	External B

Code conversion and line control required for IBM 2741's (134.5 baud) must be implemented by the user's program. The half-duplex line discipline required for the 202 dataset (1200 baud) is not supplied; full-duplex 212 datasets work fine.

The `sg_erase` and `sg_kill` fields of the argument structure specify the erase and kill characters respectively. (Defaults are DELETE and ^U.)

The `sg_flags` field of the argument structure contains several bits that determine the system's treatment of the terminal:

XTABS	0004000	Expand tabs on output
EVENP	0000200	Even parity allowed on input
ODDP	0000100	Odd parity allowed on input
RAW	0000040	Raw mode: wake up on all characters, 8-bit interface
CRMOD	0000020	Map CR into LF; output LF as CR-LF
ECHO	0000010	Echo (full duplex)
CBREAK	0000002	Return each character as soon as typed
TANDEM	0000001	Automatic inbound xon/xoff flow control

XTABS, causes tabs to be replaced by the appropriate number of spaces on output.

The flags for even and odd parity control parity checking on input and generation on output in cooked and CBREAK mode



(unless LPASS8 is enabled, see below). Even parity is generated on output unless ODDP is set and EVENP is clear, in which case odd parity is generated. Input characters with the wrong parity, as determined by EVENP and ODDP, are ignored in cooked and CBREAK mode.

RAW disables all processing save output flushing with LFLUSHO; full 8 bits of input are given as soon as it is available; all 8 bits are passed on output. A break condition in the input is reported as a null character. If the input queue overflows in raw mode all data in the input and output queues are discarded; this applies to both new and old drivers.

CRMOD causes input carriage returns to be turned into new-lines, and output and echoed new-lines to be output as a carriage return followed by a line feed.

CBREAK is a sort of half-cooked (rare?) mode. Programs can read each character as soon as typed, instead of waiting for a full line; all processing is done except the input editing: character and word erase and line kill, input reprint, and the special treatment of \ and EOT are disabled.

TANDEM mode causes the system to produce a stop character (default ^S) whenever the input queue is in danger of overflowing, and a start character (default ^Q) when the input queue has drained sufficiently. It is useful for flow control when the 'terminal' is really another computer which understands the conventions.

Note: The same ``stop'' and ``start'' characters are used for both directions of flow control; the `t_stopc` character is accepted on input as the character that stops output and is produced on output as the character to stop input, and the `t_startc` character is accepted on input as the character that restarts output and is produced on output as the character to restart input.

#### Basic ioctls

A large number of `ioctl(2)` calls apply to terminals. Some have the general form:

```
#include <sgtty.h>
```

```
ioctl(fildes, code, arg)
struct sgttyb *arg;
```

The applicable codes are:

<code>TIOCGTP</code>	Fetch the basic parameters associated with
----------------------	--

the terminal, and store in the pointed-to `sgttyb` structure.

**TIOCSETP**           Set the parameters according to the pointed-to `sgttyb` structure. The interface delays until output is quiescent, then throws away any unread characters, before changing the modes.

**TIOCSETN**           Set the parameters like **TIOCSETP** but do not delay or flush input. Input is not preserved, however, when changing to or from RAW.

With the following codes `arg` is ignored.

**TIOCEXCL**           Set "exclusive-use" mode: no further opens are permitted until the file has been closed.

**TIOCNXCL**           Turn off "exclusive-use" mode.

**TIOCHPCL**           When the file is closed for the last time, hang up the terminal. This is useful when the line is associated with an ACU used to place outgoing calls.

With the following codes `arg` is a pointer to an int.

**TIOCGETD**           `arg` is a pointer to an int into which is placed the current line discipline number.

**TIOCSETD**           `arg` is a pointer to an int whose value becomes the current line discipline number.

**TIOCFLUSH**           If the int pointed to by `arg` has a zero value, all characters waiting in input or output queues are flushed. Otherwise, the value of the int is for the **FREAD** and **FWRITE** bits defined in `<sys/file.h>`; if the **FREAD** bit is set, all characters waiting in input queues are flushed, and if the **FWRITE** bit is set, all characters waiting in output queues are flushed.

The remaining calls are not available in vanilla version 7 UNIX. In cases where arguments are required, they are described; `arg` should otherwise be given as 0.

**TIOCSTI**           the argument points to a character which the system pretends had been typed on the terminal.

TIOCSBRK           the break bit is set in the terminal.

TIOCCBRK           the break bit is cleared.

TIOCSDTR           data terminal ready is set.

TIOCCDTR           data terminal ready is cleared.

TIOCSTOP           output is stopped as if the ``stop'' character had been typed.

TIOCSTART          output is restarted as if the ``start'' character had been typed.

TIOCGPGRP          arg is a pointer to an int into which is placed the process group ID of the process group for which this terminal is the control terminal.

TIOCSPPGRP         arg is a pointer to an int which is the value to which the process group ID for this terminal will be set.

TIOCOUTQ           returns in the int pointed to by arg the number of characters queued for output to the terminal.

FIONREAD           returns in the long pointed to by arg the number of characters immediately readable from the argument descriptor. This works for files, pipes, and terminals.

### Tchars

The second structure associated with each terminal specifies characters that are special in both the old and new terminal interfaces: The following structure is defined in `<sys/ioctl.h>`, which is automatically included in `<sgtty.h>`:

```
struct tchars {
    char t_intrc; /* interrupt */
    char t_quitc; /* quit */
    char t_startc; /* start output */
    char t_stopc; /* stop output */
    char t_eofc; /* end-of-file */
    char t_brkc; /* input delimiter (like nl) */
};
```

The default values for these characters are ^C, ^\, ^Q, ^S, ^D, and -1. A character value of -1 eliminates the effect of that character. The `t_brkc` character, by default -1, acts like a new-line in that it terminates a `line,' is

echoed, and is passed to the program. The `stop' and `start' characters may be the same, to produce a toggle effect. It is probably counterproductive to make other special characters (including erase and kill) identical. The applicable ioctl calls are:

**TIOCGETC**      Get the special characters and put them in the specified structure.

**TIOCSSETC**      Set the special characters to those given in the structure.

#### Local mode

The third structure associated with each terminal is a local mode word. The bits of the local mode word are:

LCRTBS	000001	Backspace on erase rather than echoing erase
LPRTERA	000002	Printing terminal erase mode
LCRTERA	000004	Erase character echoes as backspace-space-backspace
LMDBUF	000020	Stop/start output when carrier drops
LLITOUT	000040	Suppress output translations
LTOSTOP	000100	Send SIGTTOU for background output
LFLUSHO	000200	Output is being flushed
LNOHANG	000400	Don't send hangup when carrier drops
LRTSCTS	001000	RTS/CTS flow control
LCRTKIL	002000	BS-space-BS erase entire line on line kill
LPASS8	004000	Pass all 8 bits through on input, in any mode
LCTLECH	010000	Echo input control chars as ^X, delete as ^?
LPENDIN	020000	Retype pending input at next read or input character
LDECCTQ	040000	Only ^Q restarts output after ^S, like DEC systems
LNOFLSH	100000	Inhibit flushing of pending I/O when an interrupt character is typed.

The applicable ioctl functions are:

**TIOCLBIS**      arg is a pointer to an int whose value is a mask containing the bits to be set in the local mode word.

**TIOCLBIC**      arg is a pointer to an int whose value is a mask containing the bits to be cleared in the local mode word.

**TIOCLSET**      arg is a pointer to an int whose value is stored in the local mode word.

**TIOCLGET**      arg is a pointer to an int into which the current local mode word is placed.

#### Local special chars

The final control structure associated with each terminal is the `ltchars` structure which defines control characters for the new terminal driver. Its structure is:

```
struct ltchars {
    char t_suspc; /* stop process signal */
    char t_dsuspc; /* delayed stop process signal */
    char t_rprntc; /* reprint line */
    char t_flushc; /* flush output (toggles) */
    char t_werasc; /* word erase */
    char t_lnextc; /* literal next character */
};
```

The default values for these characters are `^Z`, `^Y`, `^R`, `^O`, `^W`, and `^V`. A value of `-1` disables the character.

The applicable `ioctl` functions are:

`TIOCSLTC`      `arg` is a pointer to an `ltchars` structure which defines the new local special characters.

`TIOCG LTC`      `arg` is a pointer to an `ltchars` structure into which is placed the current set of local special characters.

#### Window/terminal sizes

Each terminal has provision for storage of the current terminal or window size in a `winsize` structure, with format:

```
struct winsize {
    unsigned short    ws_row;      /* rows, in characters */
    unsigned short    ws_col;      /* columns, in characters */
    unsigned short    ws_xpixel;   /* horizontal size, pixels */
    unsigned short    ws_ypixel;   /* vertical size, pixels */
};
```

A value of `0` in any field is interpreted as ```undefined;''` the entire structure is zeroed on final close.

The applicable `ioctl` functions are:

`TIOCGWINSZ`  
     `arg` is a pointer to a `struct winsize` into which will be placed the current terminal or window size information.

`TIOCSWINSZ`  
     `arg` is a pointer to a `struct winsize` which will be used to set the current terminal or window size information. If the new information is different than the old information, a `SIGWINCH` signal will be sent to the terminal's process group.

## FILES

/dev/tty  
/dev/tty\*  
/dev/console

## SEE ALSO

csh(1), stty(1), tset(1), ioctl(2), sigvec(2), stty(3C),  
getty(8)

## NAME

udp - Internet User Datagram Protocol

## SYNOPSIS

```
#include <sys/socket.h>
#include <netinet/in.h>
```

```
s = socket(AF_INET, SOCK_DGRAM, 0);
```

## DESCRIPTION

UDP is a simple, unreliable datagram protocol which is used to support the SOCK\_DGRAM abstraction for the Internet protocol family. UDP sockets are connectionless, and are normally used with the sendto and recvfrom calls, though the connect(2) call may also be used to fix the destination for future packets (in which case the recv(2) or read(2) and send(2) or write(2) system calls may be used).

UDP address formats are identical to those used by TCP. In particular UDP provides a port identifier in addition to the normal Internet address format. Note that the UDP port space is separate from the TCP port space (i.e. a UDP port may not be "connected" to a TCP port). In addition broadcast packets may be sent (assuming the underlying network supports this) by using a reserved "broadcast address"; this address is network interface dependent.

Options at the IP transport level may be used with UDP; see ip(4P).

## DIAGNOSTICS

A socket operation may fail with one of the following errors returned:

- [EISCONN]           when trying to establish a connection on a socket which already has one, or when trying to send a datagram with the destination address specified and the socket is already connected;
- [ENOTCONN]          when trying to send a datagram, but no destination address is specified, and the socket hasn't been connected;
- [ENOBUFS]           when the system runs out of memory for an internal data structure;
- [EADDRINUSE]        when an attempt is made to create a socket with a port which has already been allocated;
- [EADDRNOTAVAIL]     when an attempt is made to create a socket

with a network address for which no network interface exists.

SEE ALSO

getsockopt(2), recv(2), send(2), socket(2), intro(4N),  
inet(4F), ip(4P)



## NAME

vv - Proteon proNET 10 Megabit ring

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NVV  vv_controllers # V2LNI (Pronet)
```

## DESCRIPTION

The vv interface provides access to a 10 Mb/s Proteon proNET ring network.

The network address of the interface must be specified with an SIOCSIFADDR ioctl before data can be transmitted or received. It is only permissible to change the network address while the interface is marked "down".

The host's hardware address is discovered by putting the interface in digital loopback mode (not joining the ring) and sending a broadcast packet from which the hardware address is extracted.

Transmit timeouts are detected through use of a watchdog routine. Lost input interrupts are checked for when packets are sent out.

If the installation is running CTL boards which use the old broadcast address of 0 instead of the new address of 0xff, the define OLD\_BROADCAST should be specified in the driver.

The driver can use ``trailer'' encapsulation to minimize copying data on input and output. This may be disabled, on a per-interface basis, by setting the IFF\_NOTRAILERS flag with an SIOCSIFFLAGS ioctl.

## DIAGNOSTICS

vv%d: host %d. The software announces the host address discovered during autoconfiguration.

vv%d: can't initialize. The software was unable to discover the address of this interface, so it deemed "dead" will not be enabled.

vv%d: error vvcsr=%b. The hardware indicated an error on the previous transmission.

vv%d: output timeout. The token timer has fired and the token will be recreated.

vv%d: error vvicsr=%b. The hardware indicated an error in reading a packet off the ring.

en%d: can't handle af%d. The interface was handed a message with addresses formatted in an unsuitable address family; the packet was dropped.

vv%d: vs\_olen=%d. The ring output routine has been handed a message with a preposterous length. This results in an immediate panic: vs\_olen.

#### SEE ALSO

intro(4N), inet(4F)

#### BUGS

The encapsulation of trailer packets in the 4.2BSD version of this driver was incorrect (the packet type was in VAX byte order). As a result, the trailer encapsulation in this version is not compatible with the 4.2BSD VAX version.

## NAME

xp - generic SMD moving-head disk

## SYNOPSIS

```
/sys/conf/SYSTEM:
    NXPC          xp_controllers # Number of controllers
    NXPD          xp_drives  # RM02/03/05, RP04/05/06,
                          # CDC 9766, Fuji 160, etc.
    BADSECT NO          # Bad sector handling (see BUGS)

/etc/dtab:
    #Name Unit# Addr   Vector Br Handler(s)      # Comments
    xp ?      176700 254   5  xpintr      # xp driver

major device number(s):
    raw: 19
    block: 10
minor device encoding:
    bits 0007 specify partition of XP drive
    bits 0370 specify XP drive
```

## DESCRIPTION

The xp driver is a generic SMD storage module disk driver. It can be adapted to most SMD controllers although bootstrapping will not necessarily be directly possible. The drives are numbered from 0 to n on controller 0, from n+1 to m on controller 1, etc. The drives may have different geometries.

The xp driver is unique amongst 2BSD drivers in its numbering of drives. Other drivers (ra for example) number drives 0 thru 7 on controller 1, 8 thru 15 on controller 2 and so on. xp on the other hand can have drives 0 and 1 on controller 1, drives 2, 3, 4 and 5 on controller 2 and drives 6, 7 and 8 on controller 3. This is different from boot's view of the world, so if you are booting from other than unit 0 you may have to experiment a bit.

Files with minor device numbers 0 through 7 refer to various portions of drive 0; minor devices 8 through 15 refer to drive 1, etc. The standard device names begin with ``xp'' followed by the drive number and then a letter a-h for partitions 0-7 respectively. The character ? stands here for a drive number in the range 0-7.

The block files access the disk via the system's normal buffering mechanism and may be read and written without regard to physical disk records. There is also a `raw' interface which provides for direct transmission between the disk and the user's read or write buffer. A single read or write call results in exactly one I/O operation and therefore raw I/O is considerably more efficient when many words

are transmitted. The names of the raw files conventionally begin with an extra `r.'

In raw I/O the buffer must begin on a word (even) boundary, and counts should be a multiple of 512 bytes (a disk sector). Likewise lseek calls should specify a multiple of 512 bytes.

#### DISK SUPPORT

Disks must be labeled using either the standalone disklabel program on the boot tape or with the disklabel(8) program. The kernel no longer attempts to determine the drive type and geometry, instead reading this information from the disklabel. There are no partition tables coded into the xp driver, these must be placed on the drive with disklabel.

Special files should only be created for the partitions that are actually used, as the overlap in these addresses could lead to confusion otherwise. Traditionally the xp?a partition is normally used for the root file system, the xp?b partition as a swap area, and the xp?c partition for pack-pack copying (it maps the entire disk).

#### FILES

/dev/xp[0-7][a-h] block files  
/dev/rxp[0-7][a-h] raw files  
/dev/MAKEDEV script to create special files  
/dev/MAKEDEV.local script to localize special files

#### SEE ALSO

hk(4), ra(4), ram(4), rk(4), rl(4), rp(4), rx(4), si(4),  
dtab(5), autoconfig(8), newfs(8)

#### DIAGNOSTICS

xp%d%c: hard error sn%d cs2=%b erl=%b. An unrecoverable error occurred during transfer of the specified sector of the specified disk partition. The contents of the two error registers are also printed in octal and symbolically with bits decoded. The error was either unrecoverable, or a large number of retry attempts (including offset positioning and drive recalibration) could not recover the error.

xp%d: write locked. The write protect switch was set on the drive when a write was attempted. The write operation is not recoverable.

xp%d%c: soft ecc sn%d. A recoverable ECC error occurred on the specified sector of the specified disk partition. This happens normally a few times a week. If it happens more frequently than this the sectors where the errors are occurring should be checked to see if certain cylinders on the pack, spots on the carriage of the drive or heads are

indicated.

xp%d: unknown device type 0%o. The number in the drive's drive type register is unknown to the xp driver.

#### BUGS

In raw I/O read and write(2) truncate file offsets to 512-byte block boundaries, and write scribbles on the tail of incomplete blocks. Thus, in programs that are likely to access raw devices, read, write and lseek(2) should always deal in 512-byte multiples.

DEC-standard error logging should be supported.

The kernel uses partition 'h' to access the badblock information. This should have been 'c' except that almost all of the /etc/disktab entries (and thus existing systems) use 'h' for this purpose. Unless you are very careful with disklabel(8) (to make certain that no data partition overlaps the badblock area) you should probably leave BADSECT undefined in the kernel config file.